

---

# Advanced CUDA programming: asynchronous execution, memory models, unified memory

---

January 2020

Caroline Collange

Inria Rennes – Bretagne Atlantique

<https://team.inria.fr/pacap/members/collange/>

caroline.collange@inria.fr



# Agenda

---

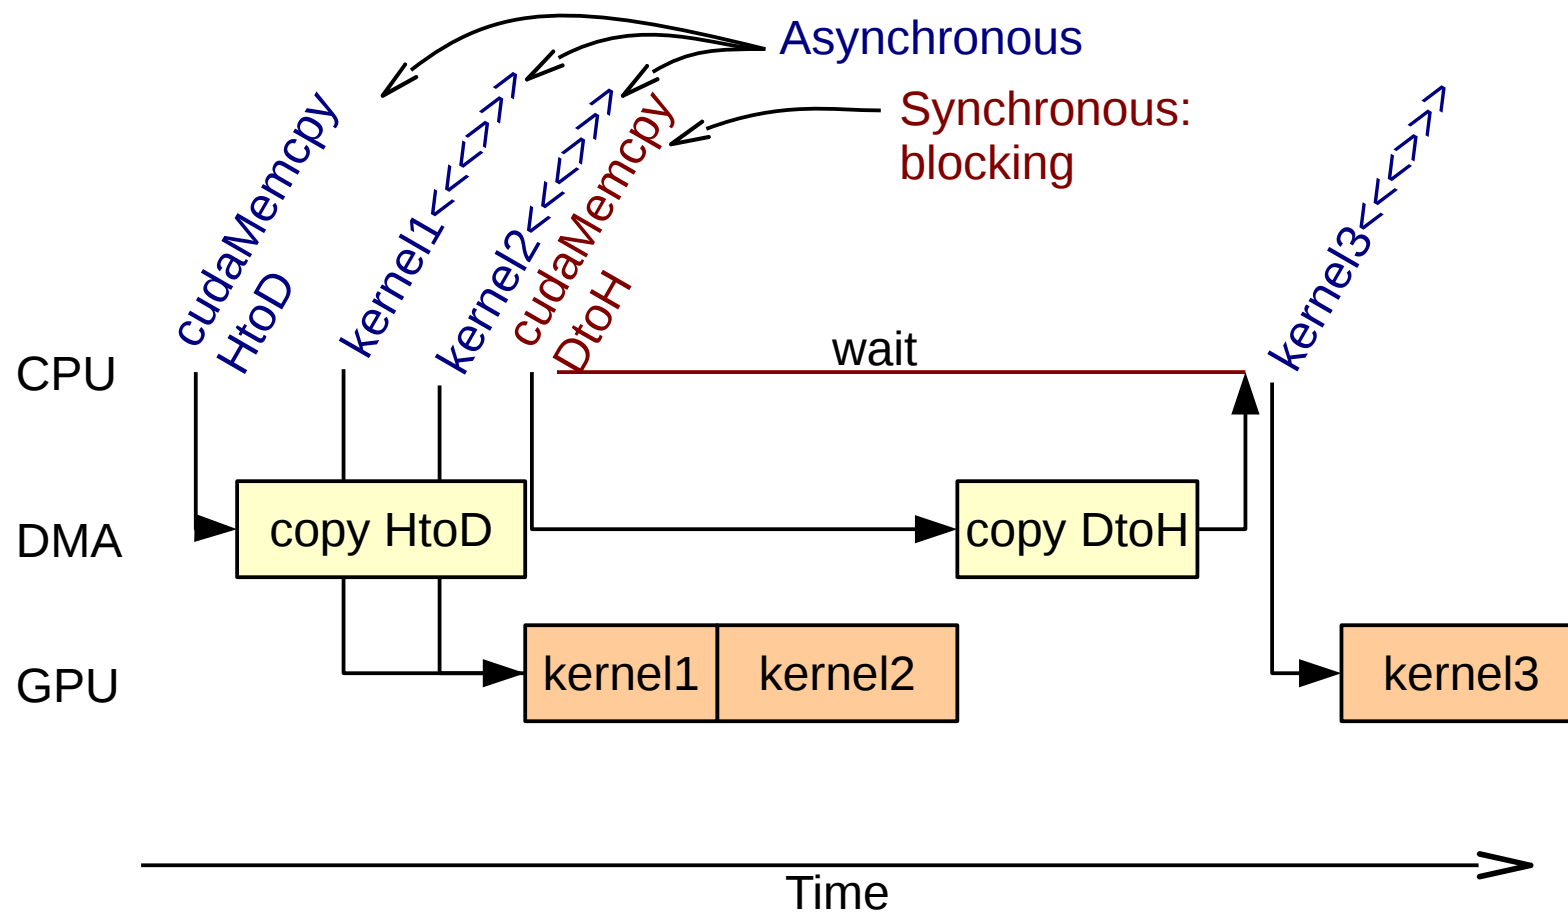
- Asynchronous execution
  - ◆ Streams
  - ◆ Task graphs
- Fine-grained synchronization
  - ◆ Atomics
  - ◆ Memory consistency model
- Unified memory
  - ◆ Memory allocation
  - ◆ Optimizing transfers

# Asynchronous execution

- By default, most CUDA function calls are *asynchronous*
  - ◆ Returns immediately to CPU code
  - ◆ GPU commands are queued and executed in-order
- Some commands are synchronous by default
  - ◆ `cudaMemcpy(..., cudaMemcpyDeviceToHost)`
  - ◆ Asynchronous version: `cudaMemcpyAsync`
- Keep it in mind when checking for errors and measuring timing!
  - ◆ Error returned by a command may be caused by an earlier command
  - ◆ Time taken by `kernel<<<>>>` launch is meaningless
- To force synchronization: `cuThreadSynchronize()`

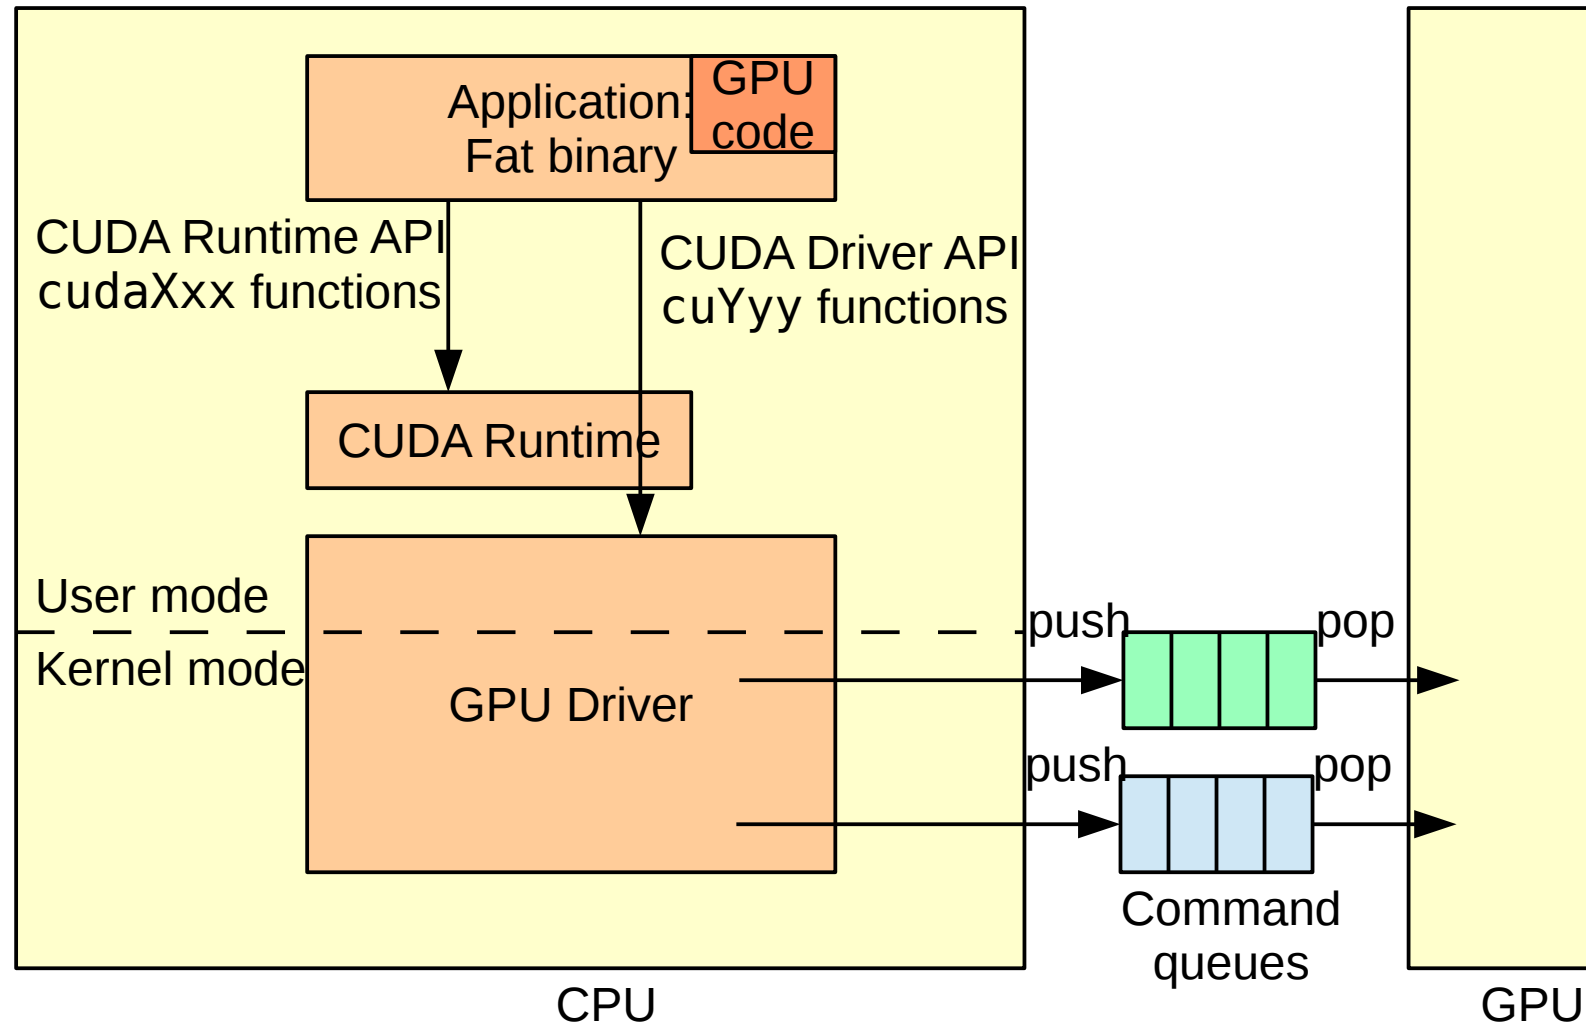
# Asynchronous transfers

- Overlap CPU work with GPU work



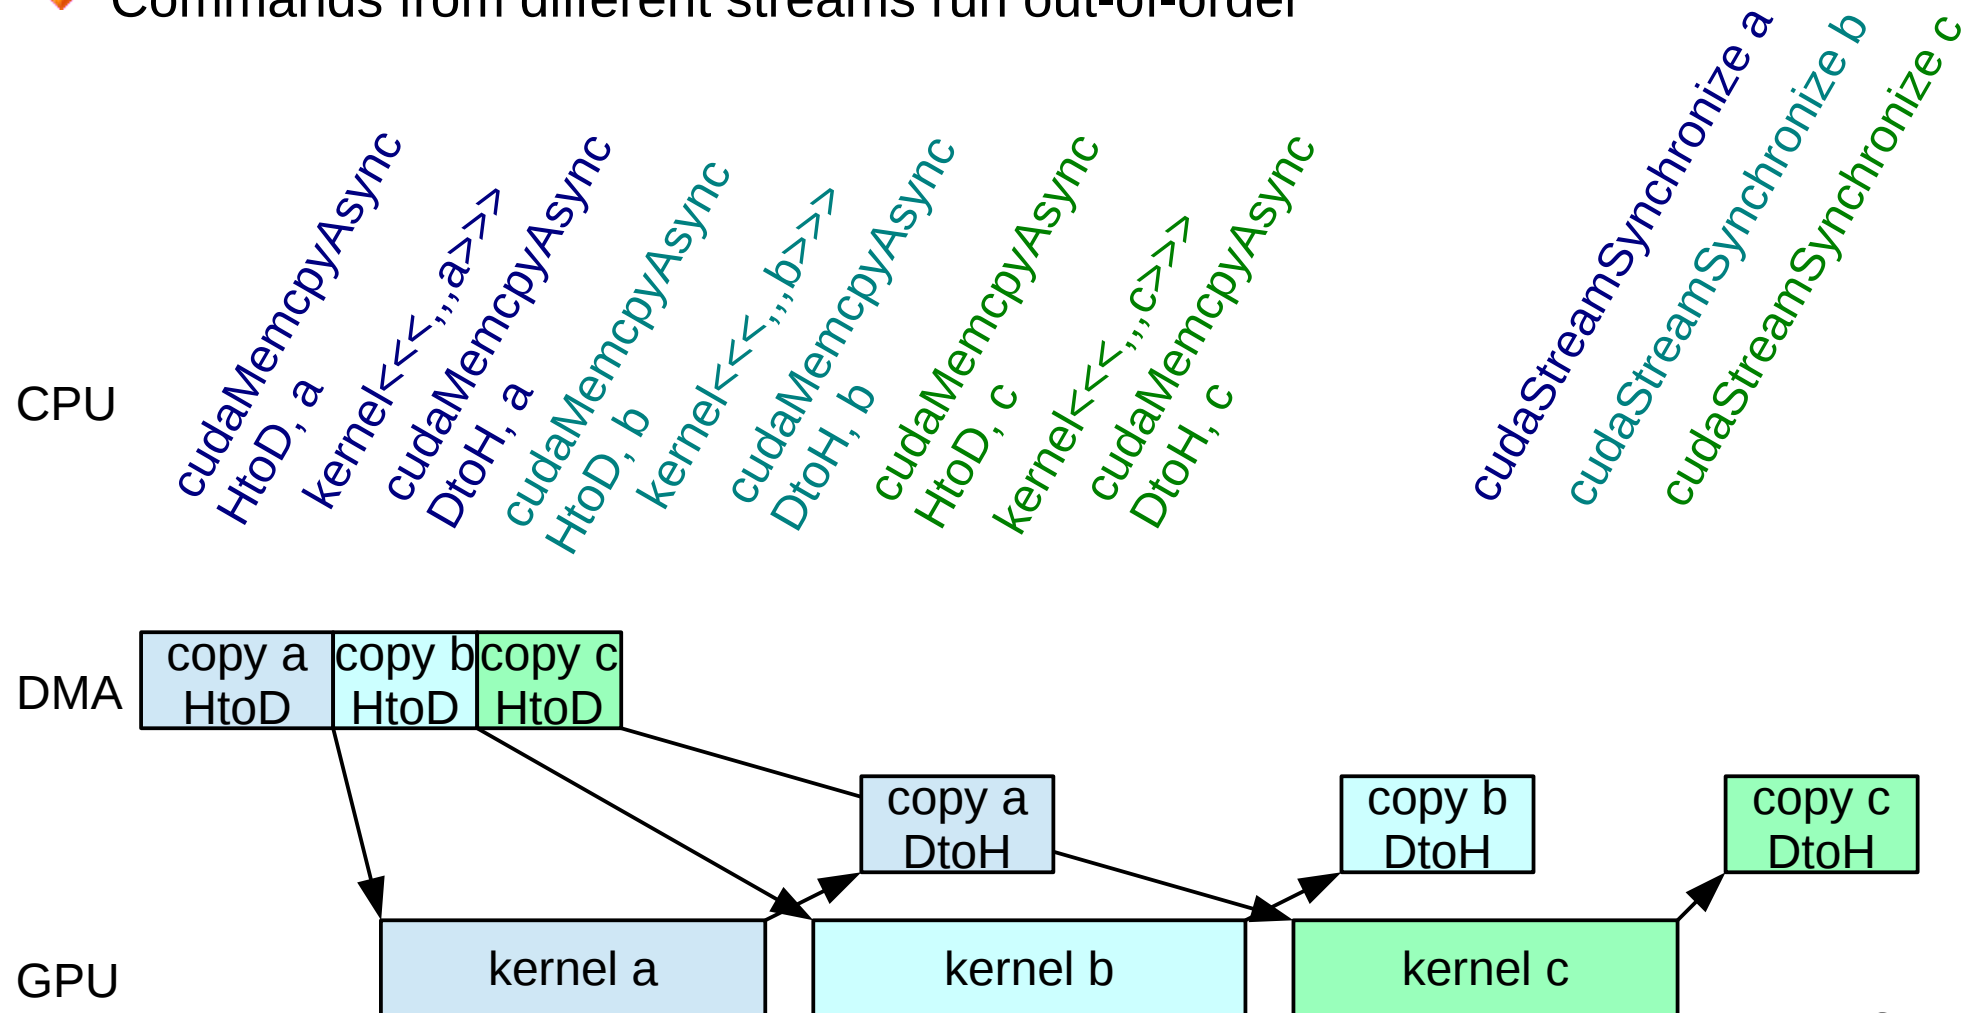
- Can we do better?

# Multiple command queues / streams



# Streams: pipelining commands

- *Command queues* in OpenCL
  - ◆ Commands from the same stream run in-order
  - ◆ Commands from different streams run out-of-order



# Streams: benefits


- Overlap CPU-GPU communication and computation:  
Direct Memory Access (DMA) copy engine  
runs CPU-GPU memory transfers in background
  - ◆ Requires **page-locked memory**
  - ◆ Some Tesla GPUs have 2 DMA engines or more:  
simultaneous send + receive + inter-GPU communication
- Concurrent kernel execution
  - ◆ Start next kernel before previous kernel finishes
  - ◆ Mitigates impact of load imbalance / tail effect

Example


Kernel<<<5,,,a>>>

Kernel<<<4,,,b>>>

Serial kernel execution



Concurrent kernel execution



# Page-locked memory

- By default, allocated memory is *pageable*
  - ◆ Can be swapped out to disk, moved by the OS...
- DMA transfers are only safe on *page-locked* memory
  - ◆ Fixed virtual → physical mapping
  - ◆ `cudaMemcpy` needs an intermediate copy:  
slower, **synchronous only**
- **cudaMallocHost** allocates page-locked memory
  - ◆ Mandatory when using streams
- Warning: page-locked memory is a limited resource!



# Streams: example

- Send data, execute, receive data

```
cudaStream_t stream[2];
for (int i = 0; i < 2; ++i)
    cudaStreamCreate(&stream[i]);

float* hostPtr;
cudaMallocHost(&hostPtr, 2 * size);

for (int i = 0; i < 2; ++i) {
    cudaMemcpyAsync(inputDevPtr + i * size, hostPtr + i * size,
                    size, cudaMemcpyHostToDevice, stream[i]);

    MyKernel <<<100, 512, 0, stream[i]>>>
        (outputDevPtr + i * size, inputDevPtr + i * size, size);

    cudaMemcpyAsync(hostPtr + i * size, outputDevPtr + i * size,
                    size, cudaMemcpyDeviceToHost, stream[i]);
}

for (int i = 0; i < 2; ++i)
    cudaStreamDestroy(stream[i]);
```

# Streams: alternative implementation

```
cudaStream_t stream[2];
for (int i = 0; i < 2; ++i)
    cudaStreamCreate(&stream[i]);
float* hostPtr;
cudaMallocHost(&hostPtr, 2 * size);

for (int i = 0; i < 2; ++i)
    cudaMemcpyAsync(inputDevPtr + i * size, hostPtr + i * size,
        size, cudaMemcpyHostToDevice, stream[i]);
for (int i = 0; i < 2; ++i)
    MyKernel<<<100, 512, 0, stream[i]>>>
        (outputDevPtr + i * size, inputDevPtr + i * size, size);
for (int i = 0; i < 2; ++i)
    cudaMemcpyAsync(hostPtr + i * size, outputDevPtr + i * size,
        size, cudaMemcpyDeviceToHost, stream[i]);


for (int i = 0; i < 2; ++i)
    cudaStreamDestroy(stream[i]);
```

- Which one is better?

# Events: synchronizing streams

- Schedule synchronization of one stream with another
  - ◆ Specify dependencies between tasks


```
cudaEvent_t e;  
cudaEventCreate(&e);  
kernel1<<<,, ,a>>>();  
cudaEventRecord(e, a);  
cudaStreamWaitEvent(b, e);  
kernel2<<<,, ,b>>>();  
  
cudaEventDestroy(e);
```



- Measure timing


```
cudaEventRecord(start, 0);  
kernel<<<>>>();  
cudaEventRecord(stop, 0);  
cudaEventSynchronize(stop);
```

```
float elapsedTime;  
cudaEventElapsedTime(&elapsedTime, start, stop);
```



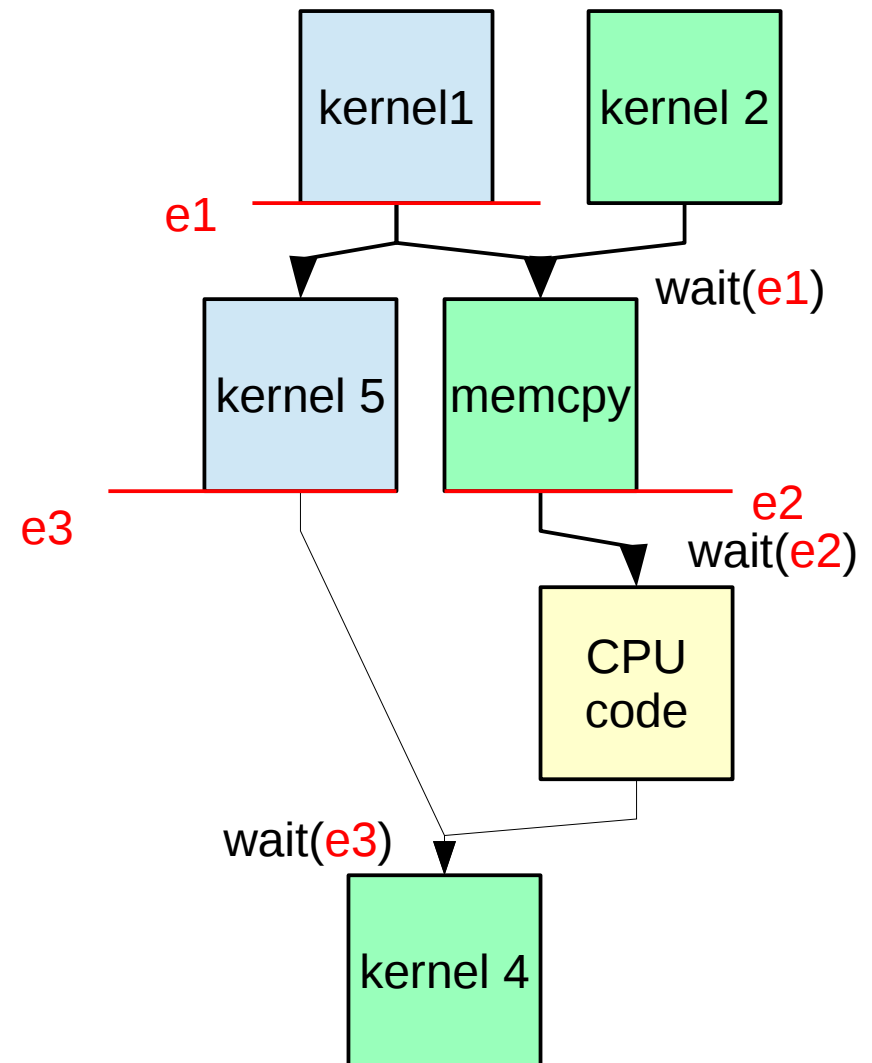
# Scheduling data dependency graphs

- With streams and events, we can express task dependency graphs
  - Equivalent to threads and events (e.g. semaphores) on CPU
- Example:
  - 2 GPU streams: a b
  - and 1 CPU thread:
  - Where should we place events?



# Scheduling data dependency graphs

```
kernel1<<<,,,a>>>();  
cudaEventRecord(e1, a);  
  
kernel2<<<,,,b>>>();  
  
cudaStreamWaitEvent(b, e1);  
  
cudaMemcpyAsync(,,,,b);  
  
cudaEventRecord(e2, b);  
  
kernel3<<<,,,a>>>();  
cudaEventRecord(e3, a);  
  
cudaEventSynchronize(e2);  
  
CPU code  
  
cudaStreamWaitEvent(b, e3);  
kernel4<<<,,,b>>>();
```



# Agenda

---

- Asynchronous execution
  - ◆ Streams
  - ◆ Task graphs
- Fine-grained synchronization
  - ◆ Atomics
  - ◆ Memory consistency model
- Unified memory
  - ◆ Memory allocation
  - ◆ Optimizing transfers

# From streams and events to task graphs

- Limitations of scheduling task graphs with streams
  - ◆ Sub-optimal scheduling : GPU runtime has no vision of tasks ahead
  - ◆ Must pay various initialization overheads when launching each task
- **New** alternative since CUDA 10.0: cudaGraph API
  - ◆ Build an in-memory representation of the dependency graph offline
  - ◆ Let the CUDA runtime optimize and schedule the task graph
  - ◆ Launch the optimized graph as needed
- Two ways we can build the dependency graph
  - ◆ Record a sequence of asynchronous CUDA calls
  - ◆ Describe the graph explicitly

# Recording asynchronous CUDA calls

- Capture a sequence of calls on a stream, **instead** of executing them

```
cudaGraph_t graph;  
cudaStreamBeginCapture(stream);
```

... CUDA calls on **stream**

```
cudaStreamEndCapture(stream, &graph);
```

- Good for converting existing asynchronous code to task graphs
  - ◆ (Almost) no code rewrite required
- Supports any number of streams (except default stream 0)
  - ◆ Follows dependencies to other streams through events
  - ◆ Capture all streams that have dependency with first captured stream
- Need all recorded calls to be asynchronous and bound to a stream
  - ◆ CPU code needs to be asynchronous to be recorded too!



# Recording asynchronous CUDA calls

- Surround asynchronous code with capture calls

```
cudaGraph_t graph;  
cudaStreamBeginCapture(a);
```

```
kernel1<<<,,,a>>>();  
cudaEventRecord(e1, a);  
kernel2<<<,,,b>>>();  
cudaStreamWaitEvent(b, e1);  
cudaMemcpyAsync(,,,b);
```

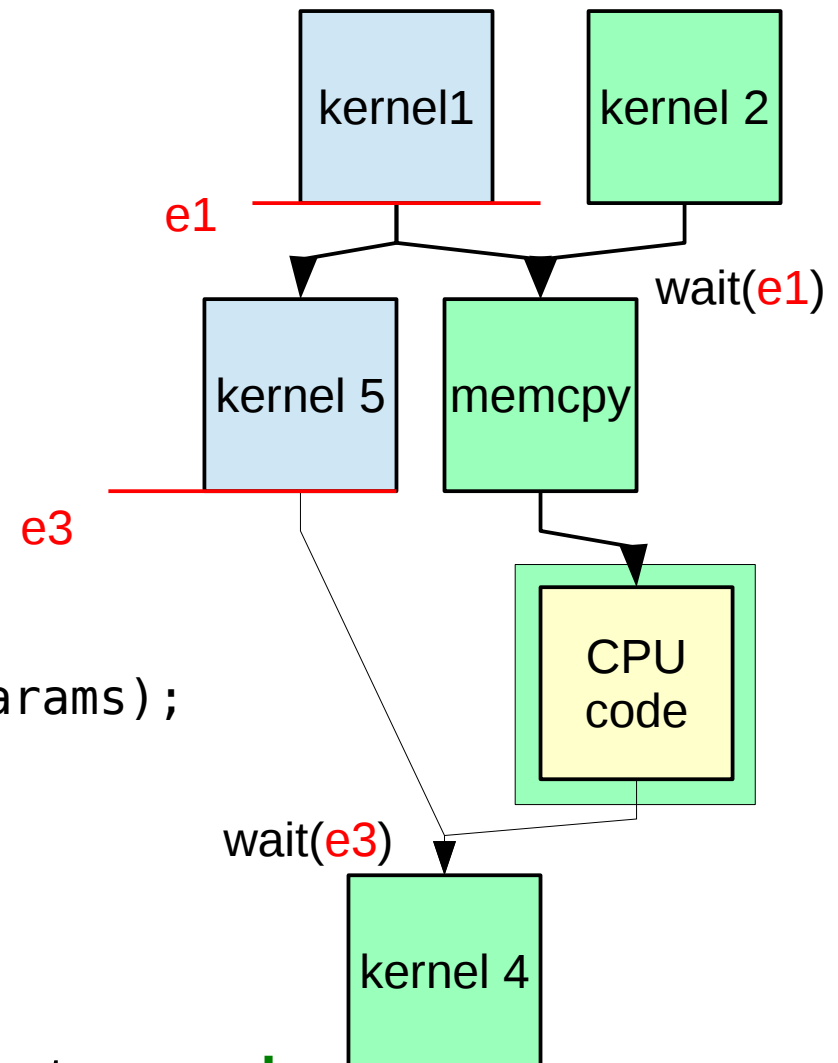
```
kernel3<<<,,,a>>>();  
cudaEventRecord(e3, a);
```

```
cudaLaunchHostFunc(b, cpucode, params);
```

```
cudaStreamWaitEvent(b, e3);  
kernel4<<<,,,b>>>();
```

```
cudaStreamEndCapture(a, &graph);
```

- Will capture stream **a**, and dependent streams: **b**



# Recording asynchronous CUDA calls

- Records only asynchronous calls: can't use immediate synchronization

```
cudaGraph_t graph;  
cudaStreamBeginCapture(a);
```

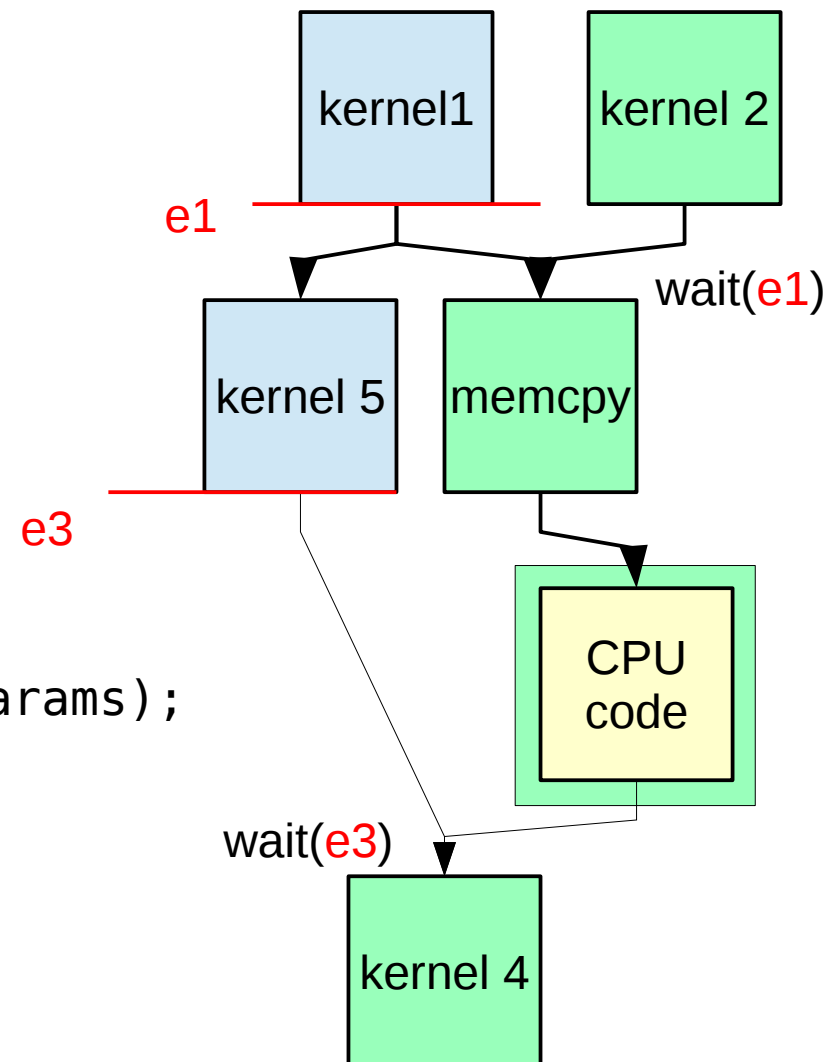
```
kernel1<<<,,,a>>>();  
cudaEventRecord(e1, a);  
kernel2<<<,,,b>>>();  
cudaStreamWaitEvent(b, e1);  
cudaMemcpyAsync(,,,b);
```

```
kernel3<<<,,,a>>>();  
cudaEventRecord(e3, a);
```

```
cudaLaunchHostFunc(b, cpucode, params);
```

```
cudaStreamWaitEvent(b, e3);  
kernel4<<<,,,b>>>();
```

```
cudaStreamEndCapture(a, &graph);
```




- Make the call to CPU code asynchronous too, on stream **b** using **cudaLaunchHostFunc**

# Describing the graph explicitly

- Add nodes to the graph: kernels, memcpy, host call...

```
cudaGraph_t graph;  
cudaGraphCreate(&graph, 0);  
  
cudaGraphNode_t k1,k2,k3,k4,mc,cpu;  
  
cudaGraphAddKernelNode(&k1, graph,  
    0, 0, // no dependency yet  
    paramsK1, 0);  
...  
cudaGraphAddKernelNode(&k4, graph,  
    0, 0, paramsK4, 0);  
  
cudaGraphAddMemcpyNode(&mc, graph,  
    0, 0, paramsMC);  
  
cudaGraphAddHostNode(&cpu, graph,  
    0, 0, paramsCPU);
```



# Passing kernel parameters

- Node creation functions take parameters as a structure pointer
- e.g. for kernel calls

```
__host__ cudaError_t
cudaGraphAddKernelNode(cudaGraphNode_t* pGraphNode,
    cudaGraph_t graph, const cudaGraphNode_t* pDependencies,
    size_t numDependencies, const cudaKernelNodeParams* pNodeParams)

struct cudaKernelNodeParams
{
    void* func;                // Kernel function
    dim3 gridDim;
    dim3 blockDim;
    unsigned int sharedMemBytes;
    void **kernelParams;      // Array of pointers to arguments
    void **extra;             // (low-level alternative to kernelParams)
};
```

- kernelParams point to memory that will contain parameters when the graph is eventually executed

# Describing the graph explicitly

- Add dependencies between nodes

```
cudaGraph_t graph;  
cudaGraphCreate(&graph, 0);
```

```
cudaGraphNode_t k1,k2,k3,k4,mc,cpu;
```

... Add nodes

```
cudaGraphAddDependencies(graph,  
    &k1, &k3, 1); // kernel1 → kernel3
```


```
cudaGraphAddDependencies(graph,  
    &k1, &mc, 1); // kernel1 → memcpy
```

```
cudaGraphAddDependencies(graph,  
    &k2, &mc, 1); // kernel2 → memcpy
```

```
cudaGraphAddDependencies(graph,  
    &mc, &cpu, 1); // memcpy → cpu
```

```
cudaGraphAddDependencies(graph,  
    &k3, &k4, 1); // kernel3 → kernel4
```

```
cudaGraphAddDependencies(graph,  
    &cpu, &k4, 1); // cpu → kernel4
```



# Instantiating and running the graph

- Instantiate the graph to create an executable graph
- Launch executable graph on a stream

```
cudaGraph_t graph;
```

```
... build or record graph
```

```
cudaGraphExec_t exec;
```

```
cudaGraphInstantiate(&exec, graph, 0, 0, 0);
```

```
cudaGraphLaunch(exec, stream);
```

```
cudaStreamSynchronize(stream);
```

- Once a graph is instantiated, its topology cannot be changed
- Kernel/memcpy/call... **parameters** can still be changed using **cudaGraphExecUpdate** or **cudaGraphExec{Kernel,Host,Memcpy,Memset}NodeSetParams**

# Agenda

---

- Asynchronous execution
  - ◆ Streams
  - ◆ Scheduling dependency graphs
- Fine-grained synchronization
  - ◆ Atomics
  - ◆ Memory consistency model
- Unified memory
  - ◆ Memory allocation
  - ◆ Optimizing transfers

# Inter-thread/inter-warp communication

- Barrier: simplest form of synchronization
  - ◆ Easy to use
  - ◆ Coarse-grained
- Atomics
  - ◆ Fine-grained
  - ◆ Can implement *wait-free* algorithms
  - ◆ May be used for blocking algorithms (locks) among **warps** of the **same block**
  - ◆ From Volta (CC $\geq$ 7.0), also among threads of the same warp
- Communication through memory
  - ◆ Beware of consistency




# Atomics

---

- Read, modify, write in one operation
  - ◆ Cannot be mixed with accesses from other thread
- Available operators
  - ◆ Arithmetic: `atomic{Add,Sub,Inc,Dec}`
  - ◆ Min-max: `atomic{Min,Max}`
  - ◆ Synchronization primitives: `atomic{Exch,CAS}`
  - ◆ Bitwise: `atomic{And,Or,Xor}`
- On global memory, and shared memory
- Performance impact in case of contention
  - ◆ Atomic operations to the same address are serialized

# Example: reduction

- After local reduction inside each block, use atomics to accumulate the result in global memory



- Complexity?
- Time including kernel launch overhead?


# Example: compare-and-swap

- Use case: perform an arbitrary associative and commutative operation atomically on a single variable
- `atomicCAS(p, old, new)` does atomically
  - ◆ if `*p == old` then assign `*p ← new`, return old
  - ◆ else return `*p`

```
__shared__ unsigned int data;
unsigned int old = data;           // Read once
unsigned int assumed;
do {
    assumed = old;
    newval = f(old, x);             // Compute
    old = atomicCAS(&data, old, newval); // Try to replace
} while(assumed != old);           // If failed, retry
```

# Memory consistency model

- Nvidia GPUs (and compiler) implement a relaxed consistency model
  - ◆ No global ordering between memory accesses
  - ◆ Threads may not see the writes/atomics in the same order



- Need to enforce explicit ordering

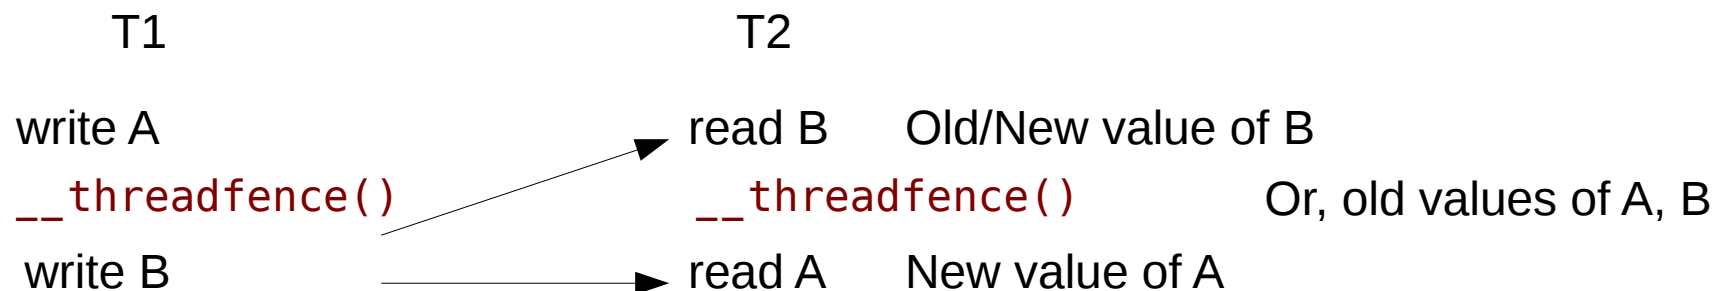
Details and precise specifications of CUDA memory consistency model:

<http://docs.nvidia.com/cuda/parallel-thread-execution/index.html#memory-consistency-model>

# Enforcing memory ordering: fences

- `__threadfence_block`  
`__threadfence`  
`__threadfence_system`

- ◆ Make writes preceding the fence appear before writes following the fence for the other threads at the block / device / system level
- ◆ Make reads preceding the fence happen after reads following the fence



- Declare shared variables as **volatile** to make writes visible to other threads (prevents compiler from removing “redundant” read/writes)
- `__syncthreads` implies `__threadfence_block`

# Floating-point atomics

- `atomicAdd` supports floating-point operands
- Remember
  - ◆ Floating-point addition is not associative
  - ◆ Thread scheduling is not deterministic
- Without FP atomics
  - ◆ Evaluation order is independent of thread scheduling
  - ◆ You should expect deterministic result for a fixed combination of GPU, driver, and runtime parameters (thread block size, etc.)
- With FP atomics
  - ◆ Evaluation order depends on thread scheduling
  - ◆ You will get different answers from one run to the other

# Agenda

---

- Asynchronous execution
  - ◆ Streams
  - ◆ Scheduling dependency graphs
- Fine-grained synchronization
  - ◆ Atomics
  - ◆ Memory consistency model
- Unified memory
  - ◆ Memory allocation
  - ◆ Optimizing transfers

# Unified memory and other features

OK, we lied

- You were told
  - ◆ CPU and GPU have distinct memory spaces
  - ◆ Blocks cannot communicate
  - ◆ You need to synchronize threads inside a block
- This is the least-common denominator across all CUDA GPUs
- This is not true (any more). We now have:
  - ◆ Device-mapped, Unified virtual address space, Unified memory
  - ◆ Global and shared memory atomics
  - ◆ Dynamic parallelism
  - ◆ Warp-synchronous programming with intra-block thread groups
  - ◆ Grid level and multi-grid level thread groups



# Unified memory

---

- Allocate memory using `cudaMallocManaged`
  - ◆ Pointer is accessible from both CPU and GPU
  - ◆ The CUDA runtime will take care of the transfers
  - ◆ No need for `cudaMemcpy` any more
  - ◆ Behaves **as if** you had a single memory space
- Suboptimal performance: on-demand page migration
  - ◆ Optimization: perform copies in advance using `cudaMemPrefetchAsync` when access patterns are known

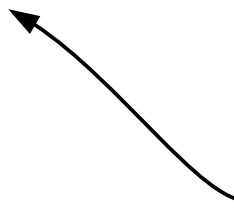
# VectorAdd using unified memory

```
int main()
{
    int numElements = 50000;
    size_t size = numElements * sizeof(float);

    float *A, *B, *C;
    cudaMallocManaged((void **)&A, size);
    cudaMallocManaged((void **)&B, size);
    cudaMallocManaged((void **)&C, size);
    Initialize(A, B);

    int blocks = numElements;
    vectorAdd2<<<blocks, 1>>>(A, B, C);
    cudaDeviceSynchronize();
    Display(C);


    cudaFree(A);
    cudaFree(B);
    cudaFree(C);
}
```



Explicit CPU-GPU synchronization  
is now mandatory! Do not forget it.  
(*We all make this mistake once*)

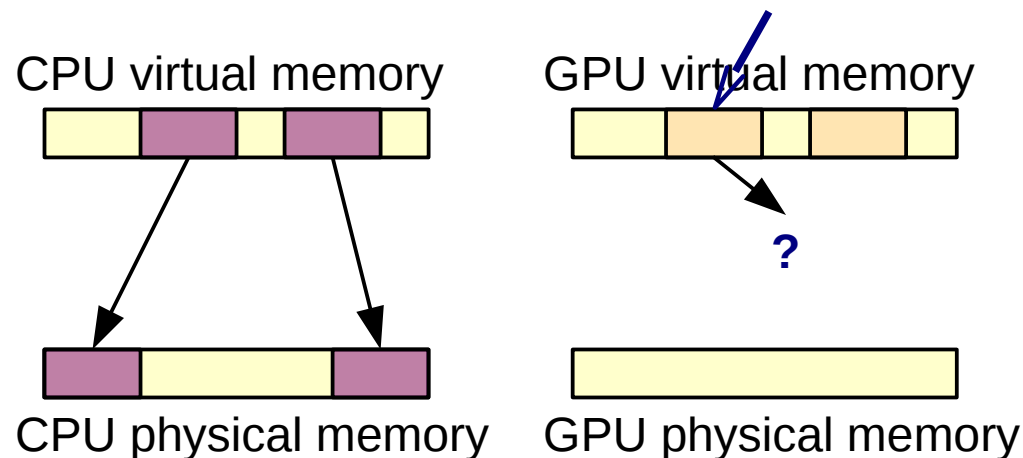
# How it works: on-demand paging

- Managed memory pages mapped in both CPU and GPU spaces
  - ♦ Same virtual address
  - ♦ Not necessarily allocated in physical memory
- Typical flow
  - ♦ **1.** Data allocated in CPU memory




# How it works: on-demand paging

- Managed memory pages mapped in both CPU and GPU spaces
  - ♦ Same virtual address
  - ♦ Not necessarily allocated in physical memory
- Typical flow
  - ♦ **1.** Data allocated in CPU memory
  - ♦ **2.** GPU code touches unallocated page, triggers page fault




# How it works: on-demand paging

- Managed memory pages mapped in both CPU and GPU spaces
  - ♦ Same virtual address
  - ♦ Not necessarily allocated in physical memory
- Typical flow
  - ♦ **1.** Data allocated in CPU memory
  - ♦ **2.** GPU code touches unallocated page, triggers page fault
  - ♦ **3.** Page fault handler allocates page in GPU mem, copies contents



# How it works: on-demand paging

- Managed memory pages mapped in both CPU and GPU spaces
  - ♦ Same virtual address
  - ♦ Not necessarily allocated in physical memory
- Typical flow
  - ♦ **1.** Data allocated in CPU memory
  - ♦ **2.** GPU code touches unallocated page, triggers page fault
  - ♦ **3.** Page fault handler allocates page in GPU mem, copies contents
  - ♦ **4.** If GPU modifies page contents, invalidate CPU copy  
Next CPU access will cause data to be copied back from GPU mem



# Prefetching data

On-demand paging has overhead

- Solution: load data in advance using **cudaMemPrefetchAsync**

```
...  
cudaStream_t stream;  
cudaStreamCreate(&stream);  
  
cudaMemPrefetchAsync(A, size, gpuId, stream);  
cudaMemPrefetchAsync(B, size, gpuId, stream);  
  
vectorAdd2<<<numElements, 1, 0, stream>>>(A, B, C);  
  
cudaMemPrefetchAsync(C, size, cudaCpuDeviceId, stream);  
  
cudaDeviceSynchronize();  
Display(C);  
...
```

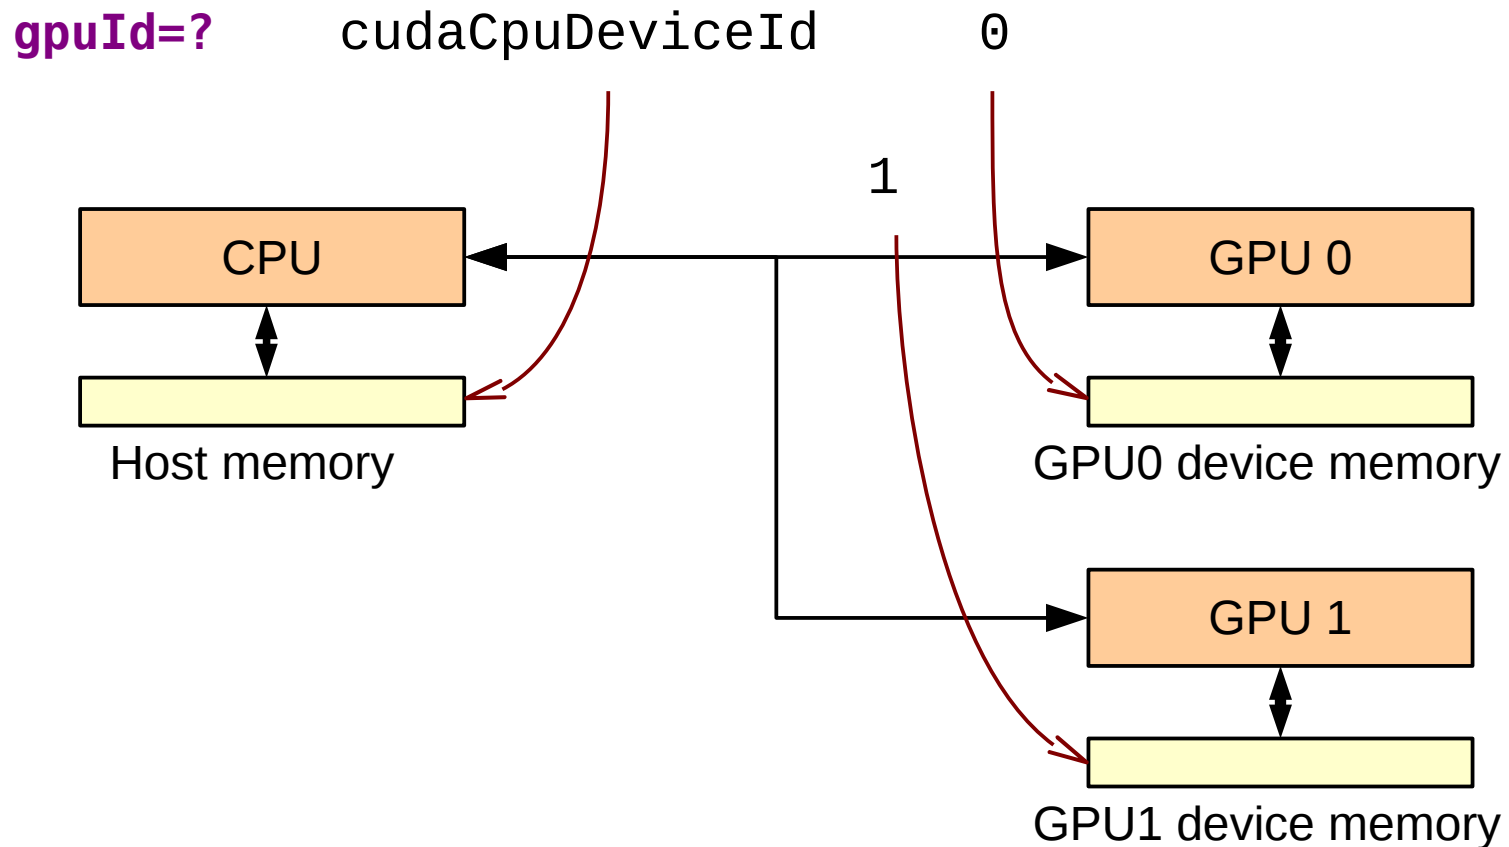
- Performance similar to manual memory management
  - ◆ Supports asynchronous copies, tolerates sloppy synchronization

# Controlling data placement

- System may have multiple GPU memory spaces

- ◆ Specify destination of prefetch

```
cudaMemPrefetchAsync(A, size, gpuId, s);
```





# References

---

- Nikolay Sakharnykh. *Maximizing Unified Memory Performance in CUDA*.  
|| $\nabla$   
<https://devblogs.nvidia.com/parallelforall/maximizing-unified-memory-performance-cuda/>