



Speaking Tensor Cores

Vijay Thakkar & Pradeep Ramani
(Representing the CUTLASS Team @NVIDIA)

June 7th, 2024



Agenda

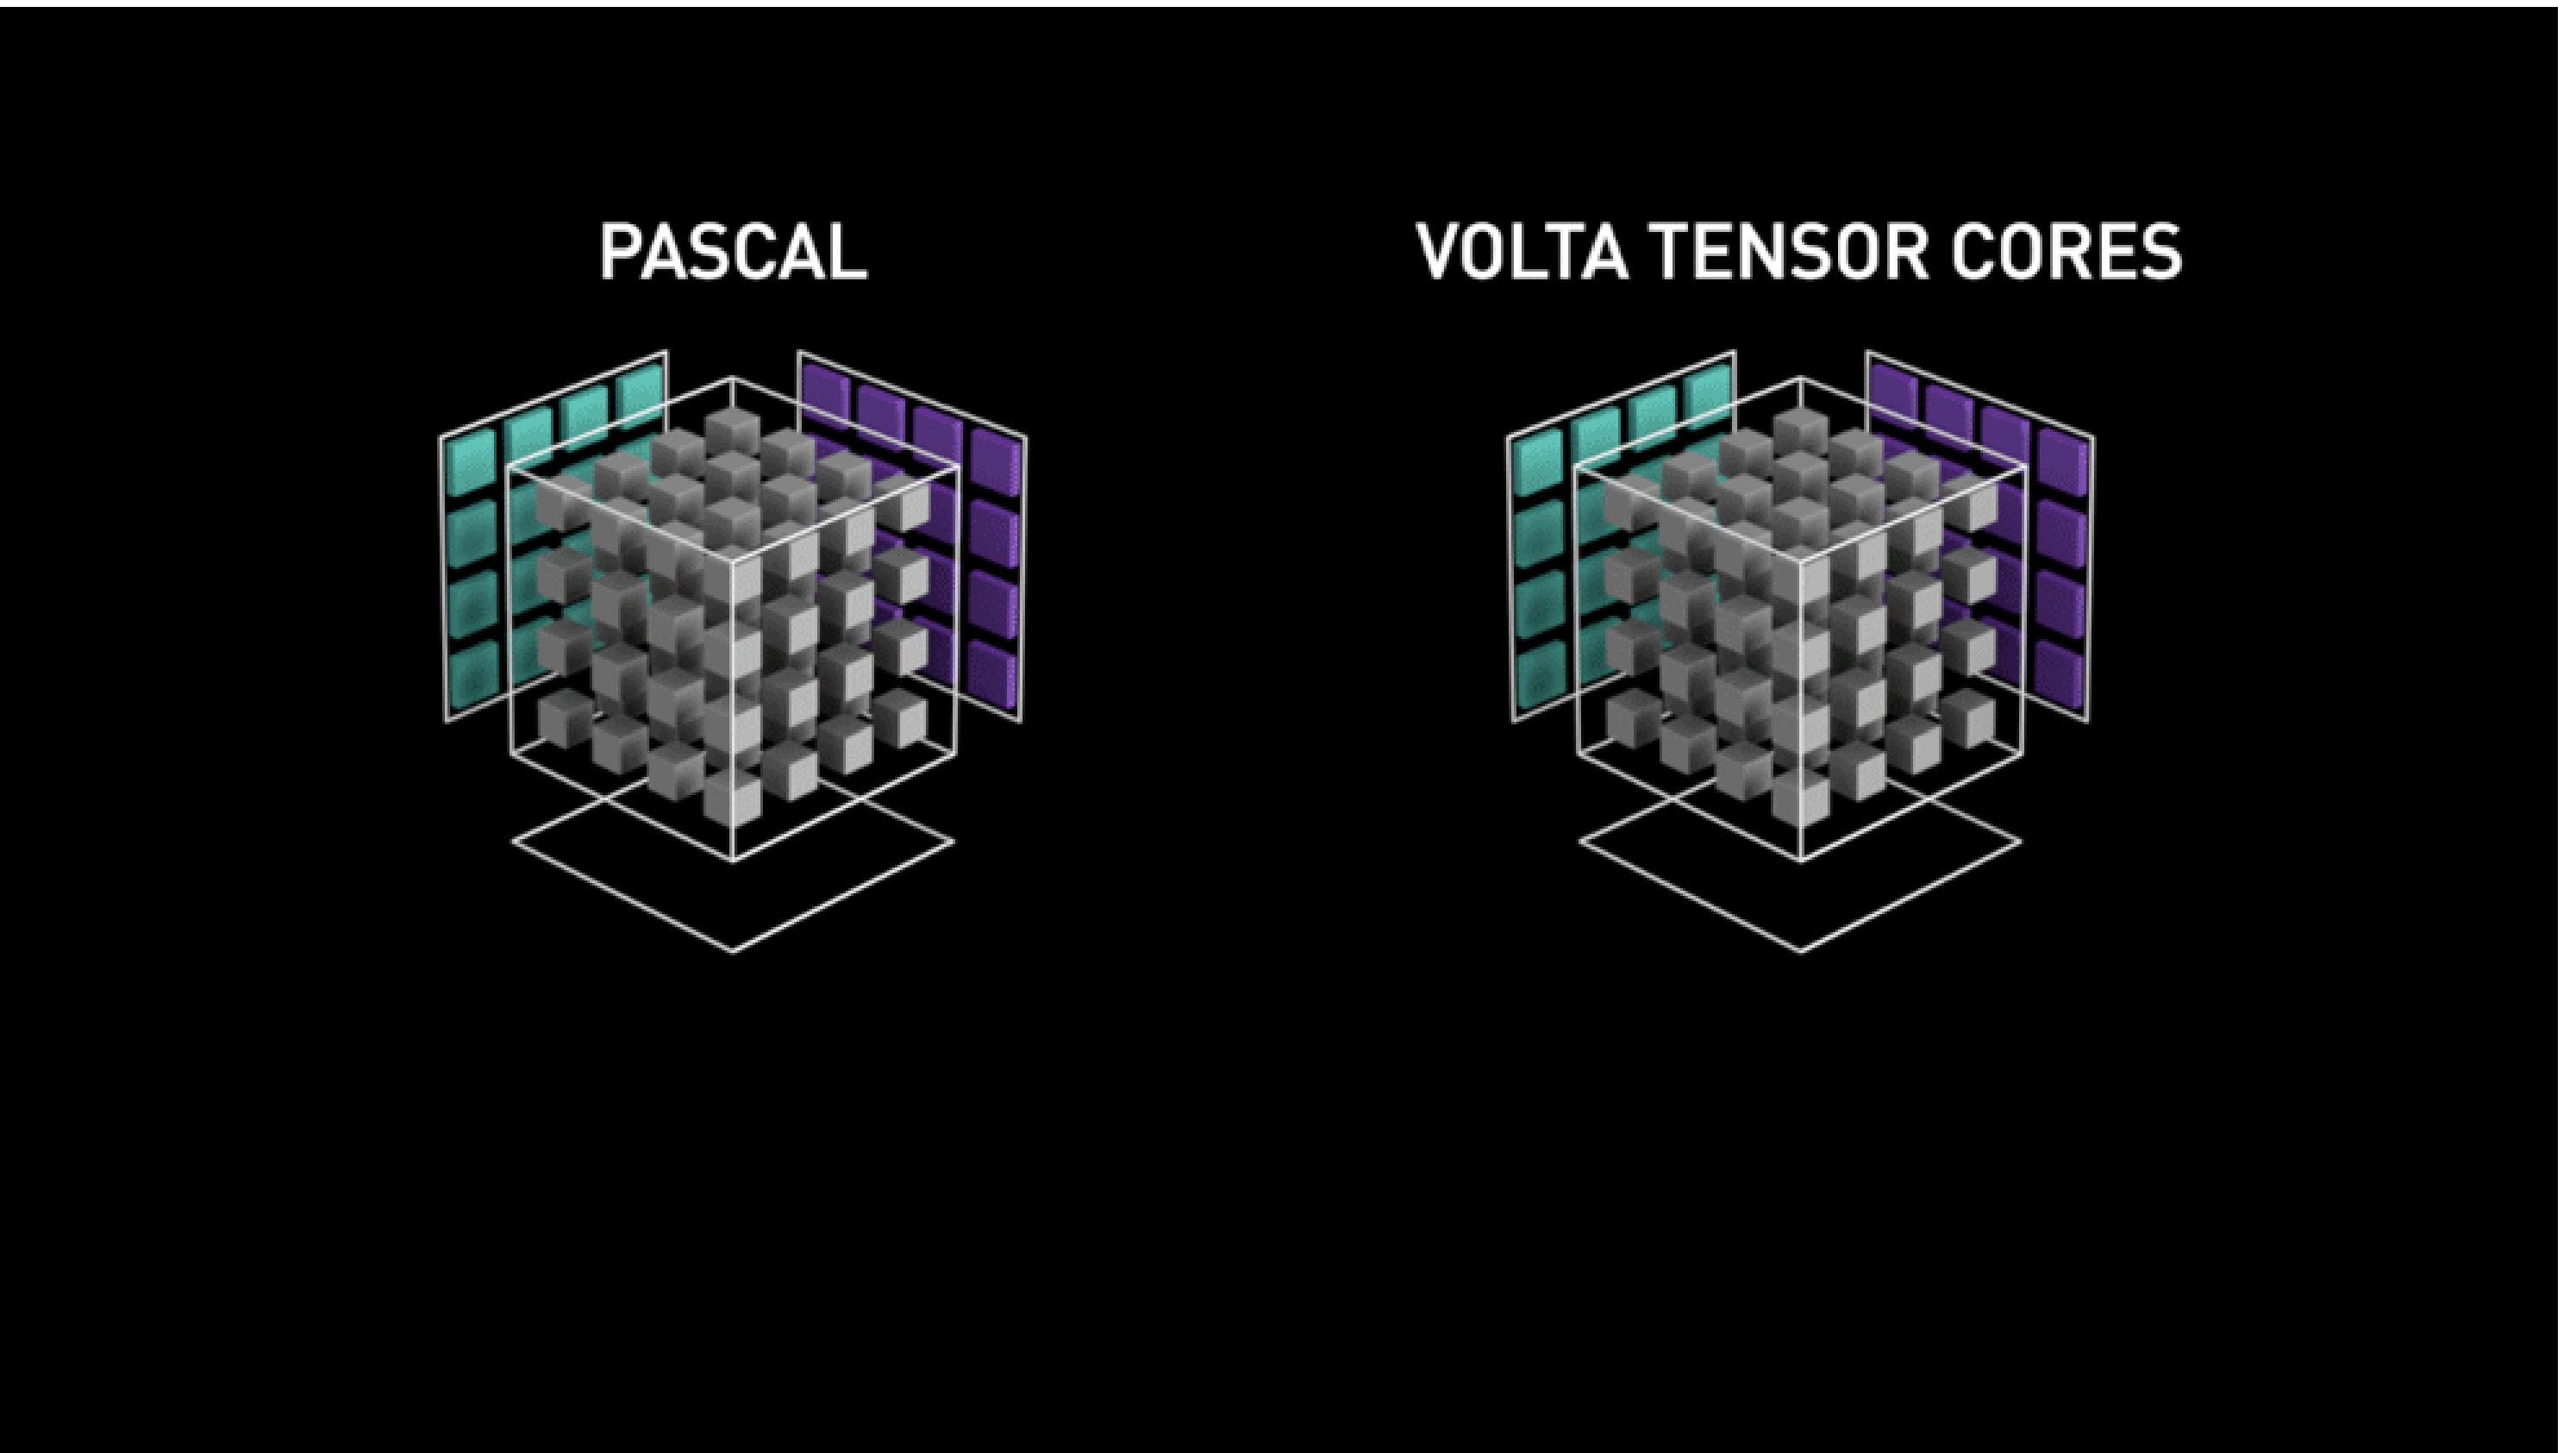
- Why Tensor Cores? What make them hard?
- CUTLASS 3.x
- Writing custom kernels with CUTLASS
- Programming Model
- Anatomy of a high performance GEMM
- Demo of a 100 line Hopper GEMM
- Q&A



What are Tensor Cores?

What is a “Tensor Core”

- Hardware block that accelerates MatMul
- FMAs accelerate vector dot products - $O(N)$ operations
- Tensor cores accelerate matrix multiplies - $O(N^3)$ operations
- Increases flop / byte ratio – more temporal and spatial reuse



Tensor Cores

MMA instruction evolution over the years

	0	1	2	3	4	5	6	7
0	T0 V0	T0 V1	T0 V2	T0 V3	T16 V0	T16 V1	T16 V2	T16 V3
1	T1 V0	T1 V1	T1 V2	T1 V3	T17 V0	T17 V1	T17 V2	T17 V3
2	T2 V0	T2 V1	T2 V2	T2 V3	T18 V0	T18 V1	T18 V2	T18 V3
3	T3 V0	T3 V1	T3 V2	T3 V3	T19 V0	T19 V1	T19 V2	T19 V3

0

	0	1	2	3
0	T0 V0	T1 V0	T2 V0	T3 V0
1	T0 V1	T1 V1	T2 V1	T3 V1
2	T0 V2	T1 V2	T2 V2	T3 V2
3	T0 V3	T1 V3	T2 V3	T3 V3
4	T16 V0	T17 V0	T18 V0	T19 V0
5	T0 V4	T0 V5	T0 V6	T2 V4
6	T0 V7	T1 V7	T2 V7	T3 V7

0

	0	0	0	0	
0	T0 V0				
1		T0 V0			
2			T0 V0		
3				T0 V0	
4					T0 V0
5					

SM70_8x8x4_F32F16F16F32_NT

0
T0 V0
T0 V1
T0 V2
T0 V3

0
T0 V0
T0 V1
T0 V2
T0 V3

FMA<double>

SM61_DP4A

0
T0 V0
T0 V1
T0 V2
T0 V3

SM80_16x8x8_F32F16F16F32_TN

0	1	2	3	4	5	6	7
T0 V0	T4 V0	T8 V0	T12 V0	T16 V0	T20 V0	T24 V0	T28 V0
T1 V1	T4 V1	T8 V1	T12 V1	T16 V1	T20 V1	T24 V1	T28 V1
T2 V1	T5 V0	T9 V0	T13 V0	T17 V0	T21 V0	T25 V0	T29 V0
T3 V1	T5 V1	T9 V1	T13 V1	T17 V1	T21 V1	T25 V1	T29 V1
T4 V0	T6 V0	T10 V0	T14 V0	T18 V0	T22 V0	T26 V0	T30 V0
T5 V0	T6 V1	T10 V1	T14 V1	T18 V1	T22 V1	T26 V1	T30 V1
T6 V0	T6 V2	T10 V2	T14 V2	T18 V2	T22 V2	T26 V2	T30 V2
T7 V0	T7 V1	T11 V1	T15 V1	T19 V1	T23 V1	T27 V1	T31 V1

0

0	1	2	3	4	5	6	7
T0 V0	T4 V0	T8 V0	T12 V0	T16 V0	T20 V0	T24 V0	T28 V0
T1 V1	T4 V1	T8 V1	T12 V1	T16 V1	T20 V1	T24 V1	T28 V1
T2 V1	T5 V0	T9 V0	T13 V0	T17 V0	T21 V0	T25 V0	T29 V0
T3 V1	T5 V1	T9 V1	T13 V1	T17 V1	T21 V1	T25 V1	T29 V1
T4 V0	T6 V0	T10 V0	T14 V0	T18 V0	T22 V0	T26 V0	T30 V0
T5 V0	T6 V1	T10 V1	T14 V1	T18 V1	T22 V1	T26 V1	T30 V1
T6 V0	T6 V2	T10 V2	T14 V2	T18 V2	T22 V2	T26 V2	T30 V2
T7 V0	T7 V1	T11 V1	T15 V1	T19 V1	T23 V1	T27 V1	T31 V1

0

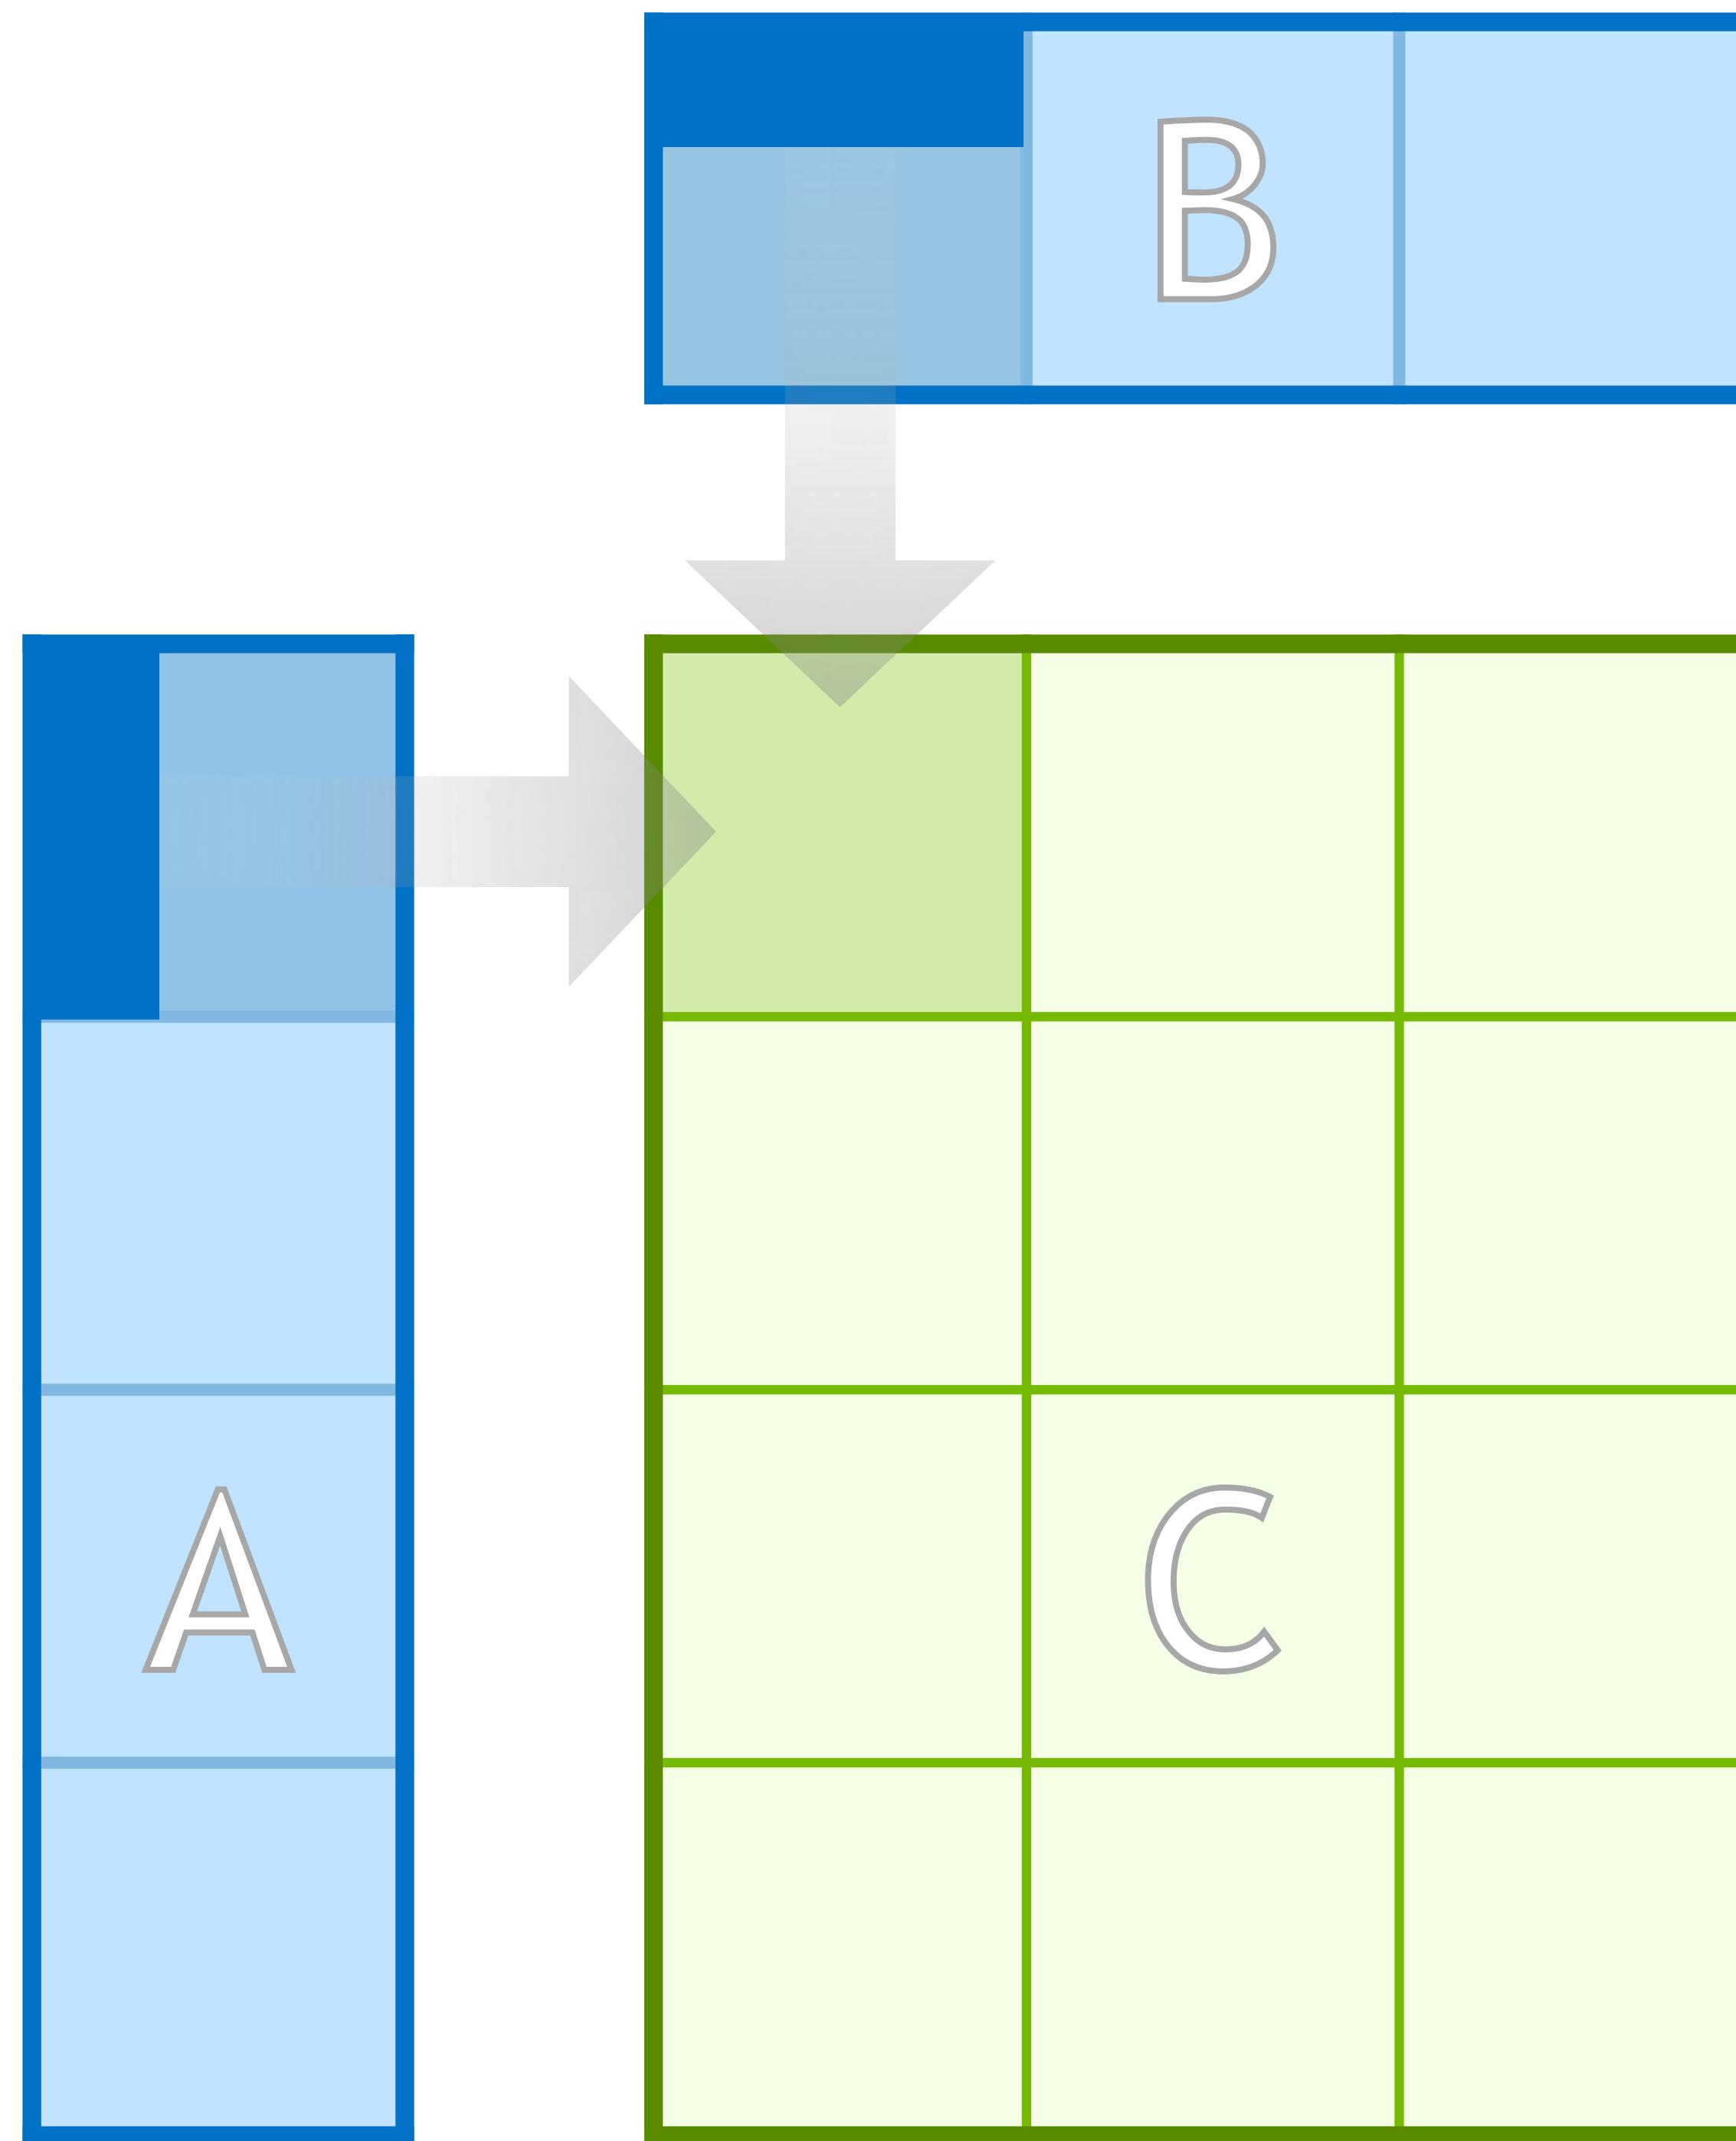
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
T0 V0	T4 V0	T8 V0	T12 V0	T16 V0	T20 V0	T24 V0	T28 V0	T32 V0	T36 V0	T40 V0	T44 V0	T48 V0	T52 V0	T56 V0	T60 V0
T1 V1	T4 V1	T8 V1	T12 V1	T16 V1	T20 V1	T24 V1	T28 V1	T32 V1	T36 V1	T40 V1	T44 V1	T48 V1	T52 V1	T56 V1	T60 V1
T2 V1	T5 V0	T9 V0	T13 V0	T17 V0	T21 V0	T25 V0	T29 V0	T33 V0	T37 V0	T41 V0	T45 V0	T49 V0	T53 V0	T57 V0	T61 V0
T3 V1	T5 V1	T9 V1	T13 V1	T17 V1	T21 V1	T25 V1	T29 V1	T33 V1	T37 V1	T41 V1	T45 V1	T49 V1	T53 V1	T57 V1	T61 V1
T4 V0	T6 V0	T10 V0	T14 V0	T18 V0	T22 V0	T26 V0	T30 V0	T34 V0	T38 V0	T42 V0	T46 V0	T50 V0	T54 V0	T58 V0	T62 V0
T5 V0	T6 V1	T10 V1	T14 V1	T18 V1	T22 V1	T26 V1	T30 V1	T34 V1	T38 V1	T42 V1	T46 V1	T50 V1	T54 V1	T58 V1	T62 V1
T6 V0	T6 V2	T10 V2	T14 V2	T18 V2	T22 V2	T26 V2	T30 V2	T34 V2	T38 V2	T42 V2	T46 V2	T50 V2	T54 V2	T58 V2	T62 V2
T7 V0	T7 V1	T11 V1	T15 V1	T19 V1	T23 V1	T27 V1	T31 V1	T35 V1	T39 V1	T43 V1	T47 V1	T51 V1	T55 V1	T59 V1	T63 V1
T8 V0	T8 V1	T12 V1	T16 V1	T20 V1	T24 V1	T28 V1	T32 V1	T36 V1	T40 V1	T44 V1	T48 V1	T52 V1	T56 V1	T60 V1	T64 V1
T9 V0	T9 V1	T13 V1	T17 V1	T21 V1	T25 V1	T29 V1	T33 V1	T37 V1	T41 V1	T45 V1	T49 V1	T53 V1	T57 V1	T61 V1	T65 V1
T10 V0	T10 V1	T14 V1	T18 V1	T22 V1	T26 V1	T30 V1	T34 V1	T38 V1	T42 V1	T46 V1	T50 V1	T54 V1	T58 V1	T62 V1	T66 V1
T11 V0	T11 V1	T15 V1	T19 V1	T23 V1	T27 V1	T31 V1	T35 V1	T39 V1	T43 V1	T47 V1	T51 V1	T55 V1	T59<br		



What's so hard about MatMuls?

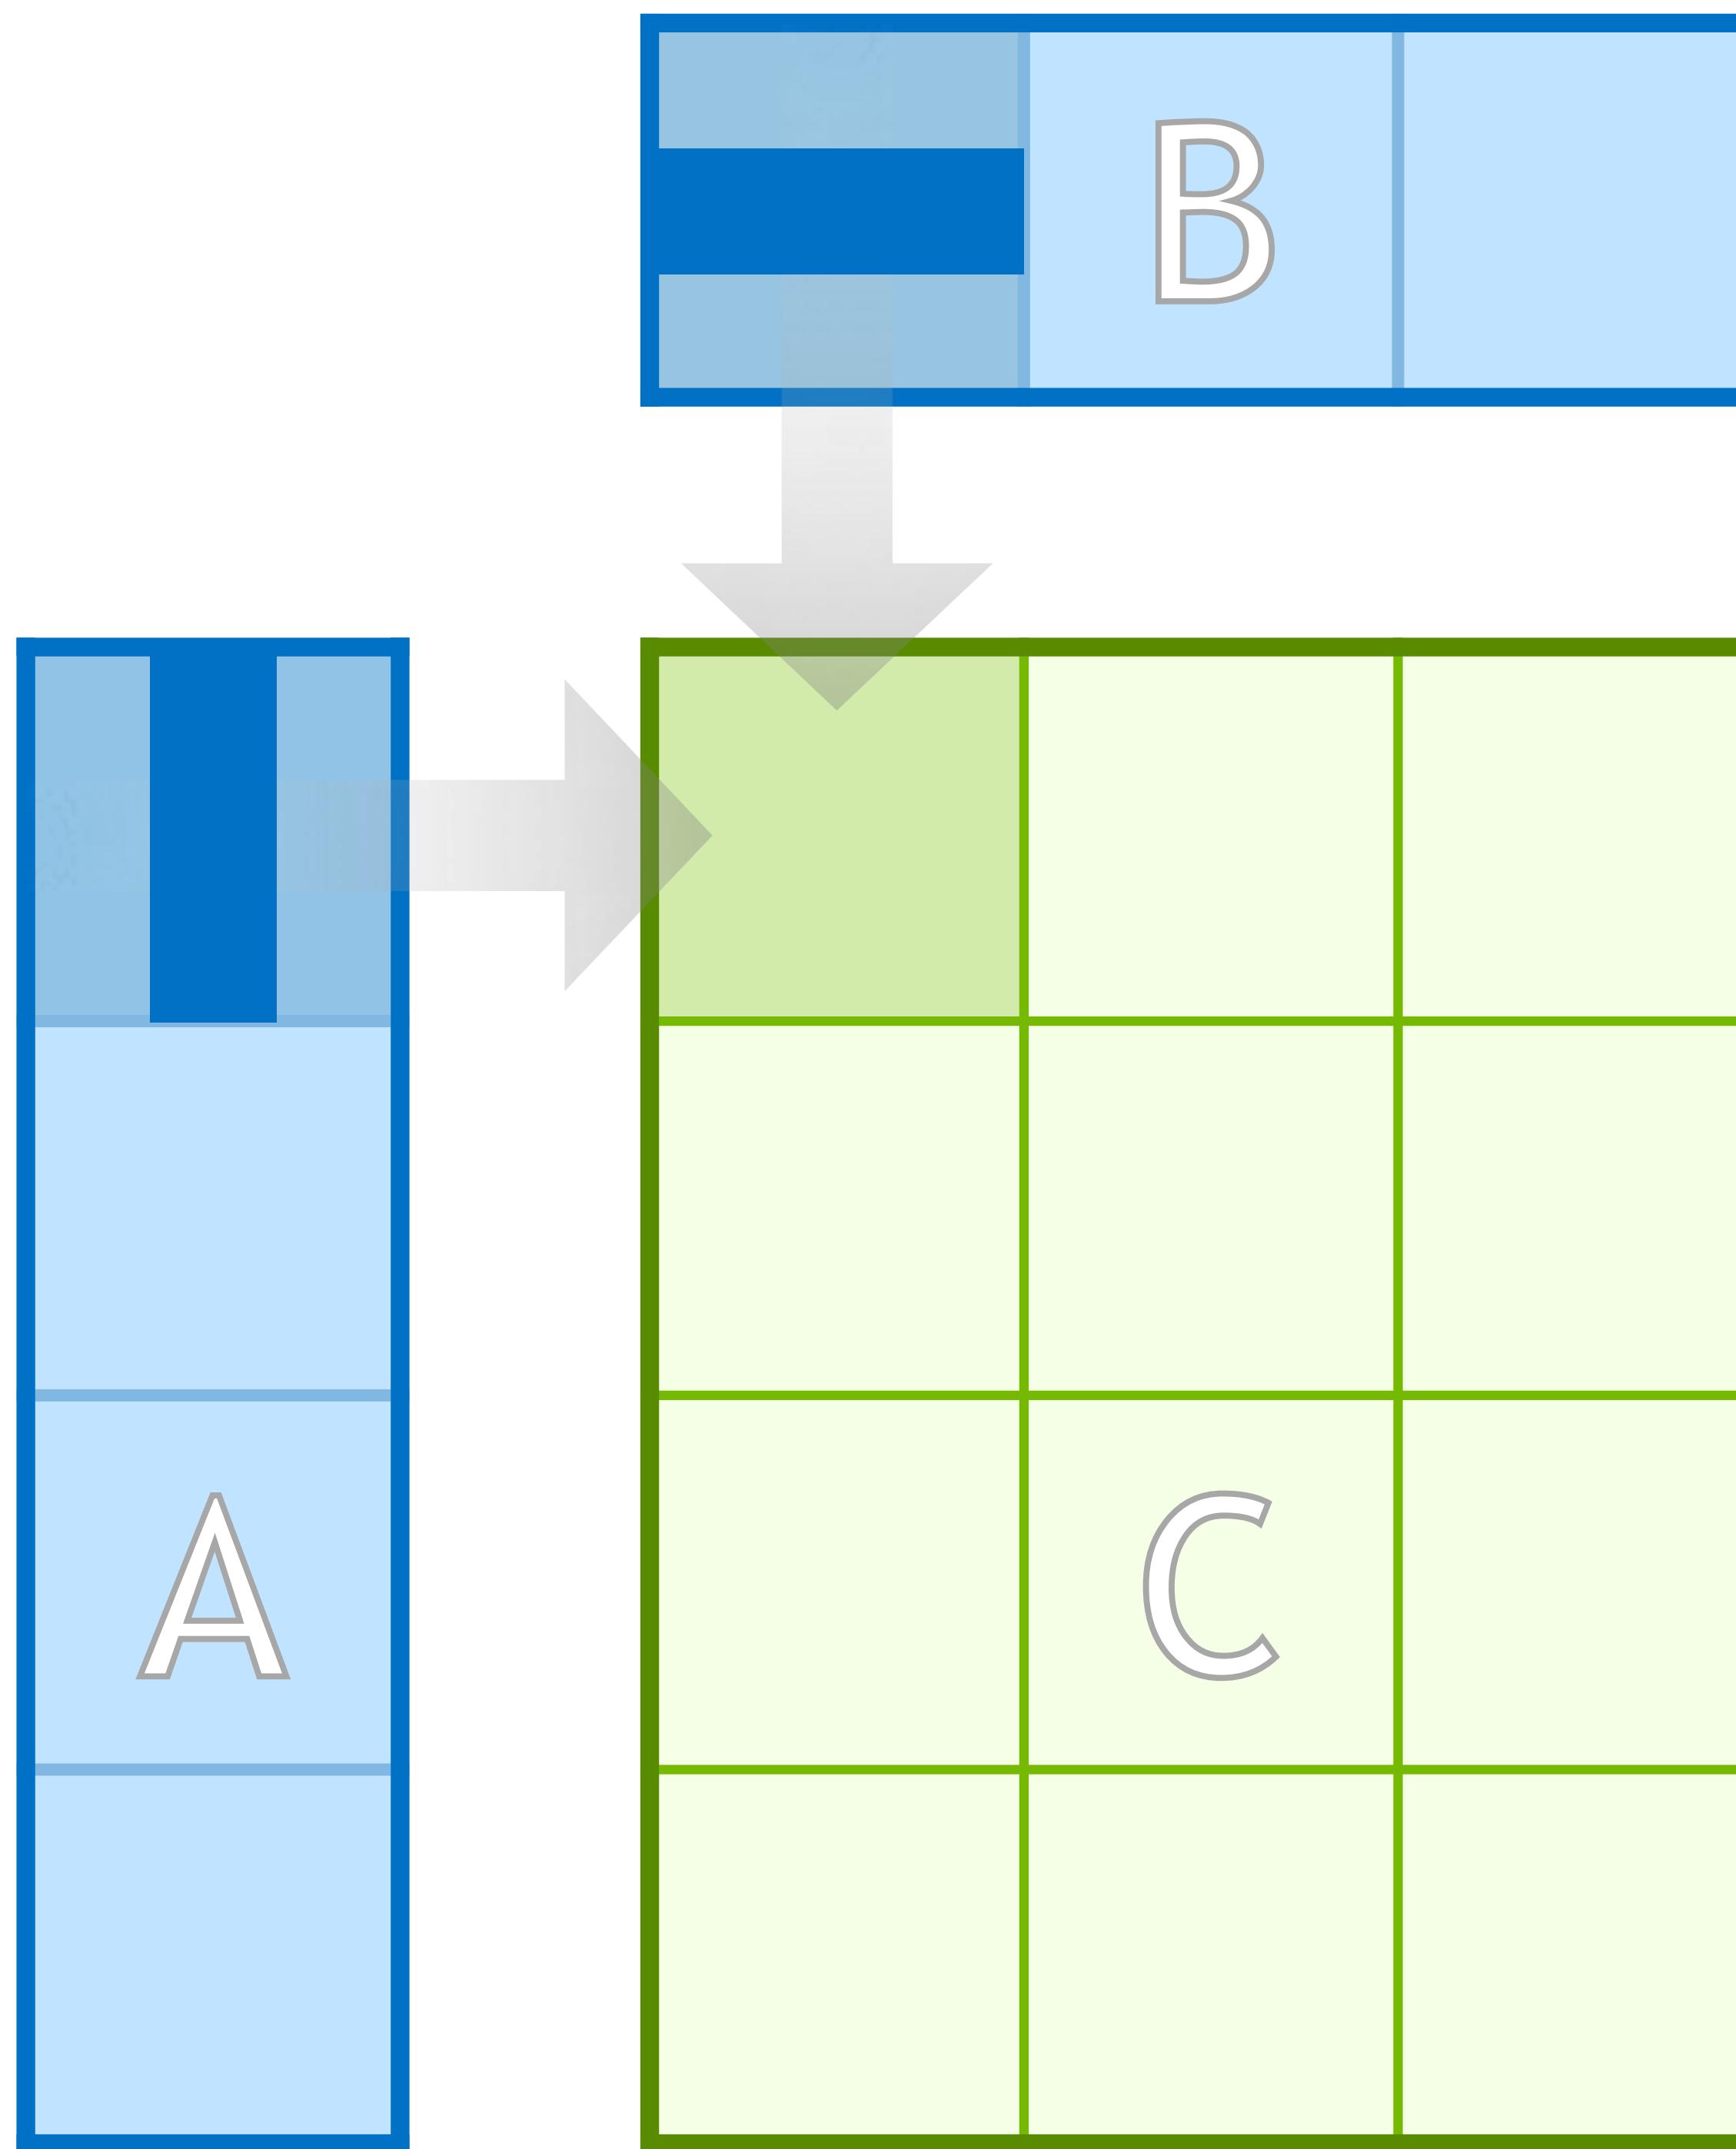
How We Think About Linear Algebra

Accumulated Inner Product of Outer Products



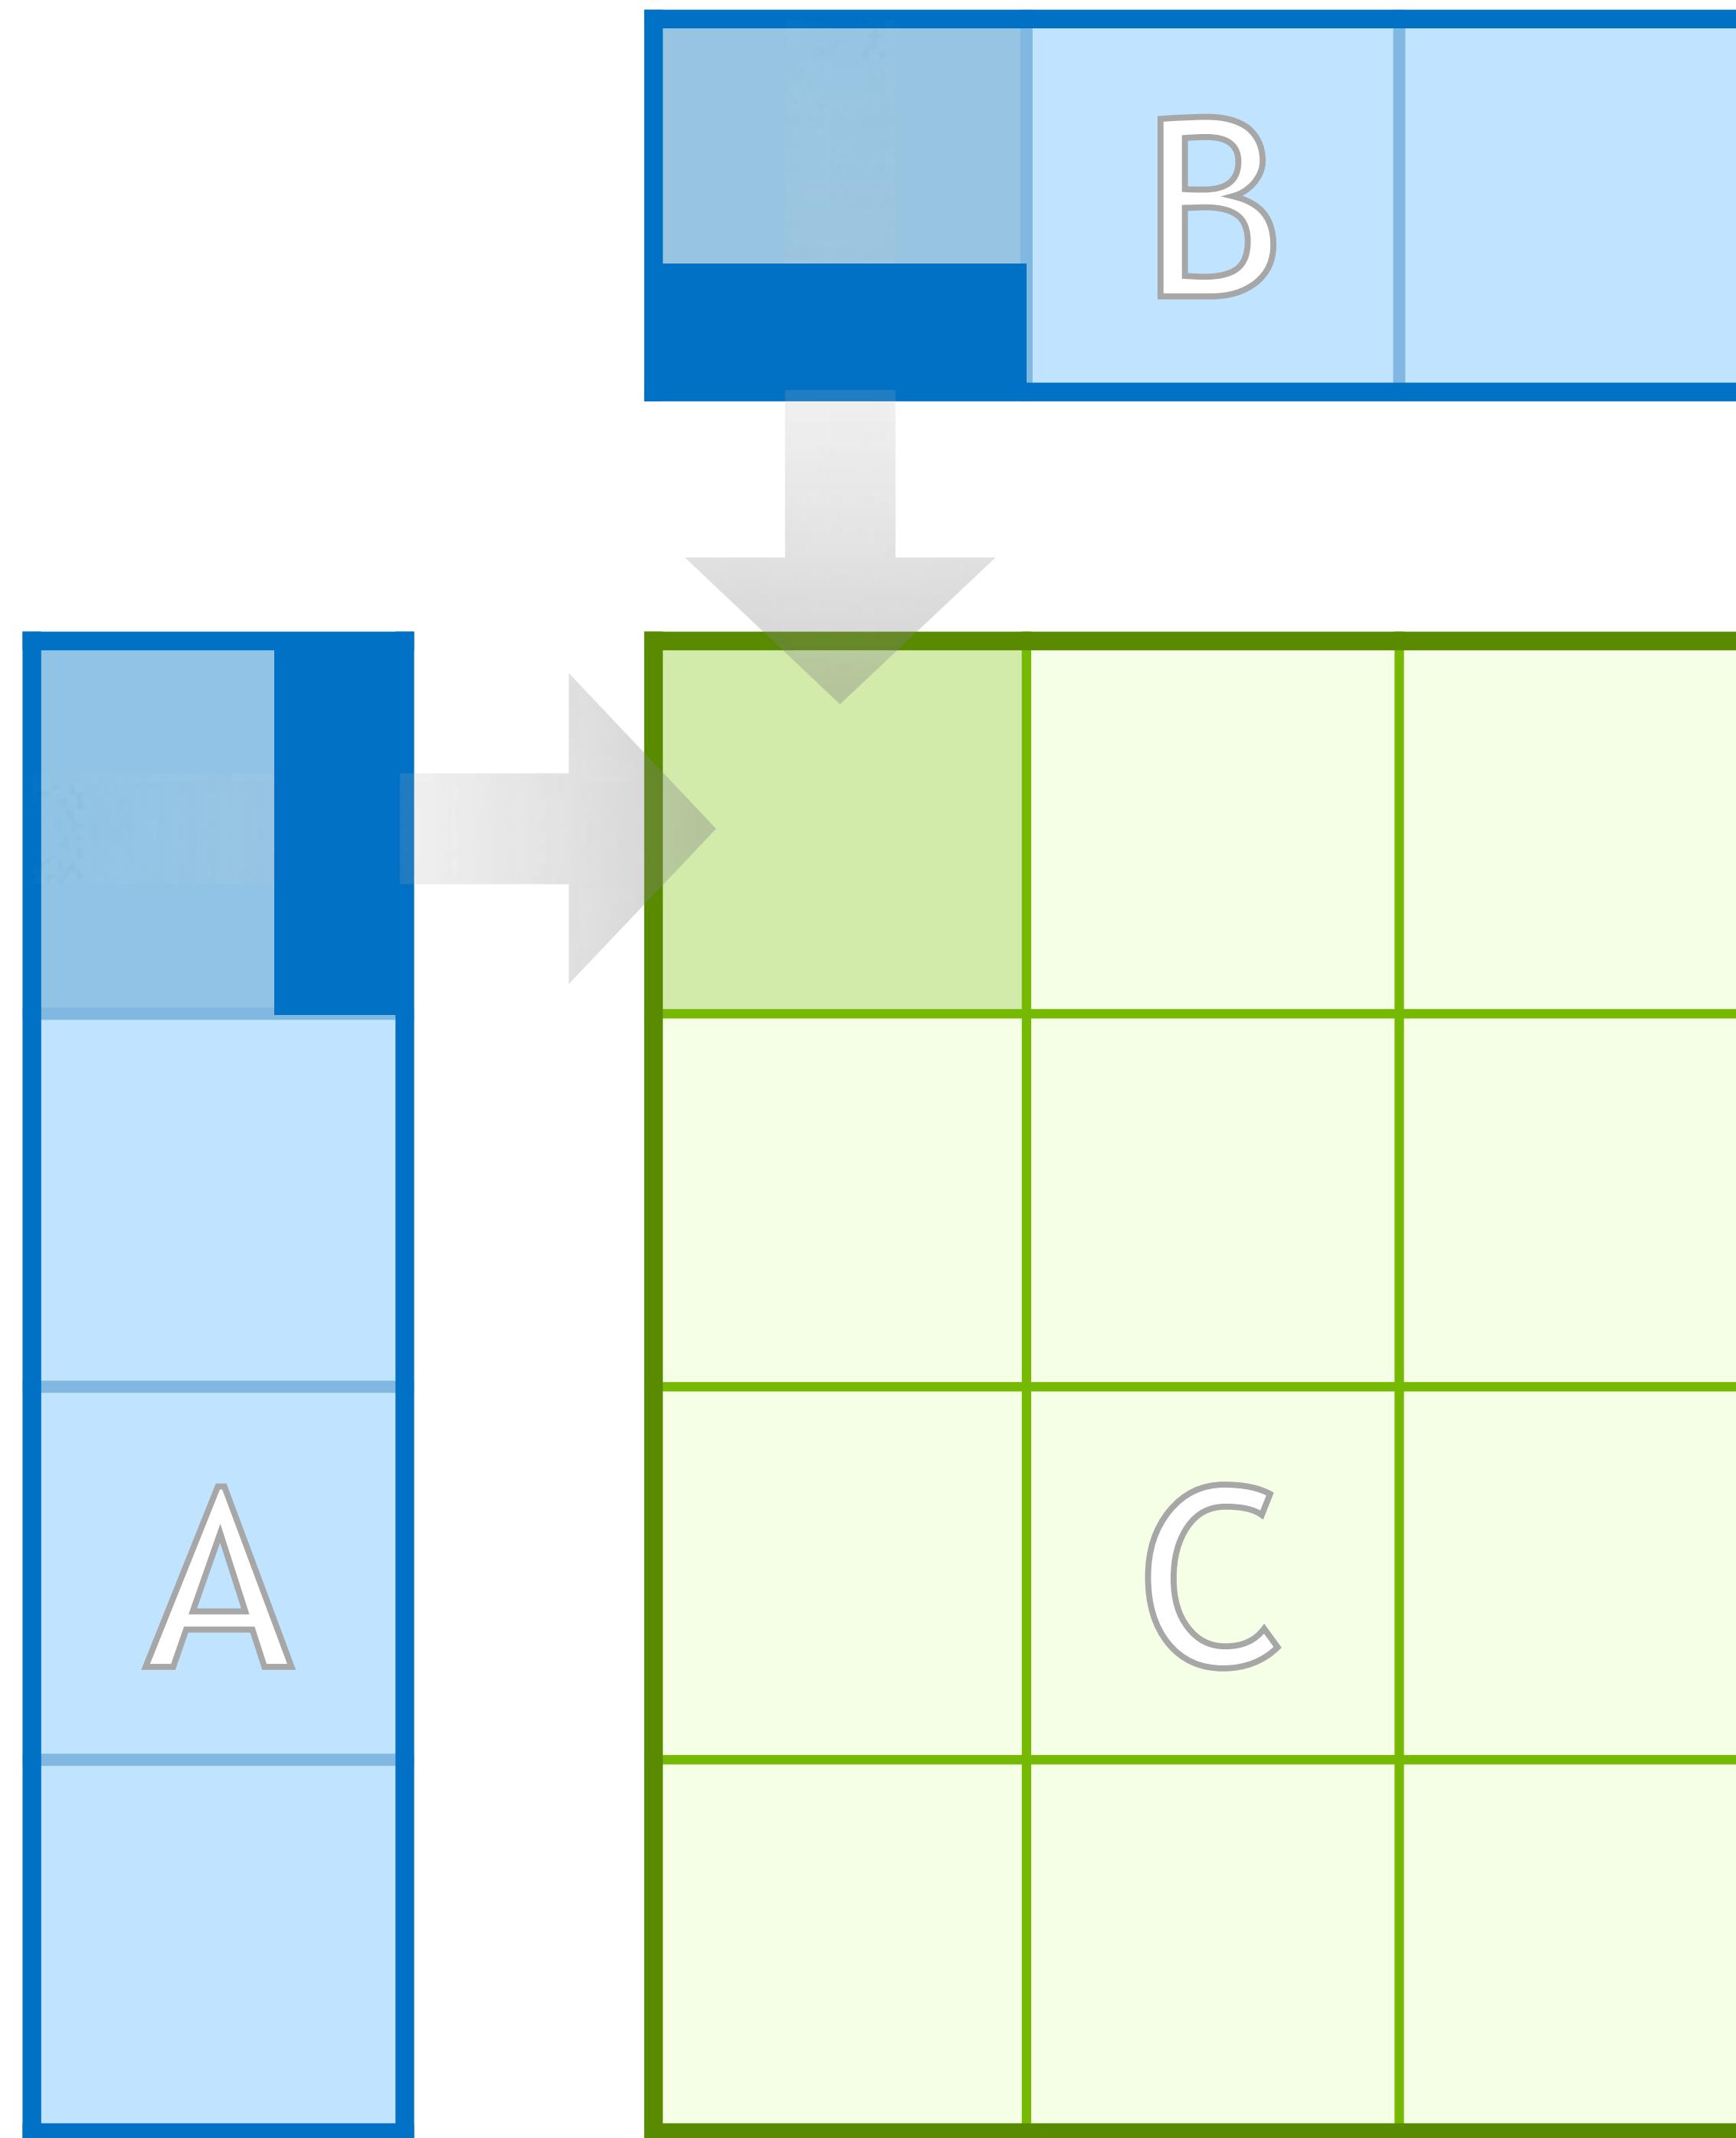
How We Think About Linear Algebra

Accumulated Inner Product of Outer Products



How We Think About Linear Algebra

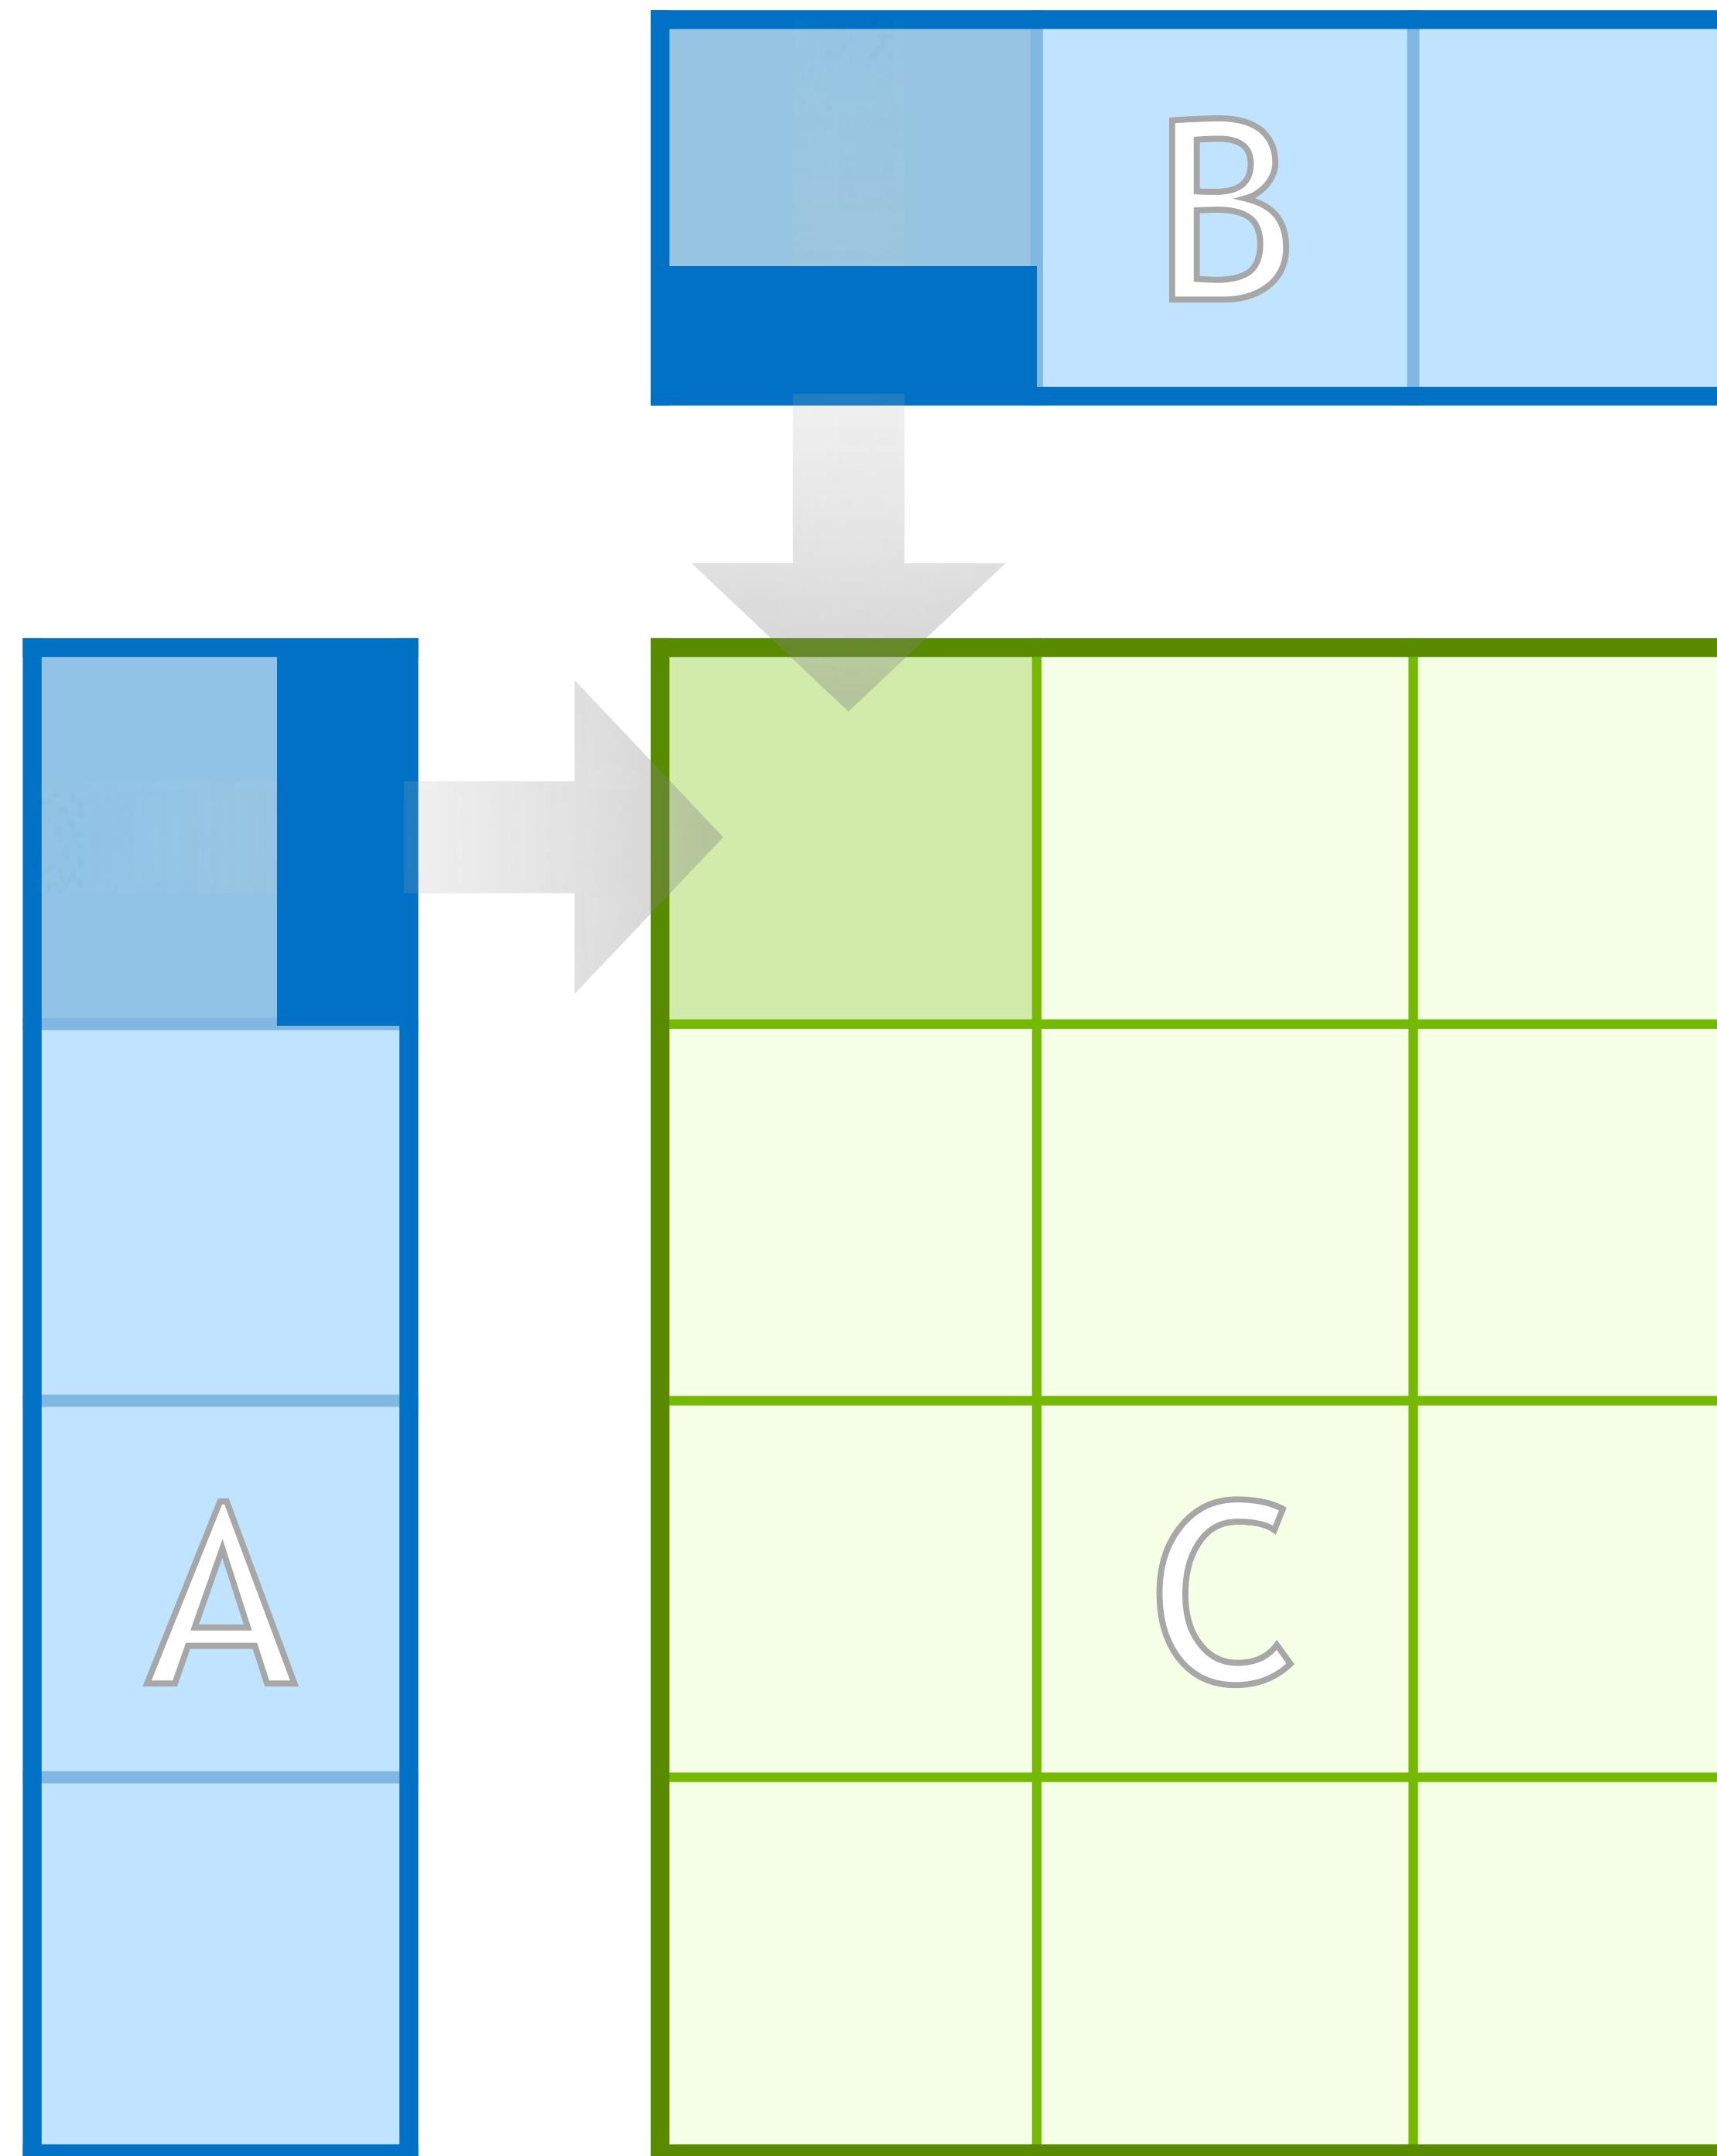
Accumulated Inner Product of Outer Products



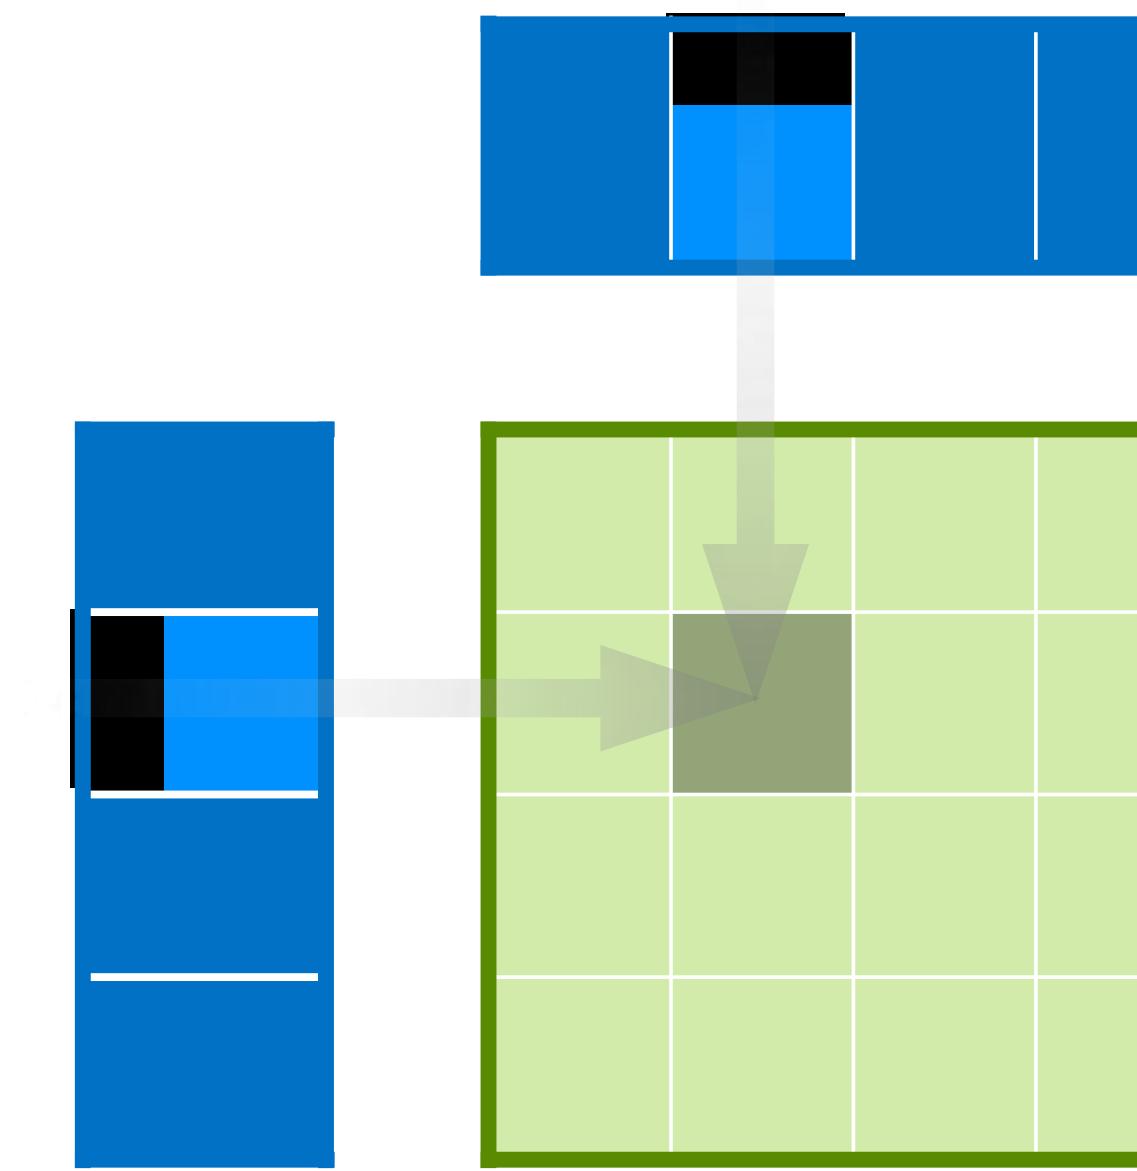
```
for (int m = 0; m < M; ++m) {  
    for (int n = 0; n < N; ++n) {  
        for (int k = 0; k < K; ++k) {  
  
            C(m, n) += A(m, k) * B(n, k);  
  
        } // end k  
    } // end n  
} // end m
```

How We Think About Linear Algebra

Parallelism and locality are hierarchical ...

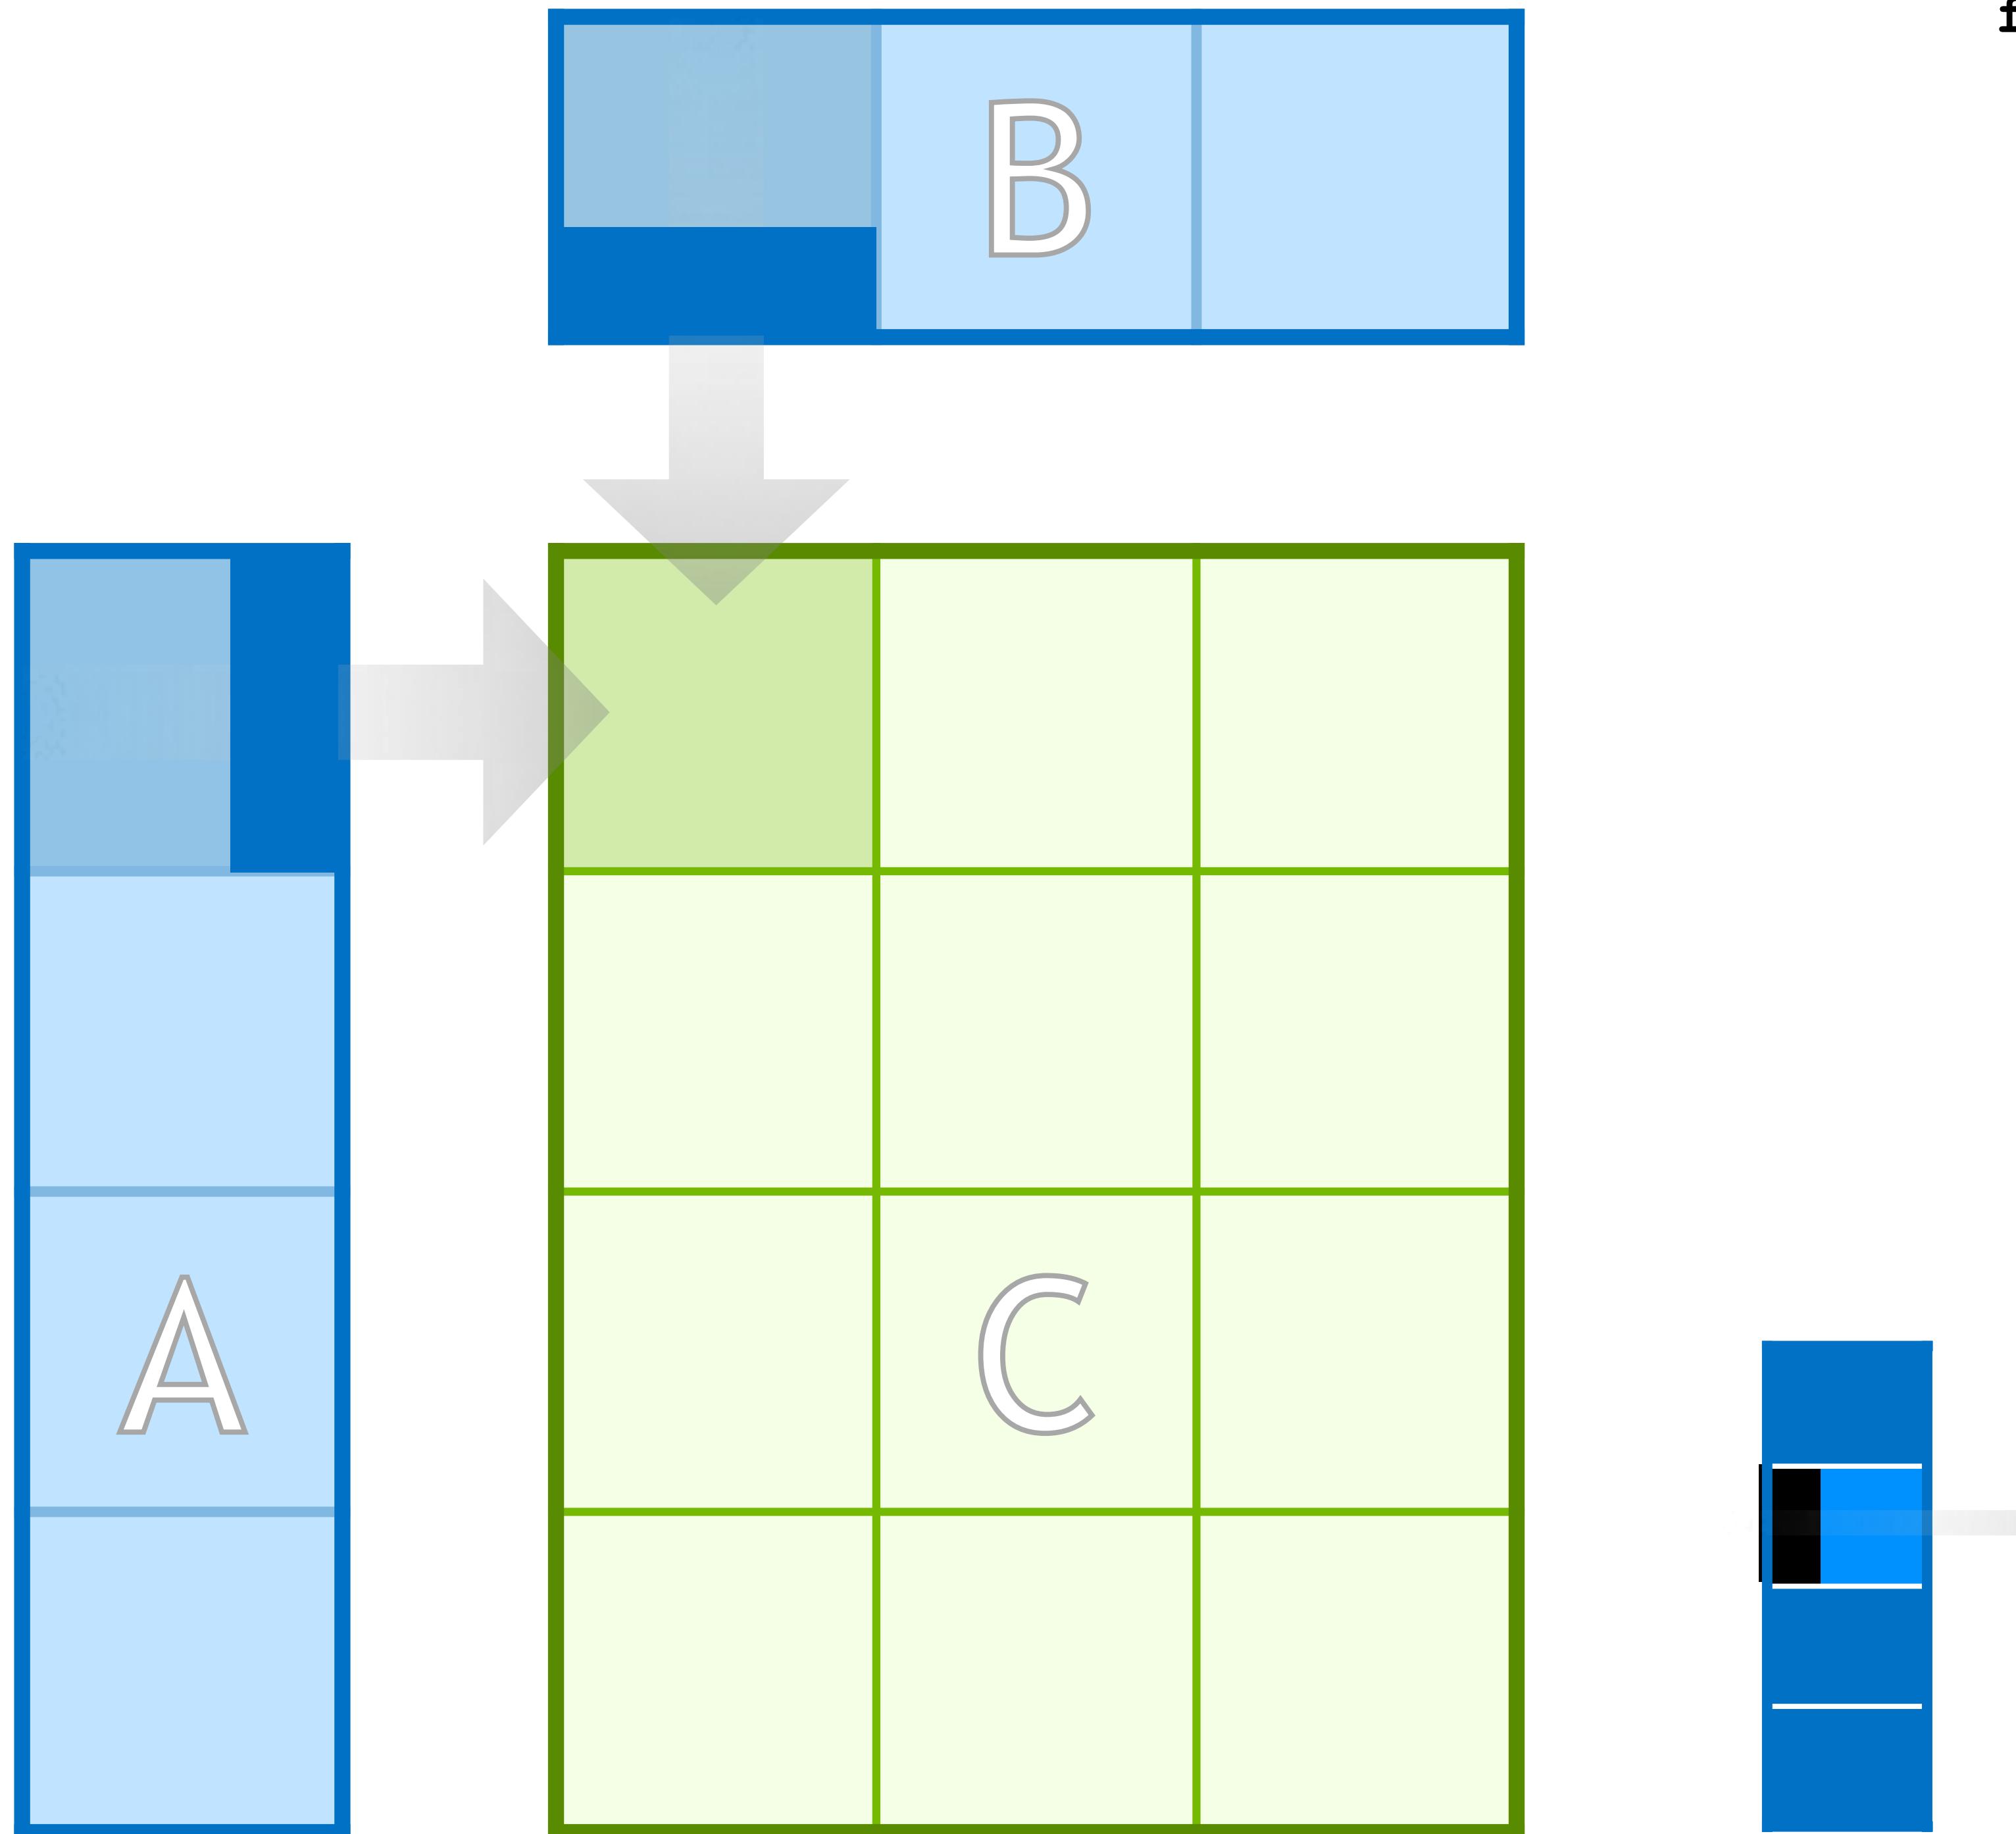


```
for (int cta_n = 0; cta_n < GemmN; cta_n += CtaN) {  
    for (int cta_m = 0; cta_m < GemmM; cta_m += CtaM) {  
        for (int cta_k = 0; cta_k < GemmK; cta_k += CtaK) {  
  
            // Copy from global to shared  
  
            for (int warp_n = 0; warp_n < CtaN; warp_n += WarpN) {  
                for (int warp_m = 0; warp_m < CtaM; warp_m += WarpM) {  
                    for (int warp_k = 0; warp_k < CtaK; warp_k += WarpK) {  
  
                        gemm(a, b, c);  
                    }  
                }  
            }  
        }  
    }  
}
```

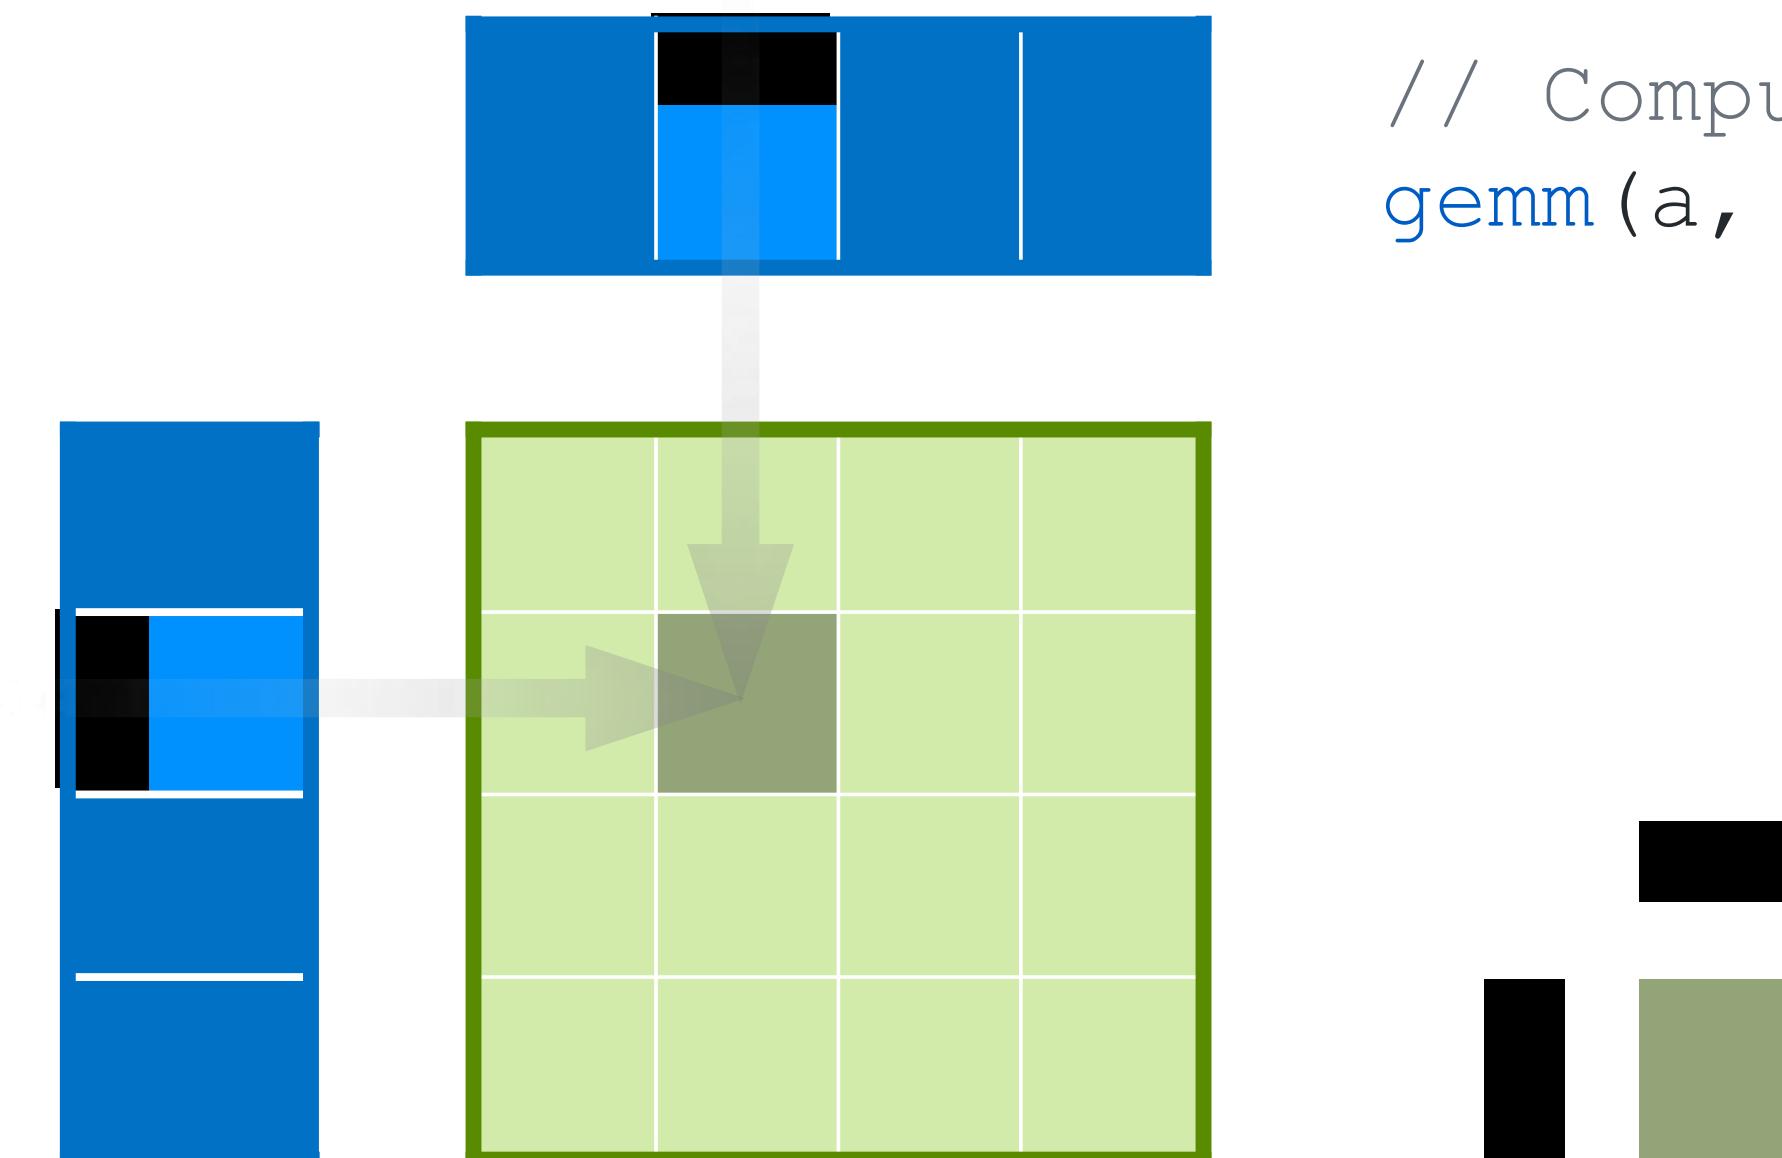


How We Think About Linear Algebra

... often with multiple nested hierarchies of both threads and data

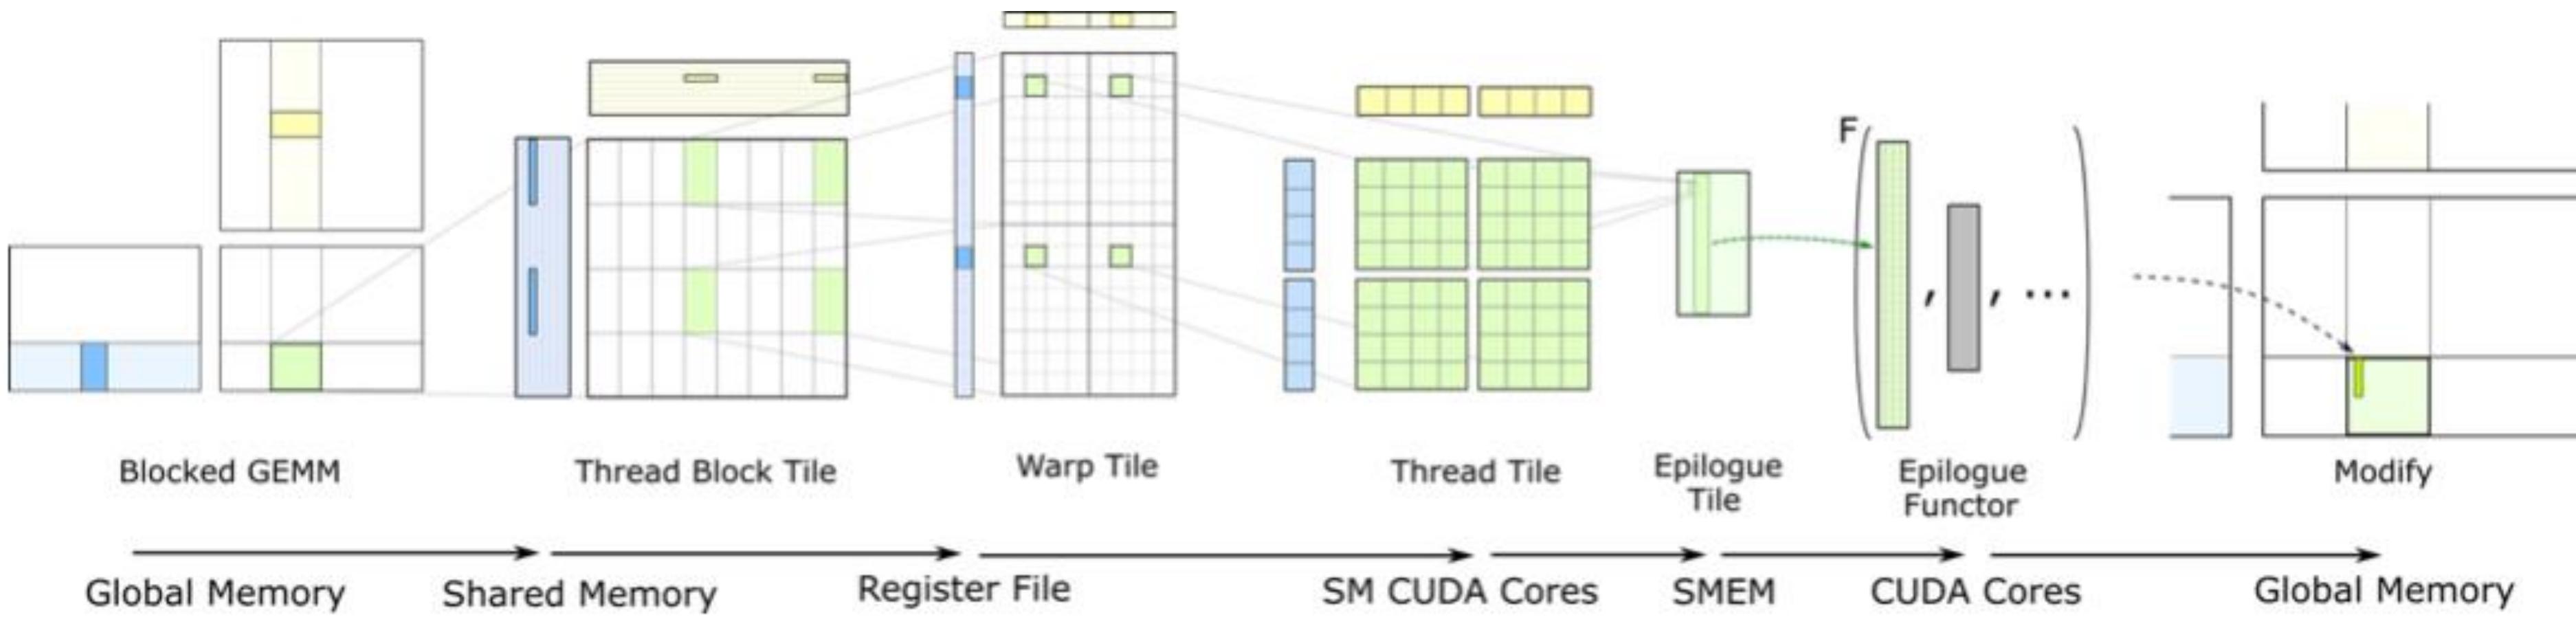


```
for (int cta_n = 0; cta_n < GemmN; cta_n += CtaN) {  
    for (int cta_m = 0; cta_m < GemmM; cta_m += CtaM) {  
        for (int cta_k = 0; cta_k < GemmK; cta_k += CtaK) {  
  
            // Copy from global to shared  
  
            for (int warp_n = 0; warp_n < CtaN; warp_n += WarpN) {  
                for (int warp_m = 0; warp_m < CtaM; warp_m += WarpM) {  
                    for (int warp_k = 0; warp_k < CtaK; warp_k += WarpK) {  
  
                        // Copy from shared to registers  
  
                        for (int mma_k = 0; mma_k < WarpK; mma_k += MmaK) {  
                            for (int mma_n = 0; mma_n < WarpN; mma_n += MmaN) {  
                                for (int mma_m = 0; mma_m < WarpM; mma_m += MmaM) {  
  
                                    // Compute in registers  
                                    gemm(a, b, c);
```



Architectural Complications ...

Push us away from coordinates to index bookkeeping



A matrix addresses	B matrix data addresses
0	0 16 32 48
16	2 3 18 19
32	4 5 20 21
48	6 7 22 23
64	8 9 24 25
80	10 11 26 27
96	12 13 28 29
112	14 15 30 31

Maxwell A access addr[tid] = addr[tid^1]
Maxwell B access addr[tid] = addr[tid^2]

Figure 3b: Maxwell thread assignments for LDS.U matching

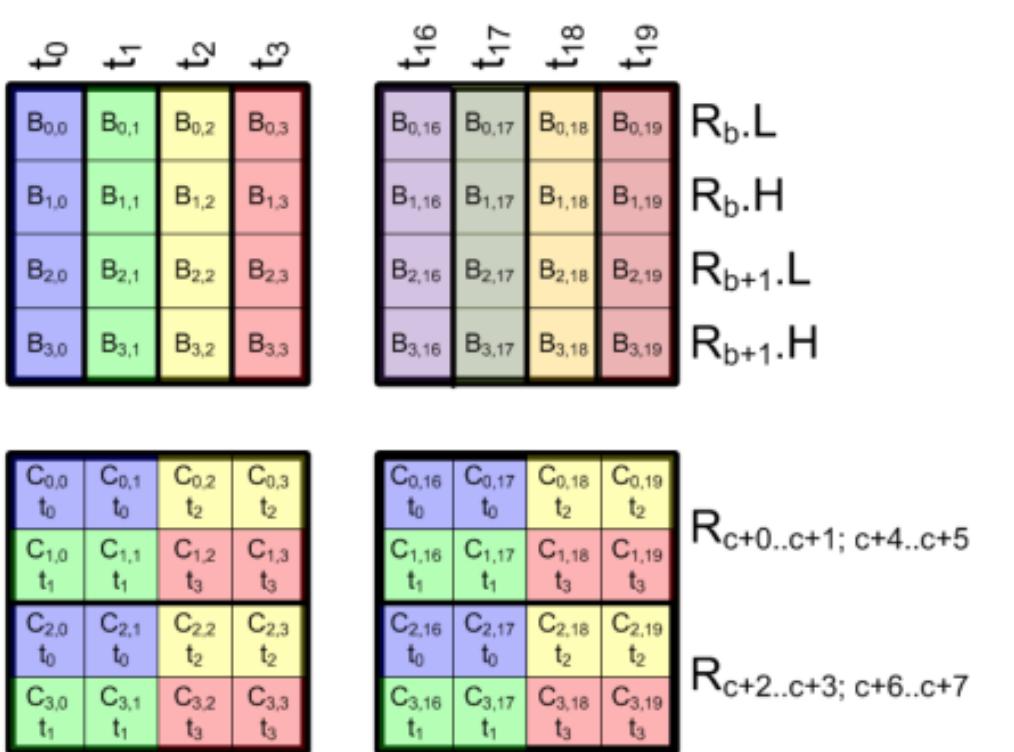


Figure 3b: Maxwell thread assignments for LDS.U matching

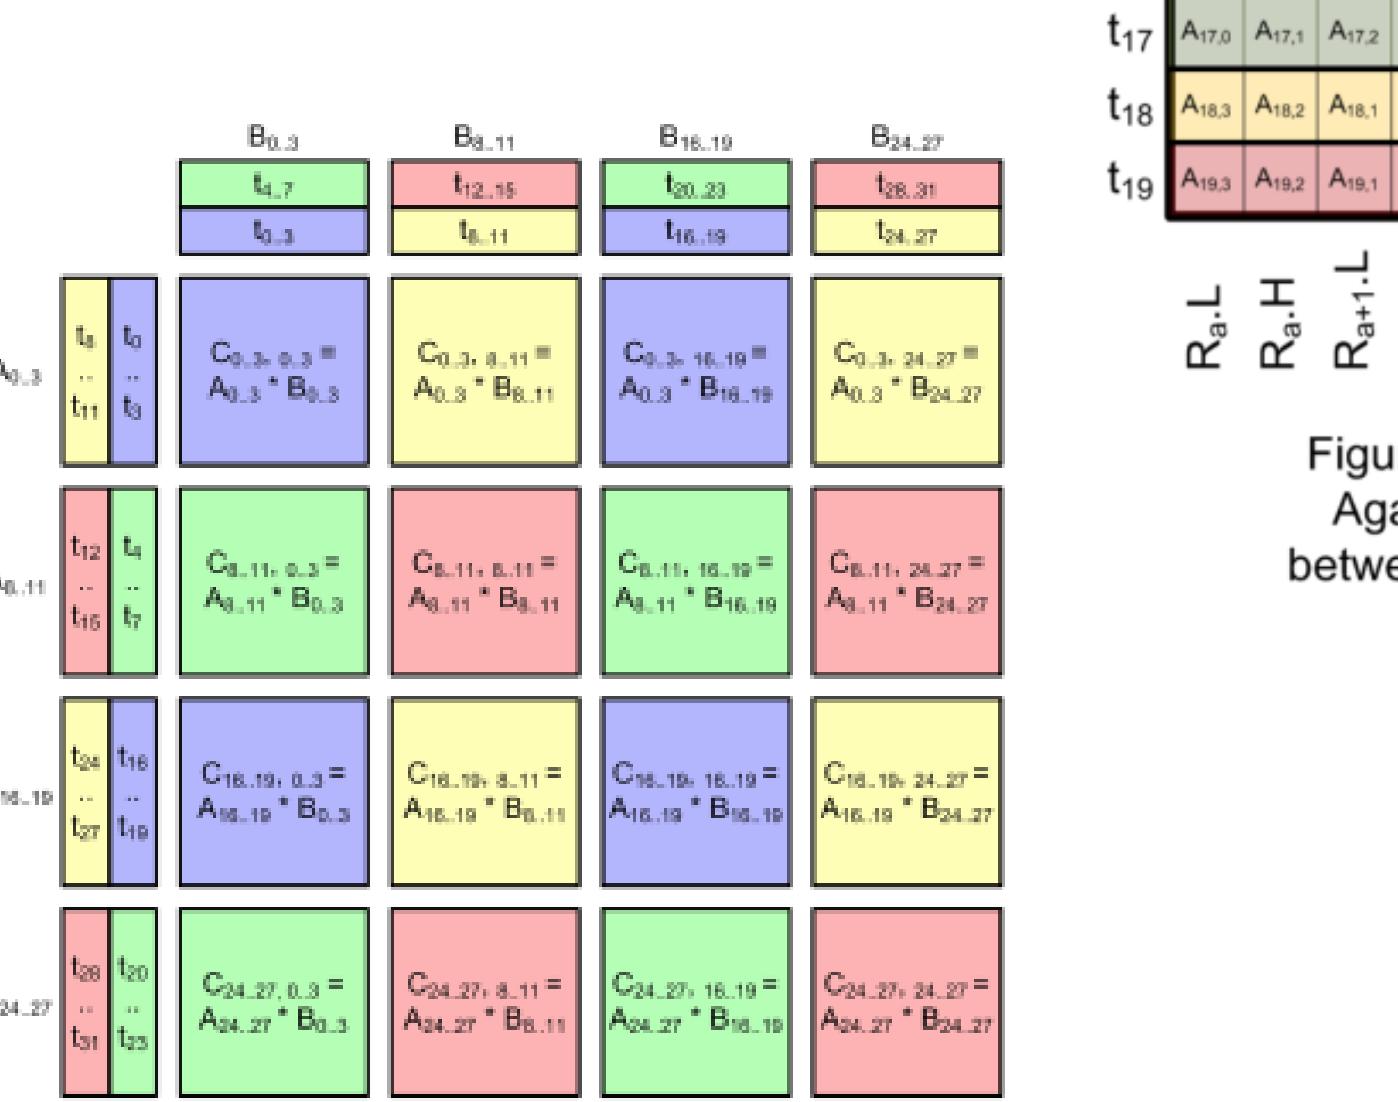
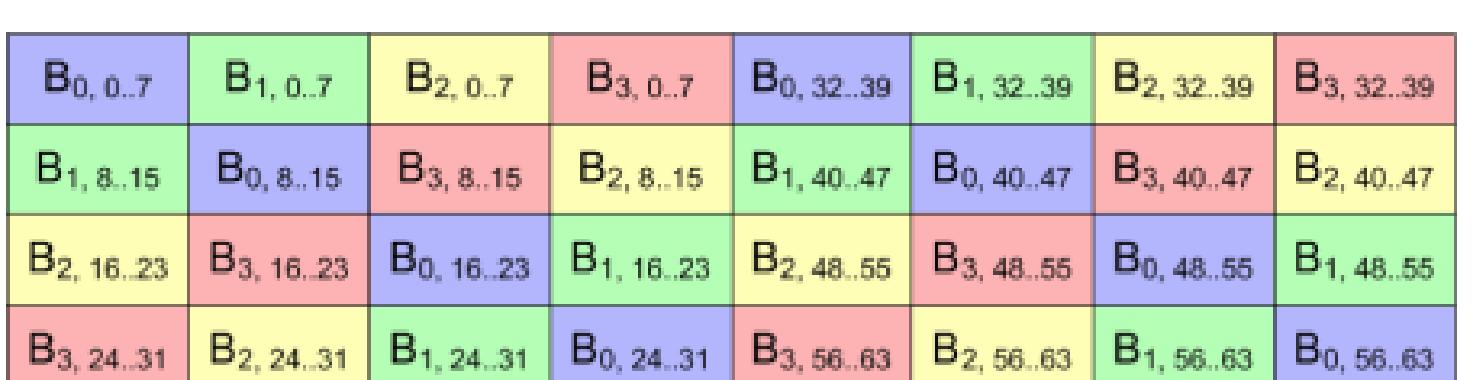
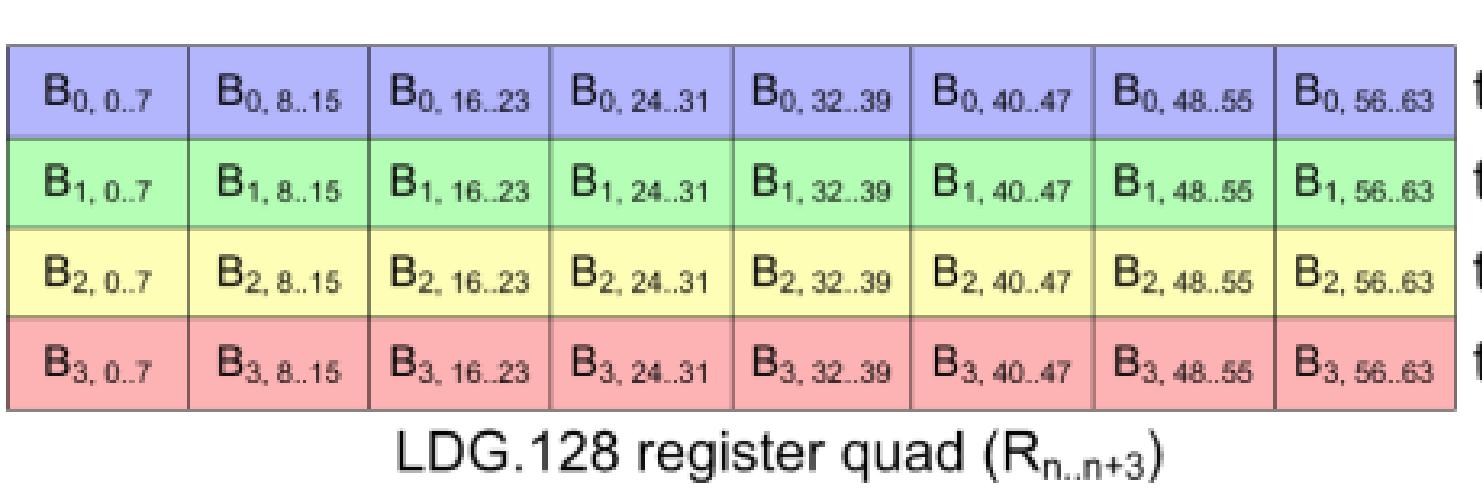


Figure 1b: HMMA.884.F32.TN data layout.
Again, ignore the apparent isomorphism between thread numbers and matrix indices.

Figure 2a: Composing sparse A[16,4] * B[4, 16] = C[16, 16]
warp-wide multiply from four A[8, 4] * B[4, 8] = C[8, 8] 8-thread multiplies.



Shared memory after STS.128 w/swizzled thread id =
(tid[1:0] << 3) | (tid & 4) | (tid[4:3] ^ tid[1:0])

```
for (int cta_n = 0; cta_n < GemmN; cta_n += CtaN) {
    for (int cta_m = 0; cta_m < GemmM; cta_m += CtaM) {
        for (int cta_k = 0; cta_k < GemmK; cta_k += CtaK) {
```

// Copy from global to shared

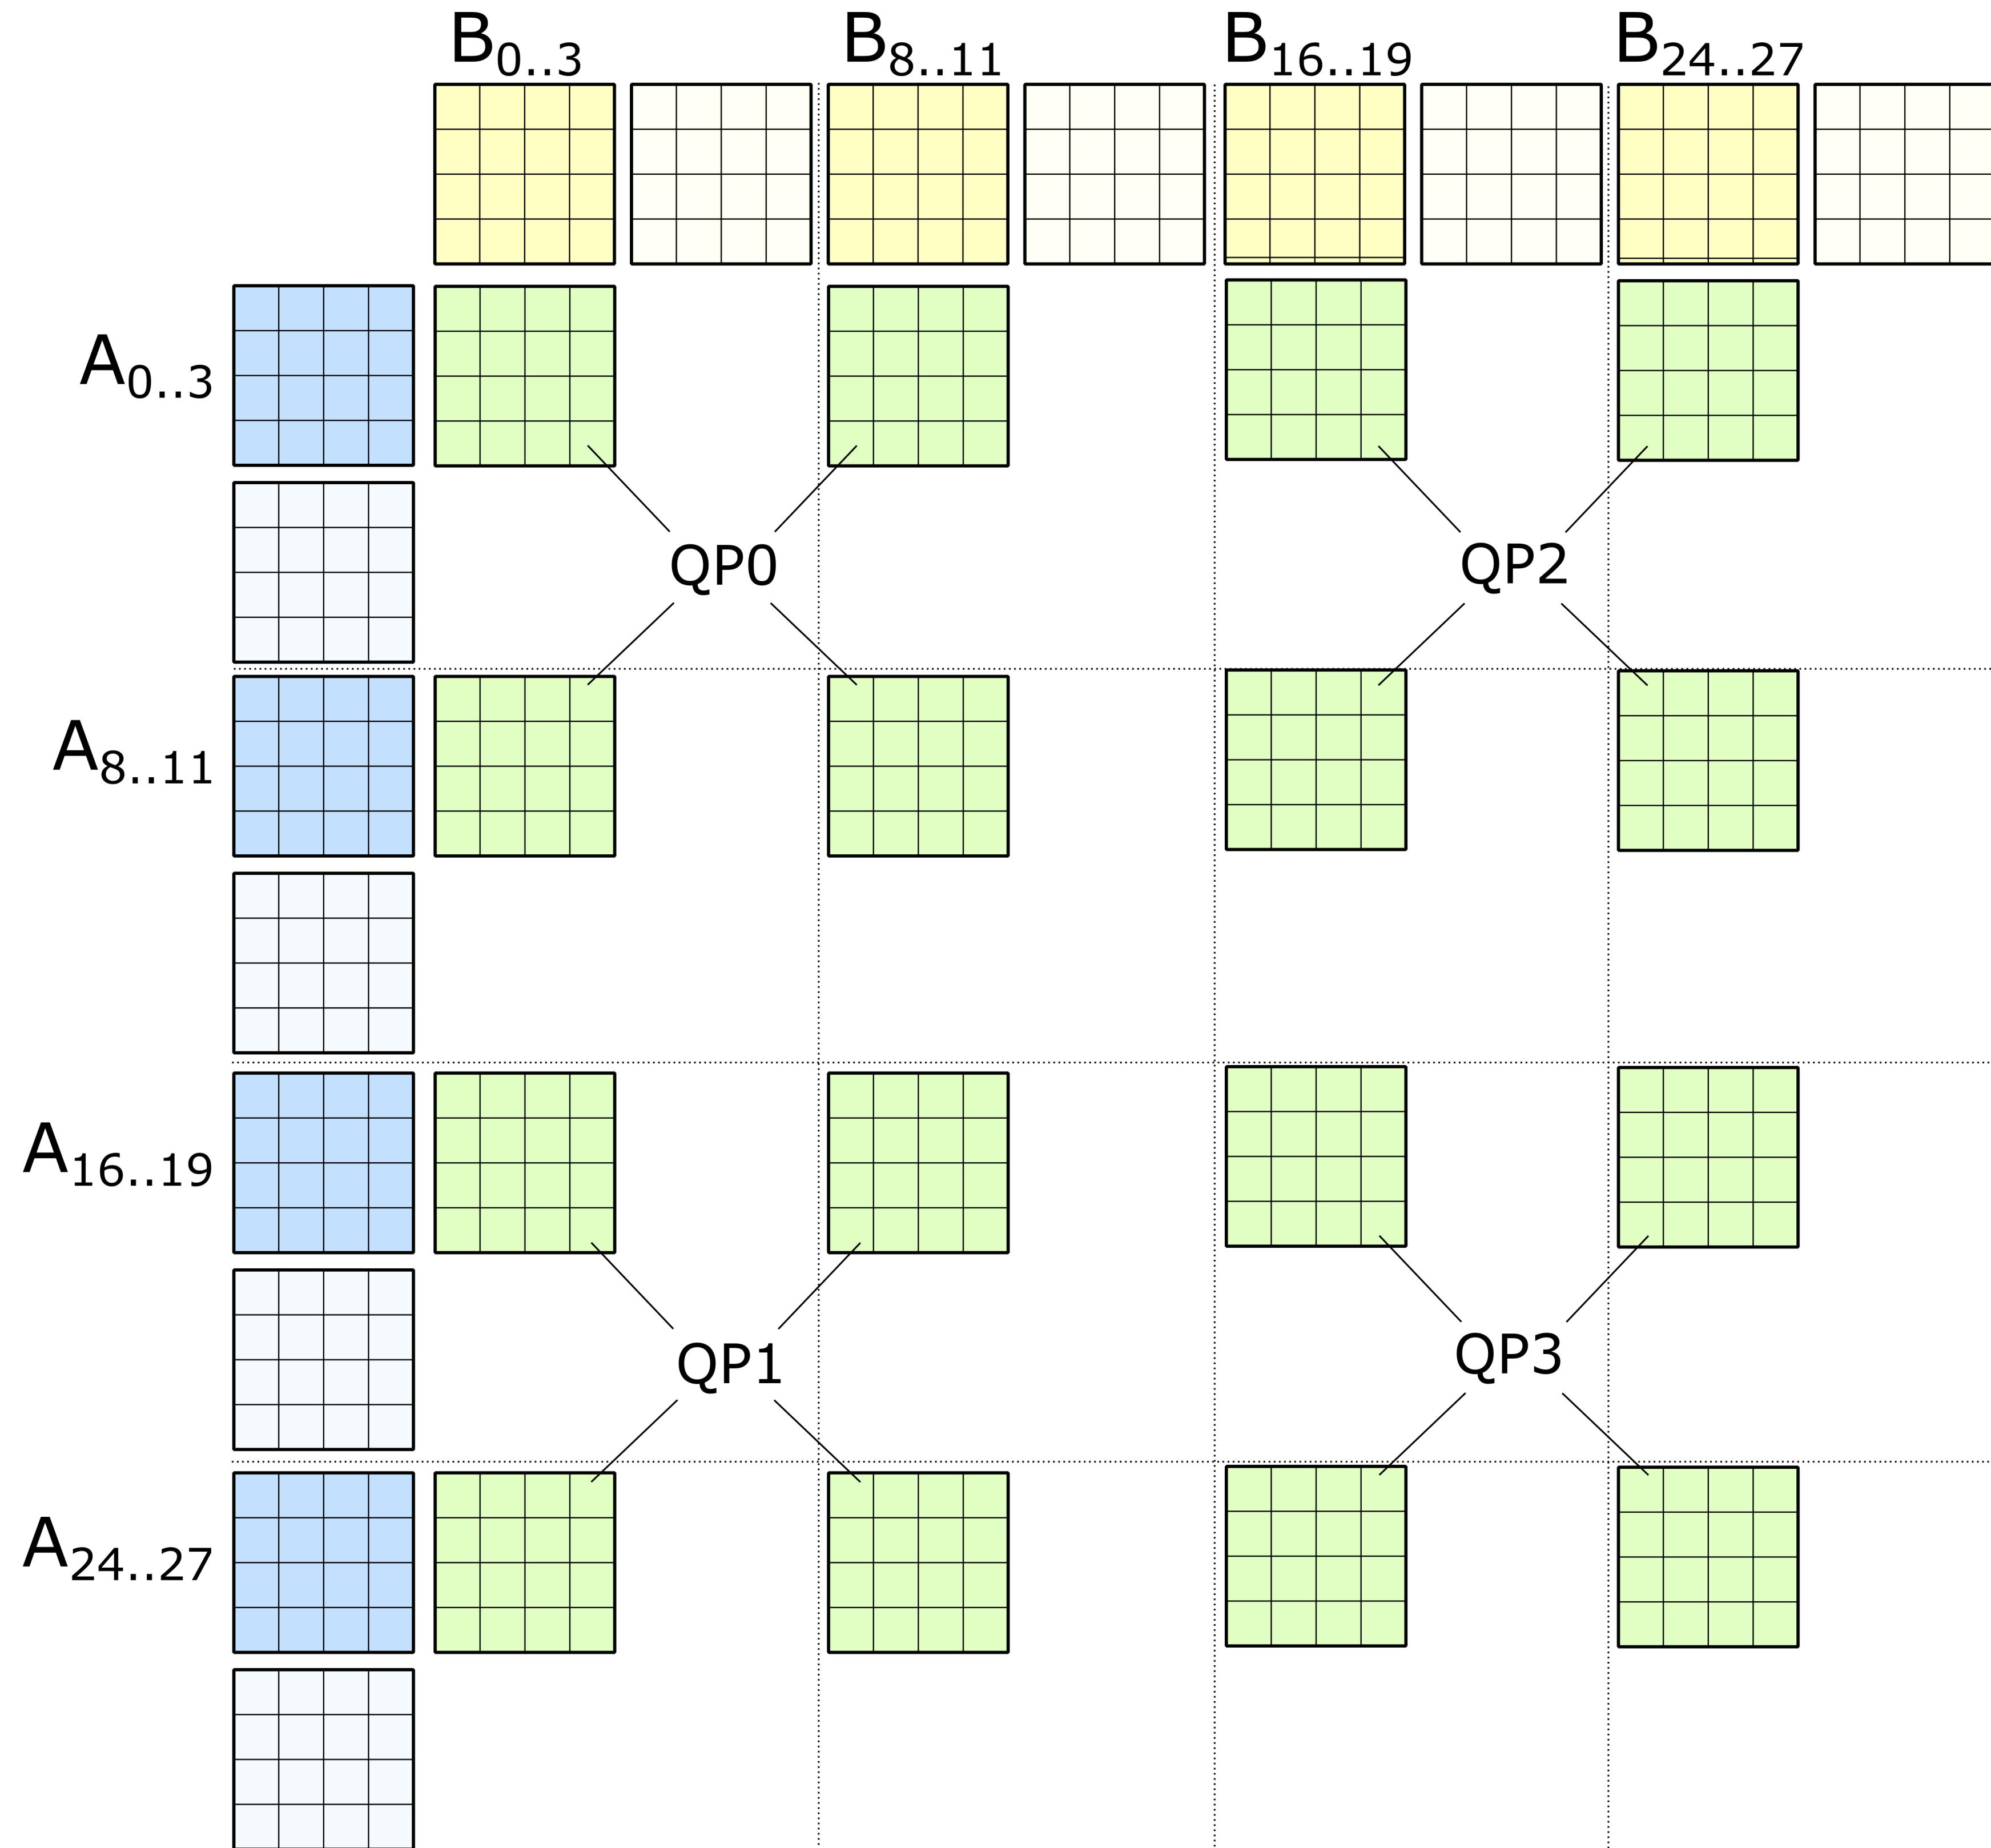
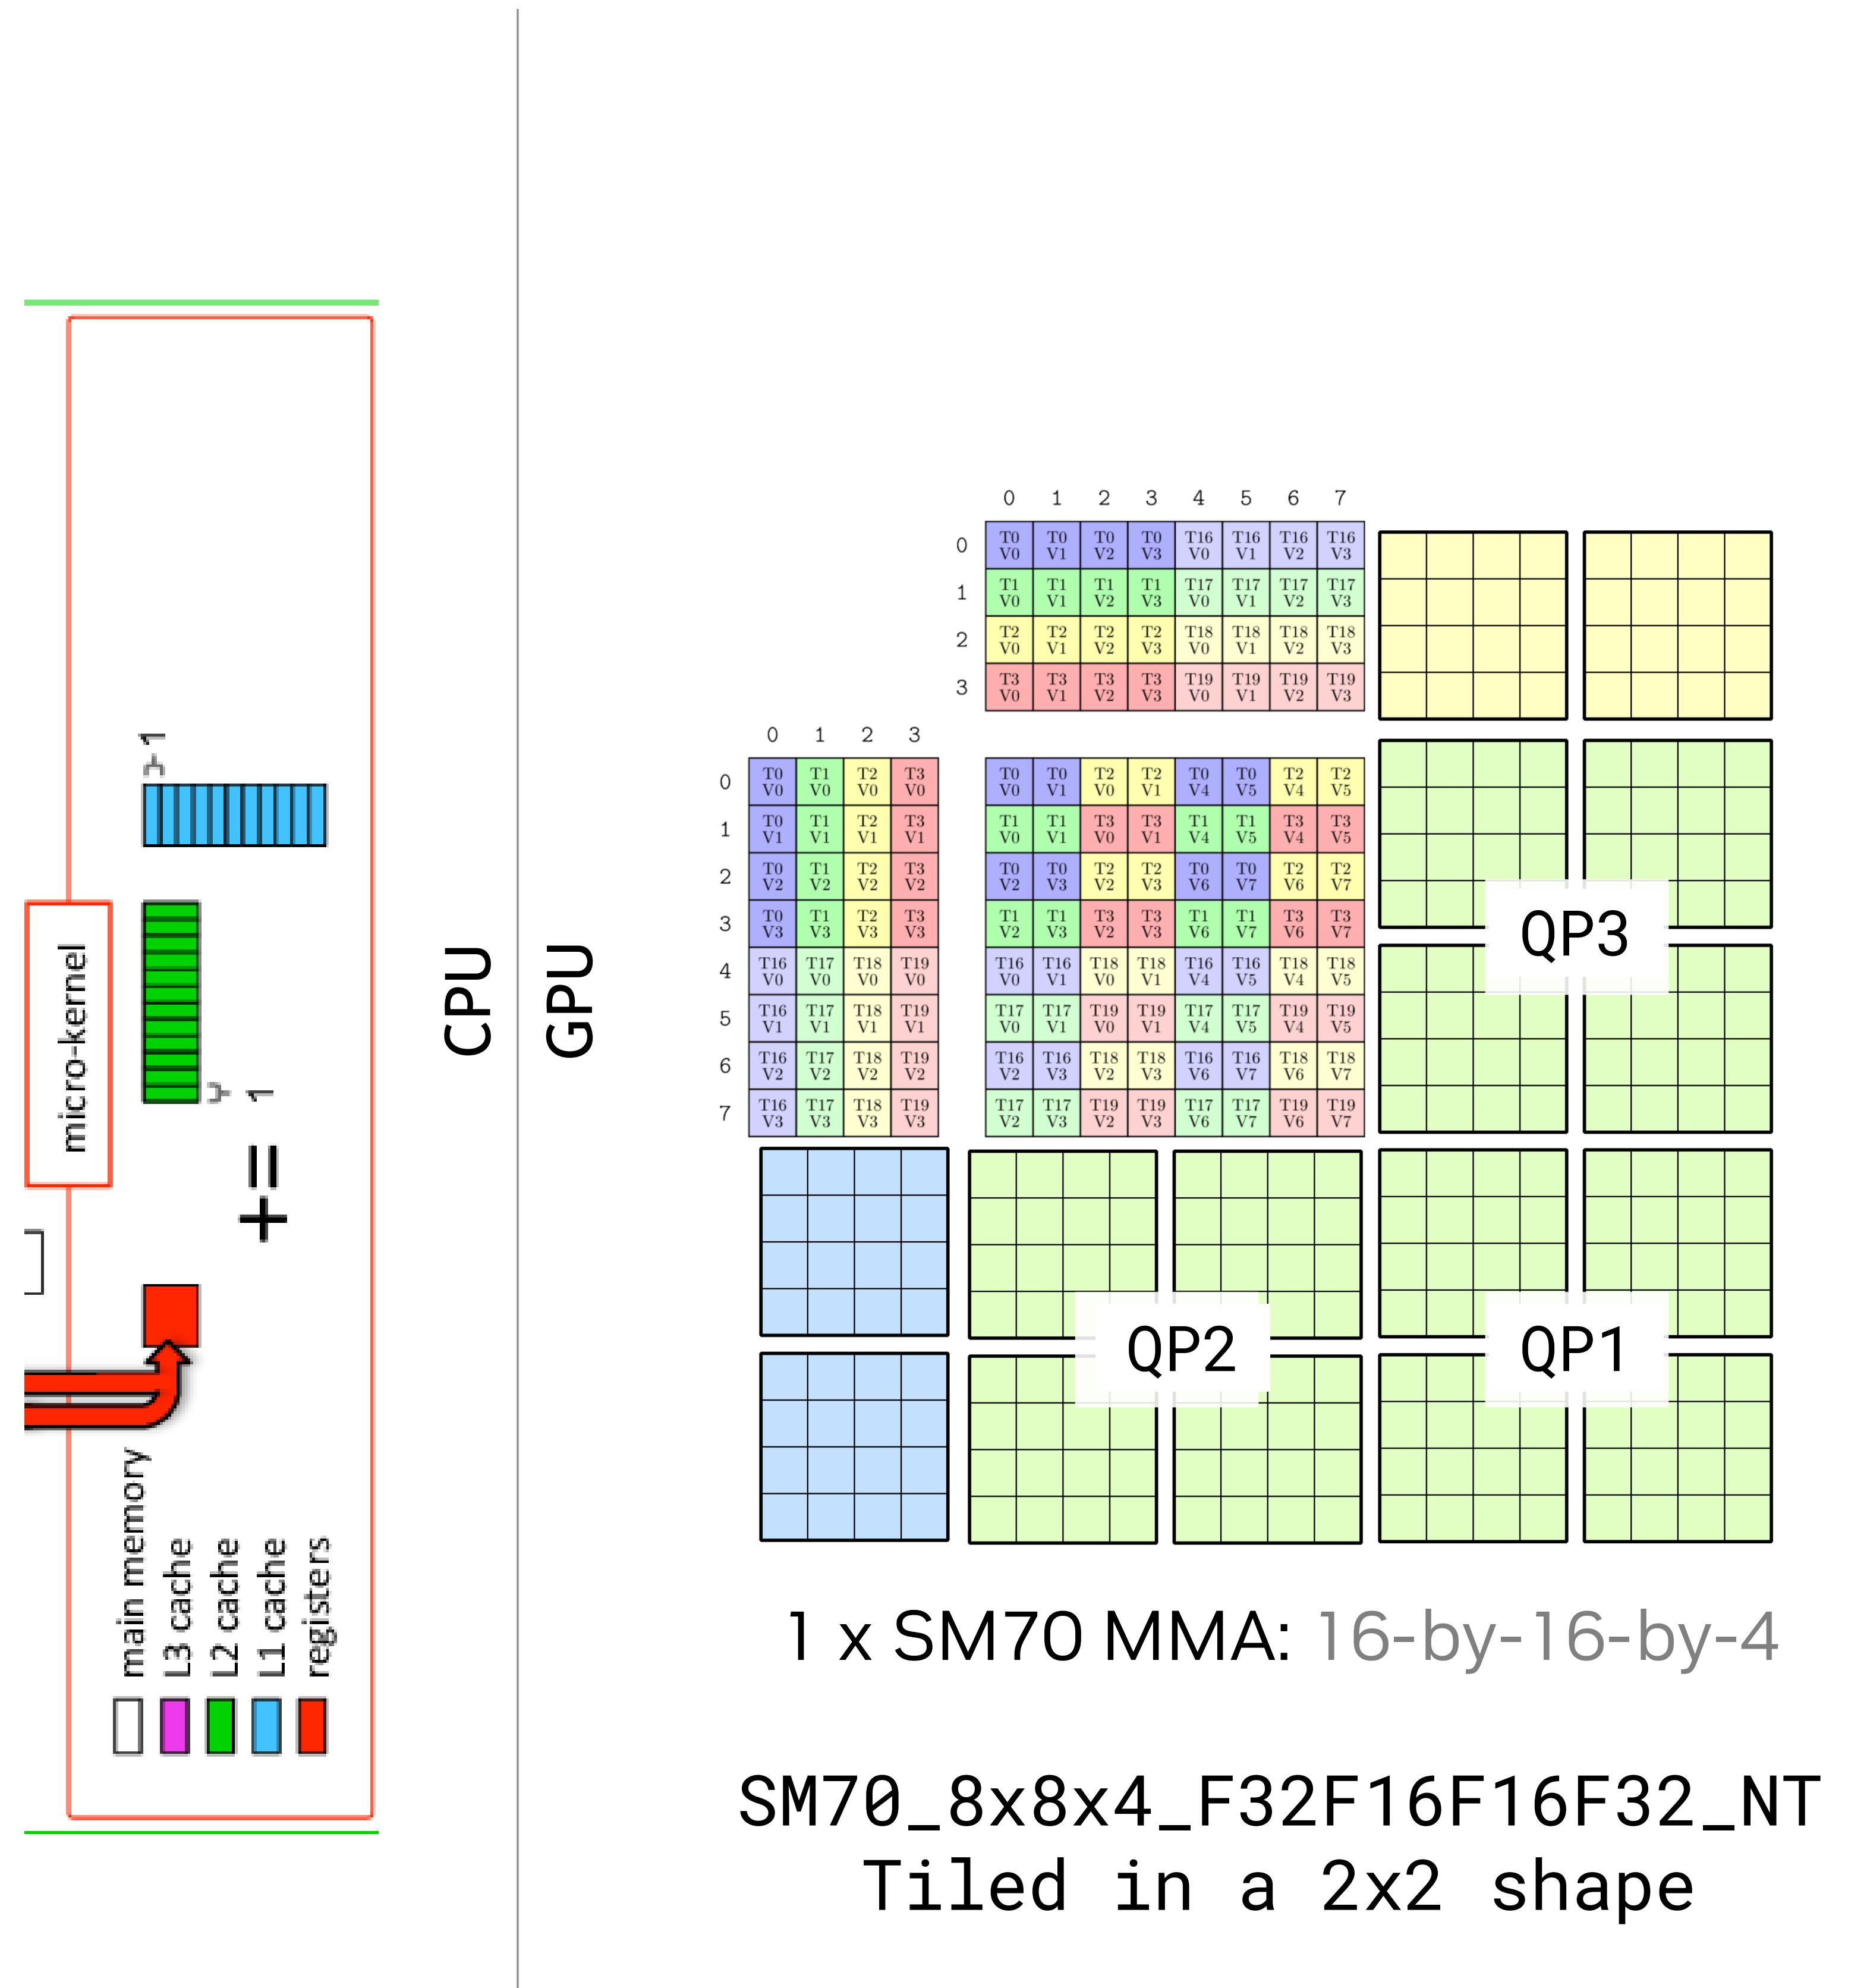
```
for (int warp_n = 0; warp_n < CtaN; warp_n += WarpN) {
    for (int warp_m = 0; warp_m < CtaM; warp_m += WarpM) {
        for (int warp_k = 0; warp_k < CtaK; warp_k += WarpK) {
```

// Copy from shared to registers

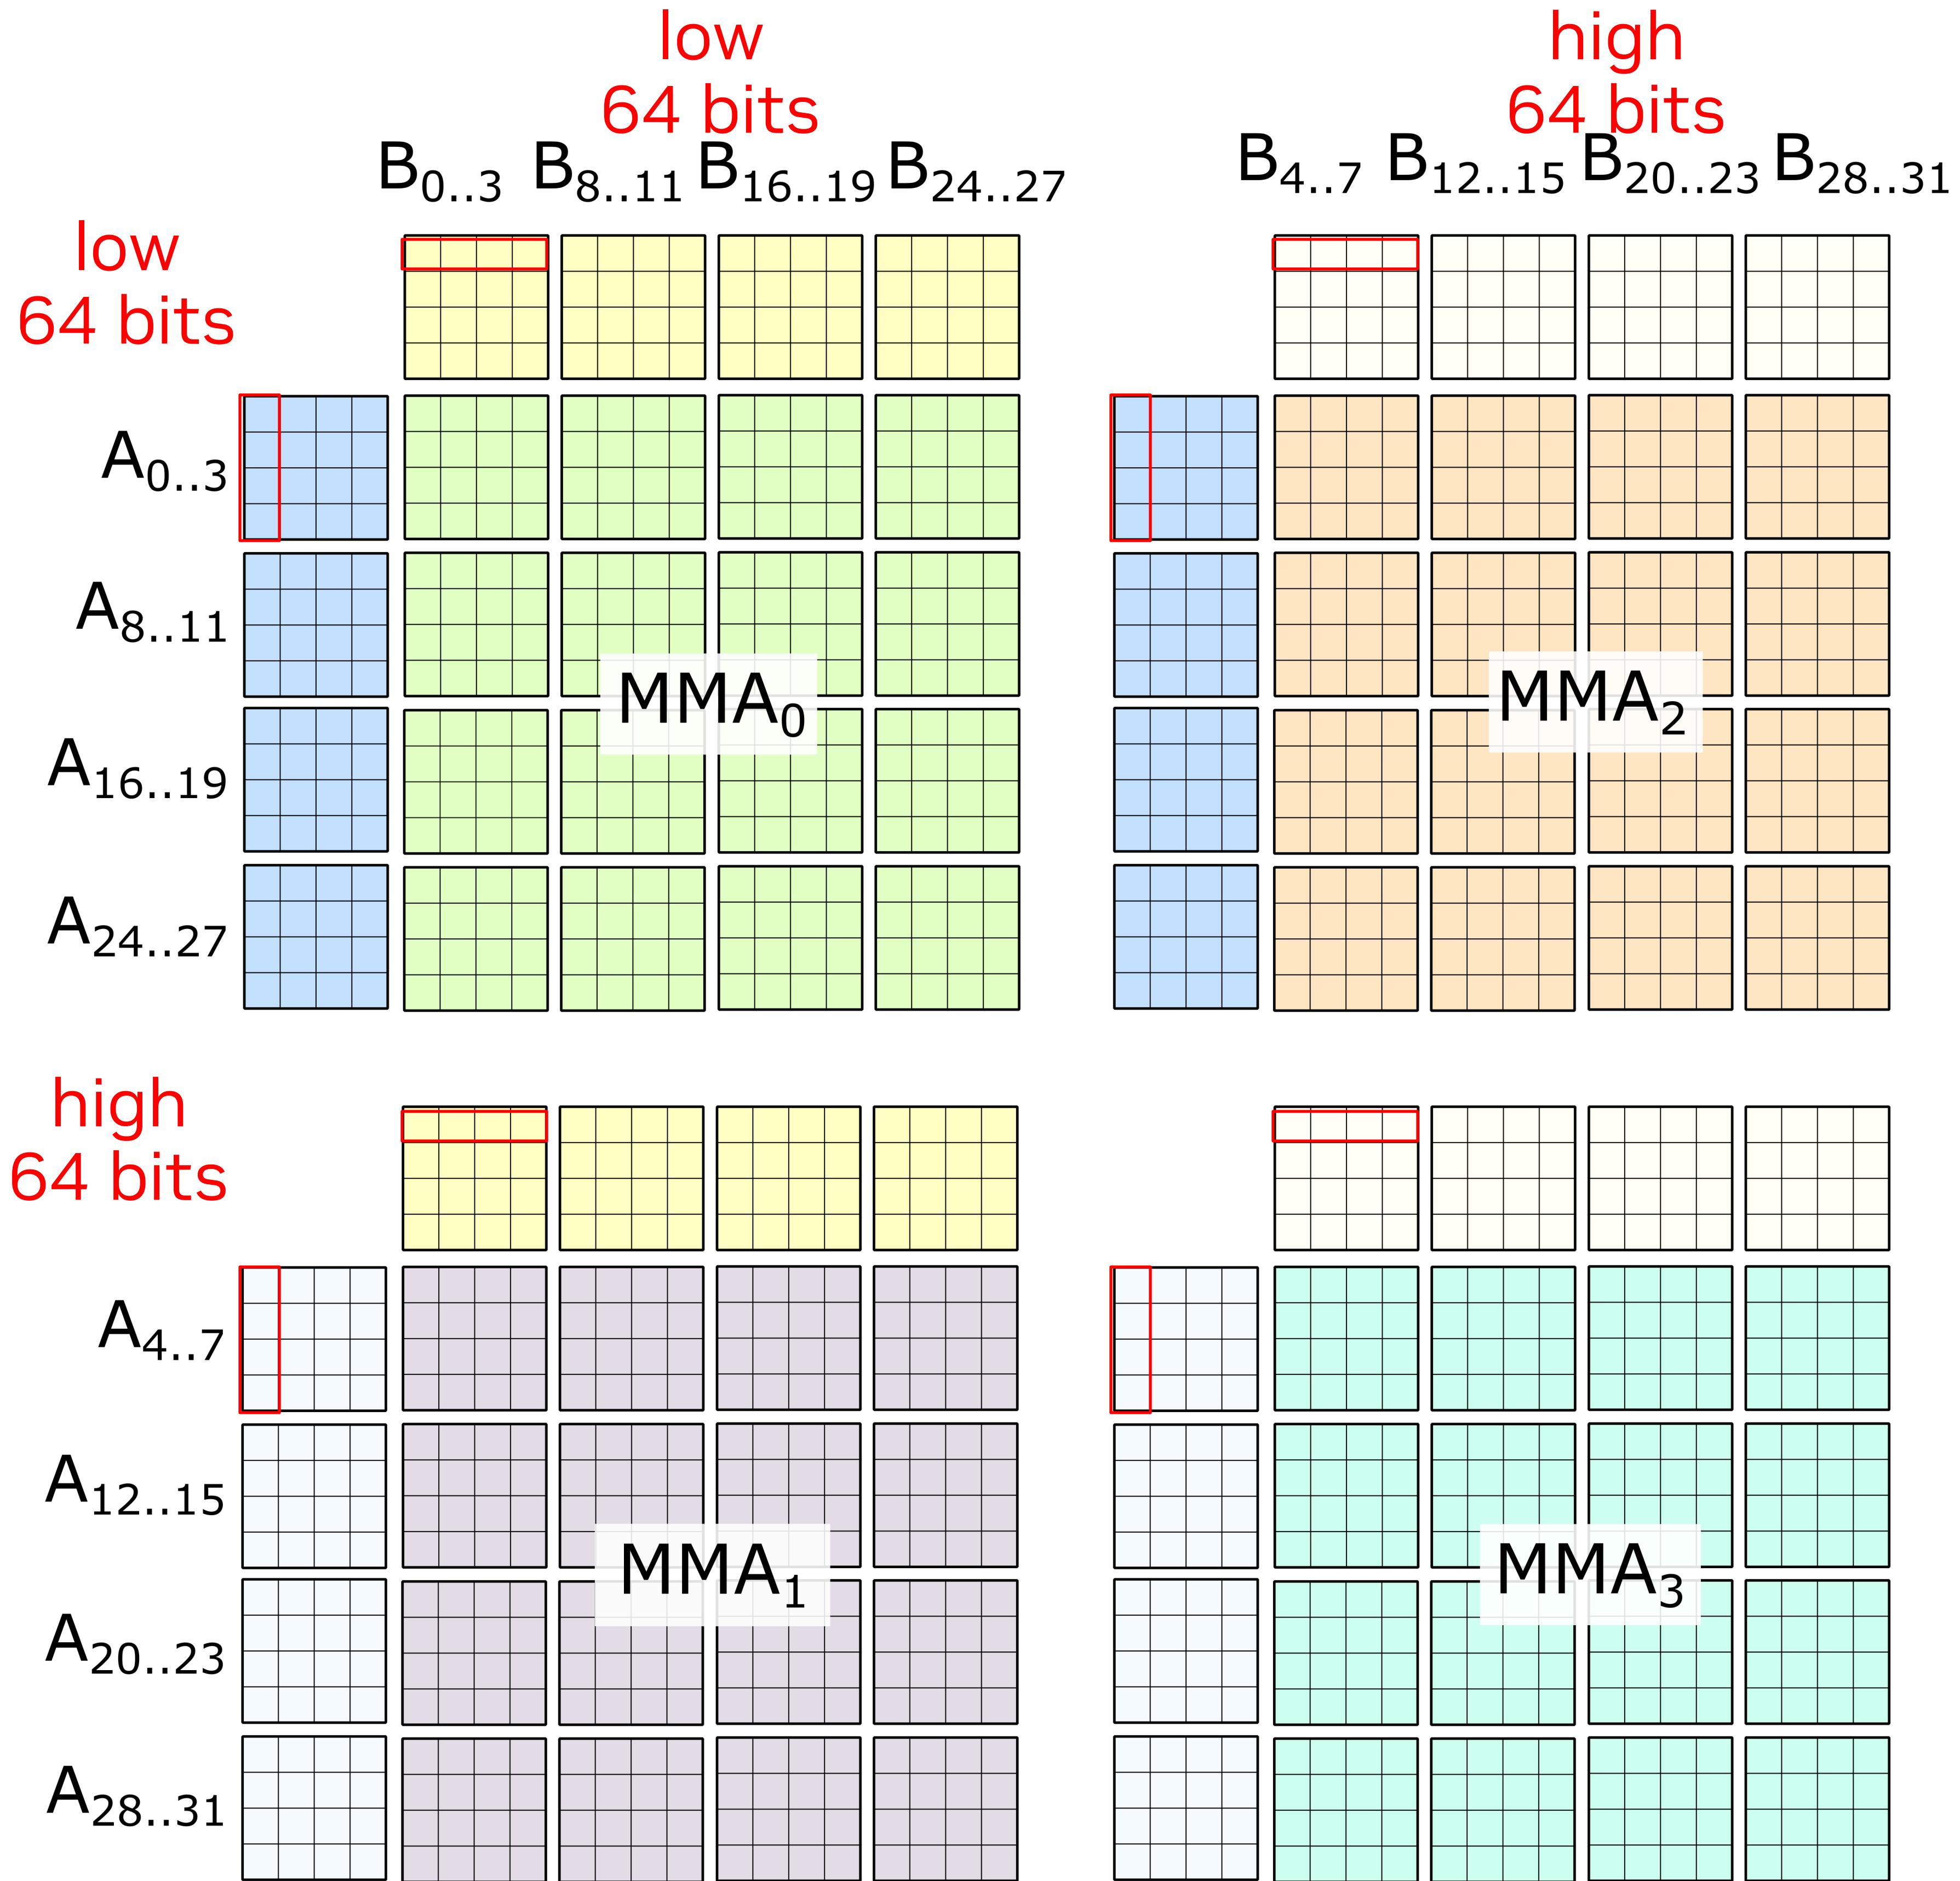
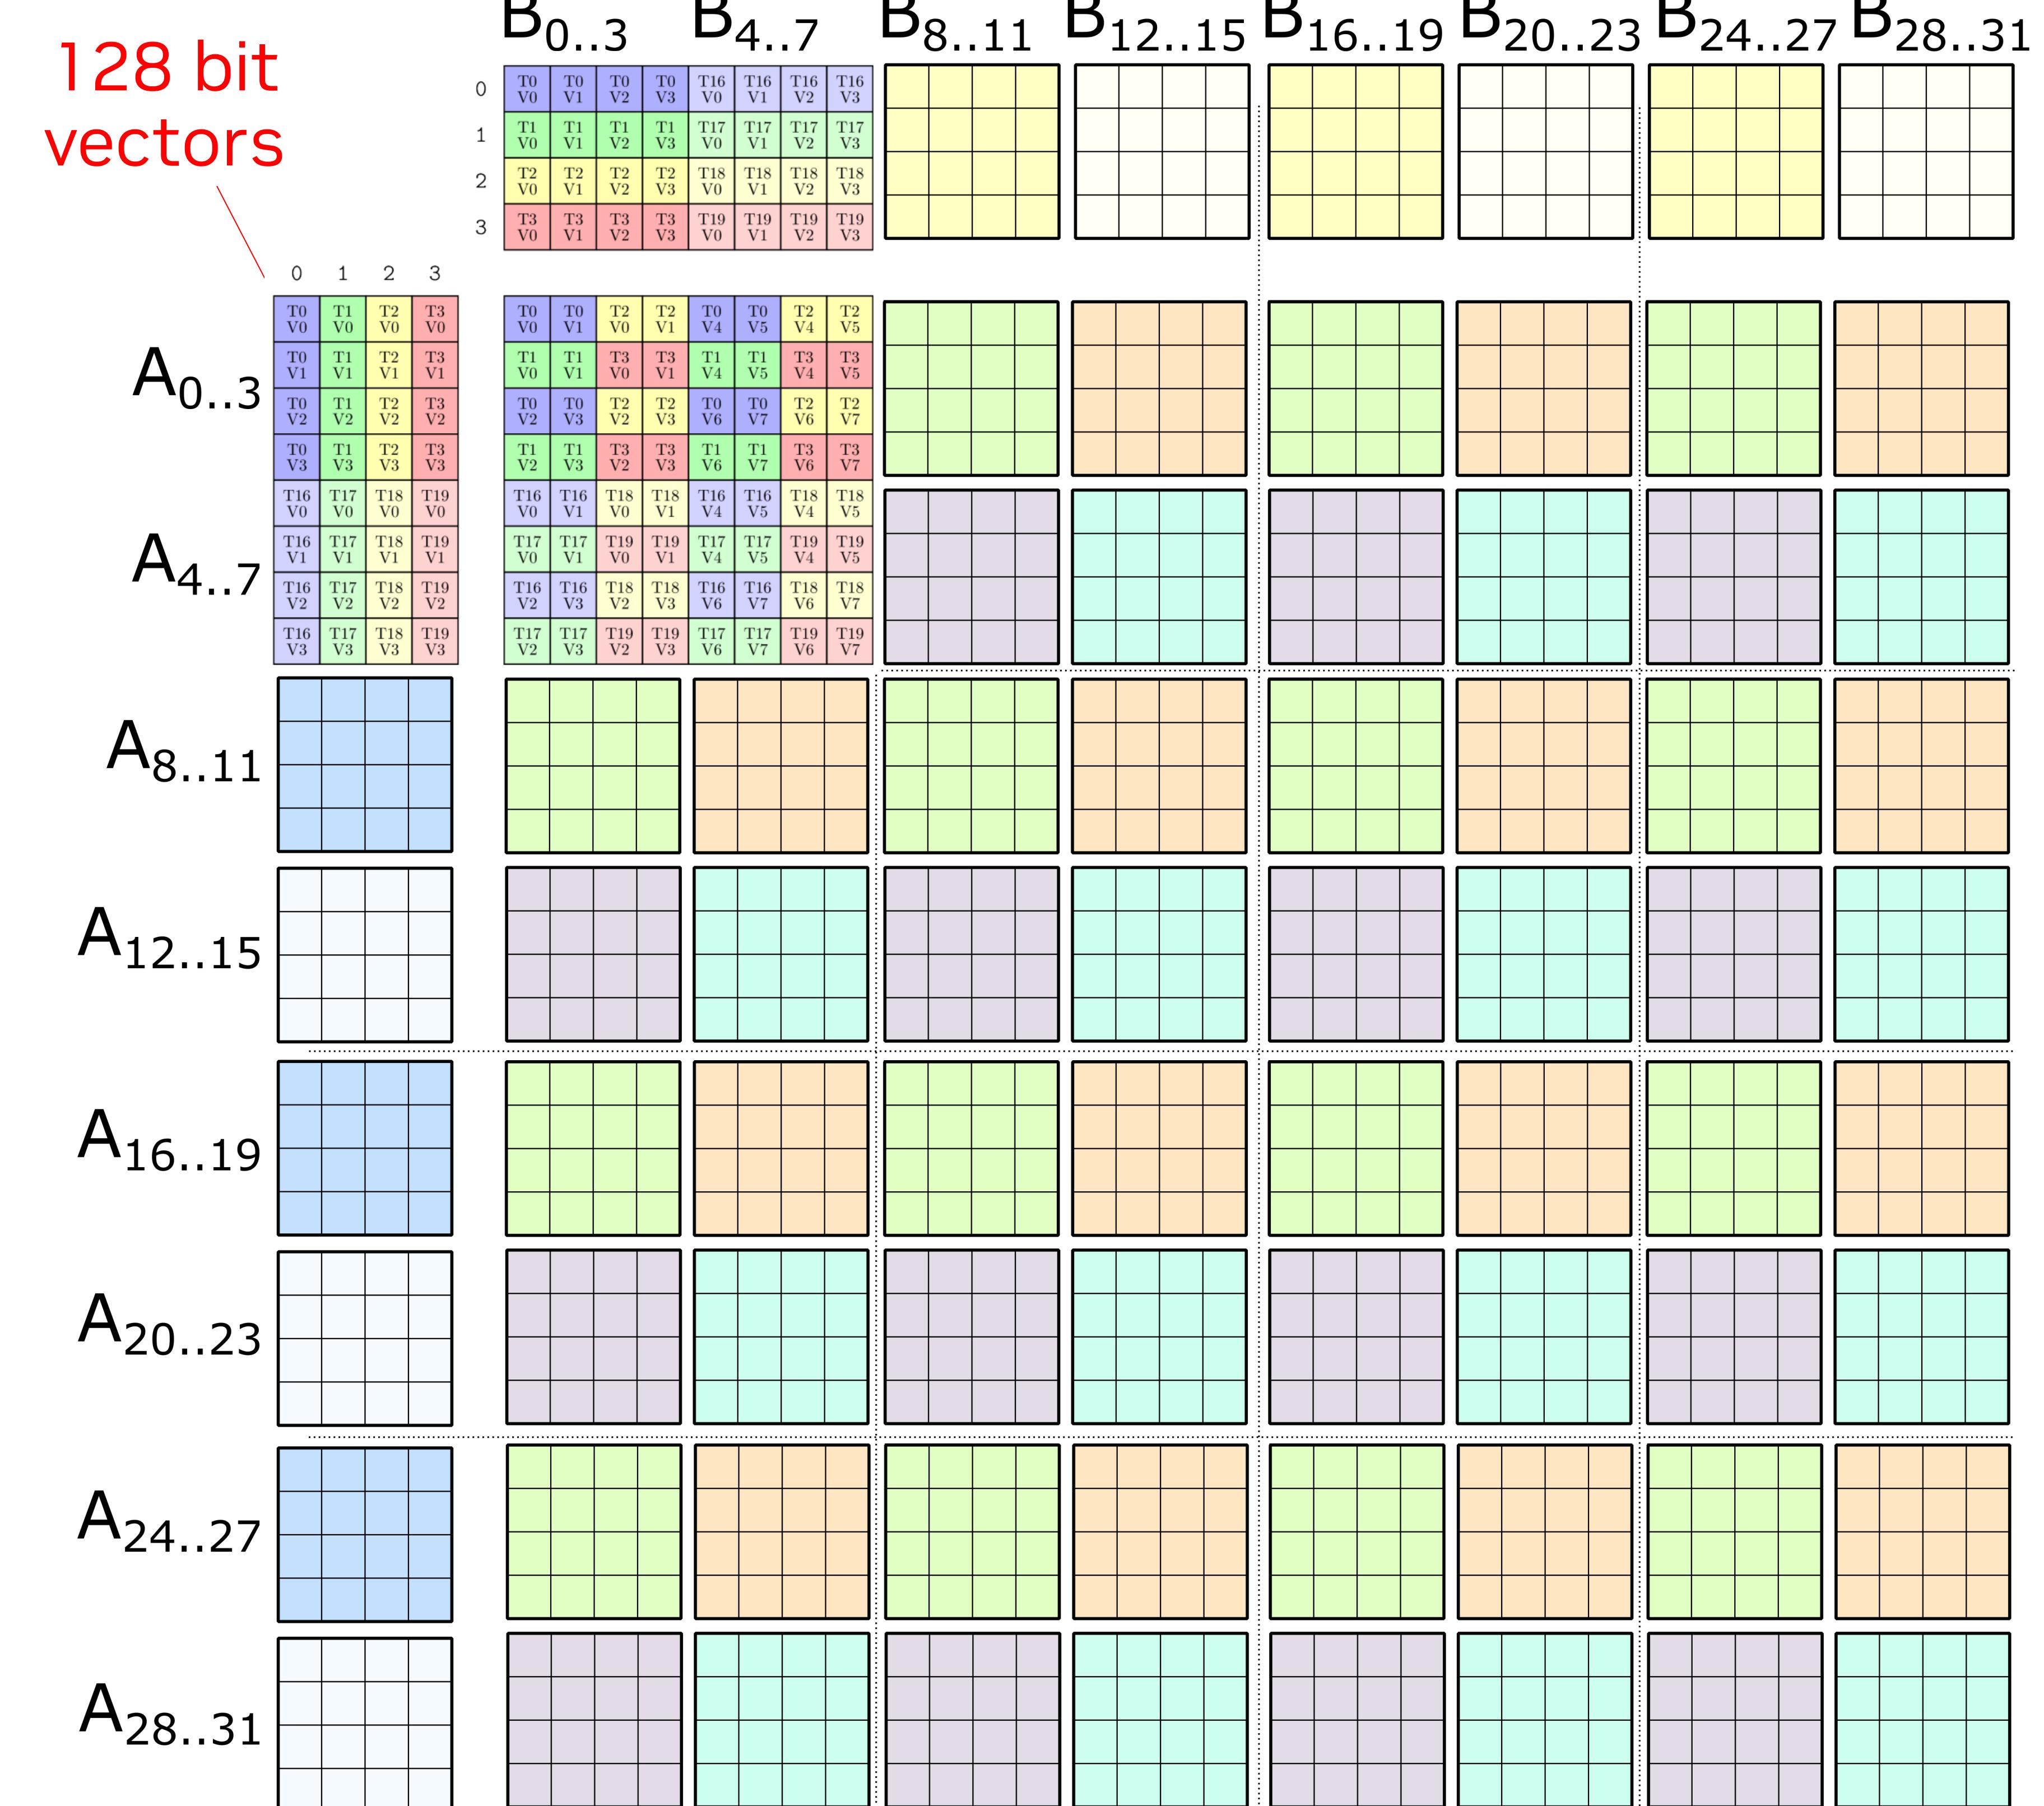
```
for (int mma_k = 0; mma_k < WarpK; mma_k += MmaK) {
    for (int mma_n = 0; mma_n < WarpN; mma_n += MmaN) {
        for (int mma_m = 0; mma_m < WarpM; mma_m += MmaM) {
```

// Compute in registers
gemm(a, b, c);

No Trivial Way To Express CPU like SIMD Loop on GPUs



Volta µKernel: Spatially Interleaved MMAs for 128 bit smem loads



2x2 Spatial interleaving allows 128 smem loads,
Allowing us to feed 2 MMAs from a single load

Ad-hoc partitioning doesn't scale

Bespoke layouts/iterators

- CUTLASS 2.x implements custom layout functors
 - Mapping logical coord -> index
- These iterators implement post-partitioned layouts
- Thread layouts are implicit in the index computations
- They cannot partition arbitrary input tensors
- Each new arch requires new layouts to be defined
- Indexing math is hard to implement, harder still to maintain

An excerpt out of CUTLASS 2.x's Volta SMEM access iterator
/// Returns the offset of a coordinate in linear memory.

```
CUTLASS_HOST_DEVICE
LongIndex operator()(TensorCoord const &coord) const {
    int vec_contiguous_idx = coord.contiguous() / kElementsPerAccess;
    int vec_strided_idx = coord.strided() / kFactor;

    int tile_contiguous_idx =
        vec_contiguous_idx / (TileShape::kContiguous / kFactor);
    int tile_contiguous_residual =
        vec_contiguous_idx % (TileShape::kContiguous / kFactor) +
        ((coord.strided() % kFactor) * (TileShape::kContiguous / kFactor));
    int tile_strided_residual = vec_strided_idx % TileShape::kStrided;

    int partition_contiguous_idx =
        tile_contiguous_residual / PartitionShape::kContiguous;
    int partition_strided_idx =
        tile_strided_residual / PartitionShape::kStrided;
    int partition_contiguous_residual =
        tile_contiguous_residual % PartitionShape::kContiguous;
    int partition_strided_residual =
        tile_strided_residual % PartitionShape::kStrided;

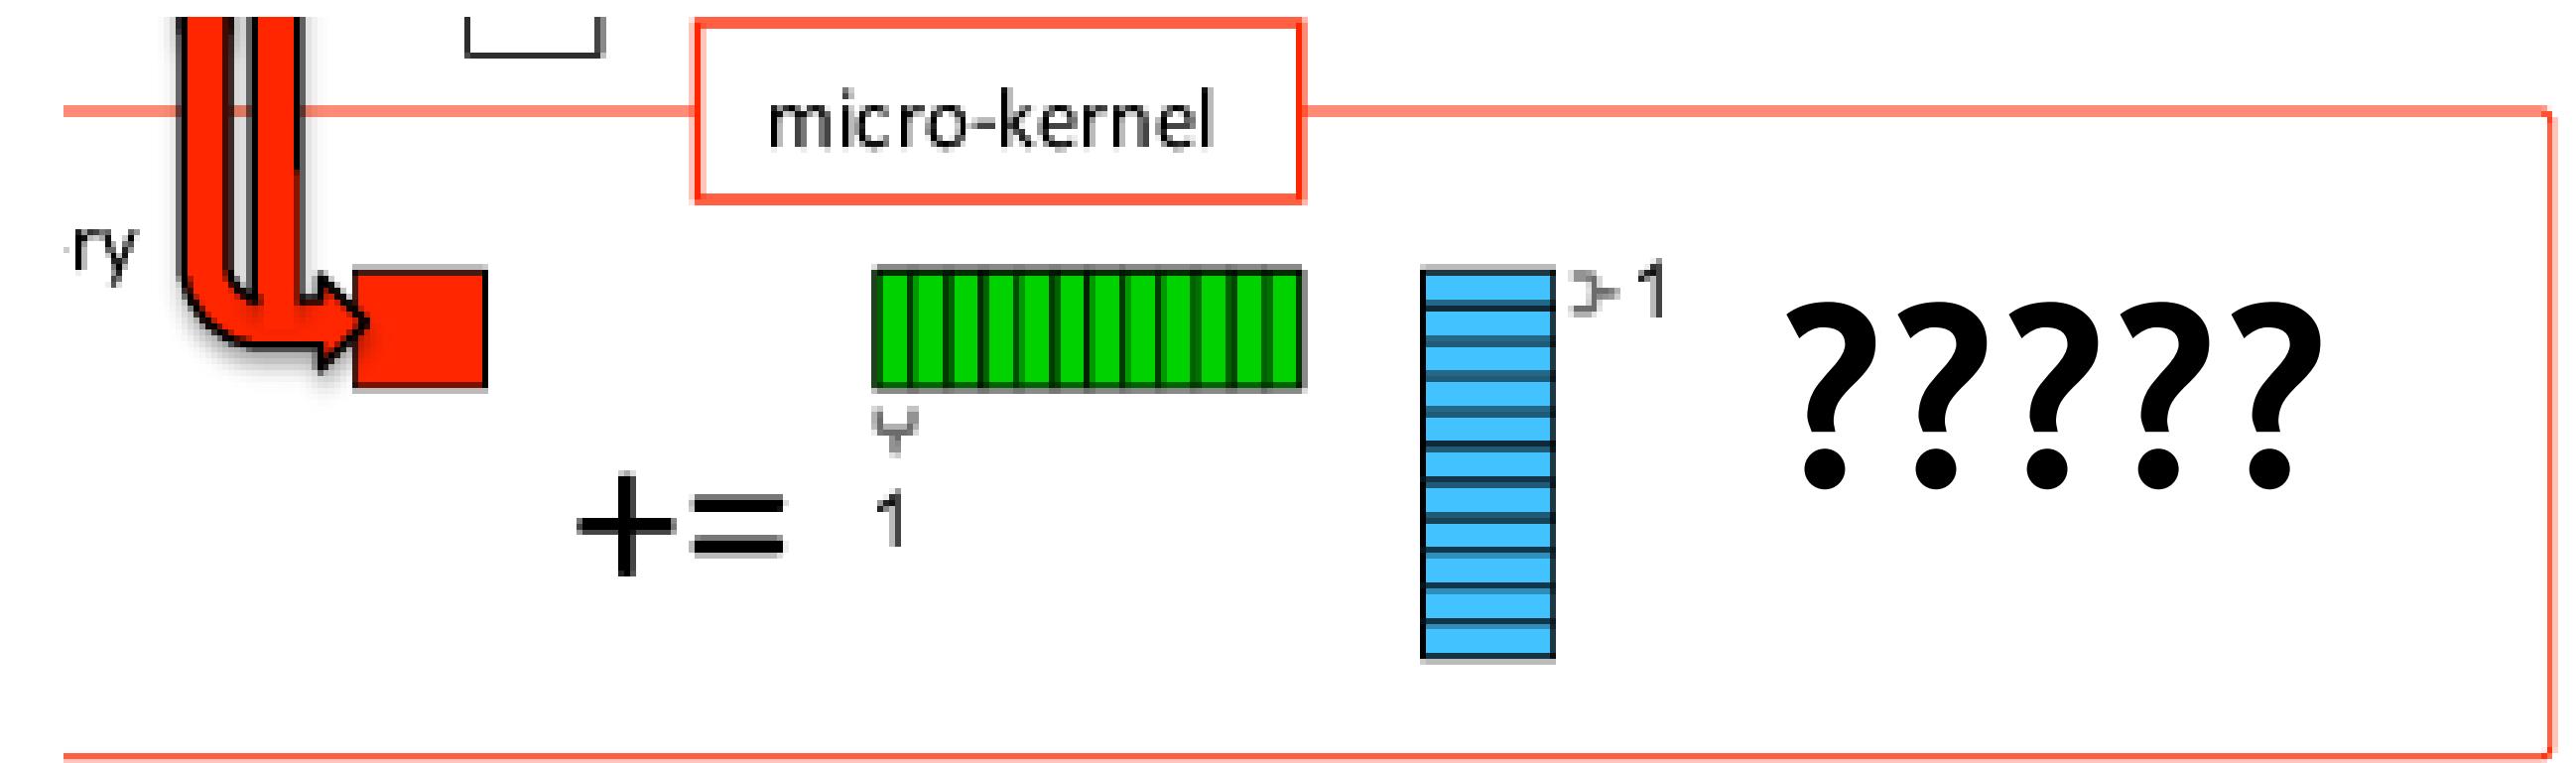
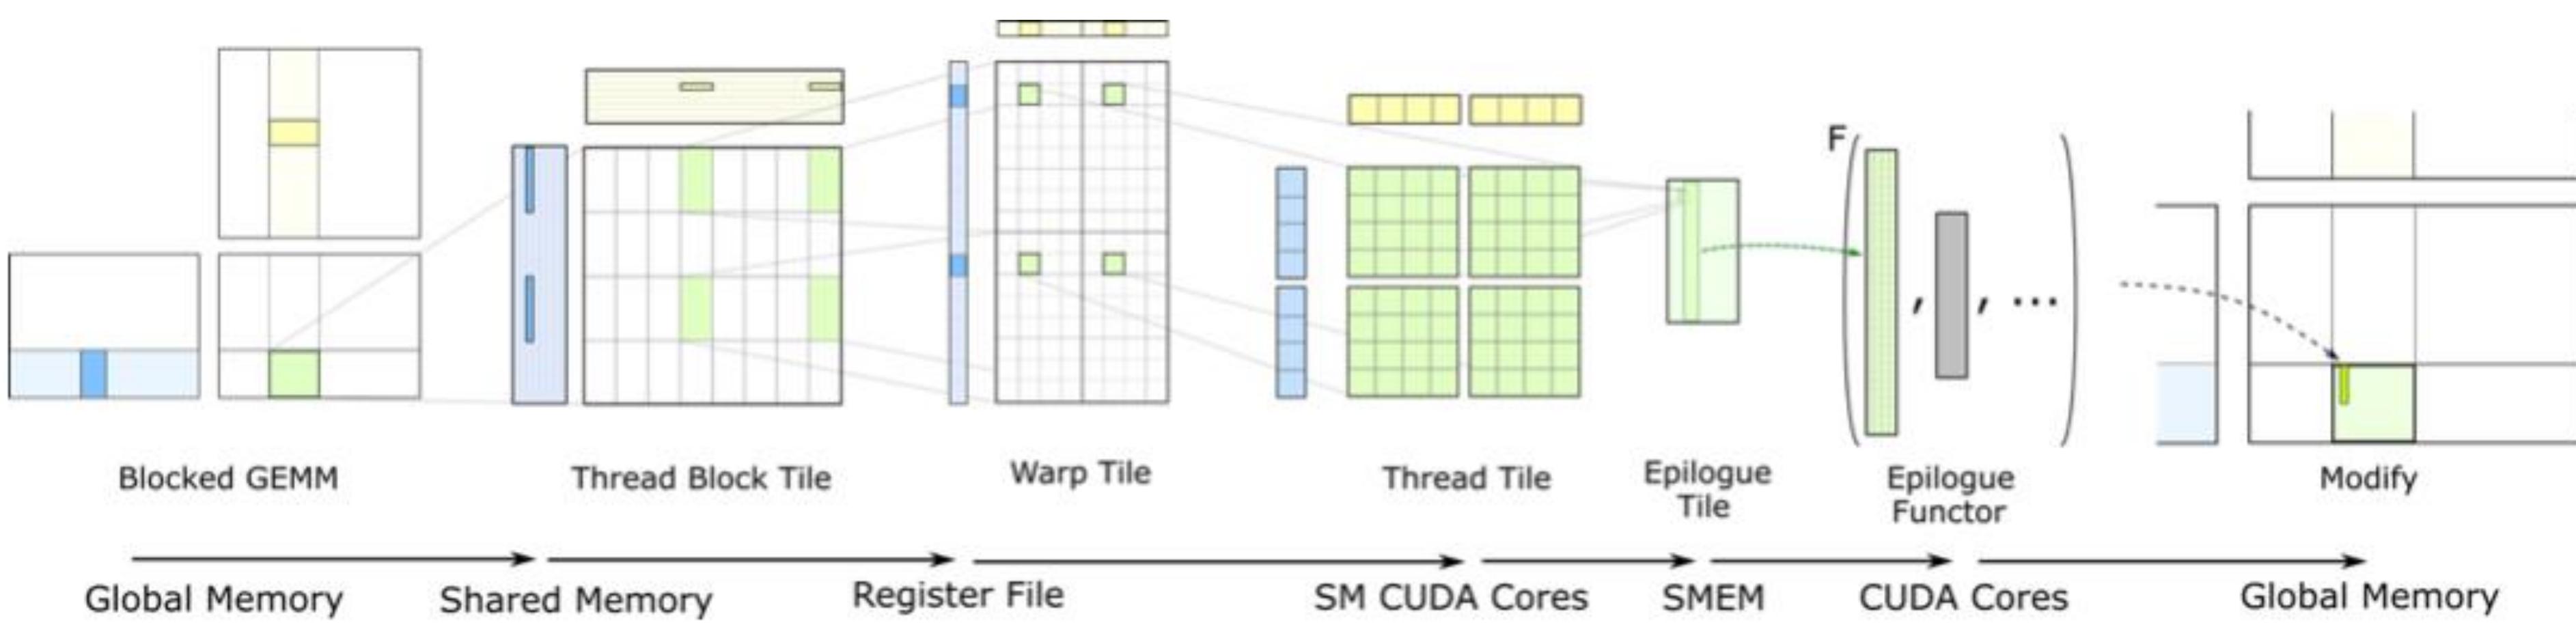
    int permuted_vec_contiguous_within_partition =
        partition_contiguous_residual ^ (partition_strided_residual % 4);
    int permuted_partition_contiguous_within_tile =
        partition_contiguous_idx ^ (partition_strided_idx % 2);

    int element_contiguous = (tile_contiguous_idx * TileShape::kContiguous +
        permuted_partition_contiguous_within_tile *
        PartitionShape::kContiguous +
        permuted_vec_contiguous_within_partition) *
        kElementsPerAccess +
        (coord.contiguous() % kElementsPerAccess);

    int element_strided = vec_strided_idx;
    return element_contiguous + element_strided * stride_[0] * kFactor;
```

Problem 1: Complicated Partitioning Patterns

Prevent us from writing canonical loops for all MMAs



```
for (int cta_n = 0; cta_n < GemmN; cta_n += CtaN) {
    for (int cta_m = 0; cta_m < GemmM; cta_m += CtaM) {
        for (int cta_k = 0; cta_k < GemmK; cta_k += CtaK) {
```

// Copy from global to shared

```
for (int warp_n = 0; warp_n < CtaN; warp_n += WarpN) {
    for (int warp_m = 0; warp_m < CtaM; warp_m += WarpM) {
        for (int warp_k = 0; warp_k < CtaK; warp_k += WarpK) {
```

// Copy from shared to registers

```
for (int mma_k = 0; mma_k < WarpK; mma_k += MmaK) {
    for (int mma_n = 0; mma_n < WarpN; mma_n += MmaN) {
        for (int mma_m = 0; mma_m < WarpM; mma_m += MmaM) {
```

// Compute in registers

gemm(a, b, c);

	B matrix data addresses								
	0	1	16	17	32	4	5	20	21
16	0	1	16	17	2	3	18	19	
32	4	5	20	21	6	7	22	23	
48	8	9	24	25	10	11	26	27	
64	10	11	26	27	12	13	28	29	
80	12	13	28	29	14	15	30	31	
96	14	15	30	31	16	17	32	33	
112	18	19	34	35	20	21	36	37	

Figure 3b: Maxwell thread assignments for LDS.U matching

Maxwell A access addr[tid] = addr[id^4]

Maxwell B access addr[tid] = addr[id^2]

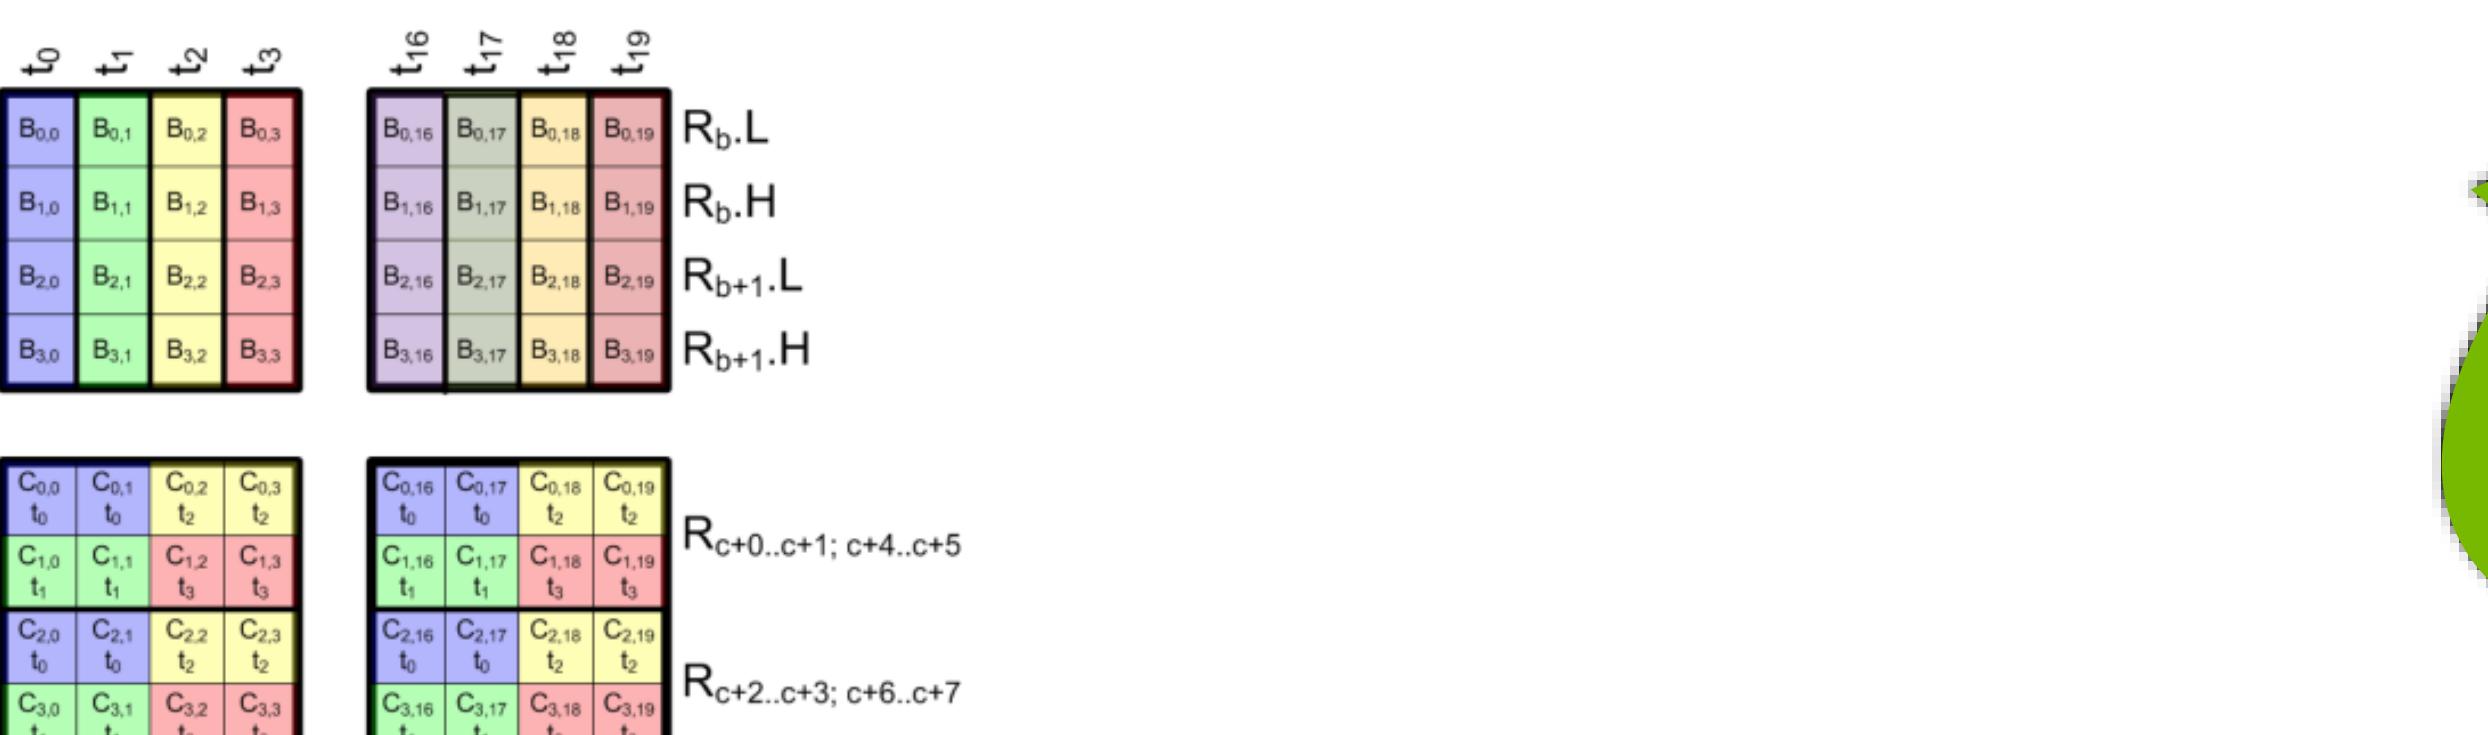


Figure 1b: HMMA.884.F32.TN data layout.
Again, ignore the apparent isomorphism
between thread numbers and matrix indices.

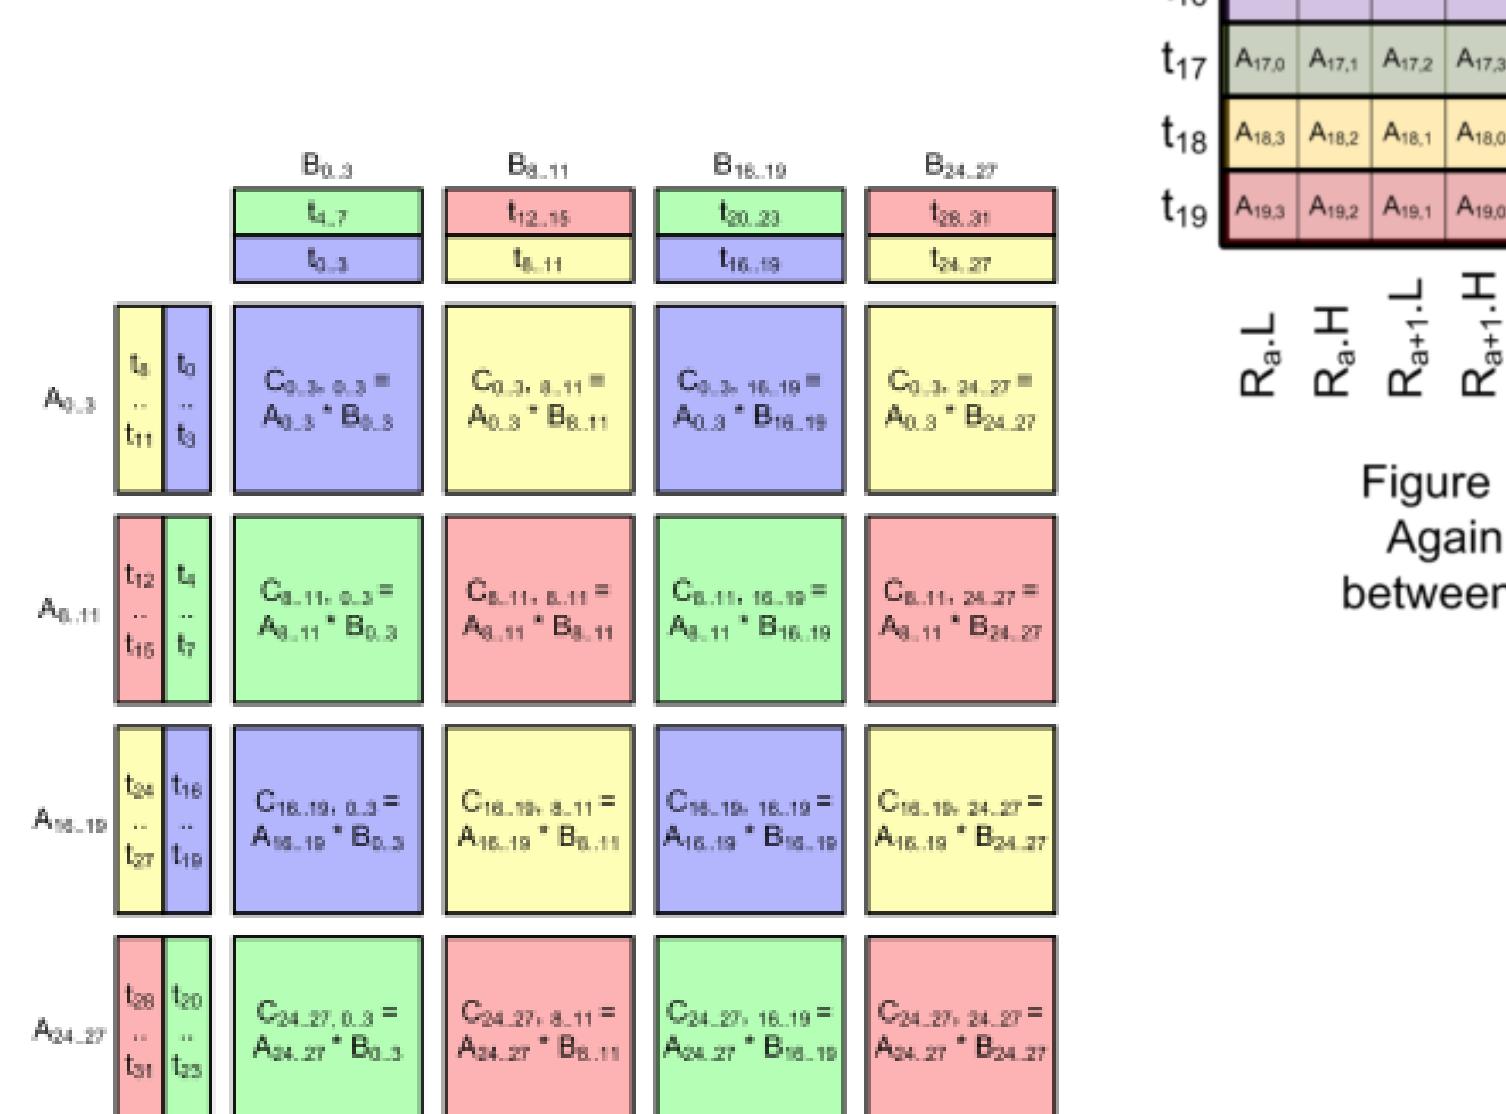
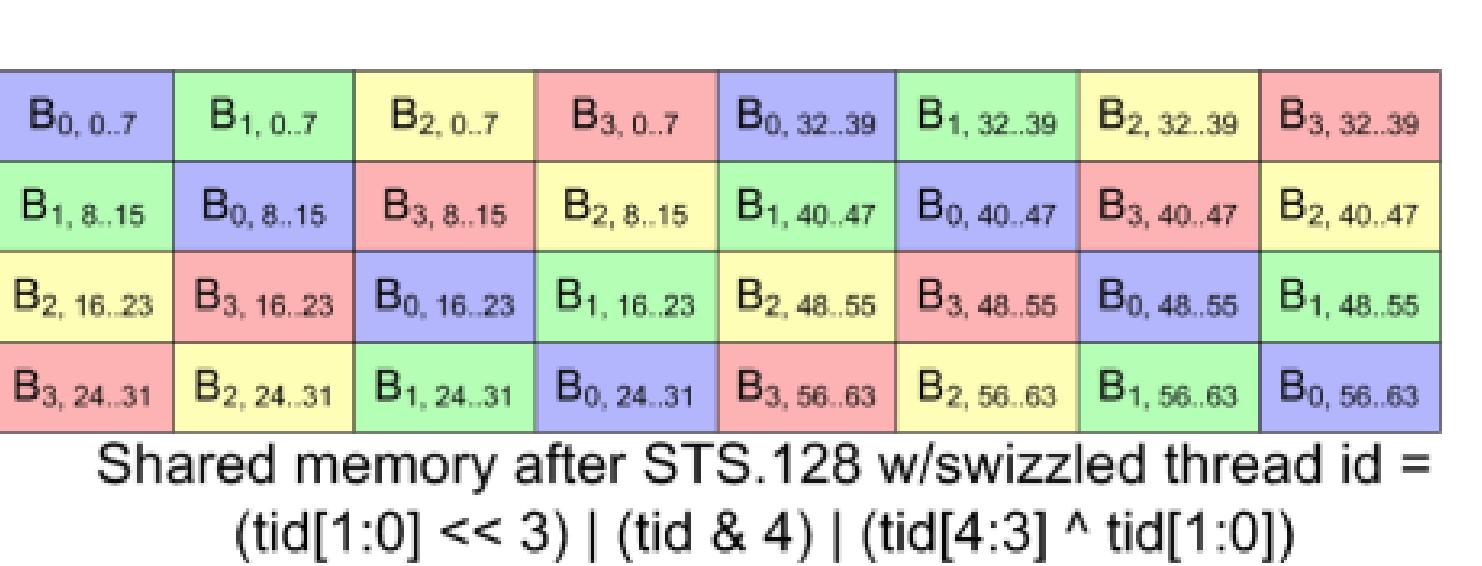


Figure 2a: Composing sparse A[16,4] * B[4, 16] = C[16, 16]
warp-wide multiply from four A[8, 4] * B[4, 8] = C[8, 8] 8-thread multiplies.



LDG.128 register quad (R_{n..n+3})

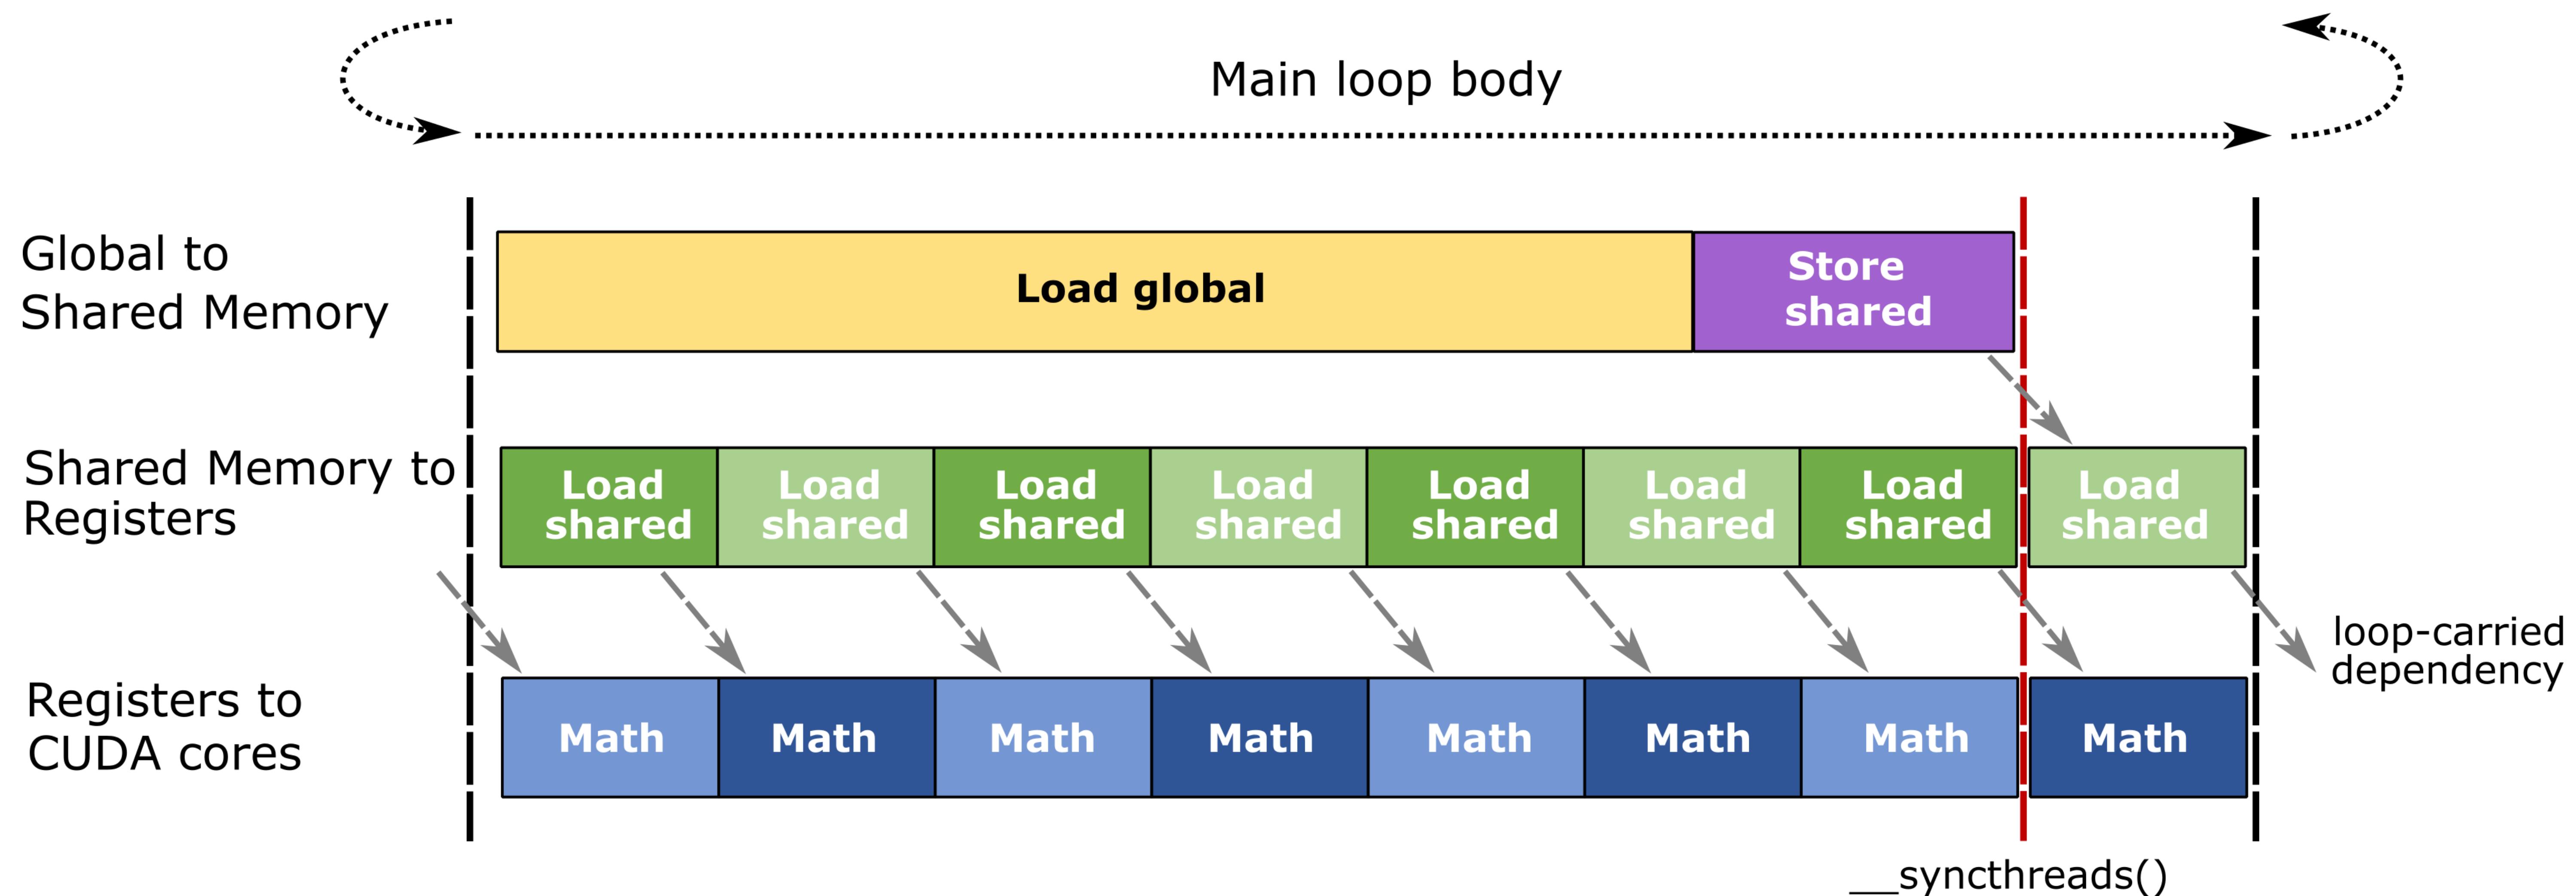
Shared memory after STS.128 w/swizzled thread id =
(tid[1:0] <> 3) | (tid & 4) | (tid[4:3] ^ tid[1:0])



Ampere Mainloop

ASYNC in SMEM and ILP in RMEM

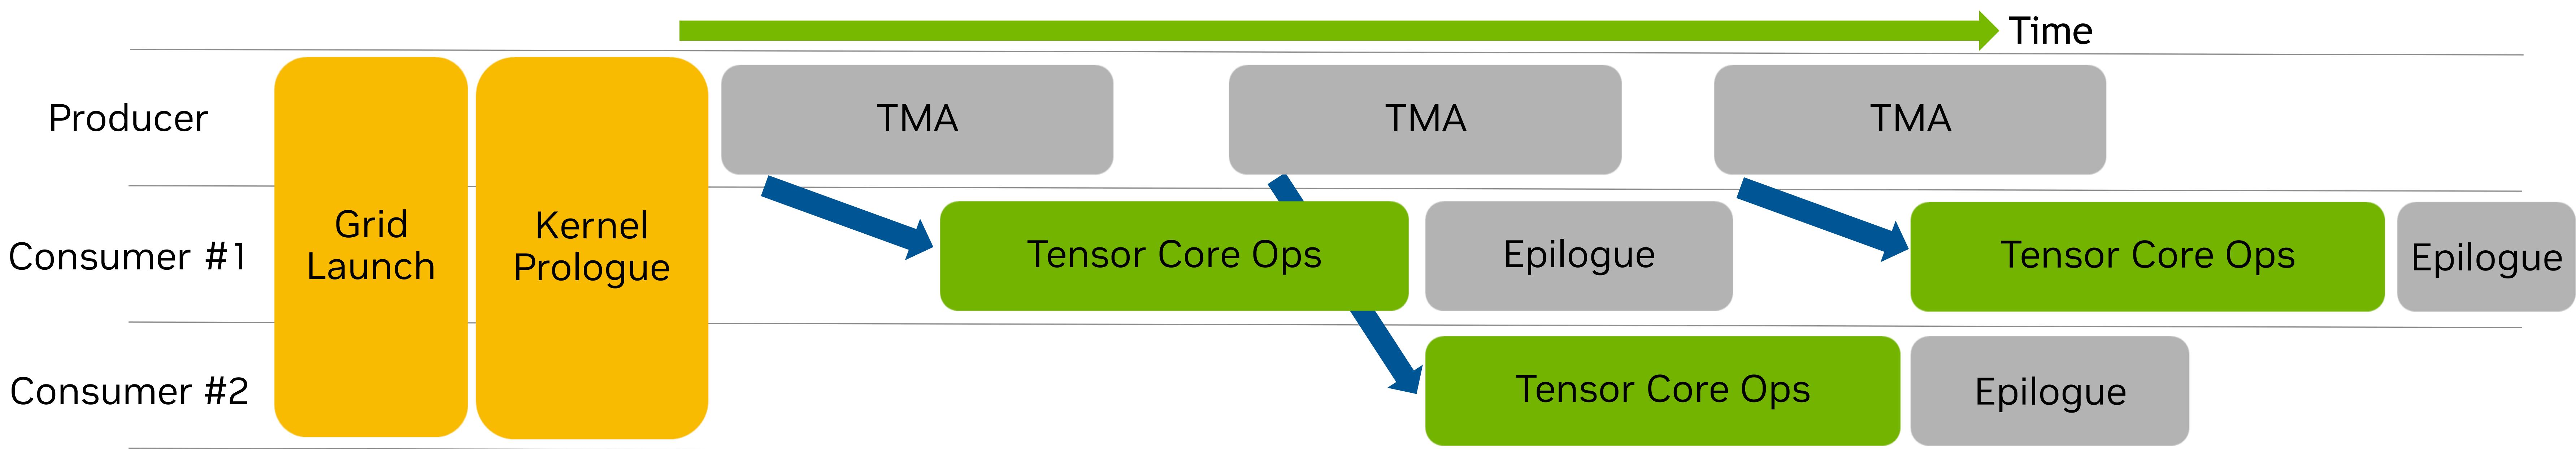
- Instruction interleaving between math and copy critical for speed of light performance
- Async pipeline feeds loads->compute
- Staged via local memories (smem and rmem)



Problem 2: Programmer Managed Asynchrony

GPUs require deeply async, managed, producer/consumer software pipelines

- Feeding the tensor cores constantly is hard – requires managing asynchrony and deep software pipelines
- With newer architectures like Hopper, even the MMA instruction is asynchronous
- Concurrency programming - Writing kernels isn't just about getting the layouts right anymore



Producer	Consumer #1	Consumer #2
<pre>PersistentTileSchedulerSm90 scheduler(problem_shape, blk_shape, cluster_shape)</pre>	<pre>// Mainloop, epilogue, and data out</pre>	
<pre>// Data in via TMA</pre> <pre>while (work_tile_info.is_valid_tile) {</pre> <pre> collective_mainloop.load()</pre> <pre> scheduler.advance_to_next_work()</pre> <pre> work_tile_info = scheduler.get_current_work()</pre> <pre>}</pre>	<pre>while (work_tile_info.is_valid_tile) {</pre> <pre> collective_mainloop.compute()</pre> <pre> scheduler.advance_to_next_work(NumConsumers)</pre> <pre> work_tile_info = scheduler.get_current_work()</pre> <pre>}</pre>	



CUTLASS 3.x

CUTLASS

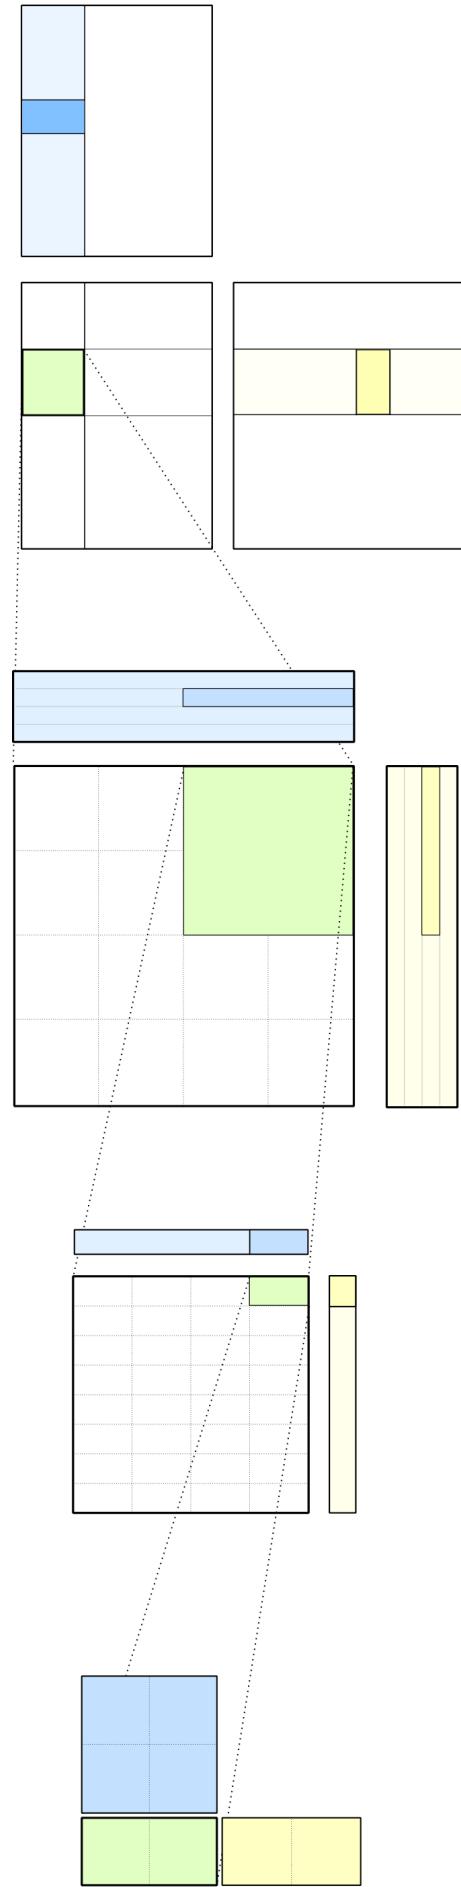
CUDA C++ Template Library for High Performance Linear Algebra



Tensor core computations at all scopes and scales, decomposed into their “moving parts”

Provides a native **tile-based programming model for GPU kernels**

Device	{ GEMM, Convolution, Reductions , BLAS3 } x { all data types } x { SIMT, Tensor Cores } x { all architectures }
Kernel	GEMM, Batched GEMM, Convolution, Reduction, Fused output operations, Fused input operations
Collective	CUTLASS temporal micro-kernels (async producer/consumer pipelines orchestrating spatial micro-kernels)
Atom	CuTe spatial micro-kernels (Tiled MMA / Copy)
Thread	Numeric conversion, <functional> operators on arrays, complex<T>, fast math algorithms
Architecture intrinsic	Templates wrapping architecture-specific PTX instructions (e.g. mma, cp.async, Idmatrix, cvt)



Open source: <https://github.com/NVIDIA/cutlass>

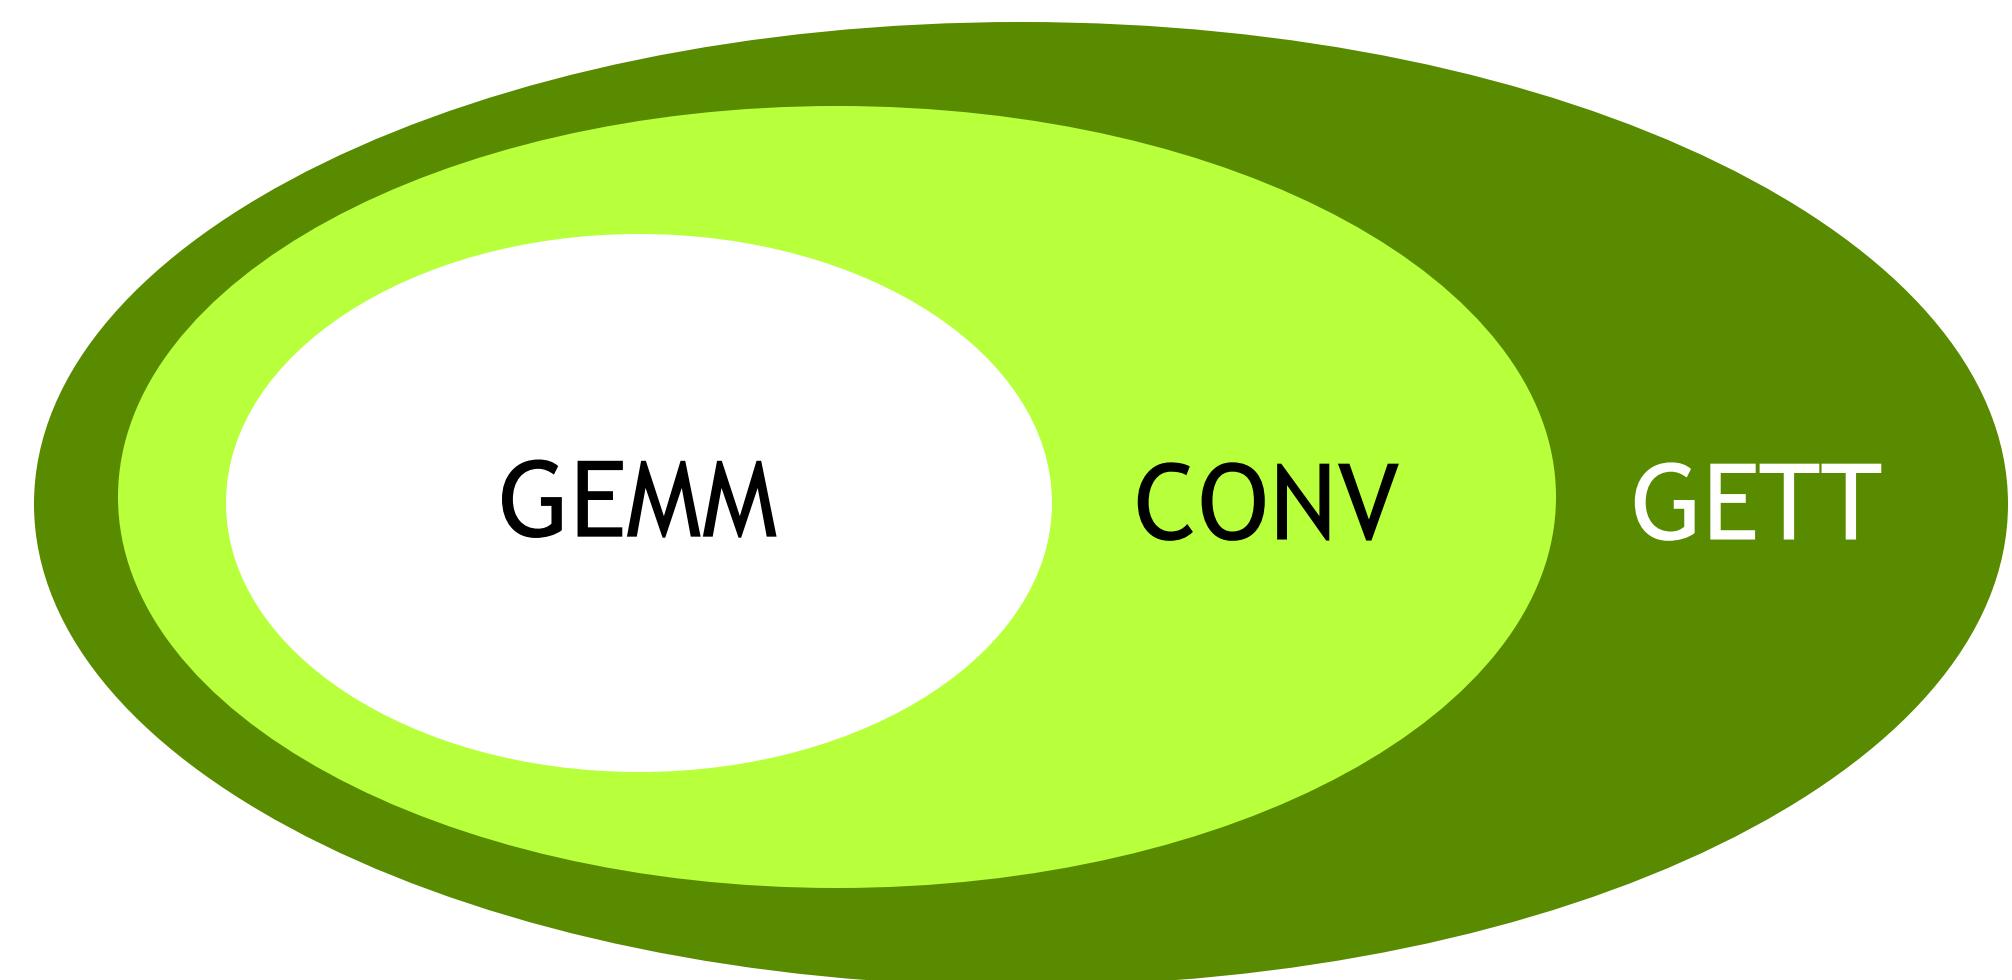
- 4.7K stars, 3M clones/month, 100+ contributors, and many active users
- Latest revision: CUTLASS 3.5
- Documentation: <https://github.com/NVIDIA/cutlass#documentation>
- Presented: [GTC'18](#), [GTC'19](#), [GTC'20](#), [GTC'21](#), [GTC'22](#), [GTC'22](#), [GTC'23](#), [GTC'24](#)
- Come join the CUTLASS channel in our discord: <https://discord.gg/CVEJqWtU>

CUTLASS 3: Design Goals

- Public Tensor Core programming model for NVIDIA GPUs
 - Serve as a production grade example for the world
- Extreme focus on developer productivity for custom kernels
 - Allow customizing any layer in the hierarchy while preserving composability with other layers
- If it compiles, it will be correct – actionable static assert messages otherwise
 - Static asserts at every layer to ensure layout and dispatch compatibilities
- Single, clear points of customization and dispatch to flatten the learning curve
 - Reduce API surface area with fewer named types

3.x Feature Highlights

Convolutions in CUTLASS 3

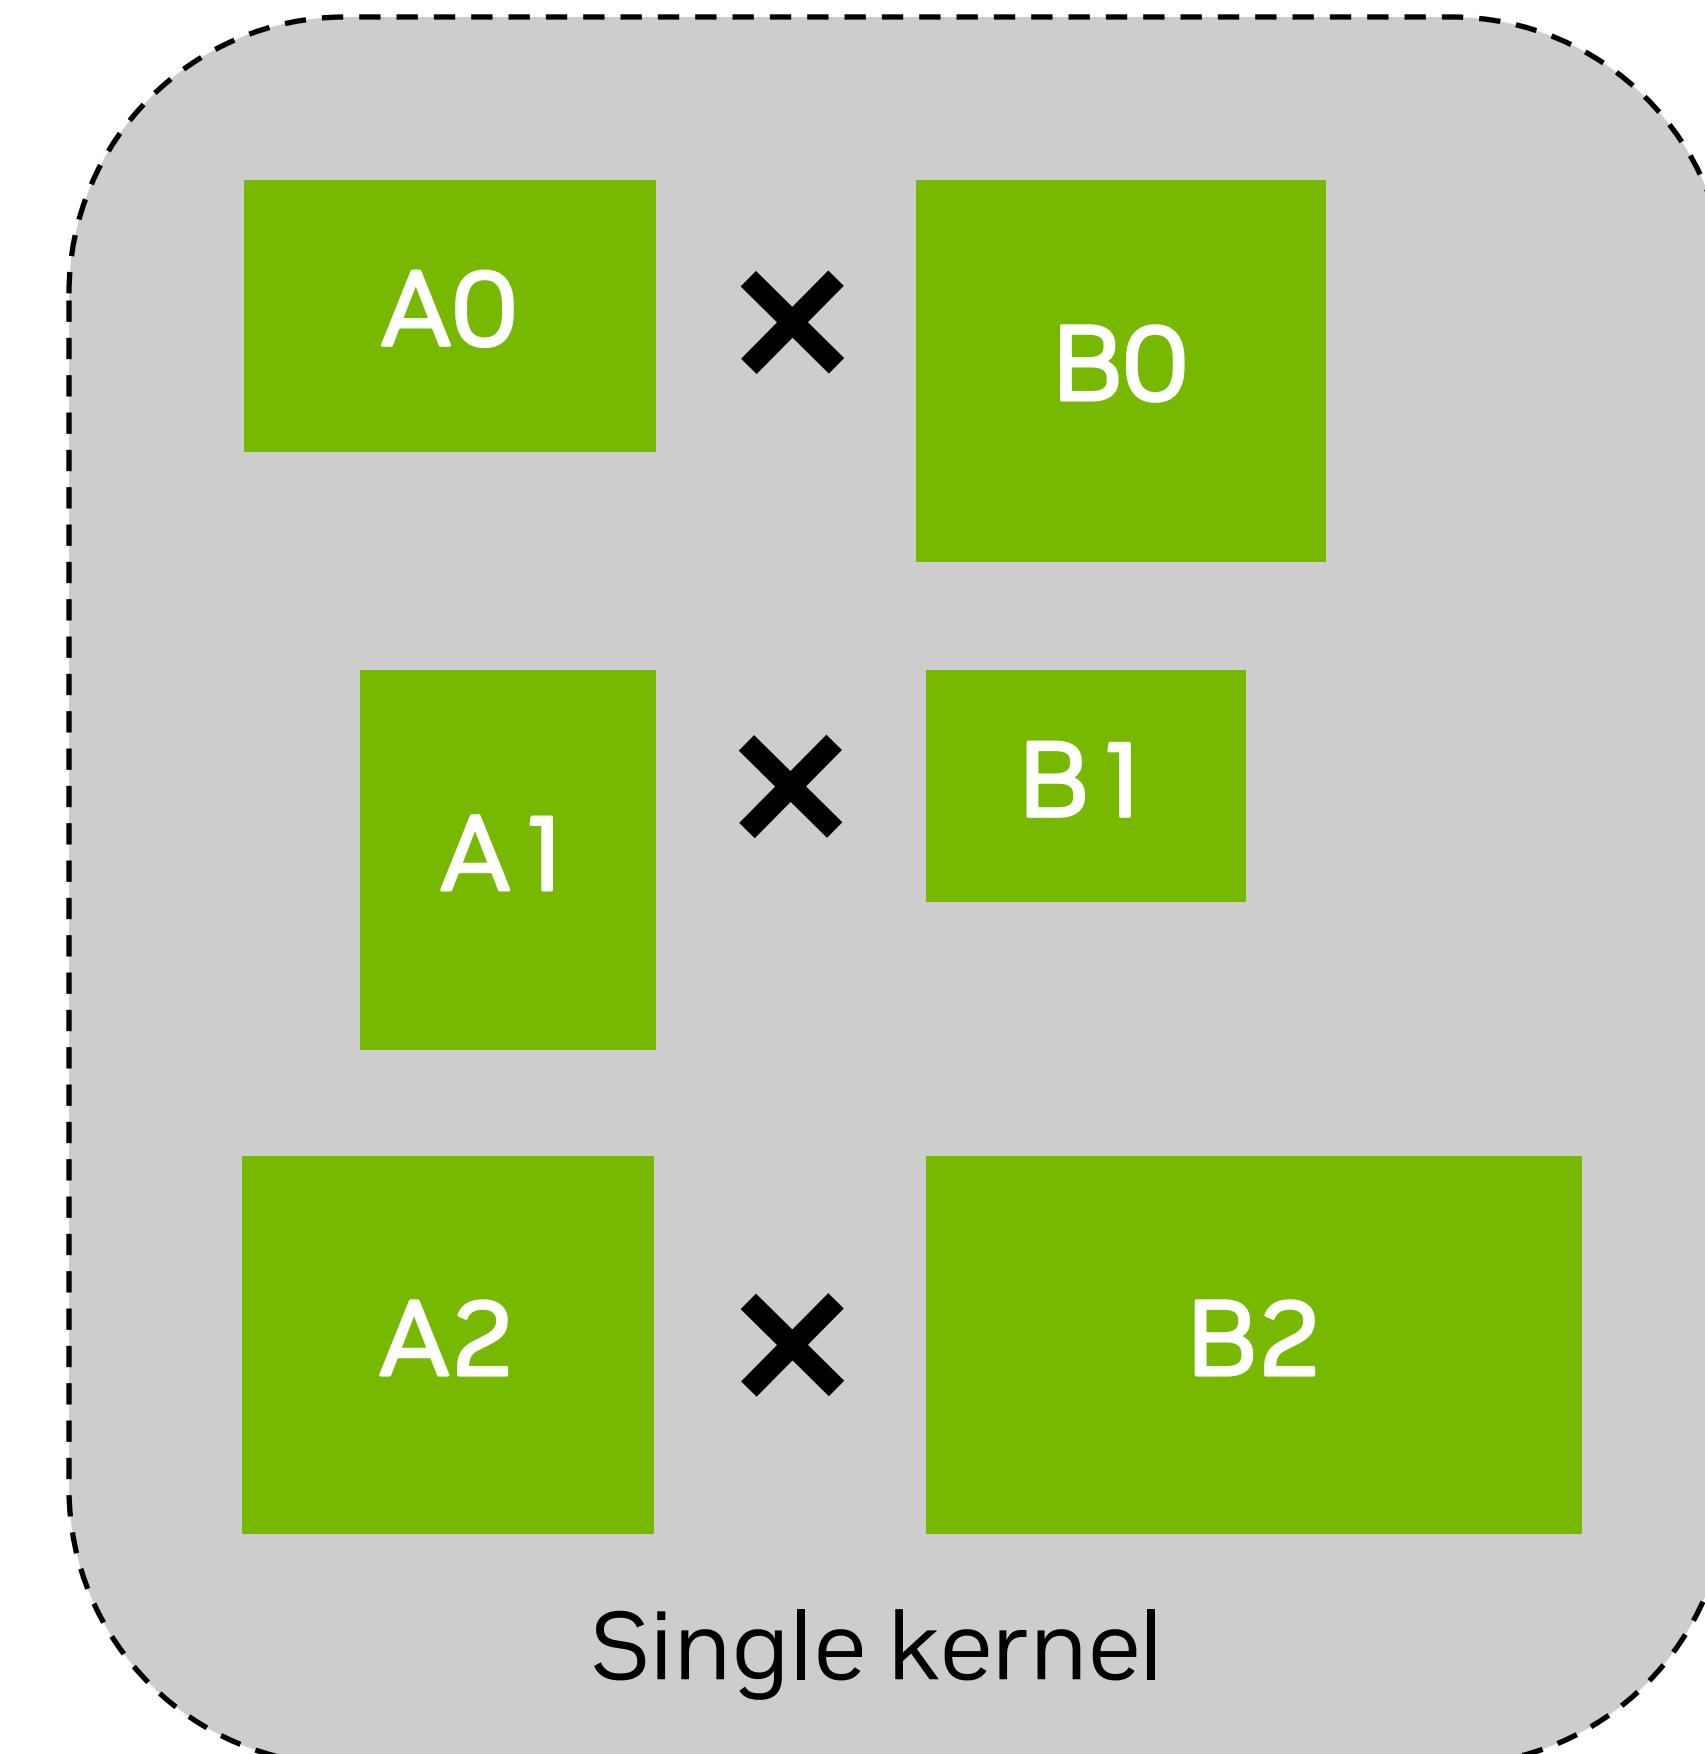


Features for LLMs

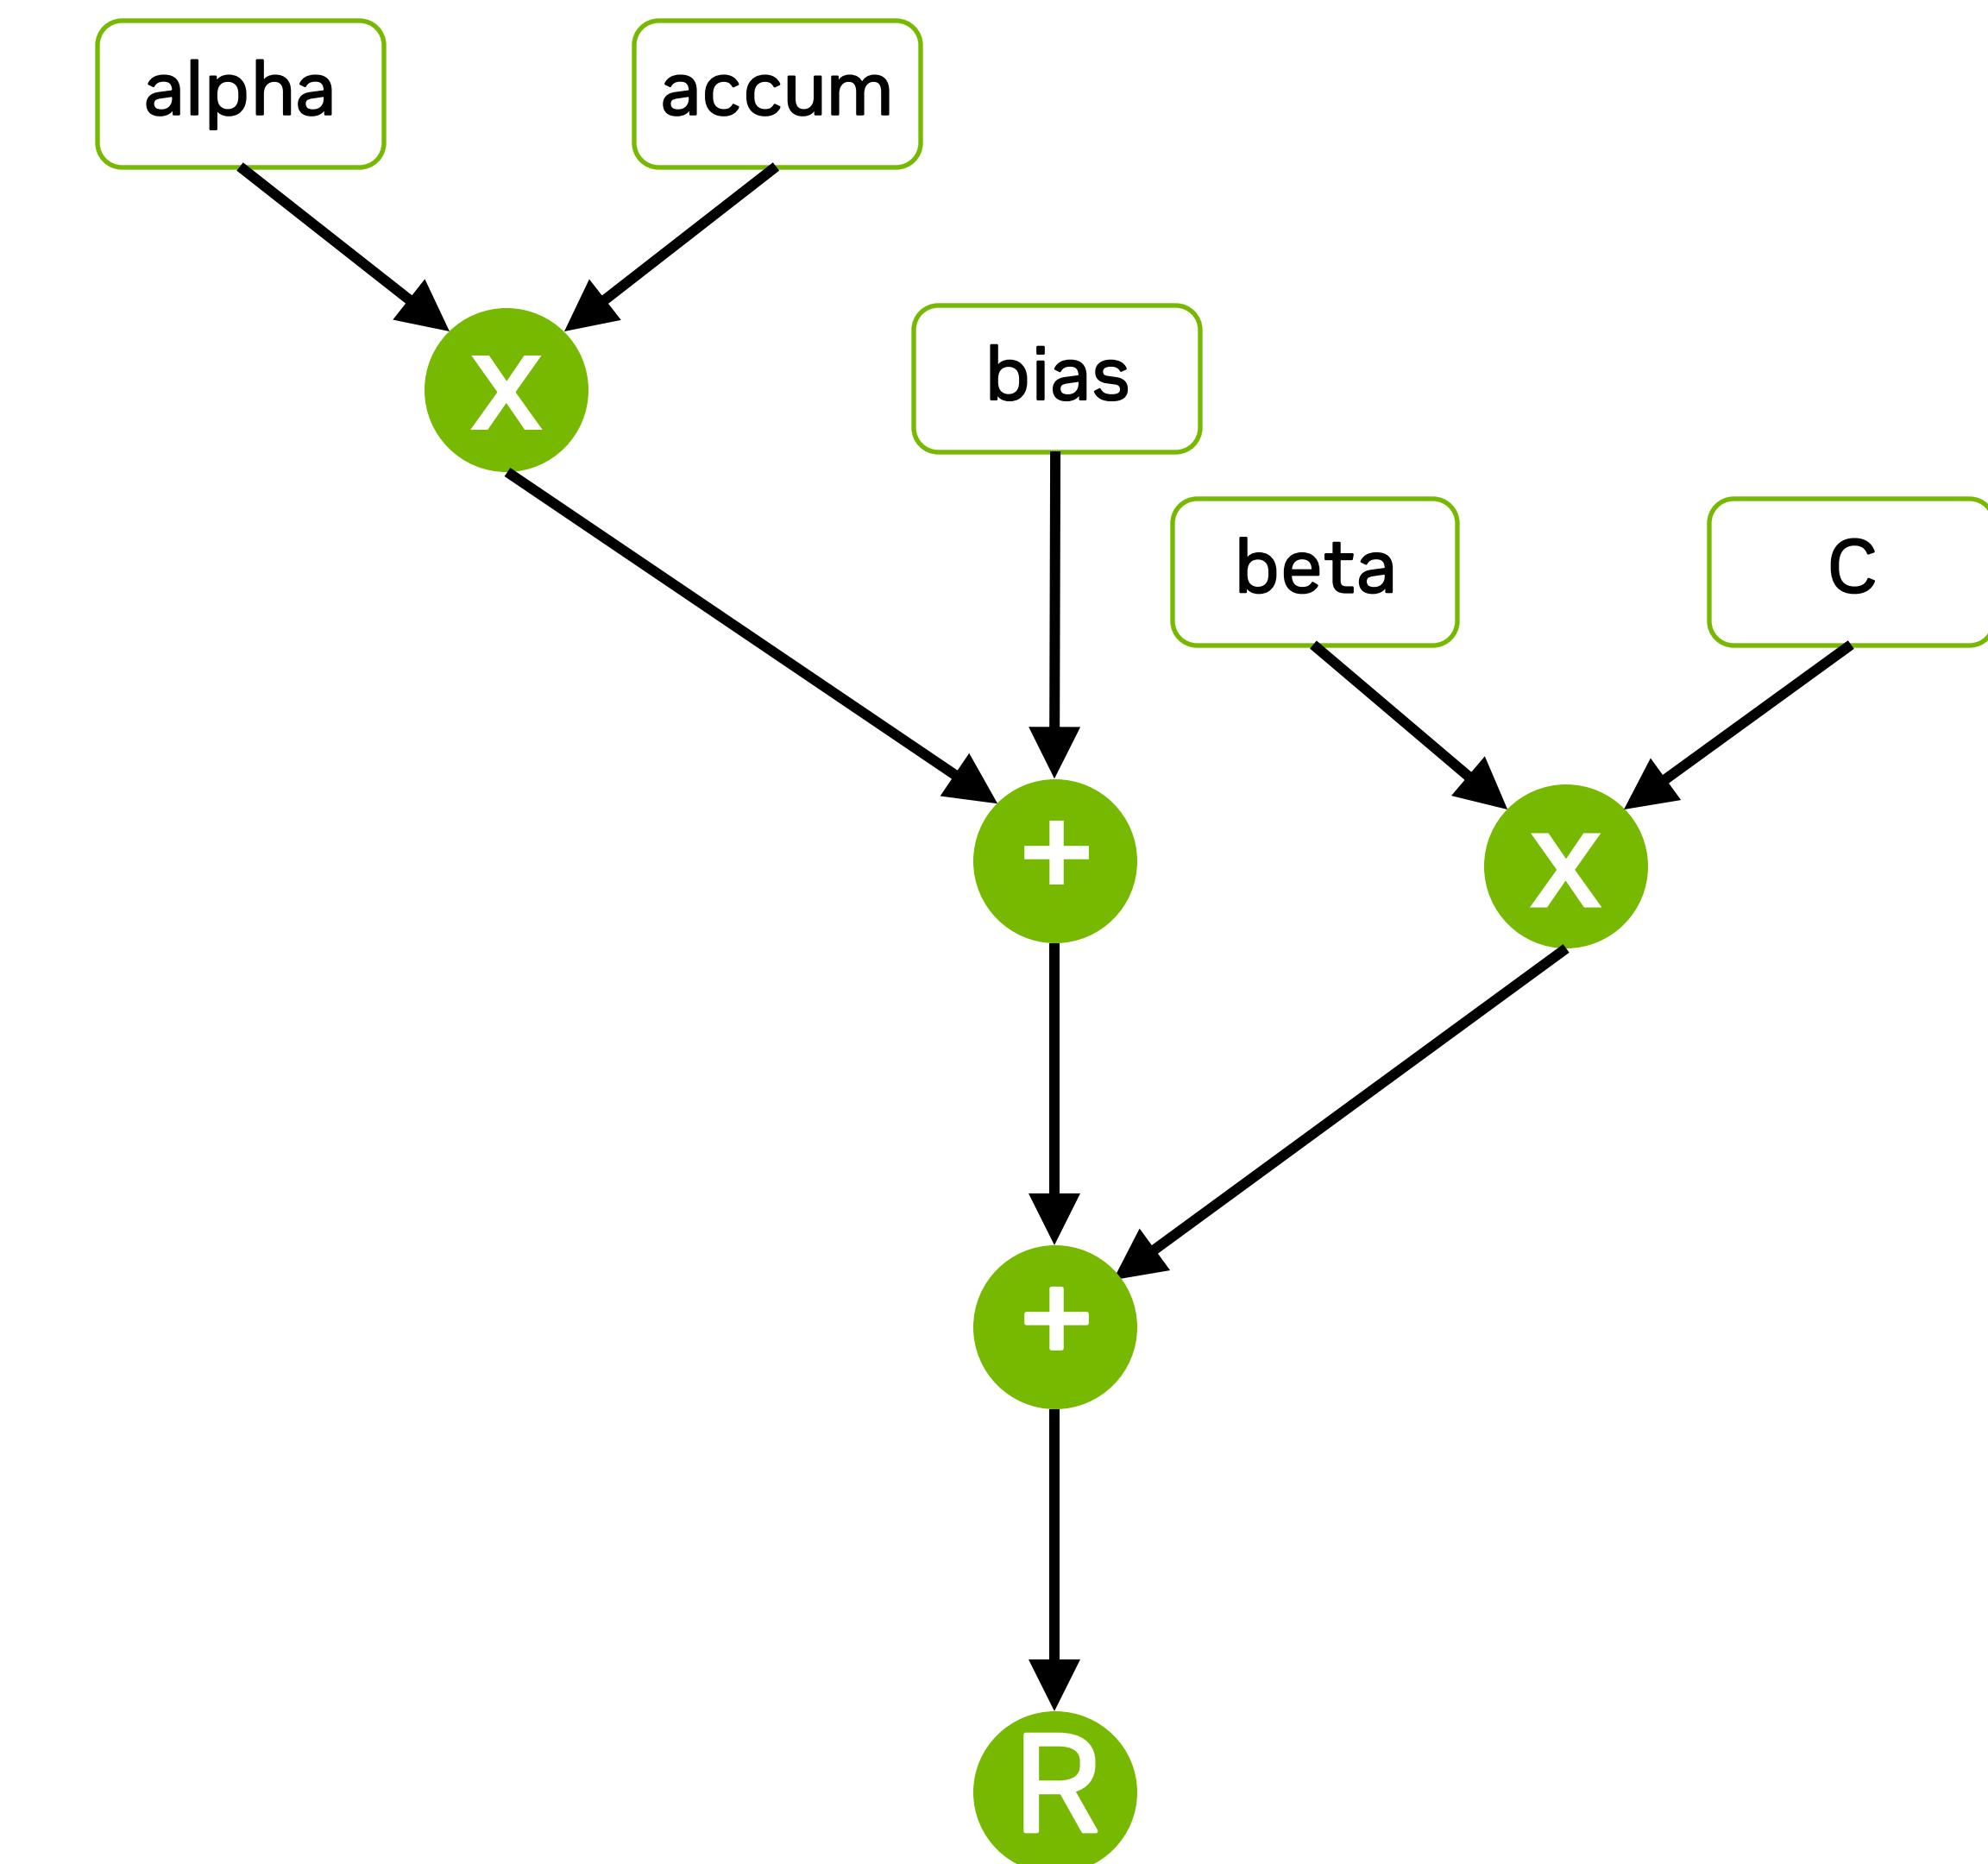
Mixed-input GEMM



Grouped GEMM

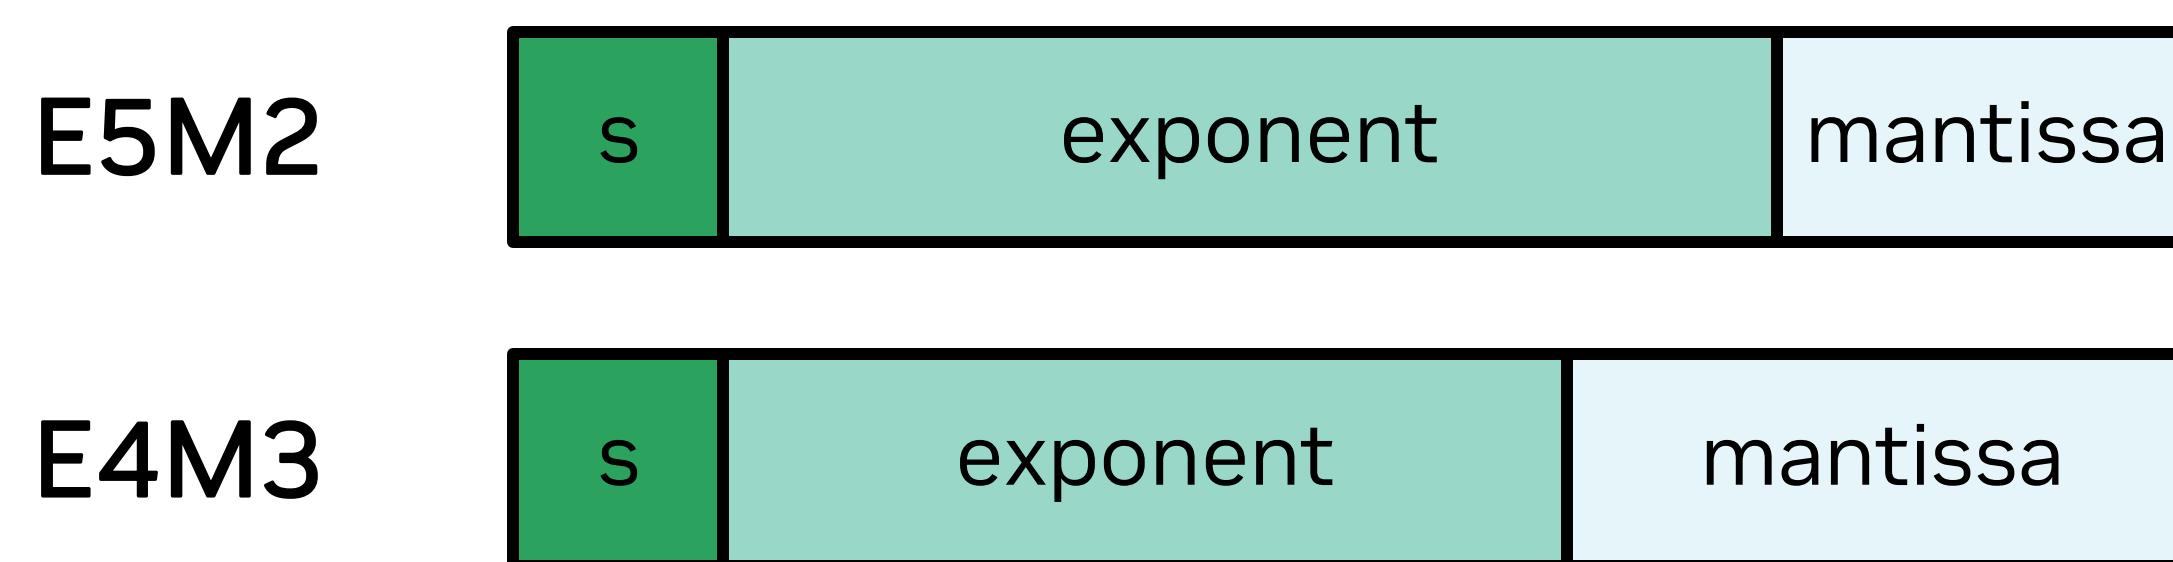


Epilogue Fusions via Visitor Tree



More CUTLASS Feature Highlights

FP8 GEMMs and convolutions for Hopper and Ada



Improved Python interface support and addition of PyPI wheel

pip install nvidia-cutlass

Tile schedulers for composable load balancing

```
using GemmKernel = cutlass::gemm::kernel::GemmUniversal<
    ProblemShape,
    CollectiveMainloop,
    CollectiveEpilogue,
    cutlass::gemm::StreamKScheduler
>;
```

More on GitHub

- Expanded compiler support
- Narrow-alignment GEMMs
- Improved documentation

CuTe Layouts and Layout Algebra

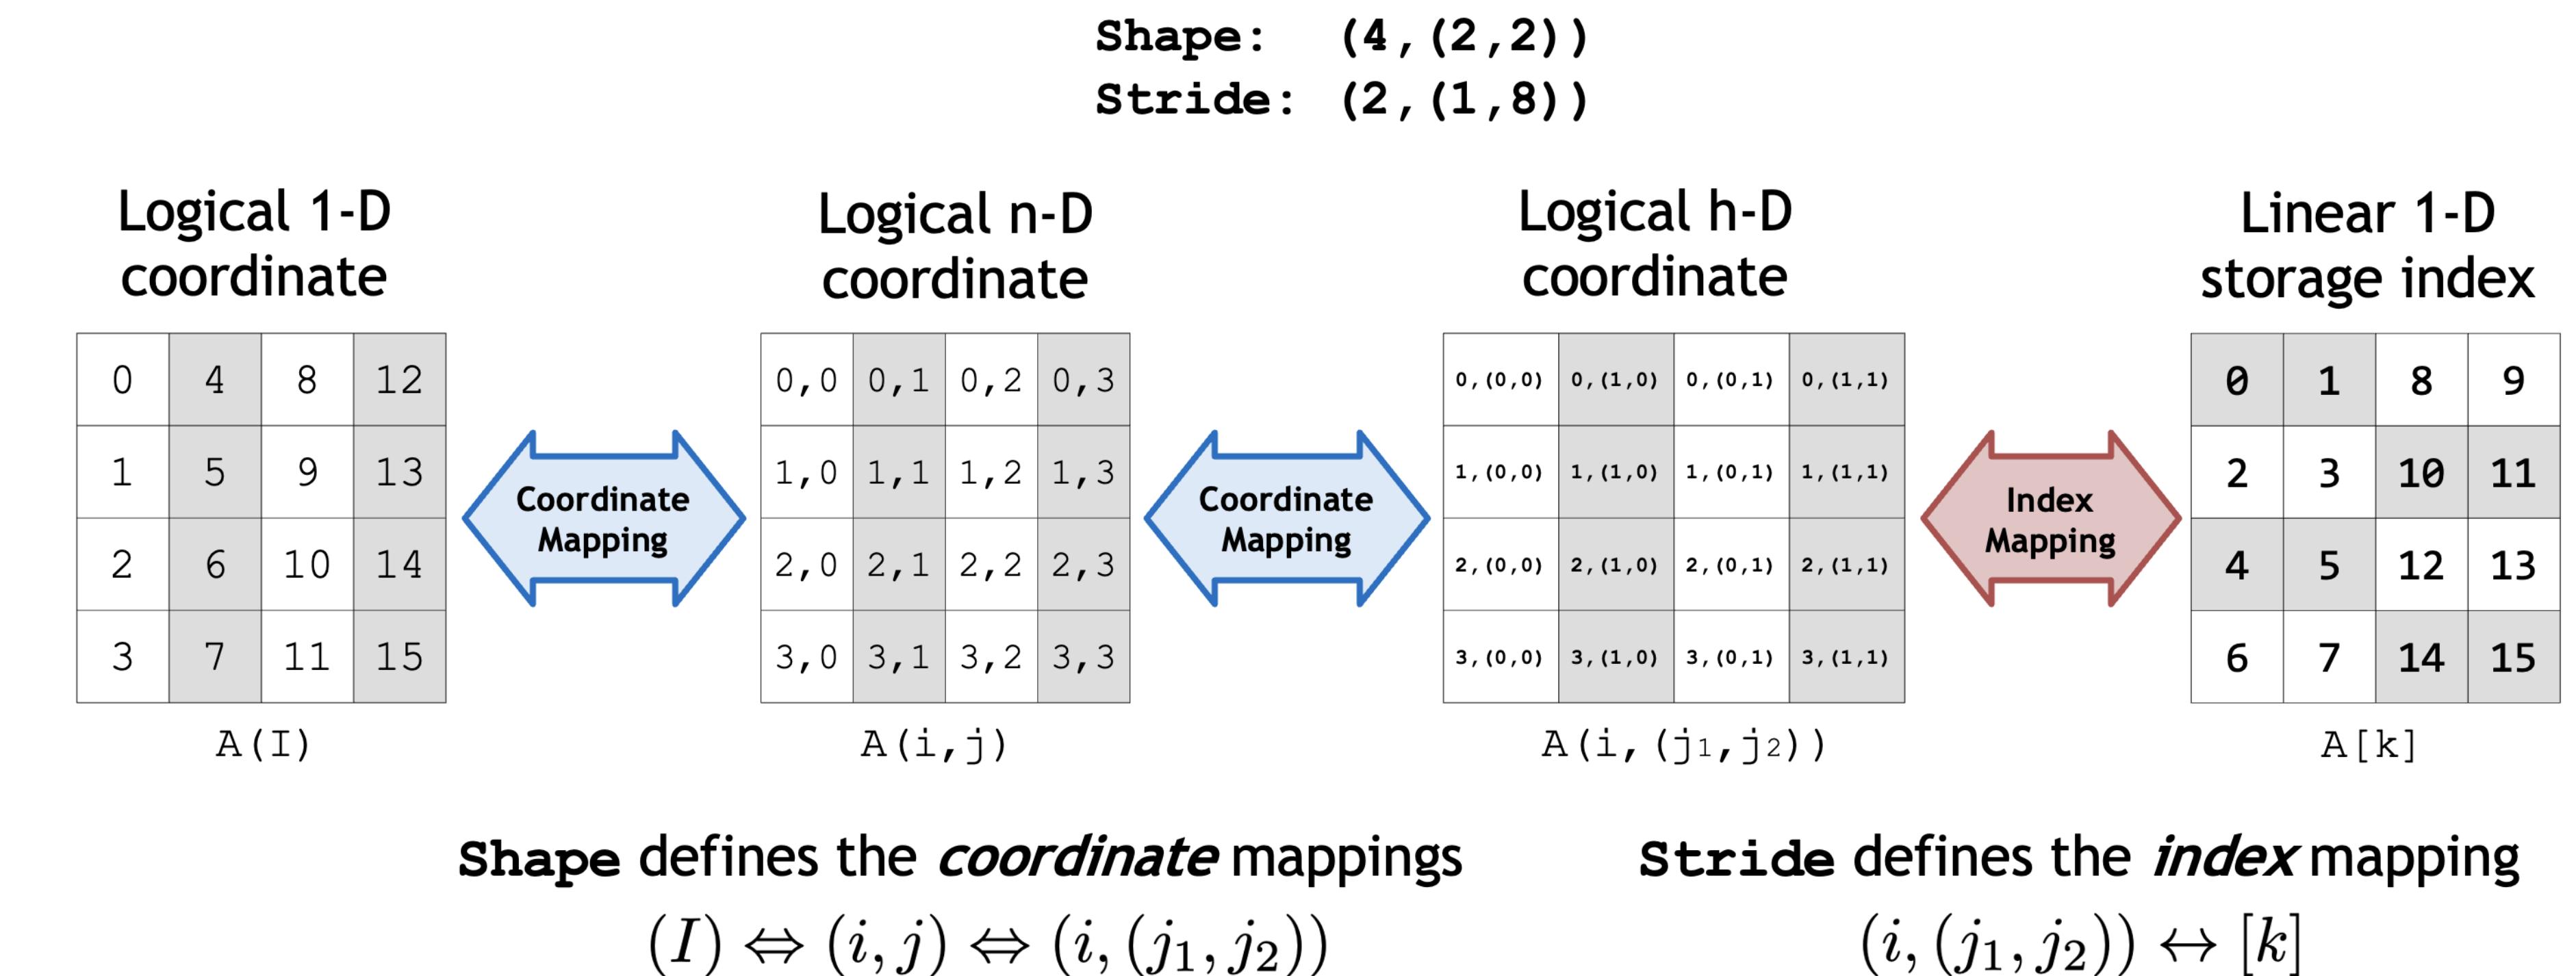
[cutlass/include/cute](#)

- **cute::Layout<Shape, Stride>**

- Compact representation of all affine layouts
- Hierarchical and multi-modal shapes & strides
- Layouts always maintain logical consistency
- Can explicitly define thread layouts
- Handles all coord->mapping
- Compose with swizzle functors

- Formalized algebra of layouts

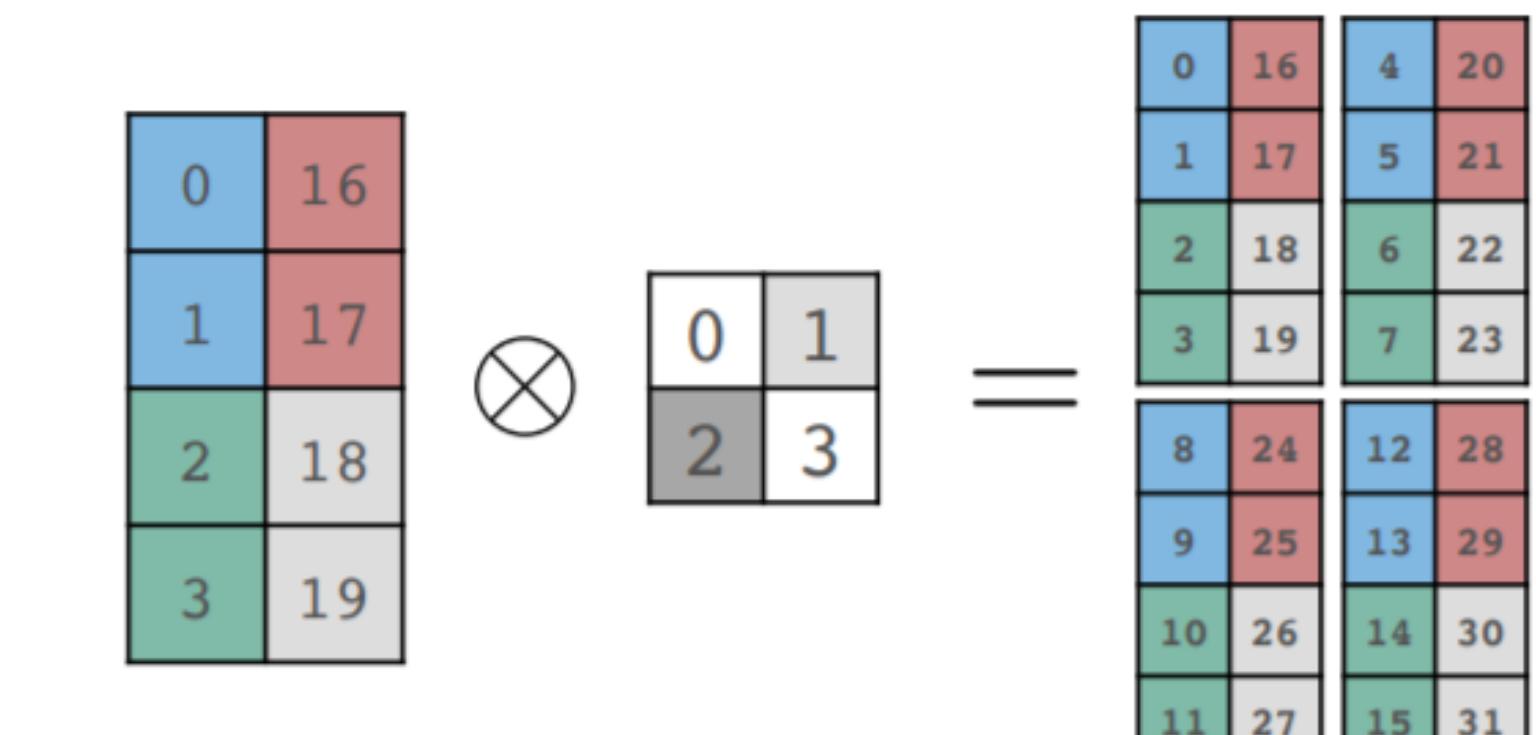
- Allows for composability at all layers
- Build complicated layouts from simple invariant layouts



Logical product:

$$f_A \otimes g_B = (f_A, f_A^* \circ g_B) \rightarrow (f_A, h_{B'})$$

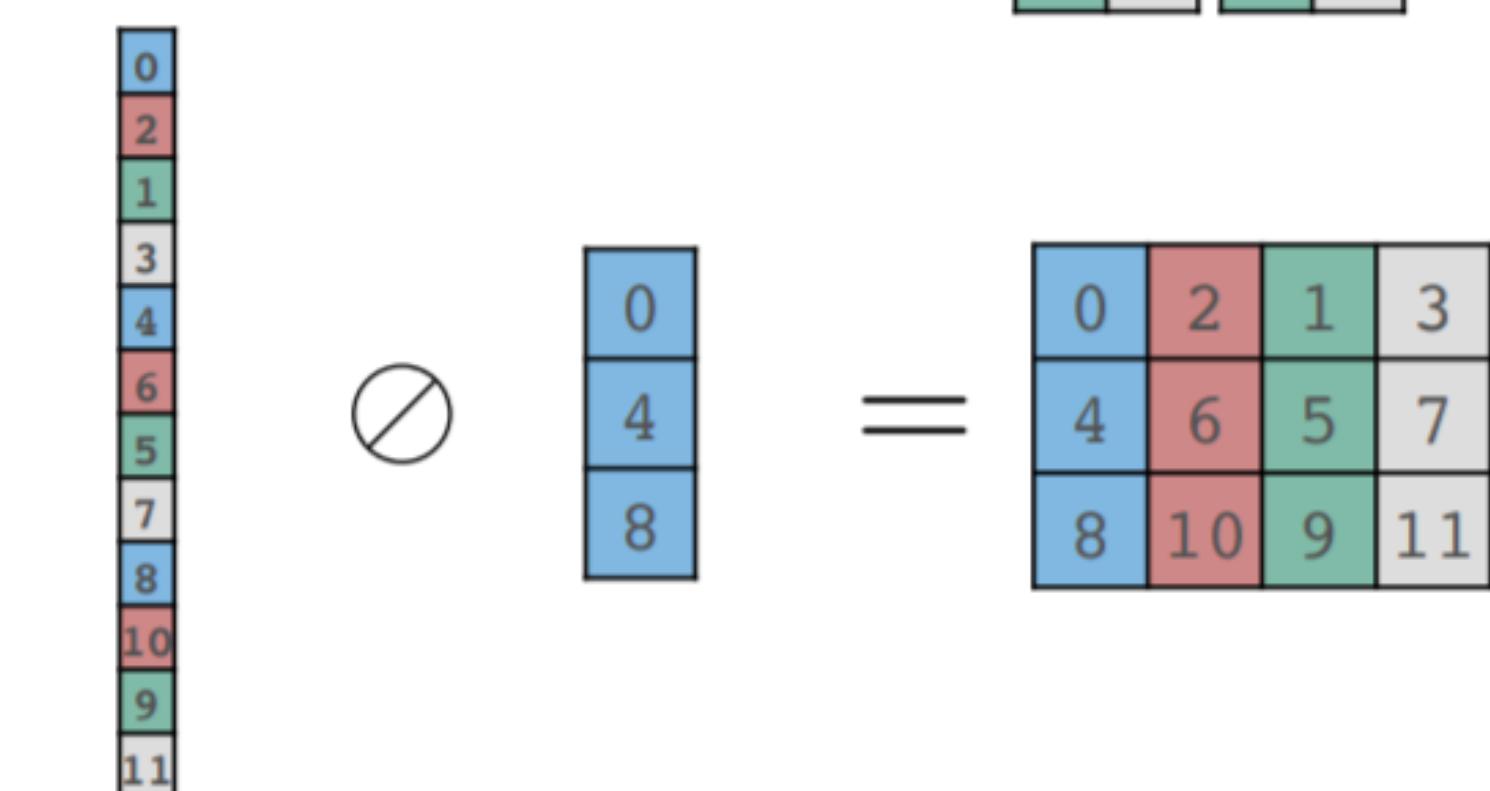
"Produce a layout where every element of layout B is a layout A."



Logical divide:

$$f_A \oslash g_B = f_A \circ (g_B, g_B^*) \rightarrow (h_{B'}, \ell_C)$$

"Produce a layout of Bs from a layout A."



Layout And Tensors

Why CuTe?

- Layout subsumes every CUTLASS-2 iterator
 - Condense complexity to a single impl. and vocab type
- Formal algebra to manipulate Layout
 - composition
 - complement
 - right_inverse, left_inverse
 - “product”
 - “divide”
- Layout for both threads and data

tridao commented 3 weeks ago

Author  ...

Amazing! Thanks a lot for the explanation @thakkarV.
I've always found smem swizzling to be the trickiest part of writing a fast gemm. It's crazy that you can replace thousands of lines of smem swizzling code with a few lines of well-designed and composable abstractions.

I'm excited to start hacking on CuTe!

  1

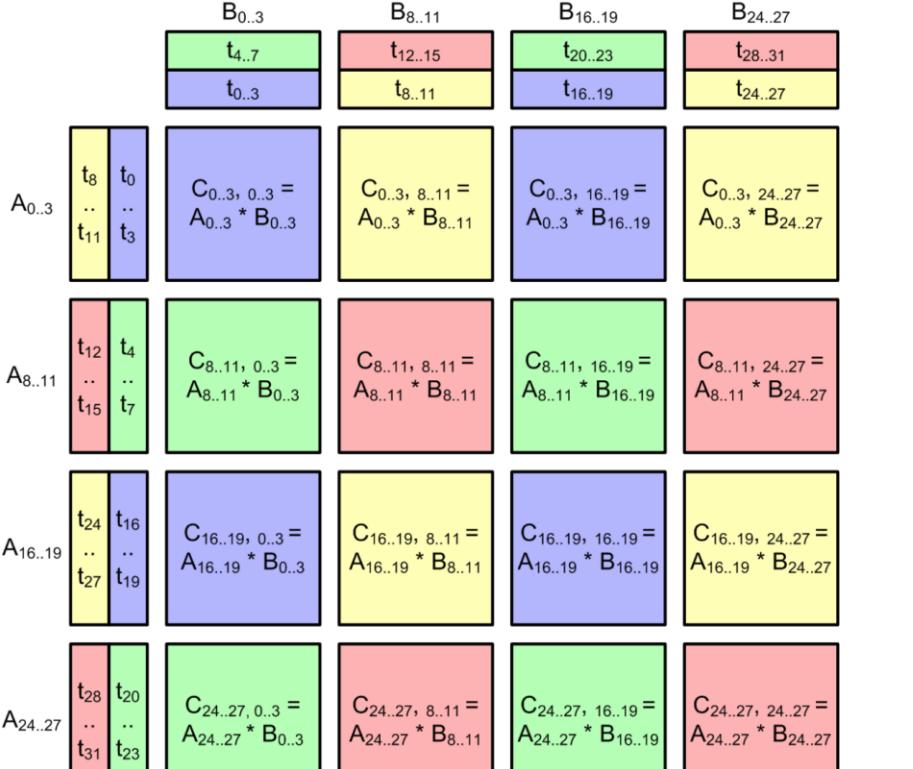


Figure 2a: Composing sparse $A[16,4] \cdot B[4, 16] = C[16, 16]$
warp-wide multiply from four $A[8, 4] \cdot B[4, 8] = C[8, 8]$ 8-thread multiples.

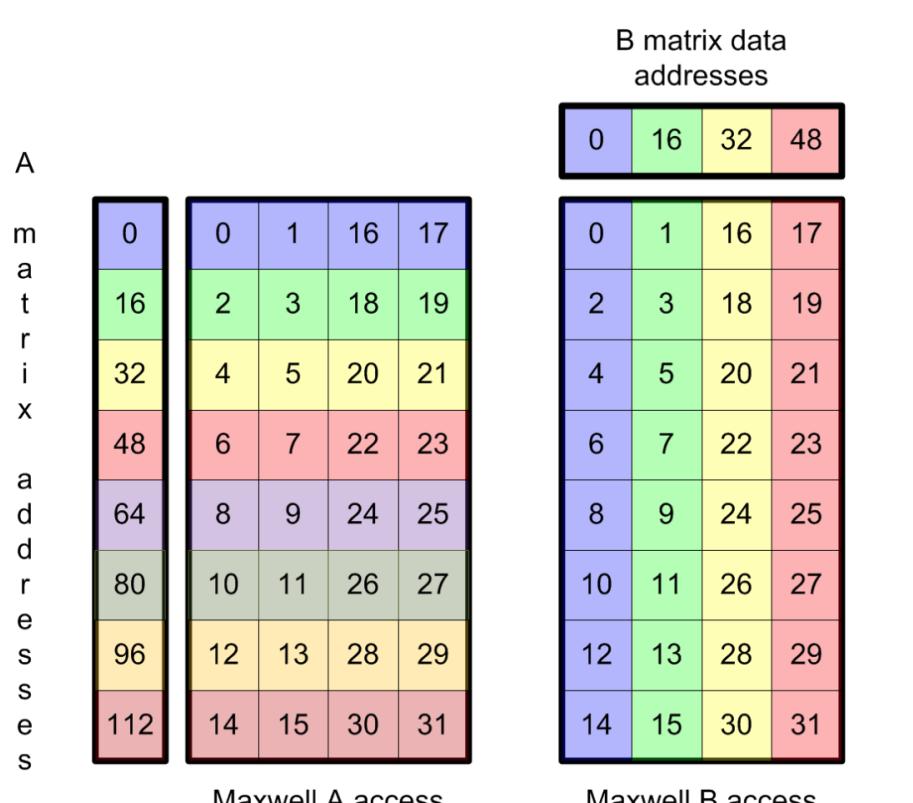


Figure 3b: Maxwell thread assignments for LDS.U matching

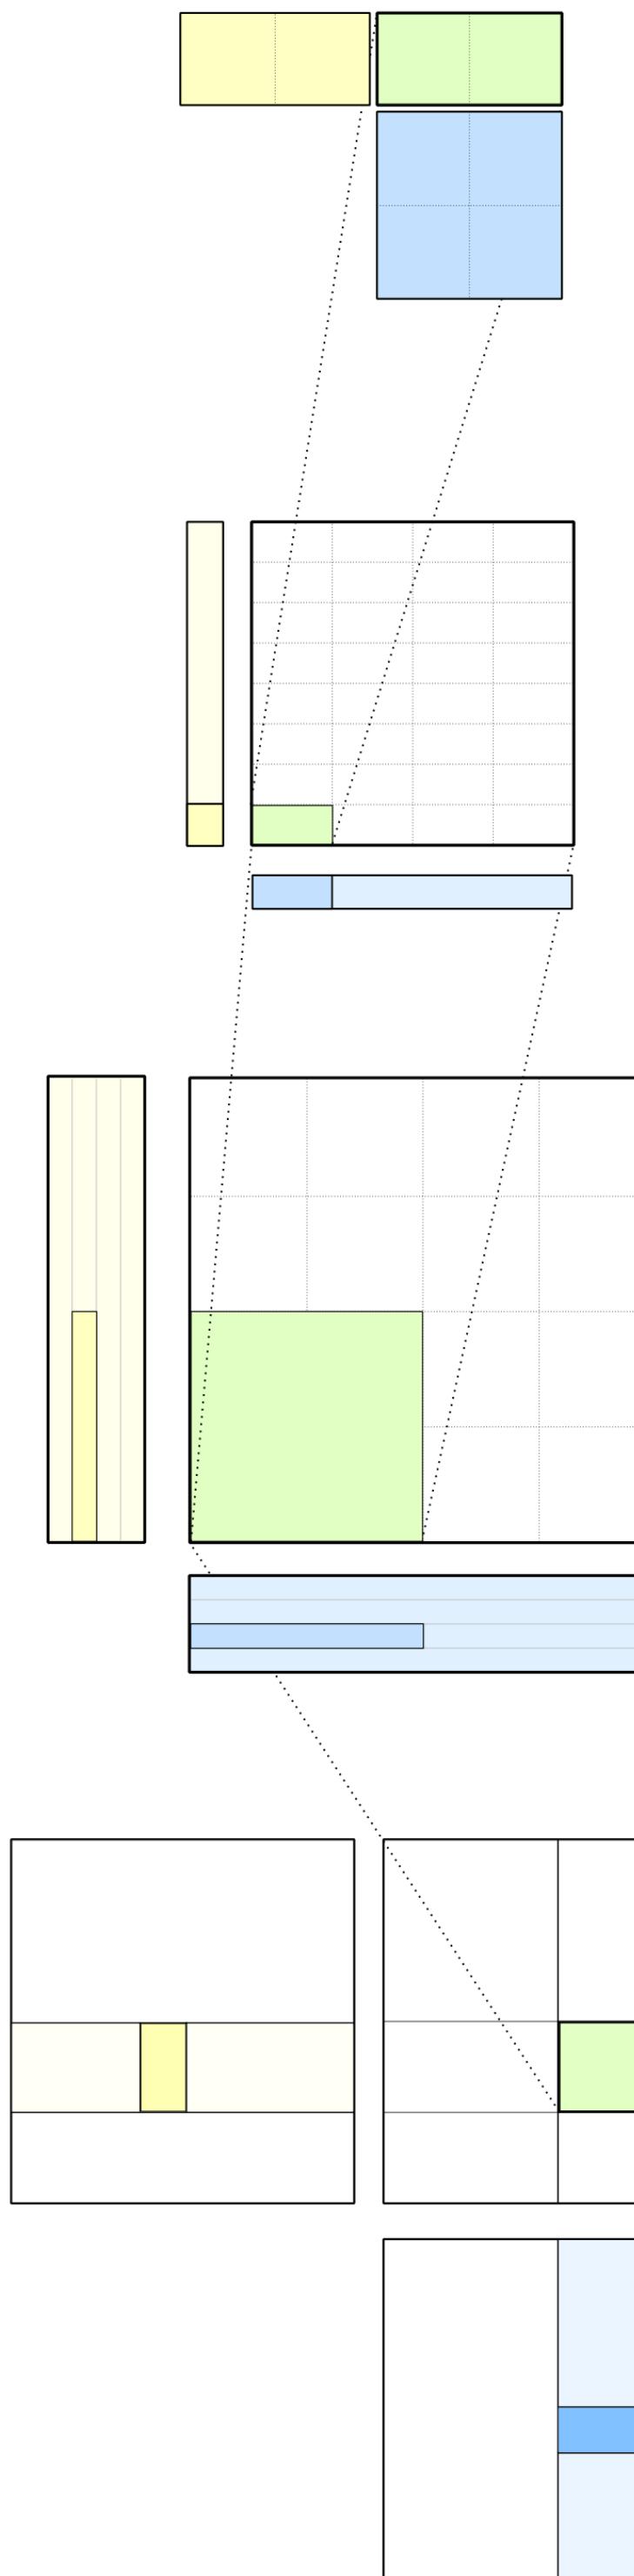
SMEM address view																xor factor																									
contiguous dim, K==>								1 0								3 2								5 4								7 6									
0	8	16	24	32	40	48	56	0	1	2	3	4	5	6	7	24	16	3	2	12	13	5	4	40	32	7	6	56	48	28	20	14	15	10	9						
0	8	9	10	11	12	13	14	15	1	2	3	4	5	6	7	19	18	10	11	12	13	14	15	40	32	7	6	56	48	28	20	14	15	10	9						
1	8	9	10	11	12	13	14	15	2	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47
2	16	17	18	19	20	21	22	23	3	24	25	26	27	28	29	30	31	32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47	48	49	50	51	52	53	54	55
3	24	25	26	27	28	29	30	31	4	32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47	48	49	50	51	52	53	54	55	56	57	58	59	60	61	62	63
4	32	33	34	35	36	37	38	39	5	40	41	42	43	44	45	46	47	48	49	50	51	52	53	54	55	56	57	58	59	60	61	62	63	64	65	66	67	68	69	70	71
5	40	41	42	43	44	45	46	47	6	48	49	50	51	52	53	54	55	64	65	66	67	68	69	70	71	72	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87
6	48	49	50	51	52	53	54	55	7	56	57	58	59	60	61	62	63	64	65	66	67	68	69	70	71	72	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87
7	56	57	58	59	60	61	62	63	8	64	65	66	67	68	69	70	71	72	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87	88	89	90	91	92	93	94	95
8	64	65	66	67	68	69	70	71	9	72	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87	88	89	90	91	92	93	94	95	96	97	98	99	100	101	102	103
9	72	73	74	75	76	77	78	79	10	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87	88	89	90	91	92	93	94	95	96	97	98	99	100	101	102	103	
10	80	81	82	83	84	85	86	87	11	88	89	90	91	92	93	94	95	96	97	98	99	100	101	102	103	104	105	106	107	108	109	110	111	112	113	114	115	116	117	118	119
11	88	89	90	91	92	93	94	95	12	96	97	98	99	100	101	102	103	104	105	106	107	108	109	110	111	112	113	114	115	116	117	118	119	120	121	122	123	124	125	126	127

SMEM address view																xor factor																									
contiguous dim, K==>								1 0								3 2								5 4								7 6									
0	8	16	24	32	40	48	56	0	1	2	3	4	5	6	7	24	16	3	2	12	13	5	4	40	32	7	6	56	48	28	20	14	15	10	9						
0	8	9	10	11	12	13	14	15	1	2	3	4	5	6	7	19	18	10	11	12	13	14	15	40	32	7	6	56	48	28	20	14	15	10	9						
1	8	9	10	11	12	13	14	15	2	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47
2	16	17	18	19	20	21	22	23	3	24	25	26	27	28																											

CUTLASS 3 Conceptual Hierarchy

CUTLASS 3.x computation hierarchy is not centered around the hardware hierarchy

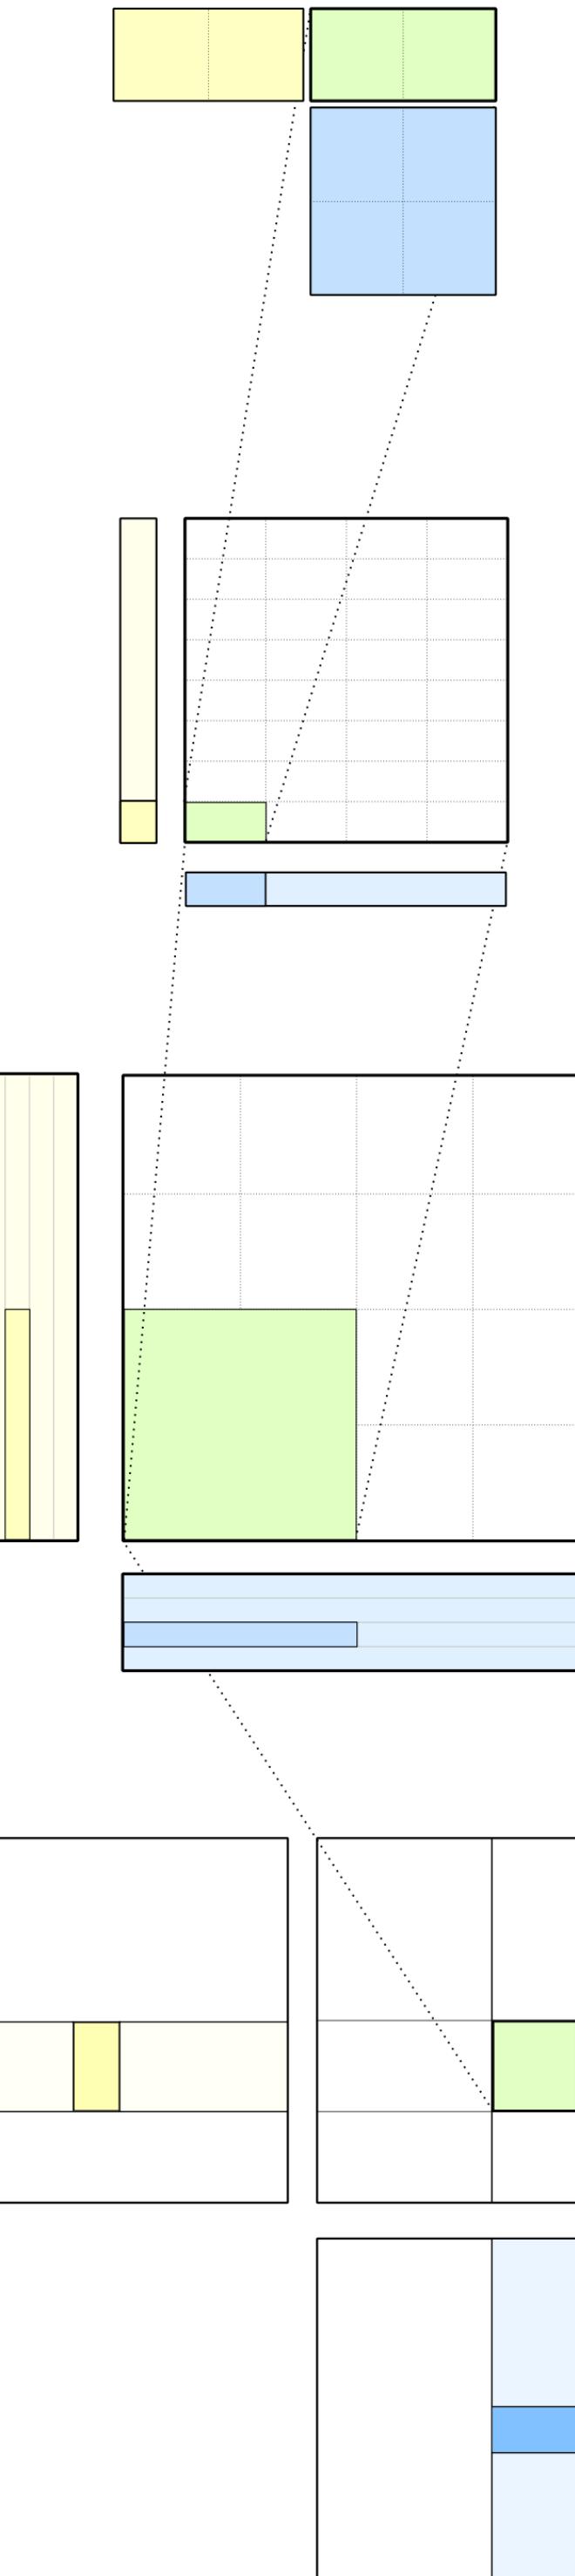
- **Atom layer:** Architecture instructions and associated meta-information
 - Smallest set of threads and values that must participate in an architecture accelerated specified math/copy op
- **Tiled MMA/Copy:** Spatial Microkernel layer
 - Describes the complete spatial tiling of a math/copy operation
 - Write canonical loops with arch specific instructions
- **Collective layer:** Temporal Microkernel layer
 - Describes the complete temporal tiling of spatial microkernels and computing one output tile
 - Abstract complex arch-specific synchronization, warp-specialization, pipelining, and instruction interleaving
- **Kernel layer:** Outermost loops around collectives
 - Conceptually: A collection of all threadblock/clusters in the grid
 - Responsible for load balancing across tiles, thread marshalling, grid planning, and arguments construction
- **Device layer:** host side setup and interface



CUTLASS 3 API

Single points of entry at each abstraction layer

- Spatial microkernels: `cute::Tiled{Mma | Copy}<>`
 - Robust representation power across a wide range GPU architectures
- Temporal Microkernels: `collective::Collective{Mma | Conv | Epilogue | Transform}<>`
 - Dispatched against by policies that also define the set of kernel schedules they can be composed with
- Kernel layer: `kernel::{Gemm | Conv}Universal<>`
 - Treats GEMM as a composition of a collective mainloop and a collective epilogue
 - **Tile Schedulers** are a first class at this level – decide which tile coords map to which program ID
- Device layer: `device::{Gemm | Conv}UniversalAdapter<>`
 - Just a handle object to the kernel
- `cutlass::Pipeline`: used for abstracting synchronization across or within the layers
- Static asserts everywhere to guard against invalid compositions or incorrect layouts





Off the shelf CUTLASS GEMM

CUTLASS 3.0 GETT

CUTLASS 3.0 Hopper GEMMs are GETTs in Disguise

```
// Build the mainloop type
using CollectiveMainloop = typename cutlass::gemm::collective::CollectiveBuilder<
    arch::Sm90, arch::OpClassTensorOp,
    ElementA, StrideA, AlignmentA,
    ElementB, StrideB, AlignmentB,
    ElementAccumulator,
    TilesShape, ClusterShape,
    cutlass::gemm::collective::StageCountAuto,
    cutlass::gemm::collective::KernelScheduleAuto
>::CollectiveOp;

// Build the epilogue type
using CollectiveEpilogue = typename cutlass::epilogue::collective::CollectiveBuilder<
    arch::Sm90, arch::OpClassTensorOp,
    TileShape, ClusterShape,
    cutlass::epilogue::collective::EpilogueTileAuto,
    ElementAccumulator, ElementCompute,
    ElementC, StrideC, AlignmentC,
    ElementD, StrideD, AlignmentD,
    cutlass::epilogue::collective::EpilogueScheduleAuto
>::CollectiveOp;

// Compose both at the kernel layer, and define the type of our problem shape tuple
using GettKernel = cutlass::gemm::kernel::GemmUniversal<
    ProblemShape_MNKL, // still a rank-4 tuple, but now is hierarchical
    CollectiveMainloop,
    CollectiveEpilogue>;

// Device layer handle to the kernel
using Gett = cutlass::gemm::device::GemmUniversalAdapter<GettKernel>;
```

CUTLASS 3: Collective API

`cutlass::collective::{Gemm|Conv|Epilogue}`

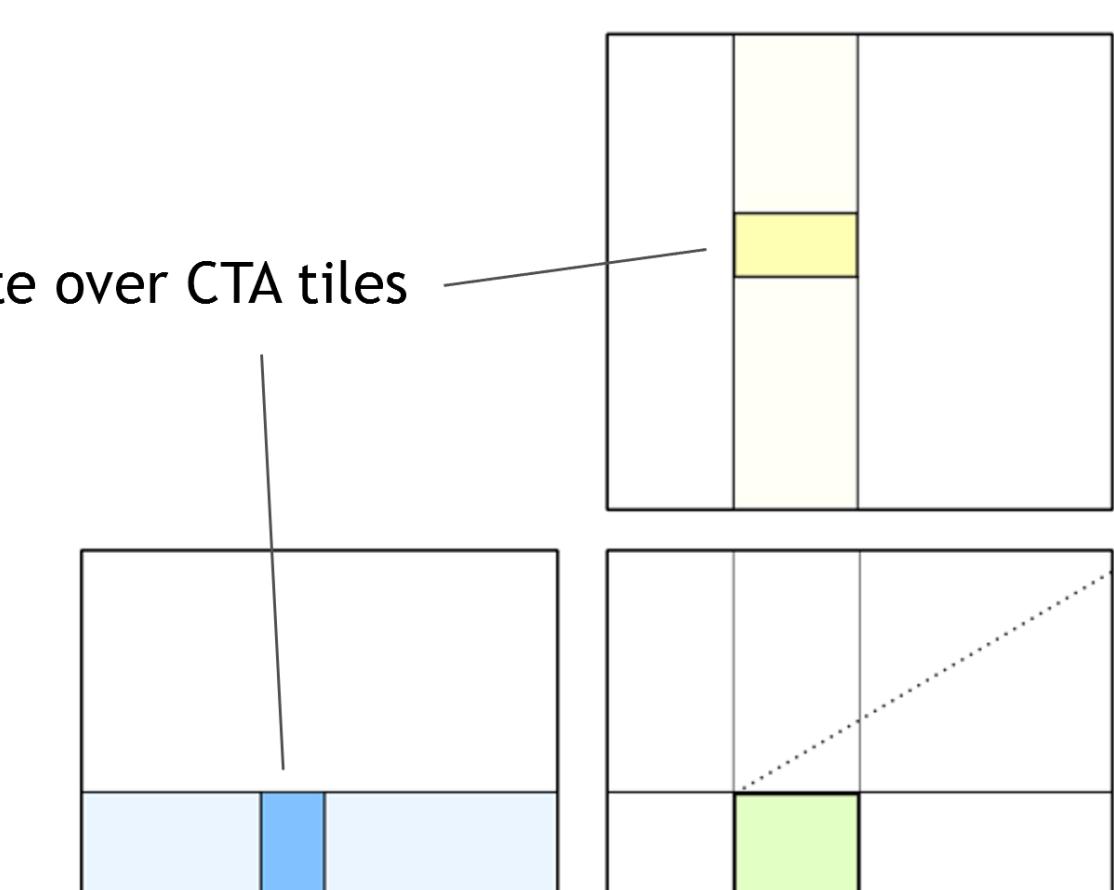
- Implements a temporal microkernel
 - Computes one output tile
 - Can be warp-specialized
 - Can have async spatial microkernels (Producer-Consumer pipelines)
- Dispatch policies used for selecting between specializations
 - One dispatch policy per mainloop variant
 - Can possibly compose with more than one kernel schedule! (next slide)

```
// Dispatch policy examples
// 2 stage pipeline through 1 stage in smem, 1 in rmem, with predicated gmem loads
struct MainloopSm70TwoStage;

// n-buffer in smem (cp.async), pipelined with registers, WITHOUT predicated gmem loads
struct MainloopSm80CpAsyncUnpredicated;

// n-buffer in smem (TMA), pipelined with Hopper GMMA and TMA, Warp specialized dynamic
// schedule
Struct MainloopSm90TmaGmmaWarpSpecialized;
```

```
template <
    class DispatchPolicy,
    class TileShape,
    class ElementA,
    class SmemLayoutA,
    class ElementB,
    class SmemLayoutB,
    class ElementC,
    class ArchTag,
    class TiledMma,
    class GmemCopyAtomA,
    class SmemCopyAtomA,
    class GmemCopyAtomB,
    class SmemCopyAtomB
>
struct CollectiveMma;
```



Multiply tiles together
Add them to output tile

CUTLASS 3: Kernel API

`cutlass::gemm::kernel::GemmUniversal<>`

- Each GEMM kernel is a composition of three parts:

- Collective mainloop (gemm or conv)
- Collective epilogue
- Tile scheduler

- Implements:

- Grid planning logic and kernel global entry point
- Thread marshalling for warp-specialized collectives
- Block/Cluster wide swizzling
- Load balancing schedules (split-K, stream-K)

- Selection via mainloop dispatch policy's `Schedule` tag

- SM90 cp.async + GMMA multistage
- SM90 TMA + GMMA multistage
- SM90 TMA + GMMA + Warp-Specialized non-persistent
- SM90 TMA + GMMA + Warp-Specialized + Pingpong Persistent
- SM90 TMA + GMMA + Warp-Specialized + Cooperative Persistent + StreamK

```
template <
    class ProblemShapeOrThreadblockMma_,
    class CollectiveMainloopOrEpilogue_,
    class CollectiveEpilogueOrThreadblockSwizzle_,
    class GridSwizzle_ = void,
    class Enable = void
>
class GemmUniversal;
```

Kernel layer back-compatible API entry point

Sensible defaults with incremental opt-ins

`cutlass::gemm::collective::CollectiveBuilder<>`

“I just want a Hopper collective”:

```
using CollectiveOp = typename collective::CollectiveBuilder<  
    arch::Sm90, arch::OpClassTensorOp,  
    half_t, LayoutA, 8,  
    half_t, LayoutB, 8,  
    float,  
    Shape<_128,_128,_64>, Shape<_1,_2,_1>,  
    gemm::collective::StageCountAuto,  
    gemm::collective::KernelScheduleAuto  
>::CollectiveOp;
```

```
template <  
    class ArchTag,  
    class OpClass,  
    class ElementA,  
    class GmemLayoutA,  
    int AlignmentA,  
    class ElementB,  
    class GmemLayoutB,  
    int AlignmentB,  
    class ElementAccumulator,  
    class TileShape_MNK,  
    class ClusterShape_MNK,  
    class StageCountType,  
    class KernelScheduleType,  
    class Enable = void  
>  
struct CollectiveBuilder;
```

“I want a Hopper collective but composed with persistent kernel and 5 stages”:

```
using CollectiveOp = typename collective::CollectiveBuilder<  
    arch::Sm90, arch::OpClassTensorOp,  
    half_t, LayoutA, 8,  
    half_t, LayoutB, 8,  
    float,  
    Shape< 128, 128, 64>, Shape< 1, 2, 1>,  
    gemm::collective::StageCount<5>,  
    gemm::KernelTmaWarpSpecializedPersistent  
>::CollectiveOp;
```

Builders do the heavy lifting ...

```
using CollectiveOp = typename collective::CollectiveBuilder<
    arch::Sm90, arch::OpClassTensorOp,
    half_t, LayoutA, 8,
    half_t, LayoutB, 8,
    float,
    Shape<_128,_128,_64>, Shape<_2,_1,_1>,
    gemm::collective::StageCountAuto,
    gemm::collective::KernelScheduleAuto
>::CollectiveOp;
```

Automatically maps down to the best mainloop config:
stage count, GMMA instruction, smem layouts, TMA mcast, kernel schedule, mainloop type ...

```
cutlass::gemm::collective::CollectiveMma<
    cutlass::gemm::MainloopSm90TmaGmmaWarpSpecialized<
        7, Shape<_2,_1,_1>, cutlass::gemm::KernelTmaWarpSpecializedCooperative>,
        Shape<_128, _128, _64>,
        half_t, Stride<int64_t, _1, int64_t>,
        half_t, Stride<int64_t, _1, int64_t>,
        TiledMMA<
            MMAAtom<SM90_64x128x16_F32F16F16_SS<GMMA::Major::K, GMMA::Major::K>,
            Layout<Shape<_2, _1, _1>>,
            SM90_TMA_LOAD,
            ComposedLayout<Swizzle<3, 4, 3>, smem_ptr_flag_bits<16>, Layout<tuple<_8, _64>, tuple<_64, _1>>>,
            void, identity,
            SM90_TMA_LOAD_MULTICAST,
            ComposedLayout<Swizzle<3, 4, 3>, smem_ptr_flag_bits<16>, Layout<tuple<_8, _64>, tuple<_64, _1>>>,
            void, identity>
```

Convolutions provide a familiar builder API

`cutlass::conv::collective::CollectiveBuilder<>`

“I just want a Hopper Fprop collective for NWC layout tensors (1-D Fprop)”:

```
using CollectiveOp = typename conv::collective::CollectiveBuilder<  
    arch::Sm90, arch::OpClassTensorOp,  
    conv::Operator::kFprop,  
    half_t, layout::TensorNWC, 8,  
    half_t, layout::TensorNWC, 8,  
    float,  
    Shape<_64,_128,Shape<_64>>, Shape<_2,_1,_1>,  
    conv::collective::StageCountAuto,  
    conv::collective::KernelScheduleAuto  
>::CollectiveOp;
```

“I want a Hopper Dgrad collective on NDHWC layout tensors (3-D Dgrad)”:

```
using CollectiveOp = typename conv::collective::CollectiveBuilder<  
    arch::Sm90, arch::OpClassTensorOp,  
    conv::Operator::kDgrad,  
    half_t, layout::TensorNDHWC, 8,  
    half_t, layout::TensorNDHWC, 8,  
    float,  
    Shape<_64,_128,Shape<_64>>, Shape<_2,_1,_1>,  
    conv::collective::StageCount<5>,  
    conv::collective::KernelScheduleAuto  
>::CollectiveOp;
```

```
template <  
    class ArchTag,  
    class OpClass,  
    conv::Operator,  
    class ElementA,  
    class GmemLayoutA,  
    int AlignmentA,  
    class ElementB,  
    class GmemLayoutB,  
    int AlignmentB,  
    class ElementAccumulator,  
    class TileShape_MNK,  
    class ClusterShape_MNK,  
    class StageCountType,  
    class KernelScheduleType,  
    class Enable = void  
>  
struct CollectiveBuilder;
```

Builders do the heavy lifting ...

```
using CollectiveOp = typename conv::collective::CollectiveBuilder<
    arch::Sm90, arch::OpClassTensorOp,
    conv::Operator::kDgrad,
    half_t, layout::TensorNDHWC, 8,
    half_t, layout::TensorNDHWC, 8,
    float,
    Shape<_64,_128,Shape<_64>>, Shape<_2,_1,_1>,
    conv::collective::StageCountAuto,
    conv::collective::KernelScheduleAuto
>::CollectiveOp;
```

Automatically maps down to the best mainloop config:
mainloop type, stage count, GMMA instruction, TMA instruction, smem layouts, kernel schedule ...

```
cutlass::conv::collective::CollectiveConv<
    conv::MainloopSm90TmaGmmaWarpSpecializedImplicitGemm<
        conv::Operator::kDgrad, 8, 3, Shape<_2,_1,_1>,
        cutlass::conv::KernelImplicitTmaWarpSpecializedSm90, 1>,
        Shape<_64,_128,Shape<_64>>,
    cutlass::half_t,
    cutlass::half_t,
    TiledMMA<
        MMA_Atom<SM90_64x128x16_F32F16F16_SS<GMMA::Major::K, GMMA::Major::MN>>,
        Layout<Shape<_1,_1,_1>>,
    cutlass::conv::collective::detail::Sm90ImplicitGemmTileTraits<
        SM90_TMA_LOAD_IM2COL,
        ComposedLayout<Swizzle<3, 4, 3>, smem_ptr_flag_bits<16>,
            Layout<Shape<_64,_64,_8>, Stride<_64,_1,_4096>>>,
        void>,
    cutlass::conv::collective::detail::Sm90ImplicitGemmTileTraits<
        SM90_TMA_LOAD_MULTICAST,
        ComposedLayout<Swizzle<3, 4, 3>, smem_ptr_flag_bits<16>,
            Layout<Shape<Shape<_64,_2>,_64,_8>, Stride<Stride<_1,_4096>,_64,_8192>>>,
        void>>
```



Ease of authoring custom kernels

Writing Custom Kernels

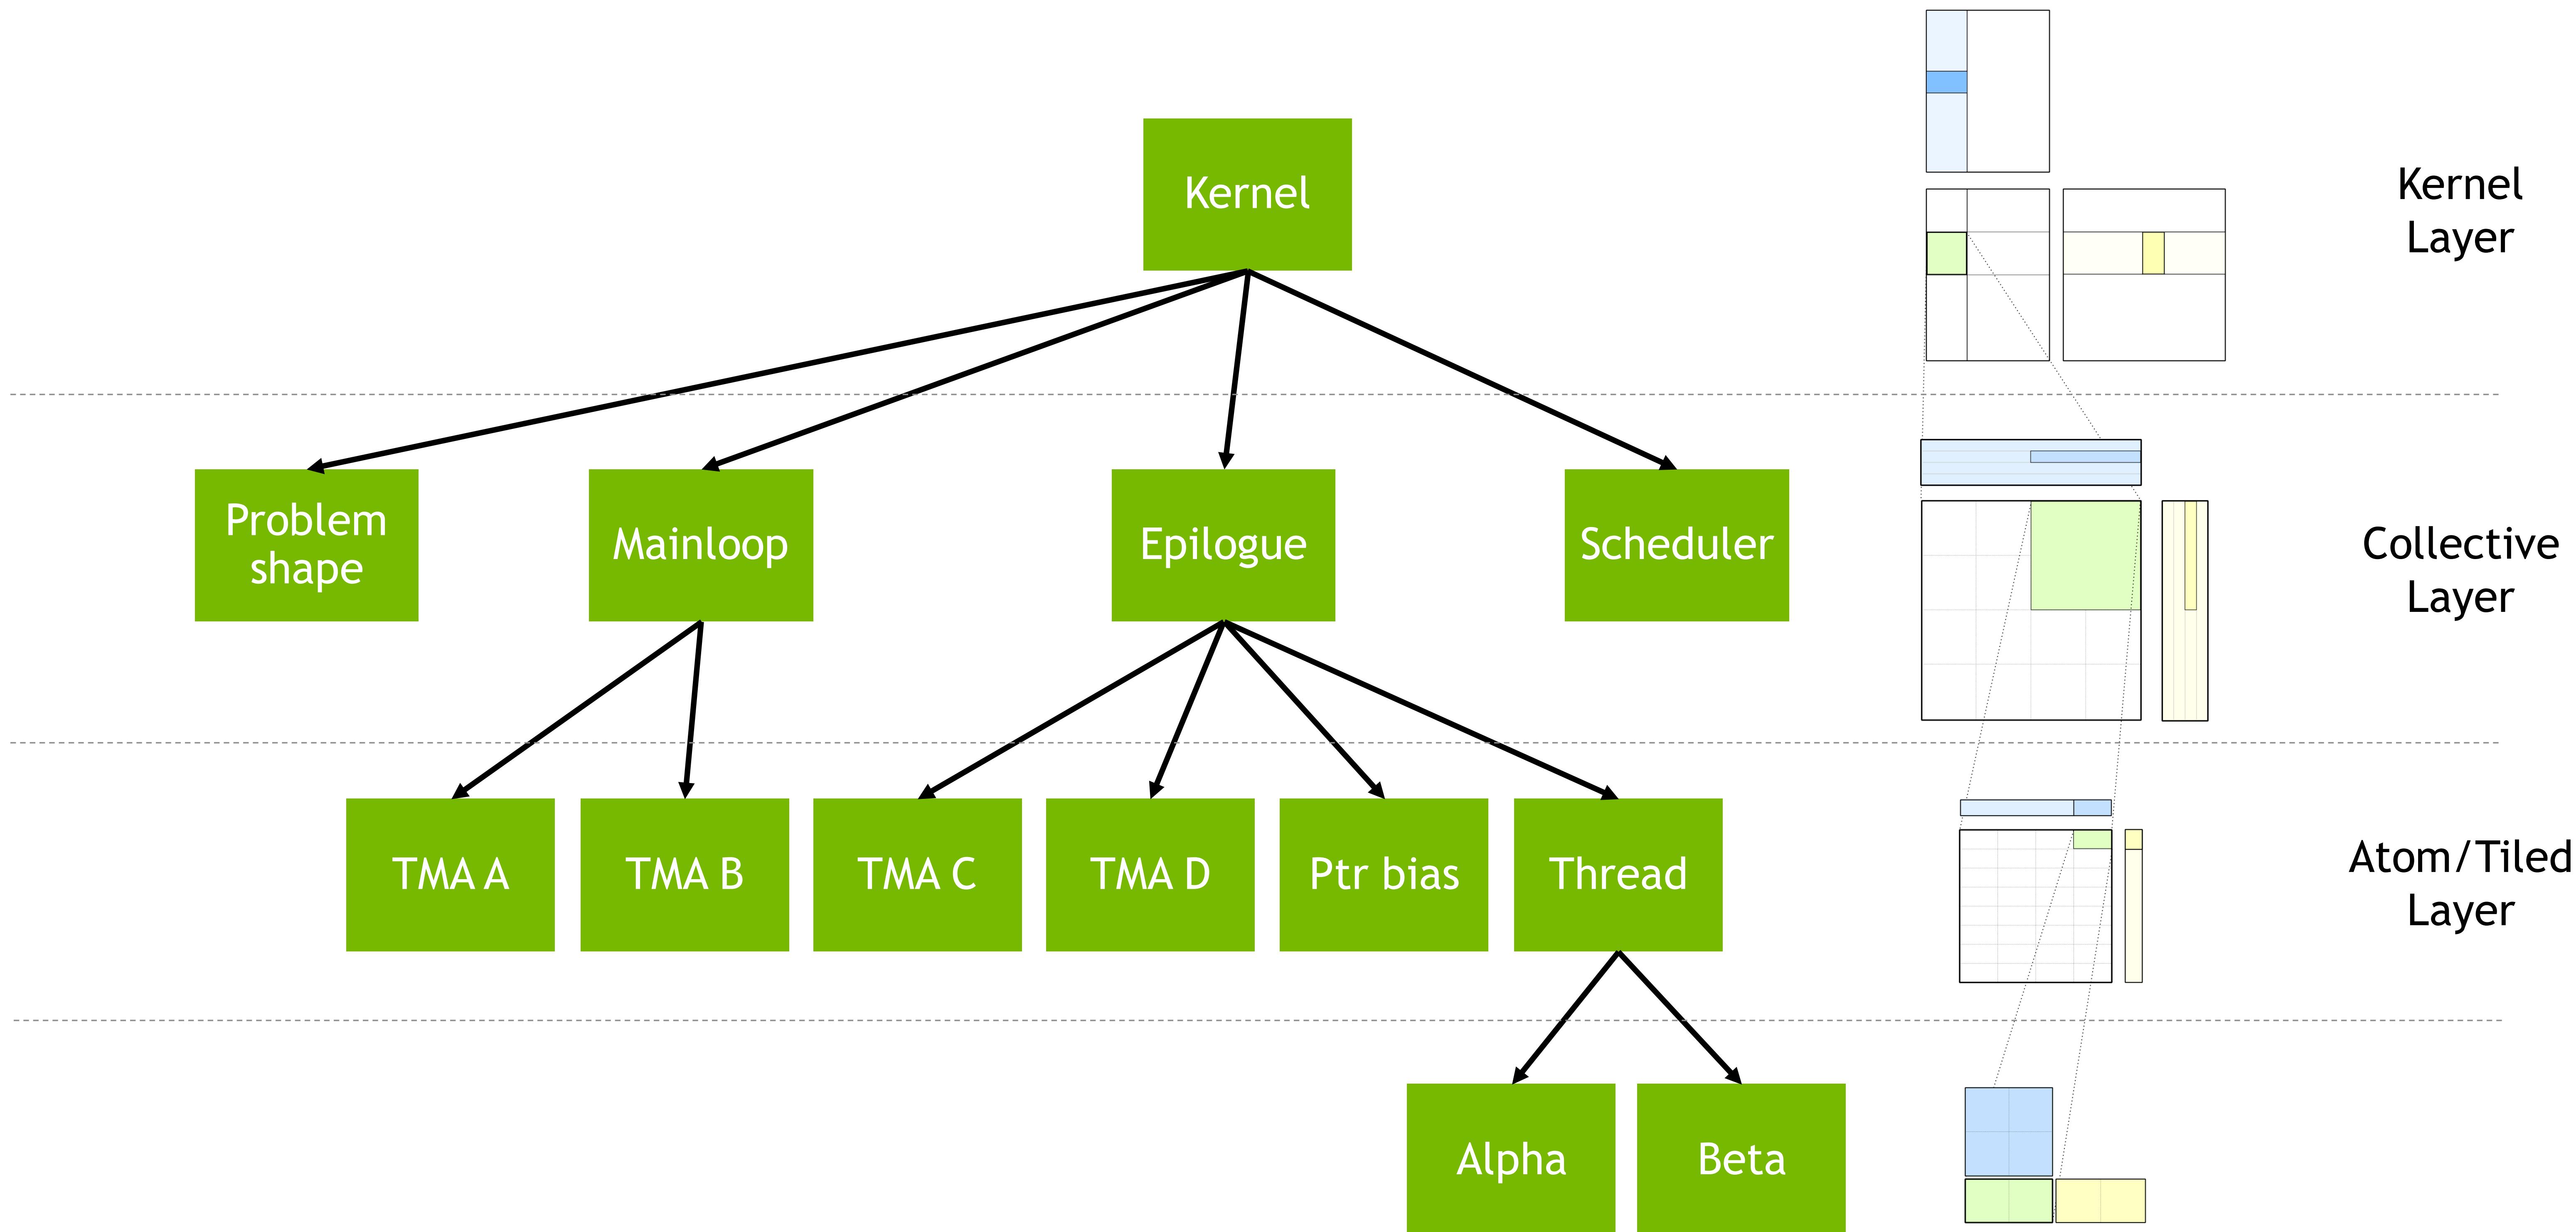
Plug and play with dispatch policies

- Write custom **mainloops** (micro-kernels), compose with existing **schedule** (outer loops) via dispatch policies
- Write custom **schedules** (outer loops), compose with existing **mainloops** (micro-kernels) via dispatch policies
- Kernel layer totally agnostic of # of in/out tensors and the semantics of the computation itself
- Kernel is a composition of a mainloop, epilogue, and tile scheduler: compose them freely

```
template<
    int Stages_,
    class ClusterShape_ = Shape<_1,_1,_1>,
    class KernelSchedule = KernelTmaWarpSpecialized // or KernelTmaWarpSpecializedPersistent
>
struct MyCustomHopperMainloopWoot {
    constexpr static int Stages = Stages_;
    using ClusterShape = ClusterShape_;
    using ArchTag = arch::Sm90;
    using Schedule = KernelSchedule;
};
```

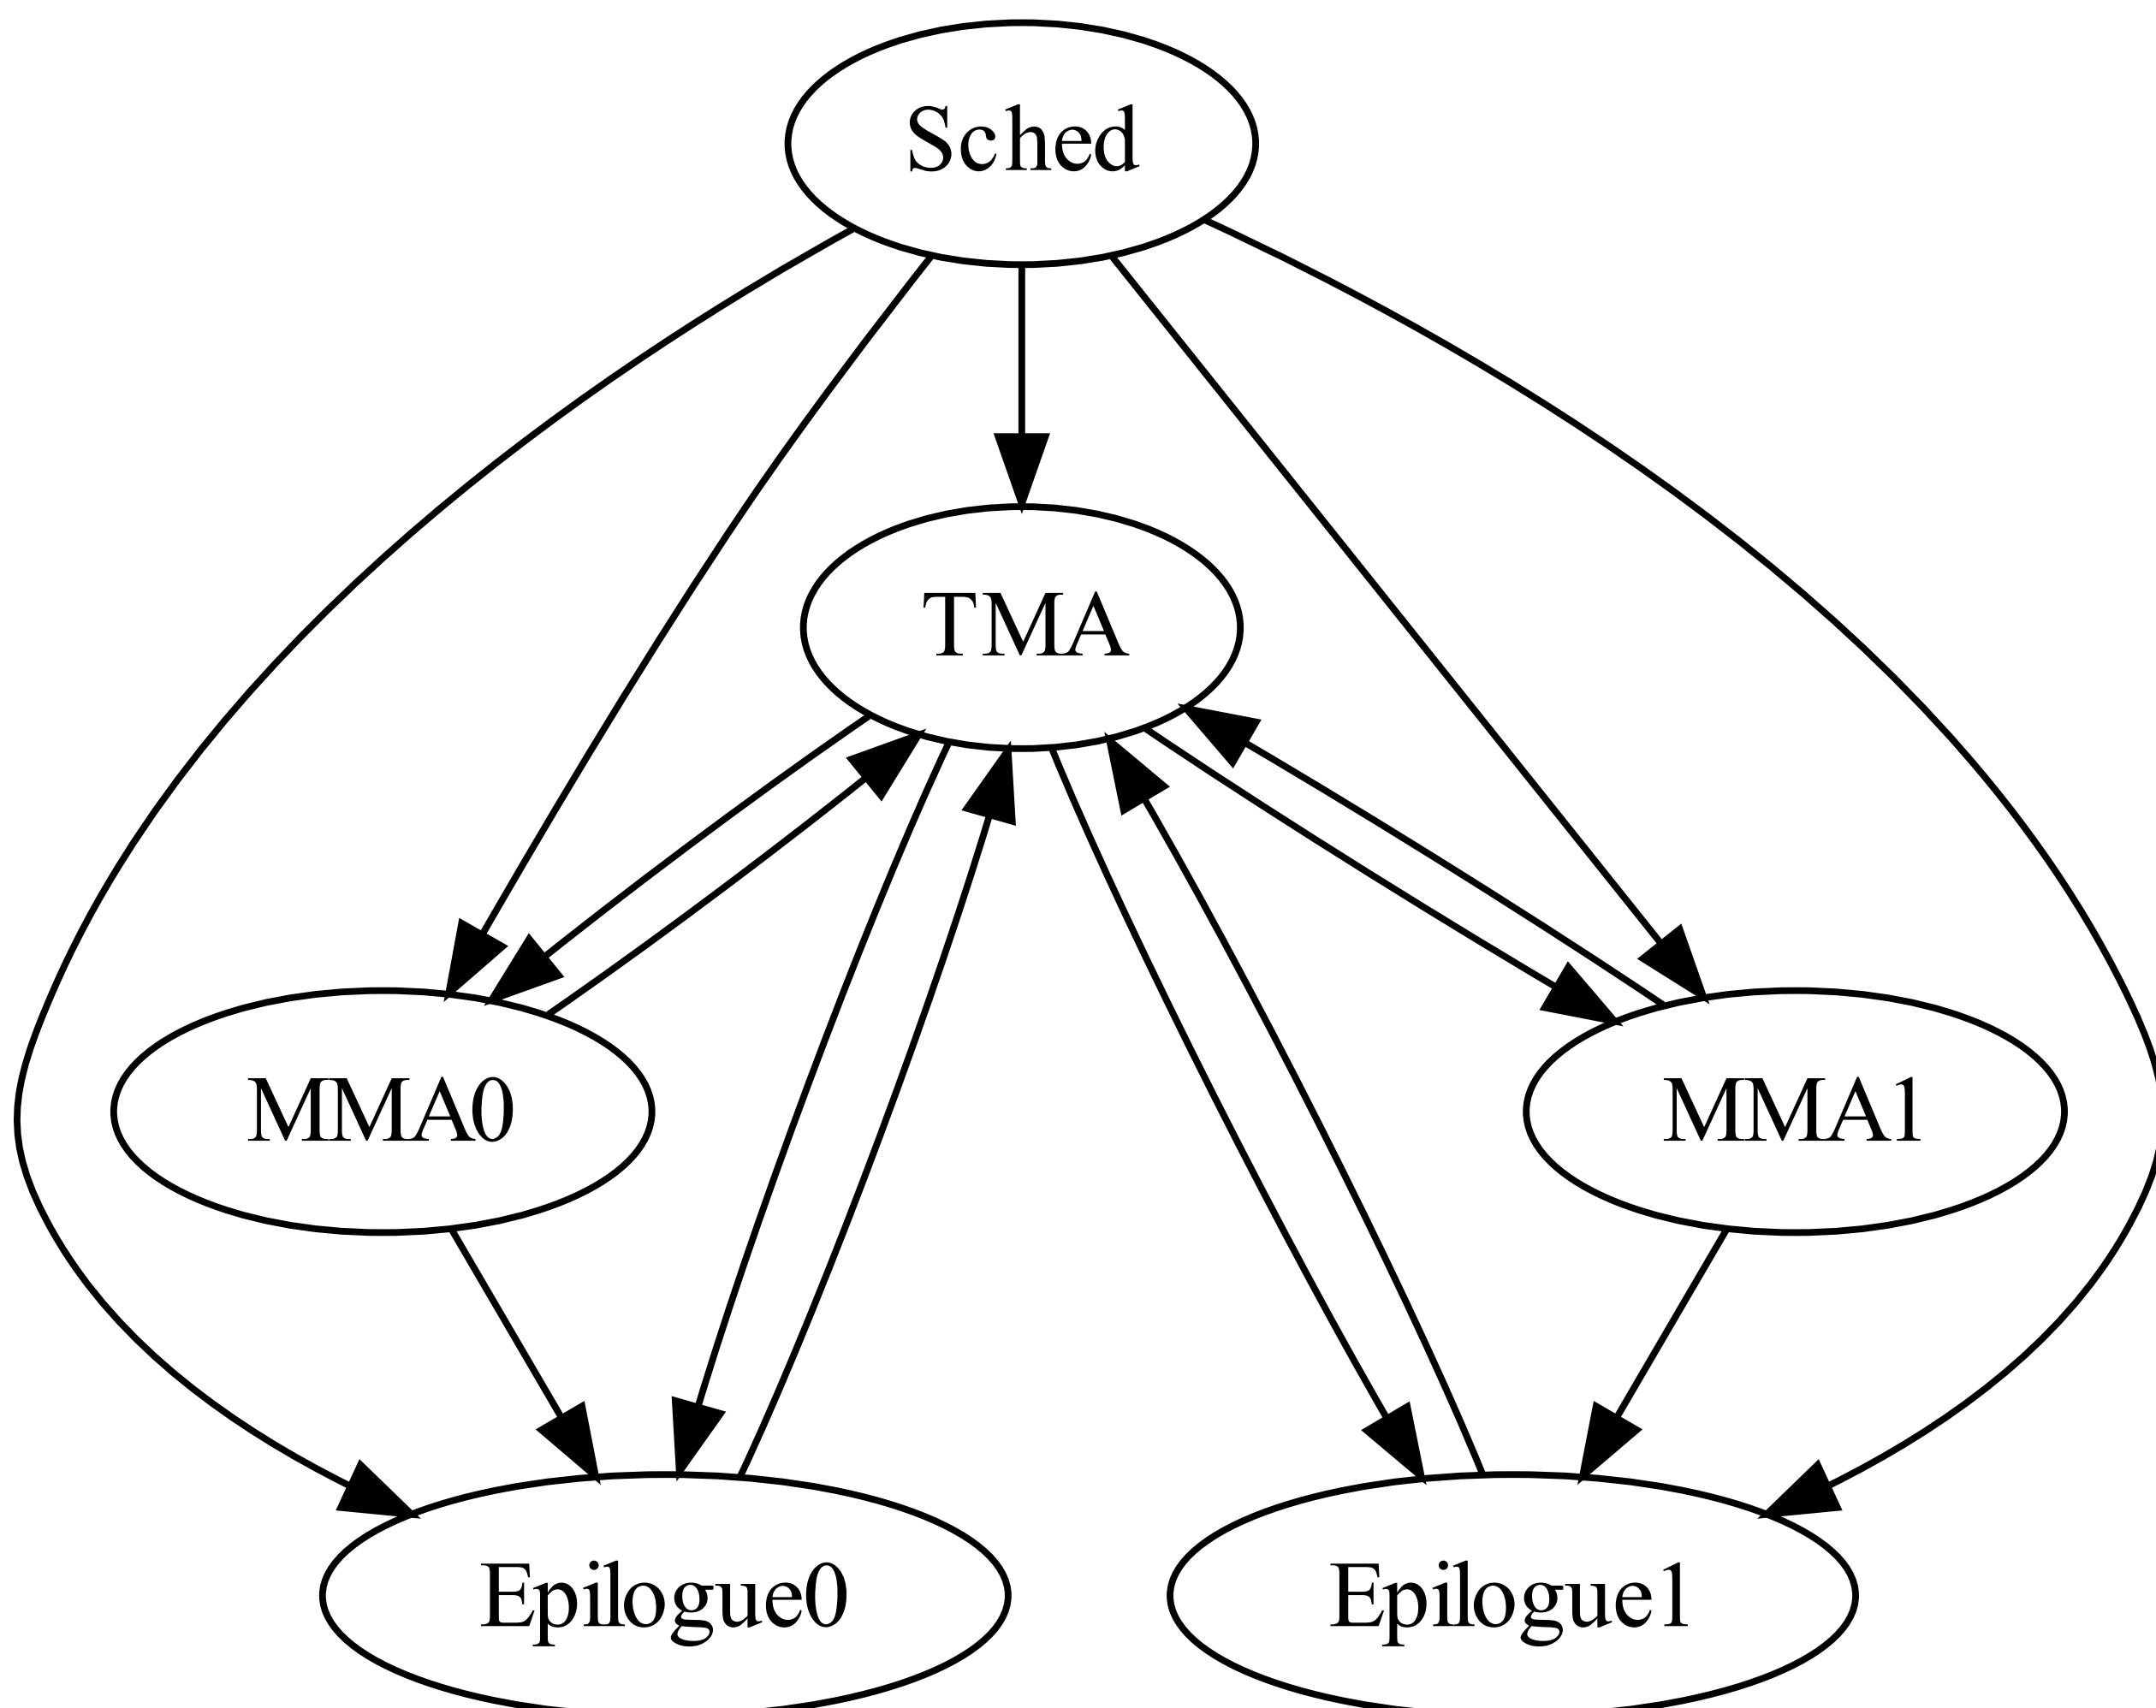
Kernel Params fully congruent with the structure of the kernel

Host Arguments -> Device Params in CUTLASS 3.x

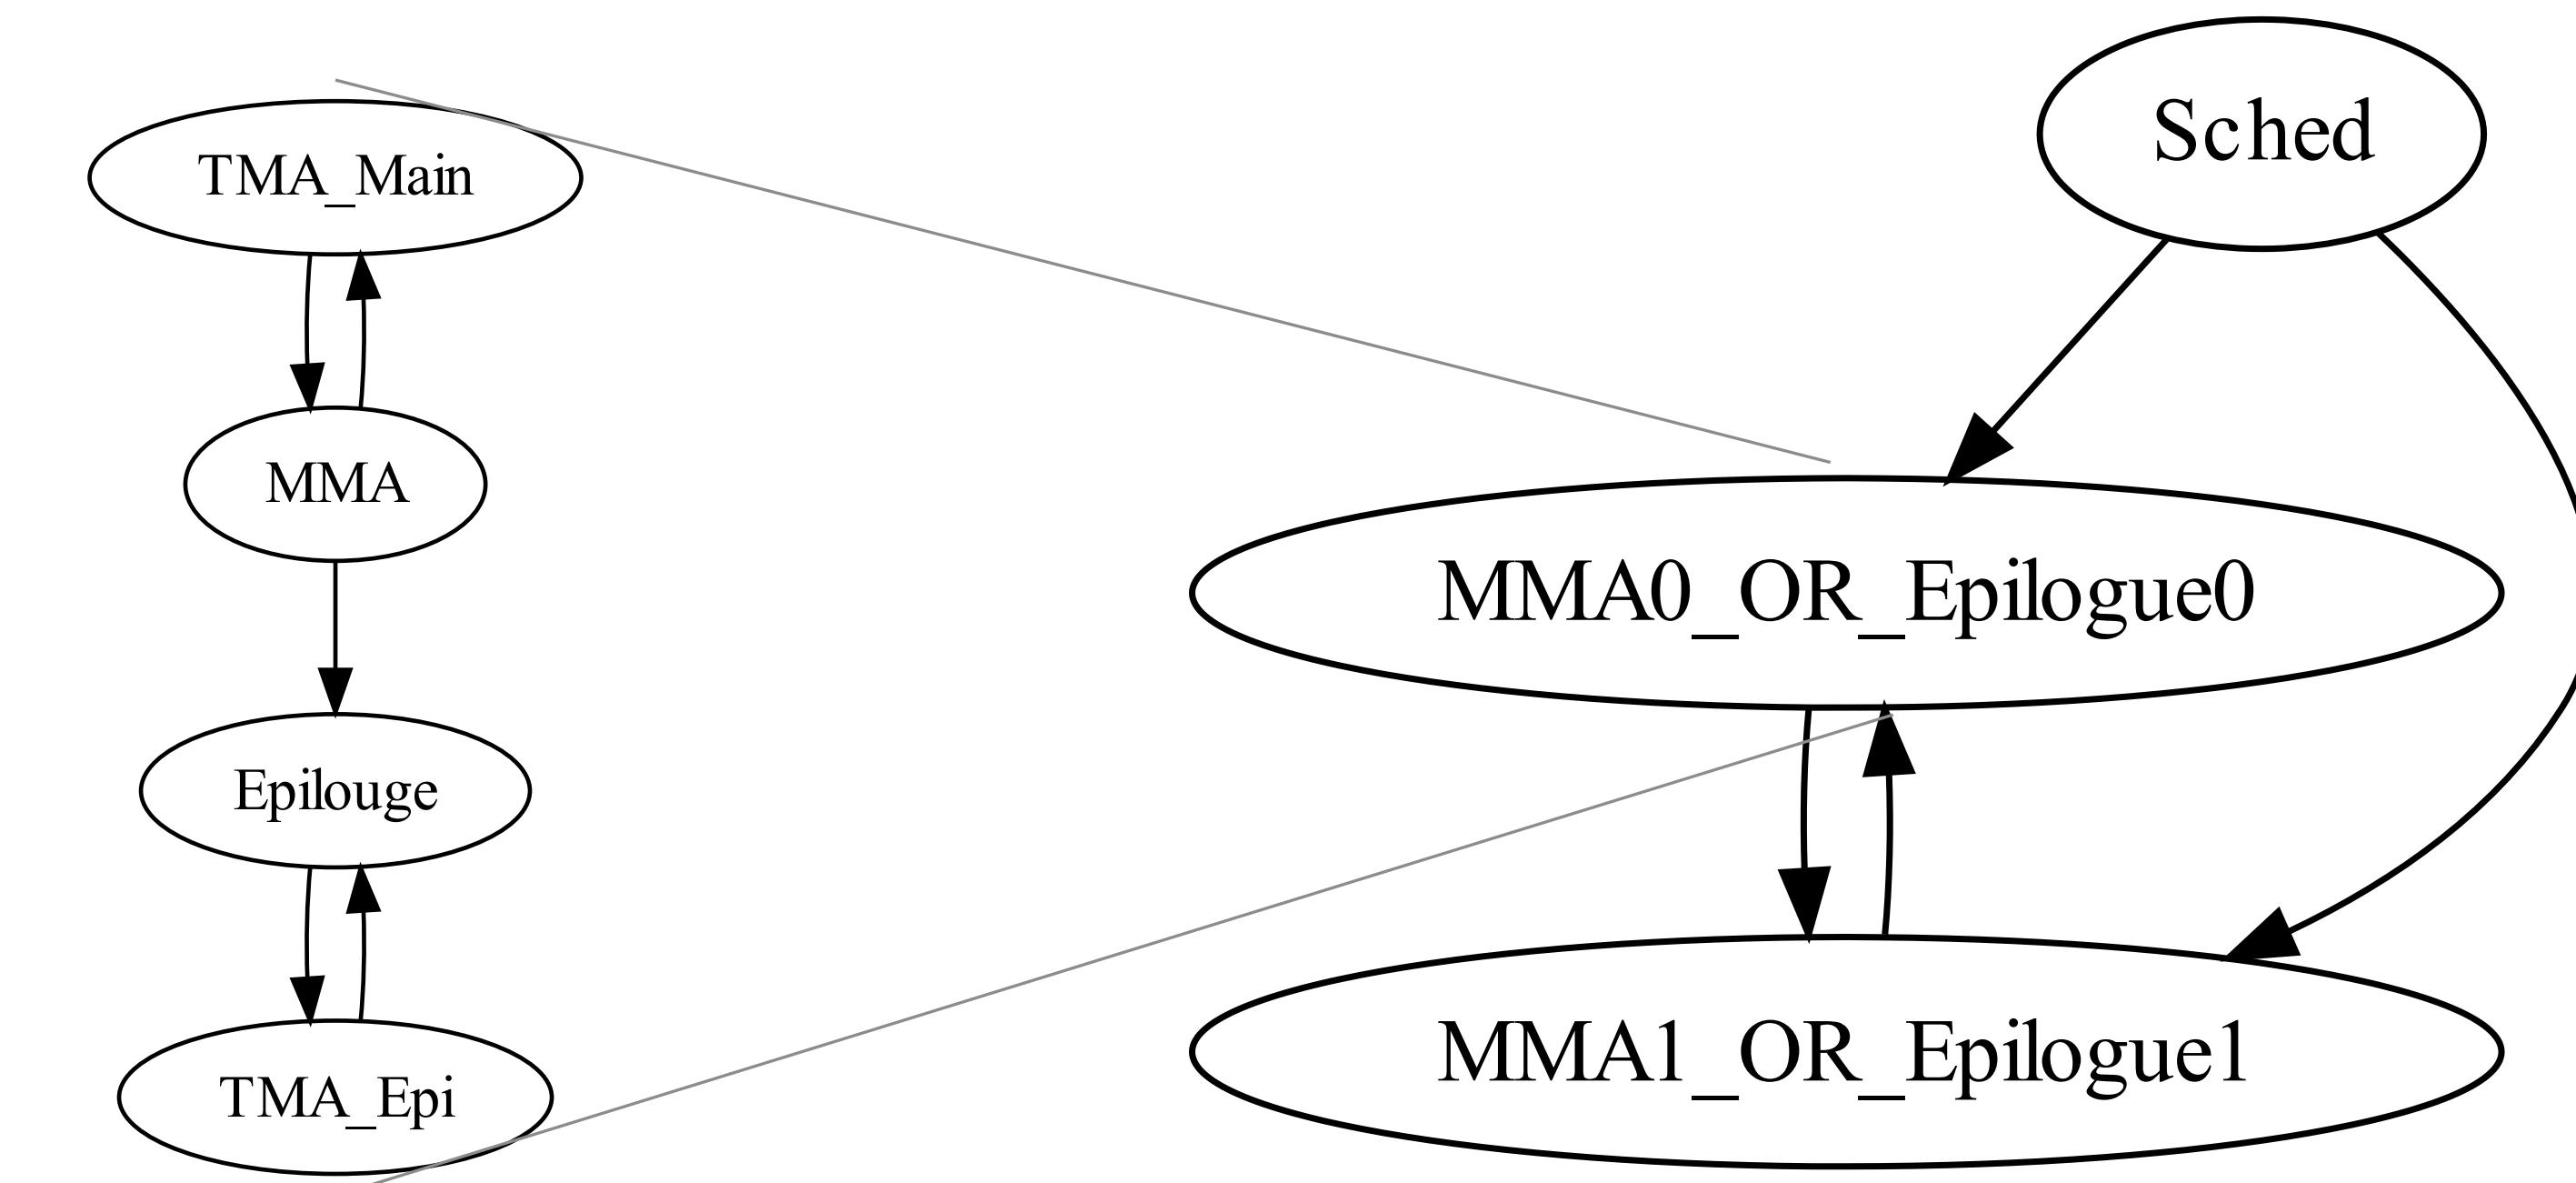


CUTLASS pipelines make synchro easy!

Collectives compose with each other and kernel schedules



Without a hierarchical pipeline model,
Hopper synchro model is too complex to reason
about



Pipelines greatly simplify synchronization by
grouping the synchro graph nodes hierarchically



Custom kernel rules of thumb

What changes when?

- Kernel fusion?
 - Write a custom kernel schedule that compose multiple existing collectives together
 - Get tile schedulers, pipelines, and highly optimized collectives for free
- Mainloop fusion?
 - Customize an existing collective MMA out of CUTLASS
 - Define a new dispatch policy and compose with an existing kernel layer
 - Hopper Mixed Input GEMMs reuse the vanilla kernel layer! (get all outer loop optimizations for free!)
 - Prefer doing mainloop fusions on A tensor only (due to MMA feature supporting A from registers)
- Epilogue fusion?
 - Use epilogue visitor tree for to make your life extremely easy
- Custom load balancing, L2 locality, or other outer loop shenanigans?
 - Write a custom tile scheduler, compose with existing mainloops, epilogue, and entire kernel layers
 - This is how StreamK is implemented in CUTLASS 3.x – it does not need a new kernel layer

Support for a variety of mixed-input combinations

```
using CollectiveMainloop =  
    typename cutlass::gemm::collective::CollectiveBuilder<  
        cutlass::arch::Sm90, cutlass::arch::OpClassTensorOp,  
        ElementA, LayoutA, AlignmentA,  
        ElementB, LayoutB, AlignmentB,  
        ElementAccumulator,  
        TileShape, ClusterShape,  
        StageCount,  
        KernelScheduleType  
>::CollectiveOp;
```

Requirements:

- Tensor Core instruction exists for wider data type
- Operands meet TMA requirements

FP16 X INT8

FP16 X INT4

BF16 X INT8

FP8 X INT4

INT8 X FP16

INT8 X BF16

INT4 X FP8

★ INT2 X FP16

★ INT1 X FP8

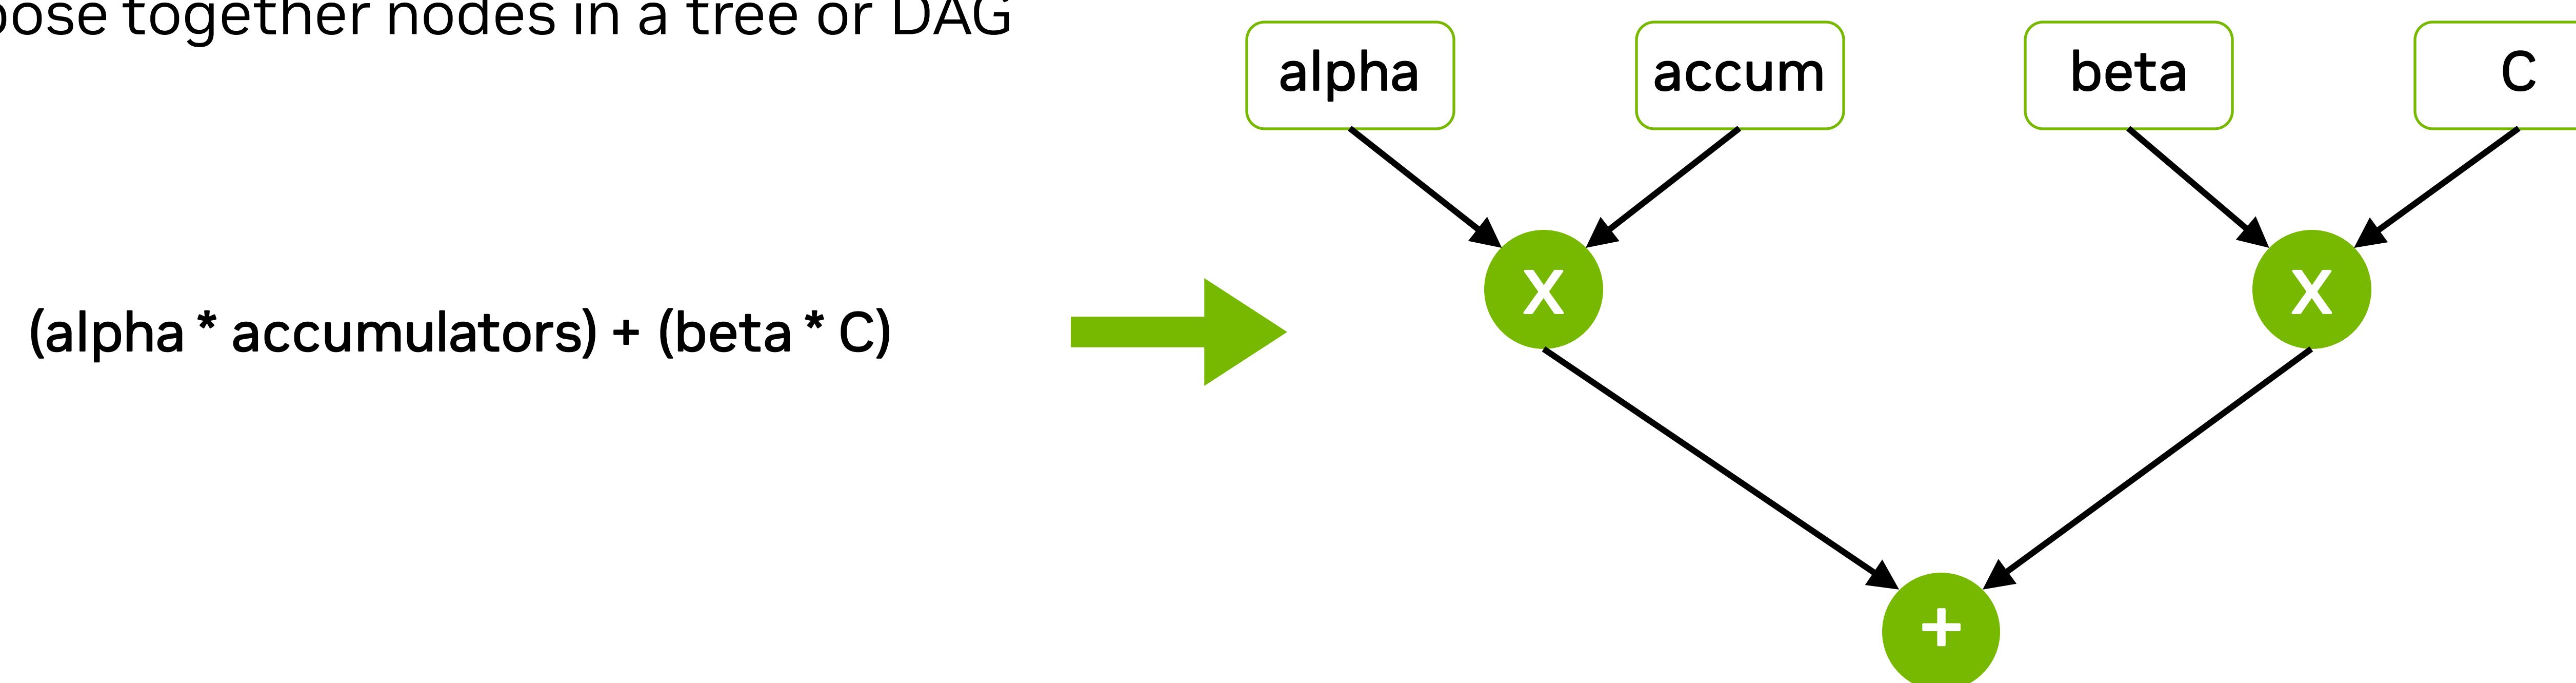
(and more)

Epilogue Visitor Tree (EVT): building blocks for composing fused epilogues

- Set of primitive nodes that can be composed together to build complex epilogues

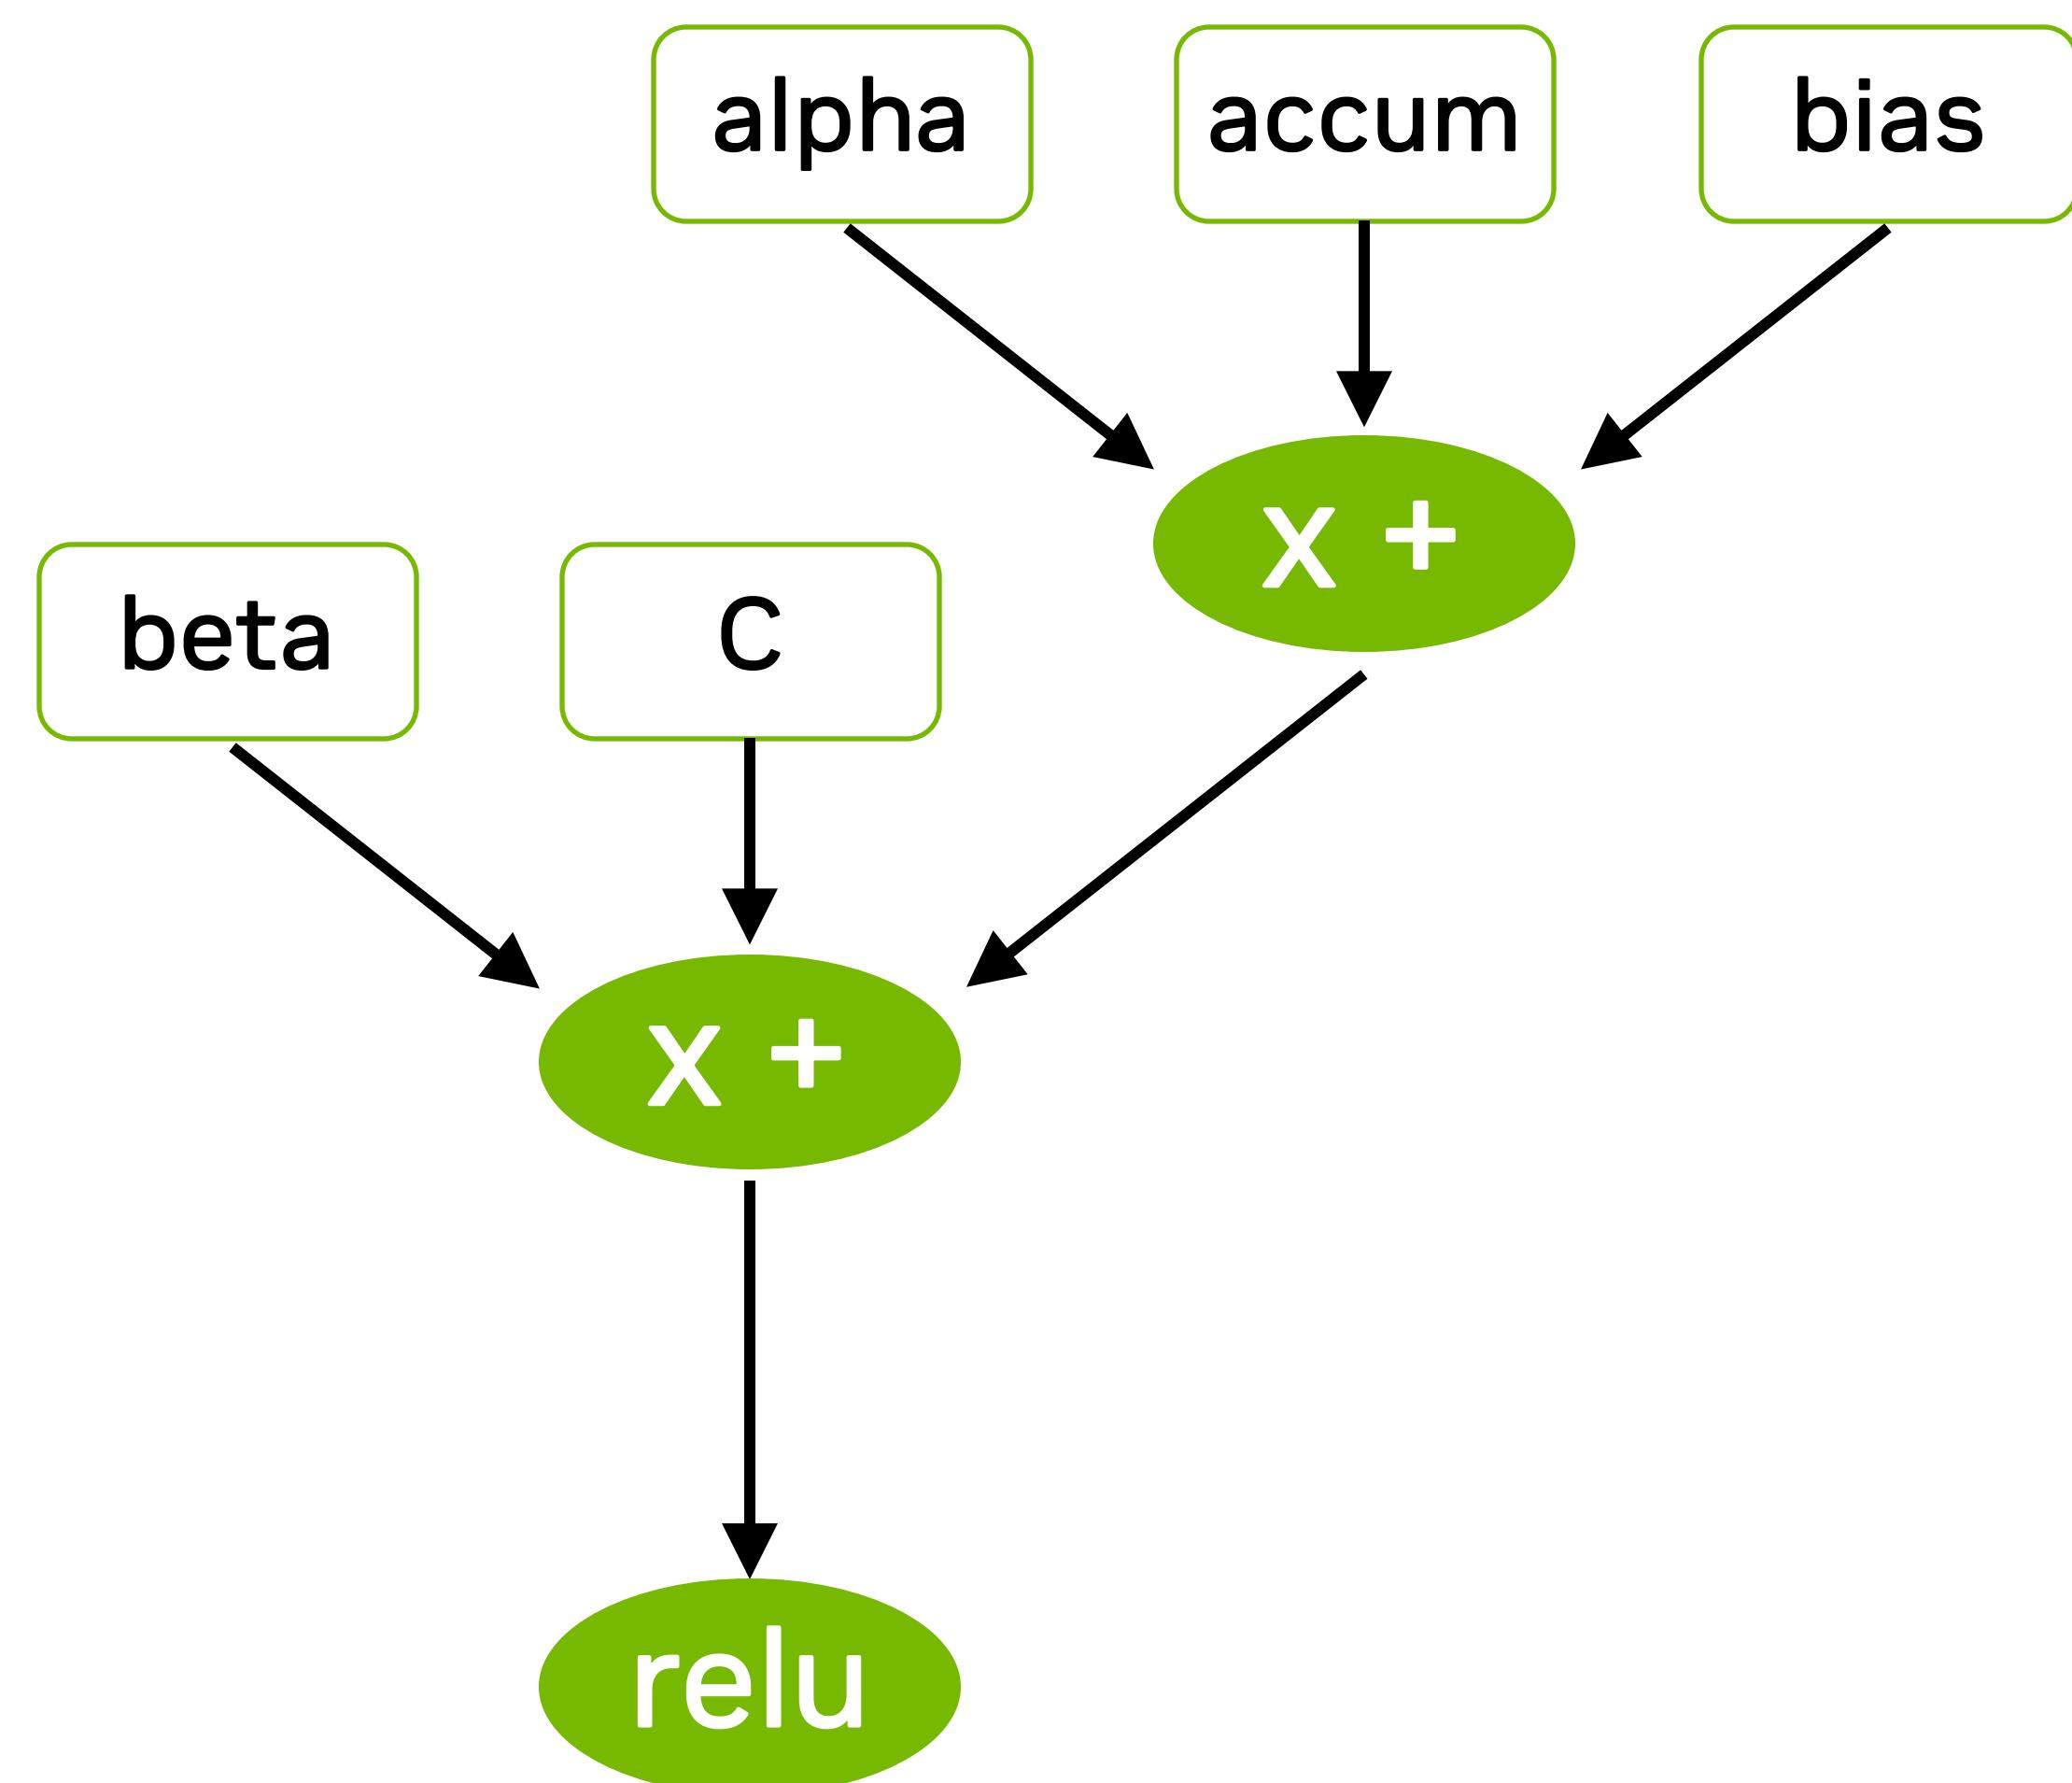
Load	Compute	Store
Accumulator	Compute (elementwise, binary, ternary)	Aux tensor
Aux tensor		Row reduction
Row broadcast		Column reduction
Column broadcast		Scalar reduction
Scalar broadcast		

- Compose together nodes in a tree or DAG



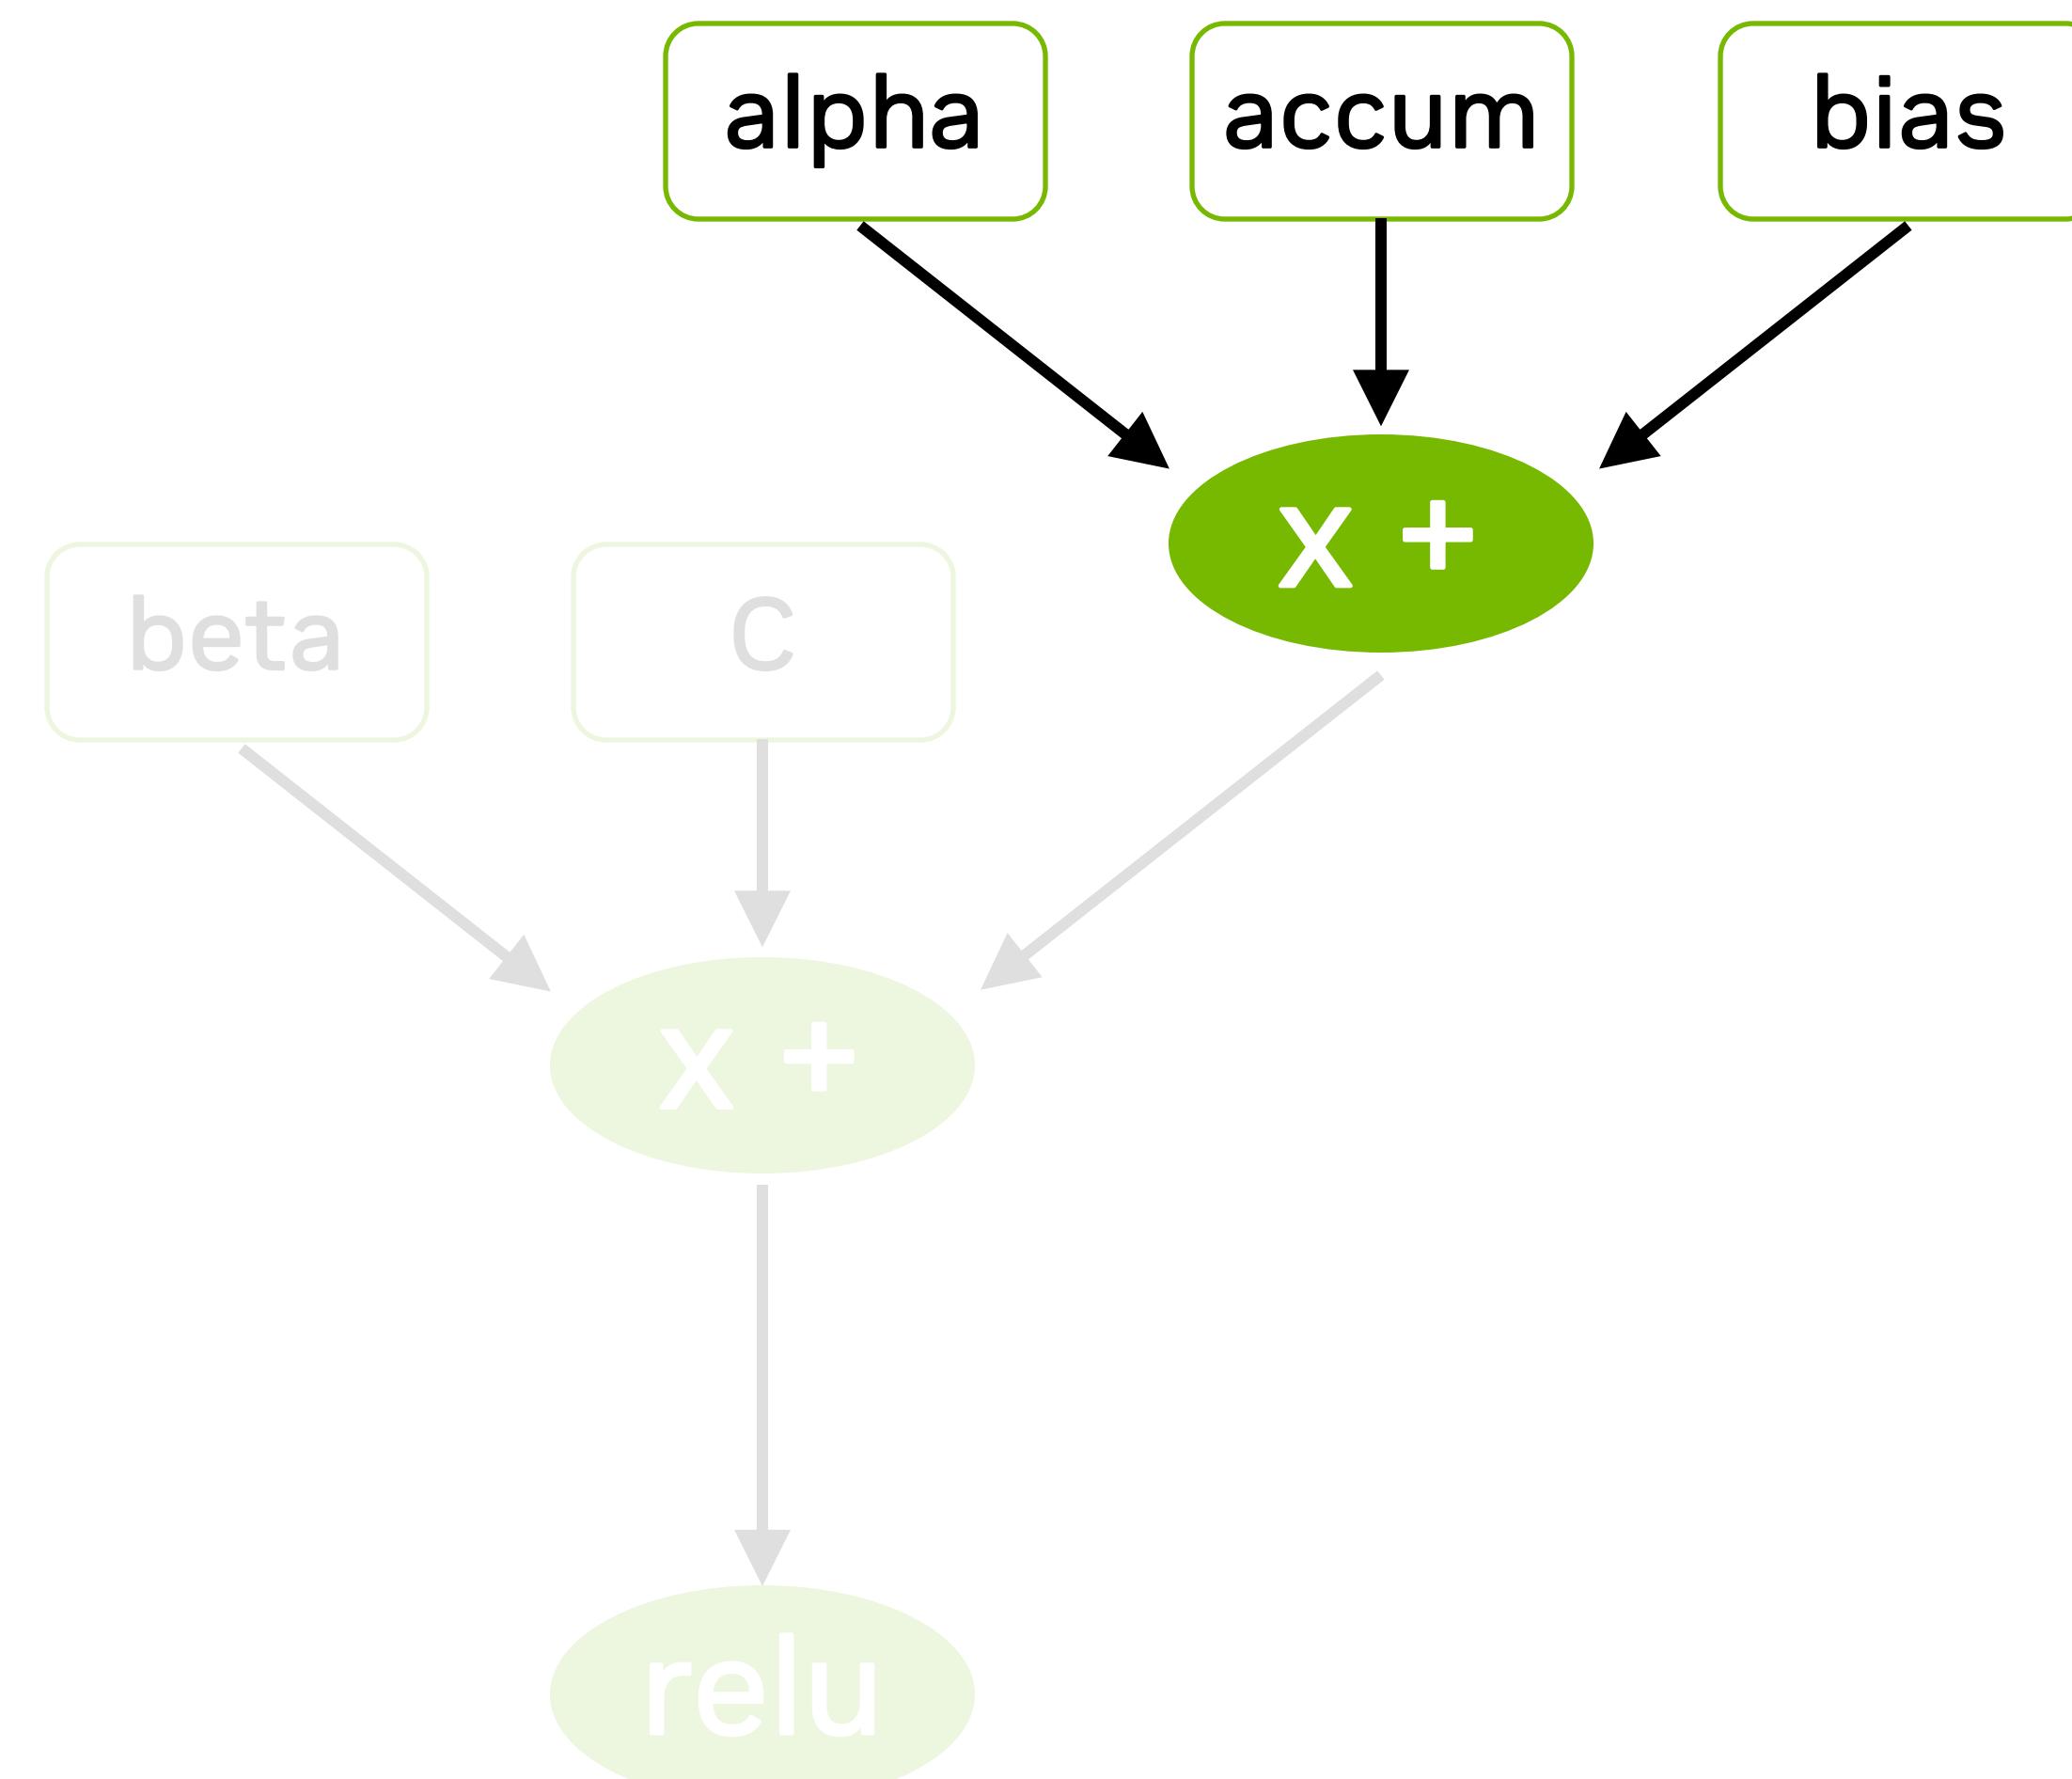
Adding an epilogue visitor tree to a GEMM in CUTLASS 3.x (C++)

$\text{ReLU}(\text{alpha} * \text{accumulators} + \text{bias} + (\text{beta} * \text{C}))$



Adding an epilogue visitor tree to a GEMM in CUTLASS 3.x (C++)

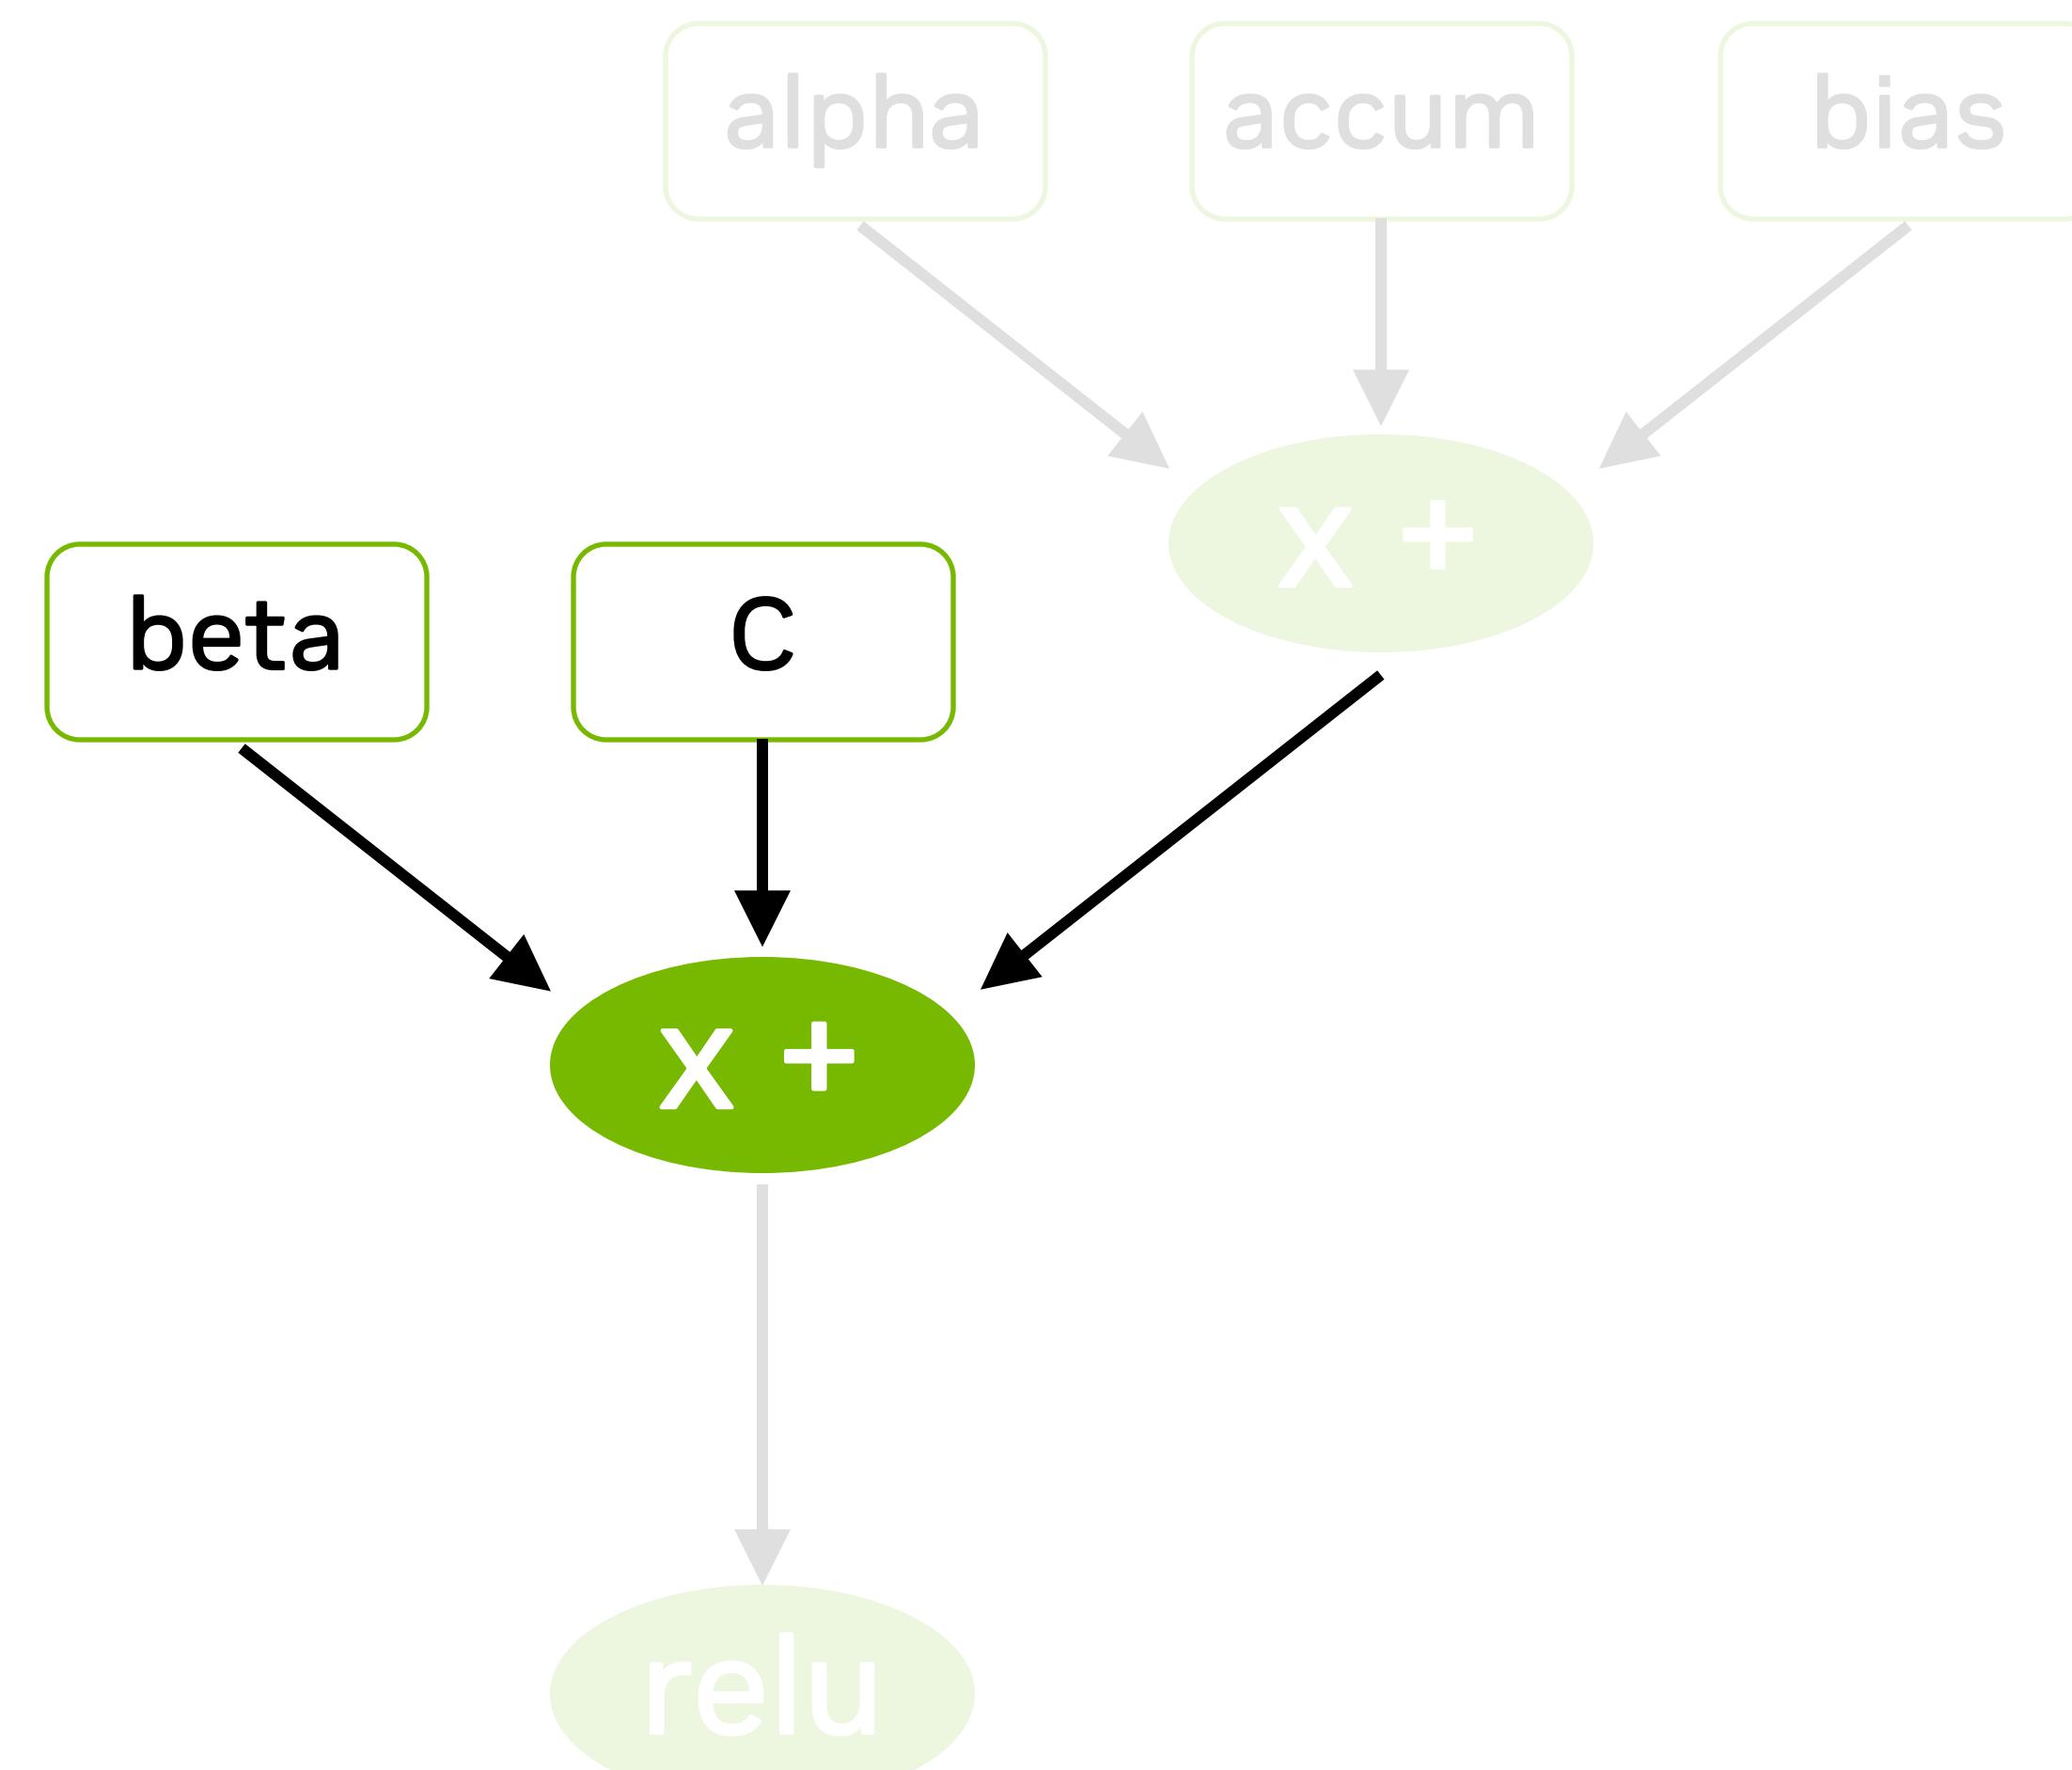
ReLU($\text{alpha} * \text{accumulators} + \text{bias} + (\text{beta} * \text{C})$)



```
using Alpha = Sm90ScalarBroadcast<ElementScalar>;  
using Accum = Sm90AccFetch;  
using Bias = Sm90ColBroadcast<  
    0, TileShape, ElementBias, Stride<_1,_0,int>, AlignmentBias>;  
using MultiplyAdd = Sm90Compute<  
    multiply_add, ElementCompute, ElementCompute, RoundStyle>;  
using EVTCompute0 = Sm90EVT<MultiplyAdd, Alpha, Accum, Bias>;
```

Adding an epilogue visitor tree to a GEMM in CUTLASS 3.x (C++)

ReLU($\text{alpha} * \text{accumulators} + \text{bias} + (\text{beta} * \mathbf{C})$)

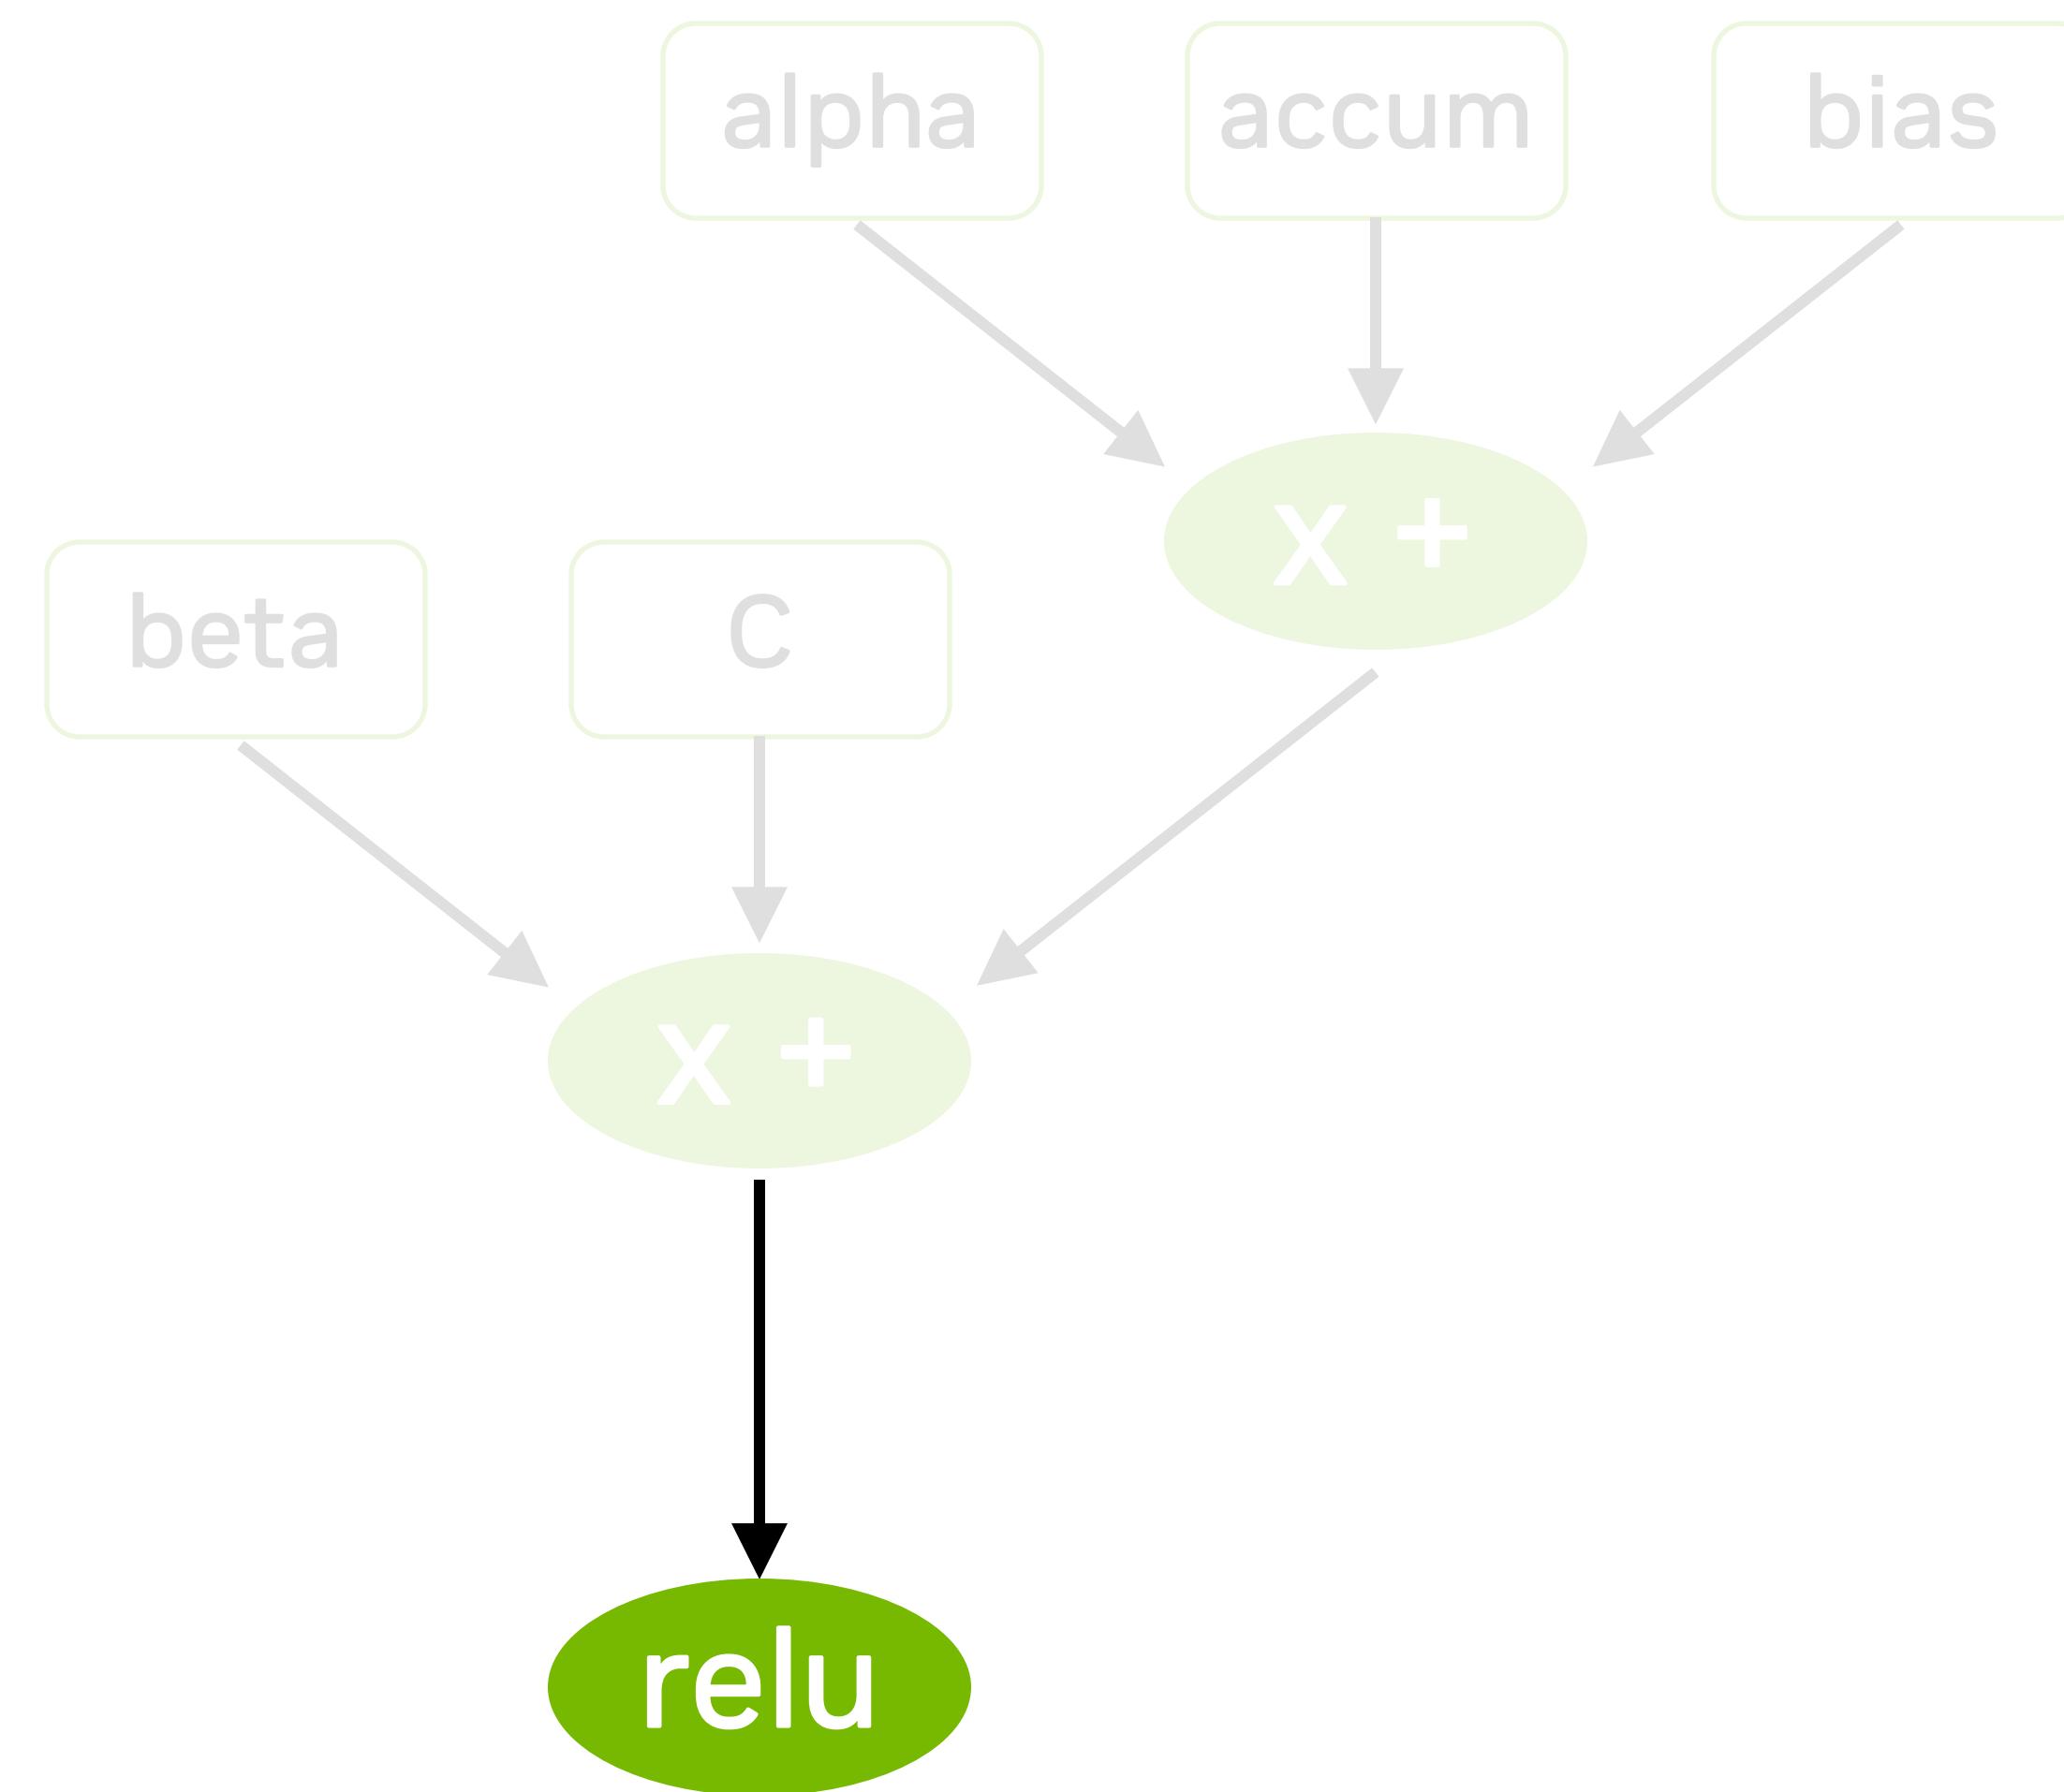


```
using Alpha = Sm90ScalarBroadcast<ElementScalar>;  
using Accum = Sm90AccFetch;  
using Bias = Sm90ColBroadcast<  
    0, TileShape, ElementBias, Stride<_1,_0,int>, AlignmentBias>;  
using MultiplyAdd = Sm90Compute<  
    multiply_add, ElementCompute, ElementCompute, RoundStyle>;  
using EVTCompute0 = Sm90EVT<MultiplyAdd, Alpha, Accum, Bias>;
```

```
using Beta = Sm90ScalarBroadcast<ElementScalar>;  
using C = Sm90SrcFetch<ElementSource>;  
using EVTCompute1 = Sm90EVT<MultiplyAdd, Beta, C, EVTCompute0>;
```

Adding an epilogue visitor tree to a GEMM in CUTLASS 3.x (C++)

ReLU($\text{alpha} * \text{accumulators} + \text{bias} + (\text{beta} * \mathbf{C})$)



```
using Alpha = Sm90ScalarBroadcast<ElementScalar>;
using Accum = Sm90AccFetch;
using Bias = Sm90ColBroadcast<0, TileShape, ElementBias, Stride<_1,_0,int>, AlignmentBias>;
using MultiplyAdd = Sm90Compute<multiply_add, ElementCompute, ElementCompute, RoundStyle>;
using EVTCompute0 = Sm90EVT<MultiplyAdd, Alpha, Accum, Bias>;
```



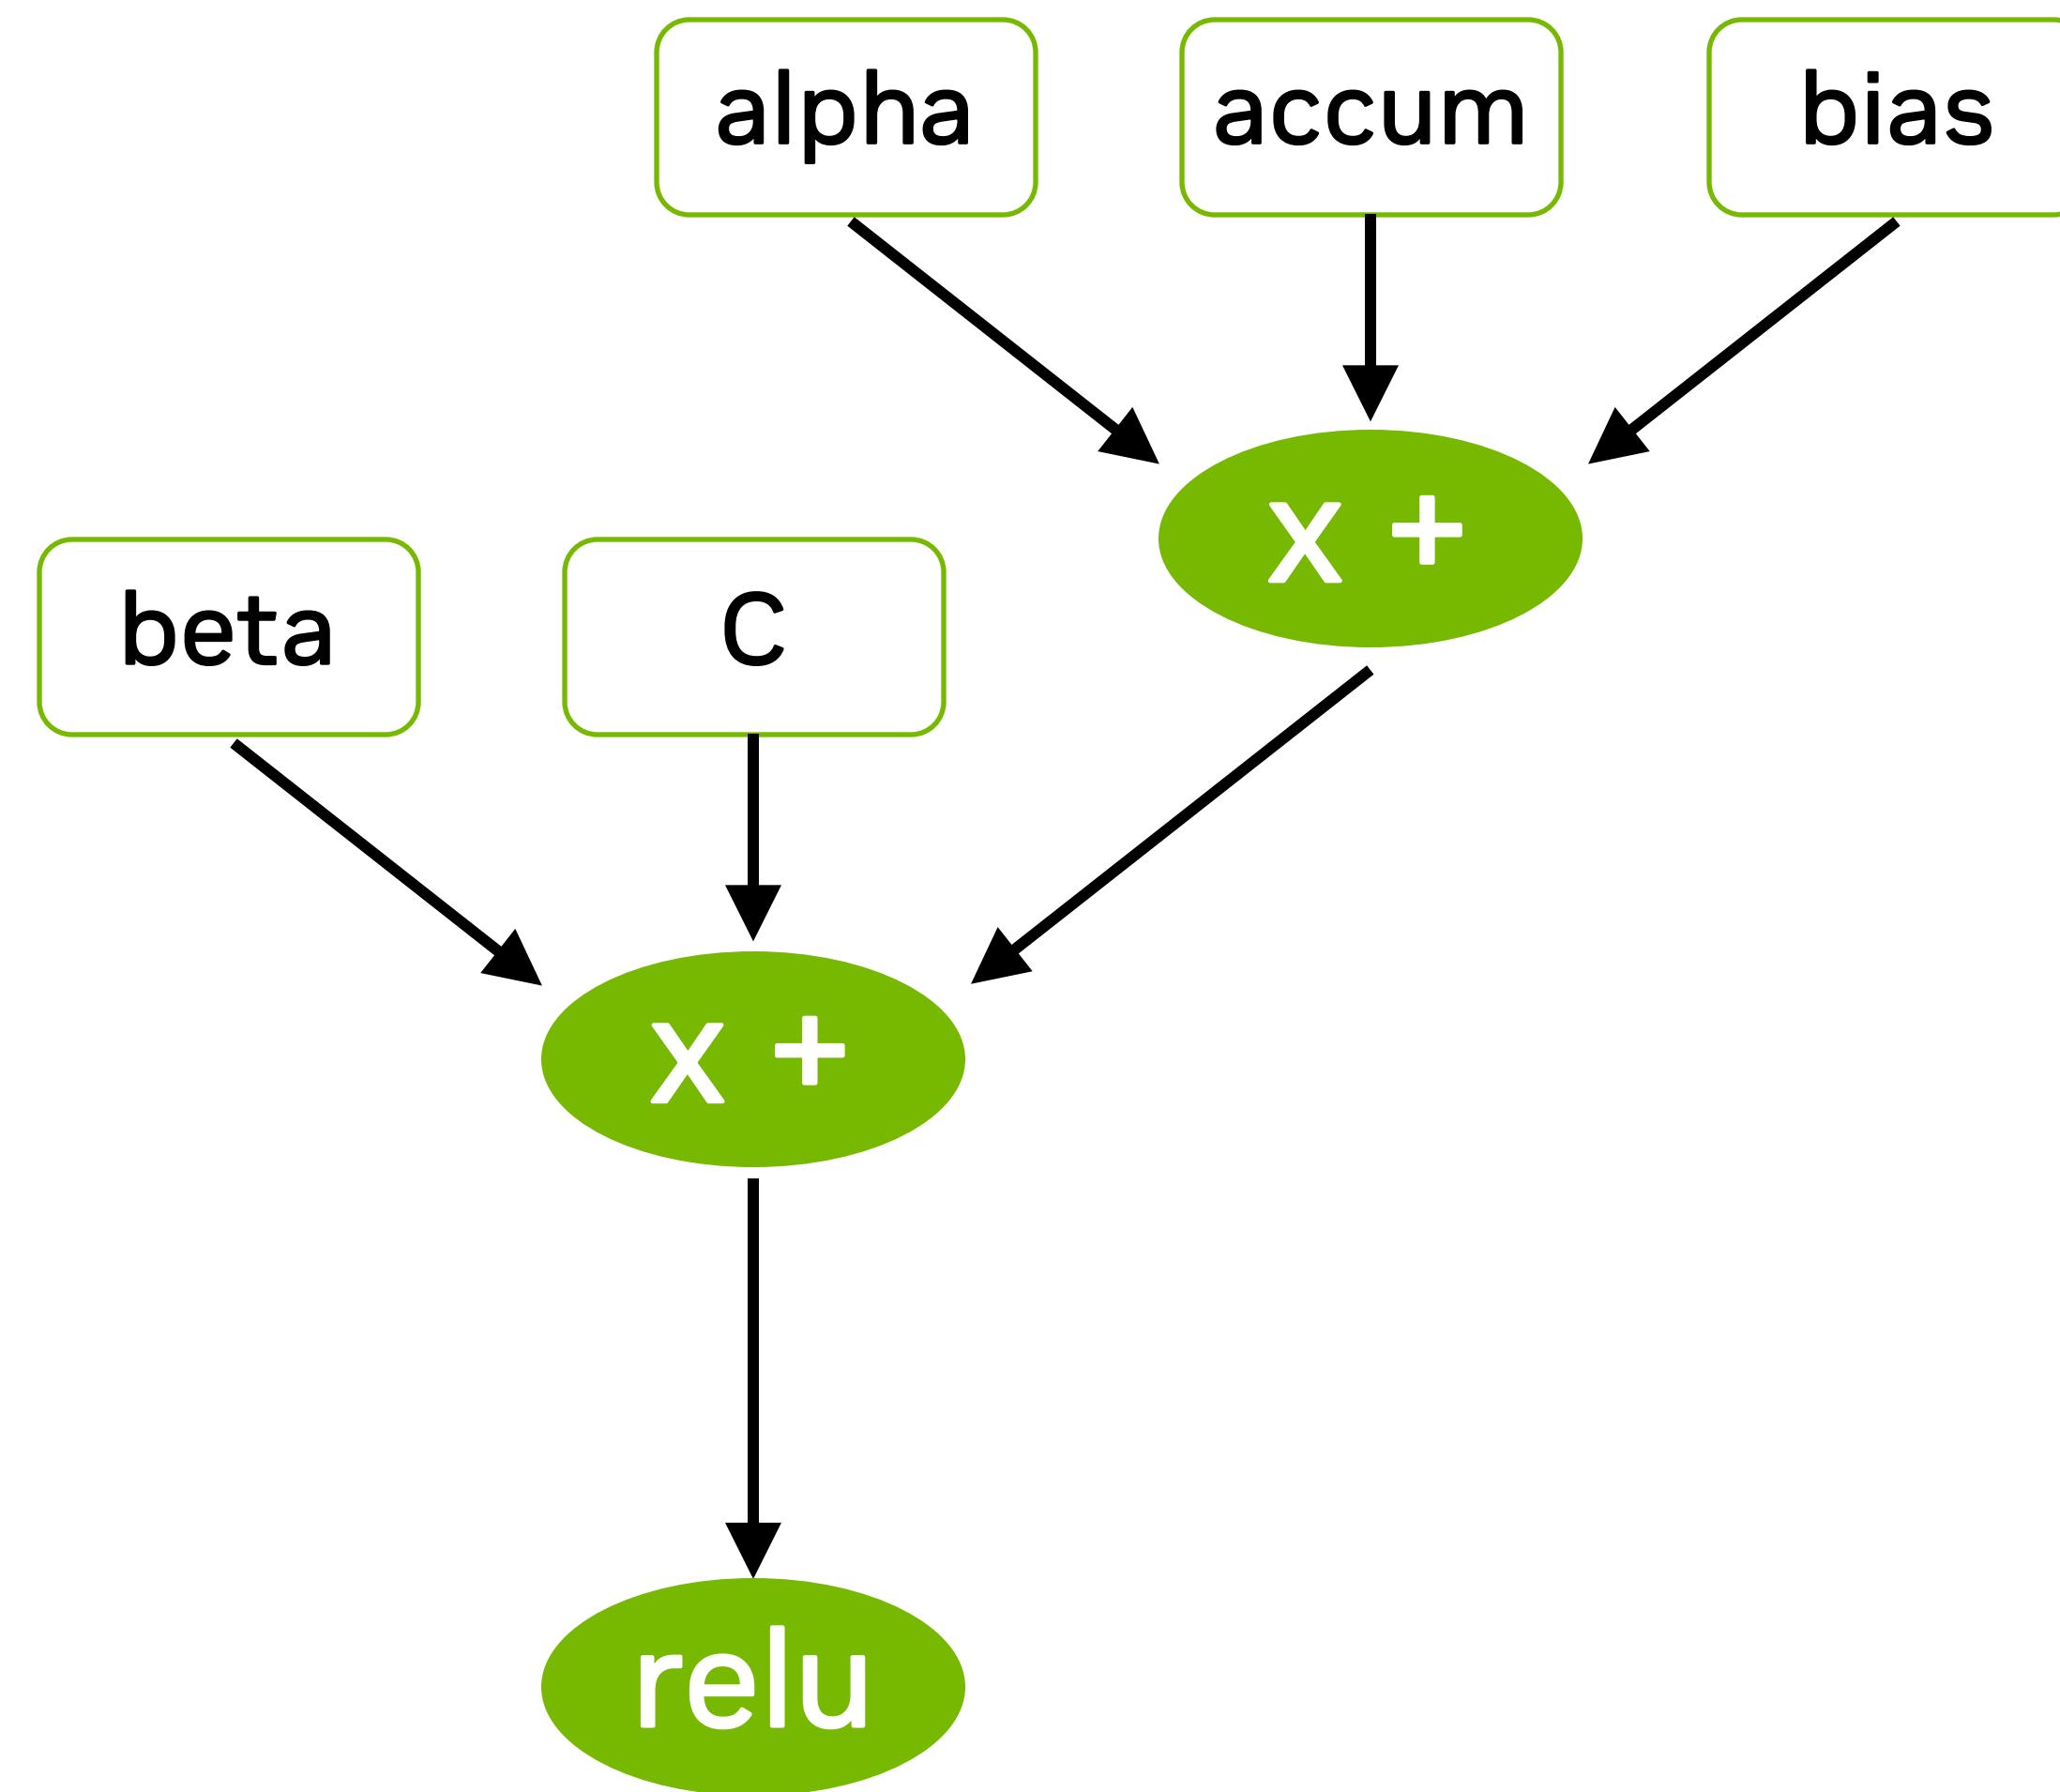
```
using Beta = Sm90ScalarBroadcast<ElementScalar>;
using C = Sm90SrcFetch<ElementSource>;
using EVTCompute1 = Sm90EVT<MultiplyAdd, Beta, C, EVTCompute0>;
```



```
using ReLUAct = Sm90Compute<ReLU, ElementOutput, ElementCompute, RoundStyle>;
using EVTOutput = Sm90EVT<ReLUAct, EVTCompute1>;
```

Adding an epilogue visitor tree to a GEMM in CUTLASS 3.x (C++)

ReLU($\text{alpha} * \text{accumulators} + \text{bias} + (\text{beta} * \text{C})$)



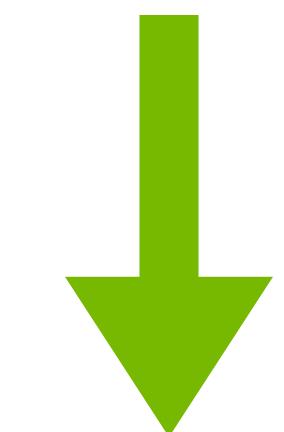
```
using EVTOutput = Sm90EVT<ReLUAct, EVTCompute1>;  
  
using CollectiveEpilogue = typename  
cutlass::epilogue::collective::CollectiveBuilder<  
    cutlass::arch::Sm90, cutlass::arch::OpClassTensorOp,  
    TileShape, ClusterShape,  
    cutlass::epilogue::collective::EpilogueTileAuto,  
    ElementAccumulator, ElementCompute,  
    ElementC, LayoutC, AlignmentC,  
    ElementD, LayoutD, AlignmentD,  
    EpilogueScheduleType,  
    EVTOutput  
>::CollectiveOp;
```

Pre-baked aliases for common patterns

```
using Alpha = Sm90ScalarBroadcast<ElementScalar>;
using Accum = Sm90AccFetch;
using Bias = Sm90ColBroadcast<
    0, TileShape, ElementBias, Stride<_1,_0,int>, AlignmentBias>;
using MultiplyAdd = Sm90Compute<
    multiply_add, ElementCompute, ElementCompute, RoundStyle>;
using EVTCompute0 = Sm90EVT<MultiplyAdd, Alpha, Accum, Bias>;
```

```
using Beta = Sm90ScalarBroadcast<ElementScalar>;
using C = Sm90SrcFetch<ElementSource>;
using EVTCompute1 = Sm90EVT<MultiplyAdd, Beta, C, EVTCompute0>;
```

```
using ReLUAct = Sm90Compute<ReLU, ElementOutput, ElementCompute, RoundStyle>;
using EVTOutput = Sm90EVT<ReLUAct, EVTCompute1>;
```



```
using EVTOutput = Sm90LinCombPerRowBiasEltAct<
    TileShape, ReLU, ElementOutput, ElementCompute>;
```

Epilogue visitor tree in Python

1. Declare a basic GEMM

```
import cutlass
plan = cutlass.op.Gemm(
    element=torch.float32,
    layout=cutlass.LayoutType.RowMajor)
```

2. Define an epilogue as a Python function

```
def my_epilogue(accum, alpha, C, beta, bias):
    D = relu(alpha * accum + beta * C + bias)
    return D
```

3. Define types and shapes of each EVT operand/output

```
empty_mn = torch.empty(size=(m, n), dtype=torch.float32)
empty_bias = torch.empty(size=(m, 1), dtype=torch.float32)
examples_inputs = {
    "accum": empty_mn, "C": empty_mn, "D": empty_mn,
    "alpha": 1.0, "beta": 1.0, "bias": empty_bias,
}
```

4. Construct the EVT and assign it to the GEMM

```
plan.epilogue_visitor = cutlass.epilogue.trace(
    my_epilogue, examples_inputs)
```

5. Compile and run the kernel

```
A, B, C, D, bias = ...
visitor_args = {
    "alpha": 2.0, "beta": 0.0,
    "C": C, "D": D, "bias": bias
}
plan.run(A, B, C, D, visitor_args=visitor_args)
```

EVT handles writing optimized epilogue loops for you

Consumer store warpgroup pseudocode

```
for (int i = 0; i < NumEpiSubtiles; ++i) {
    Tensor tRS_rAcc_frg_mn = tRS_rAcc_frg(_,i);

    copy(tiled_s2r, ...); // Copy C from smem to registers

    // Compute epilogue on accumulator fragments
    tRS_rD_frg(epi_v) = compute(tRS_rAcc_frg_mn, ...);

    copy(tiled_r2s, ...); // Copy results from register to smem

    copy(params.tma_store_d, ...); // Copy D smem to gmem

}
```

EVT handles writing optimized epilogue loops for you

Consumer store warpgroup pseudocode

```
callbacks.begin();
for (int i = 0; i < NumEpiSubtiles; ++i) {
    Tensor tRS_rAcc_frg_mn = tRS_rAcc_frg(_,i);

    copy(tiled_s2r, ...); // Copy C from smem to registers

    callbacks.previsit(...);

    // Compute epilogue on accumulator fragments
    tRS_rD_frg(epi_v) = callbacks.visit(tRS_rAcc_frg_mn, ...);

    callbacks.reduce(...); // Smem reduction callback entry

    copy(tiled_r2s, ...); // Copy results from register to smem

    callbacks.postreduce(...);

    copy(params.tma_store_d, ...); // Copy D smem to gmem

    callbacks.tma_store(..);
}

callbacks.end();
```

Using CUTLASS Profiler

Whether you are a C++ user or targeting PTX

- CUTLASS has a python based kernel emitter and a manifest to hold a bunch of kernels
- Autotuning strategy is to stamp out a set of candidates kernels and then ...
- Use the CUTLASS profiler to pick the best kernel for your problems of interest
- It is also possible to dump ptx of the best performing kernel with `cuobjdump` or -DCUTLASS_NVCC_KEEP



Hopper Programming Model

Courtesy : Stephen Jones

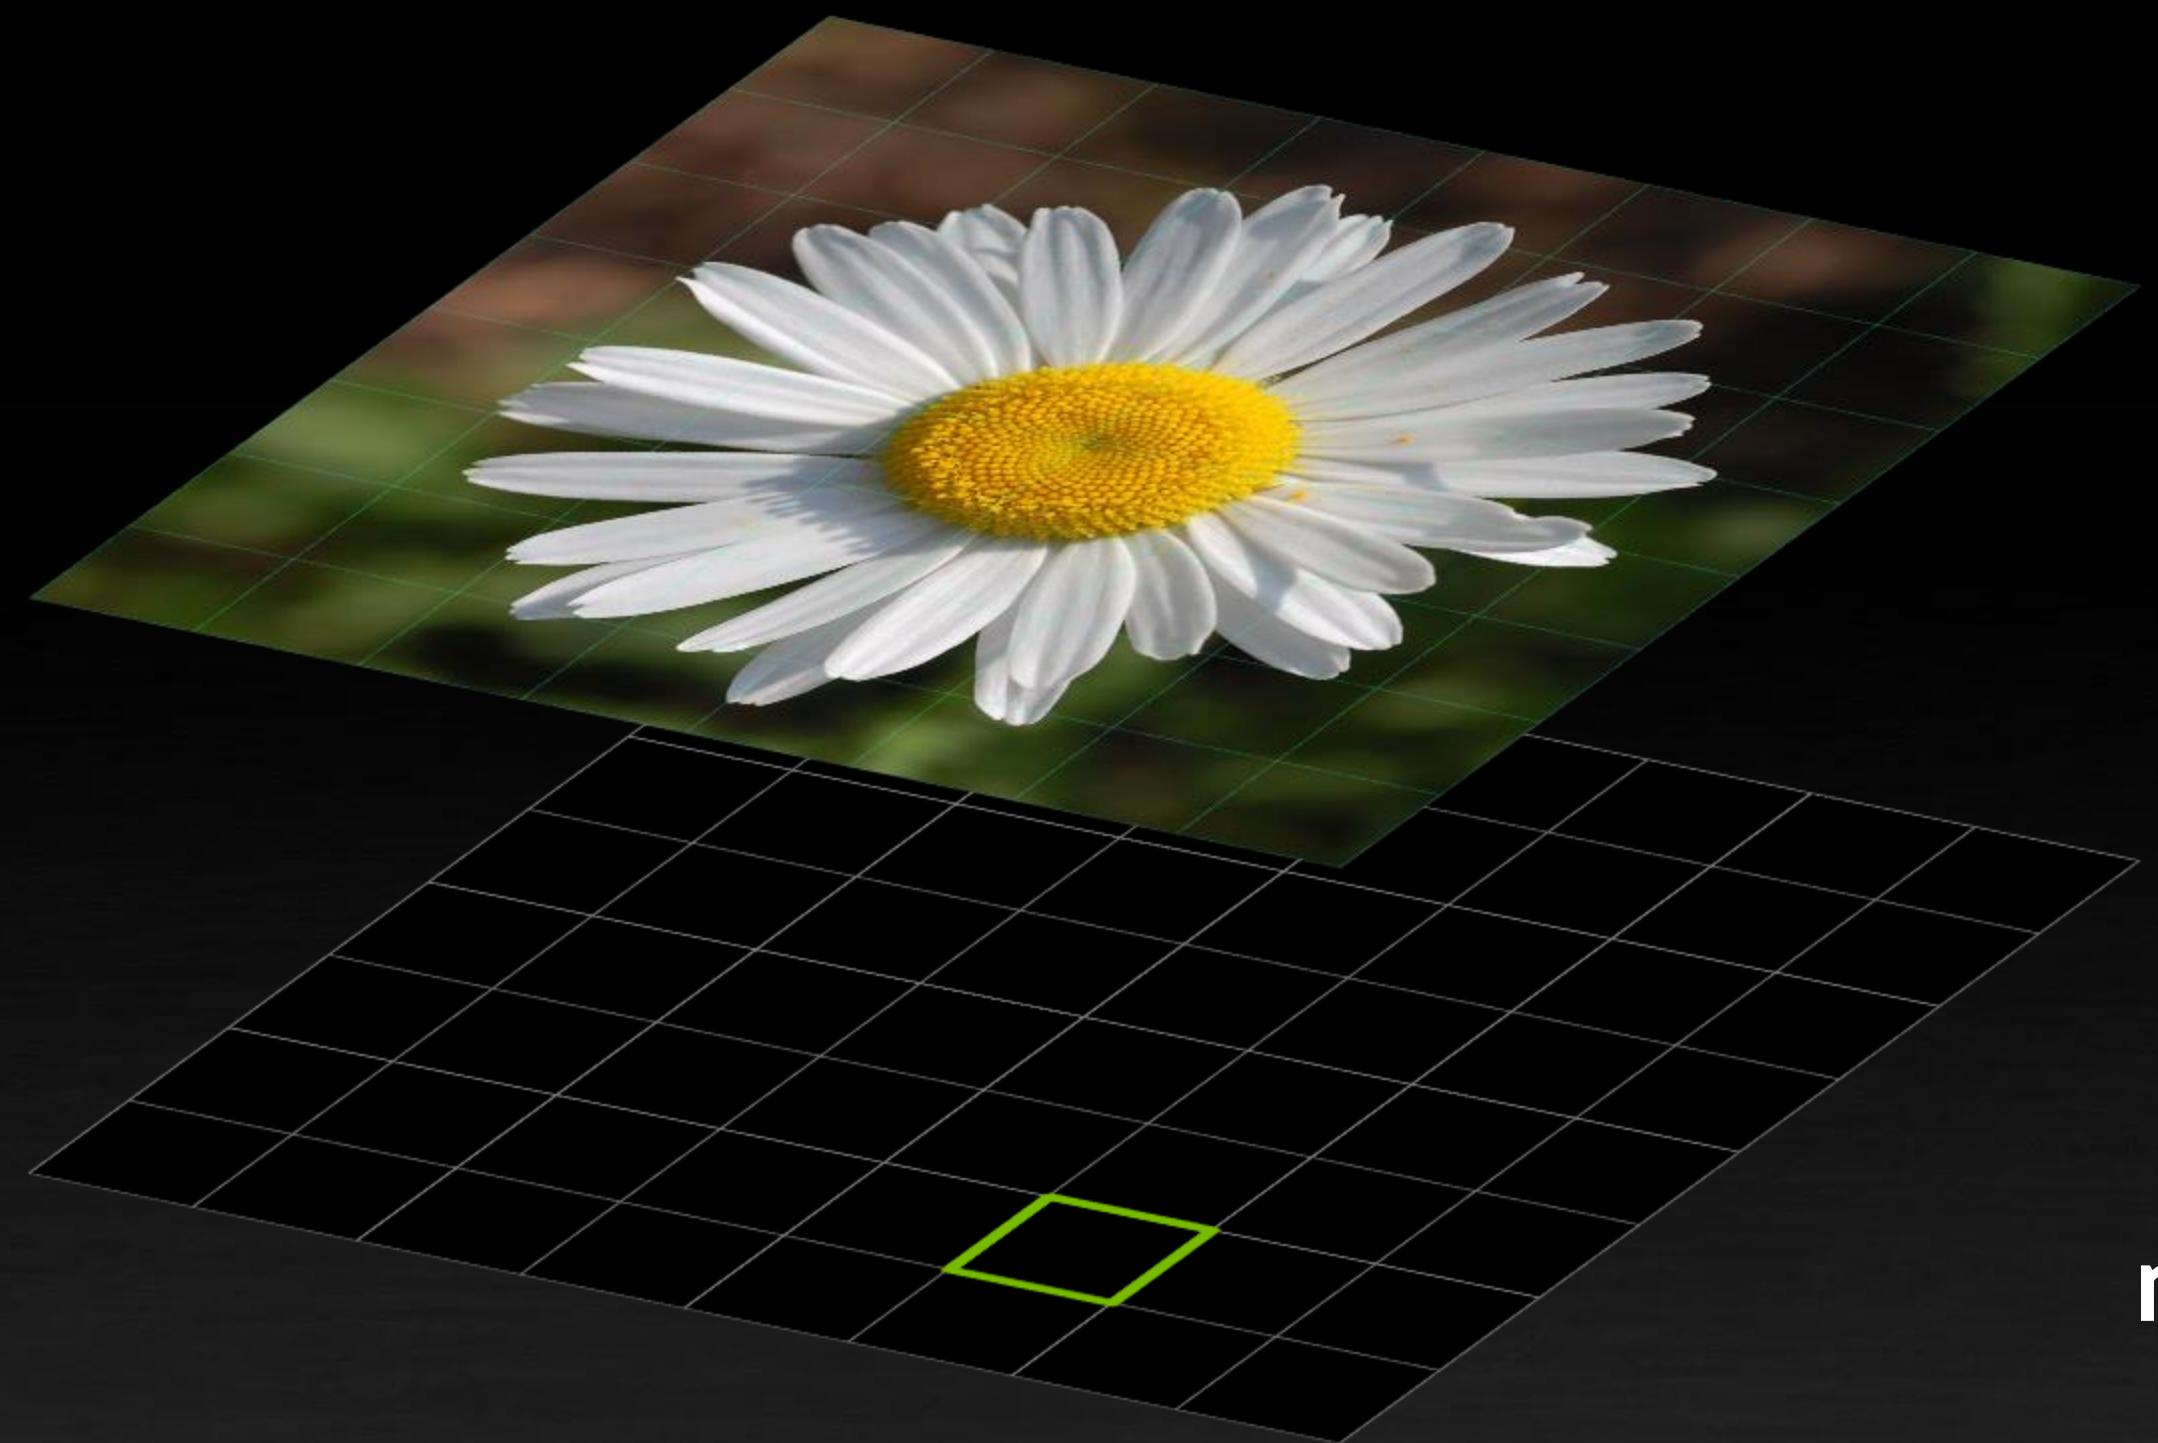
THE CUDA PROGRAMMING MODEL: GRID → BLOCKS → THREADS



Grid
of work

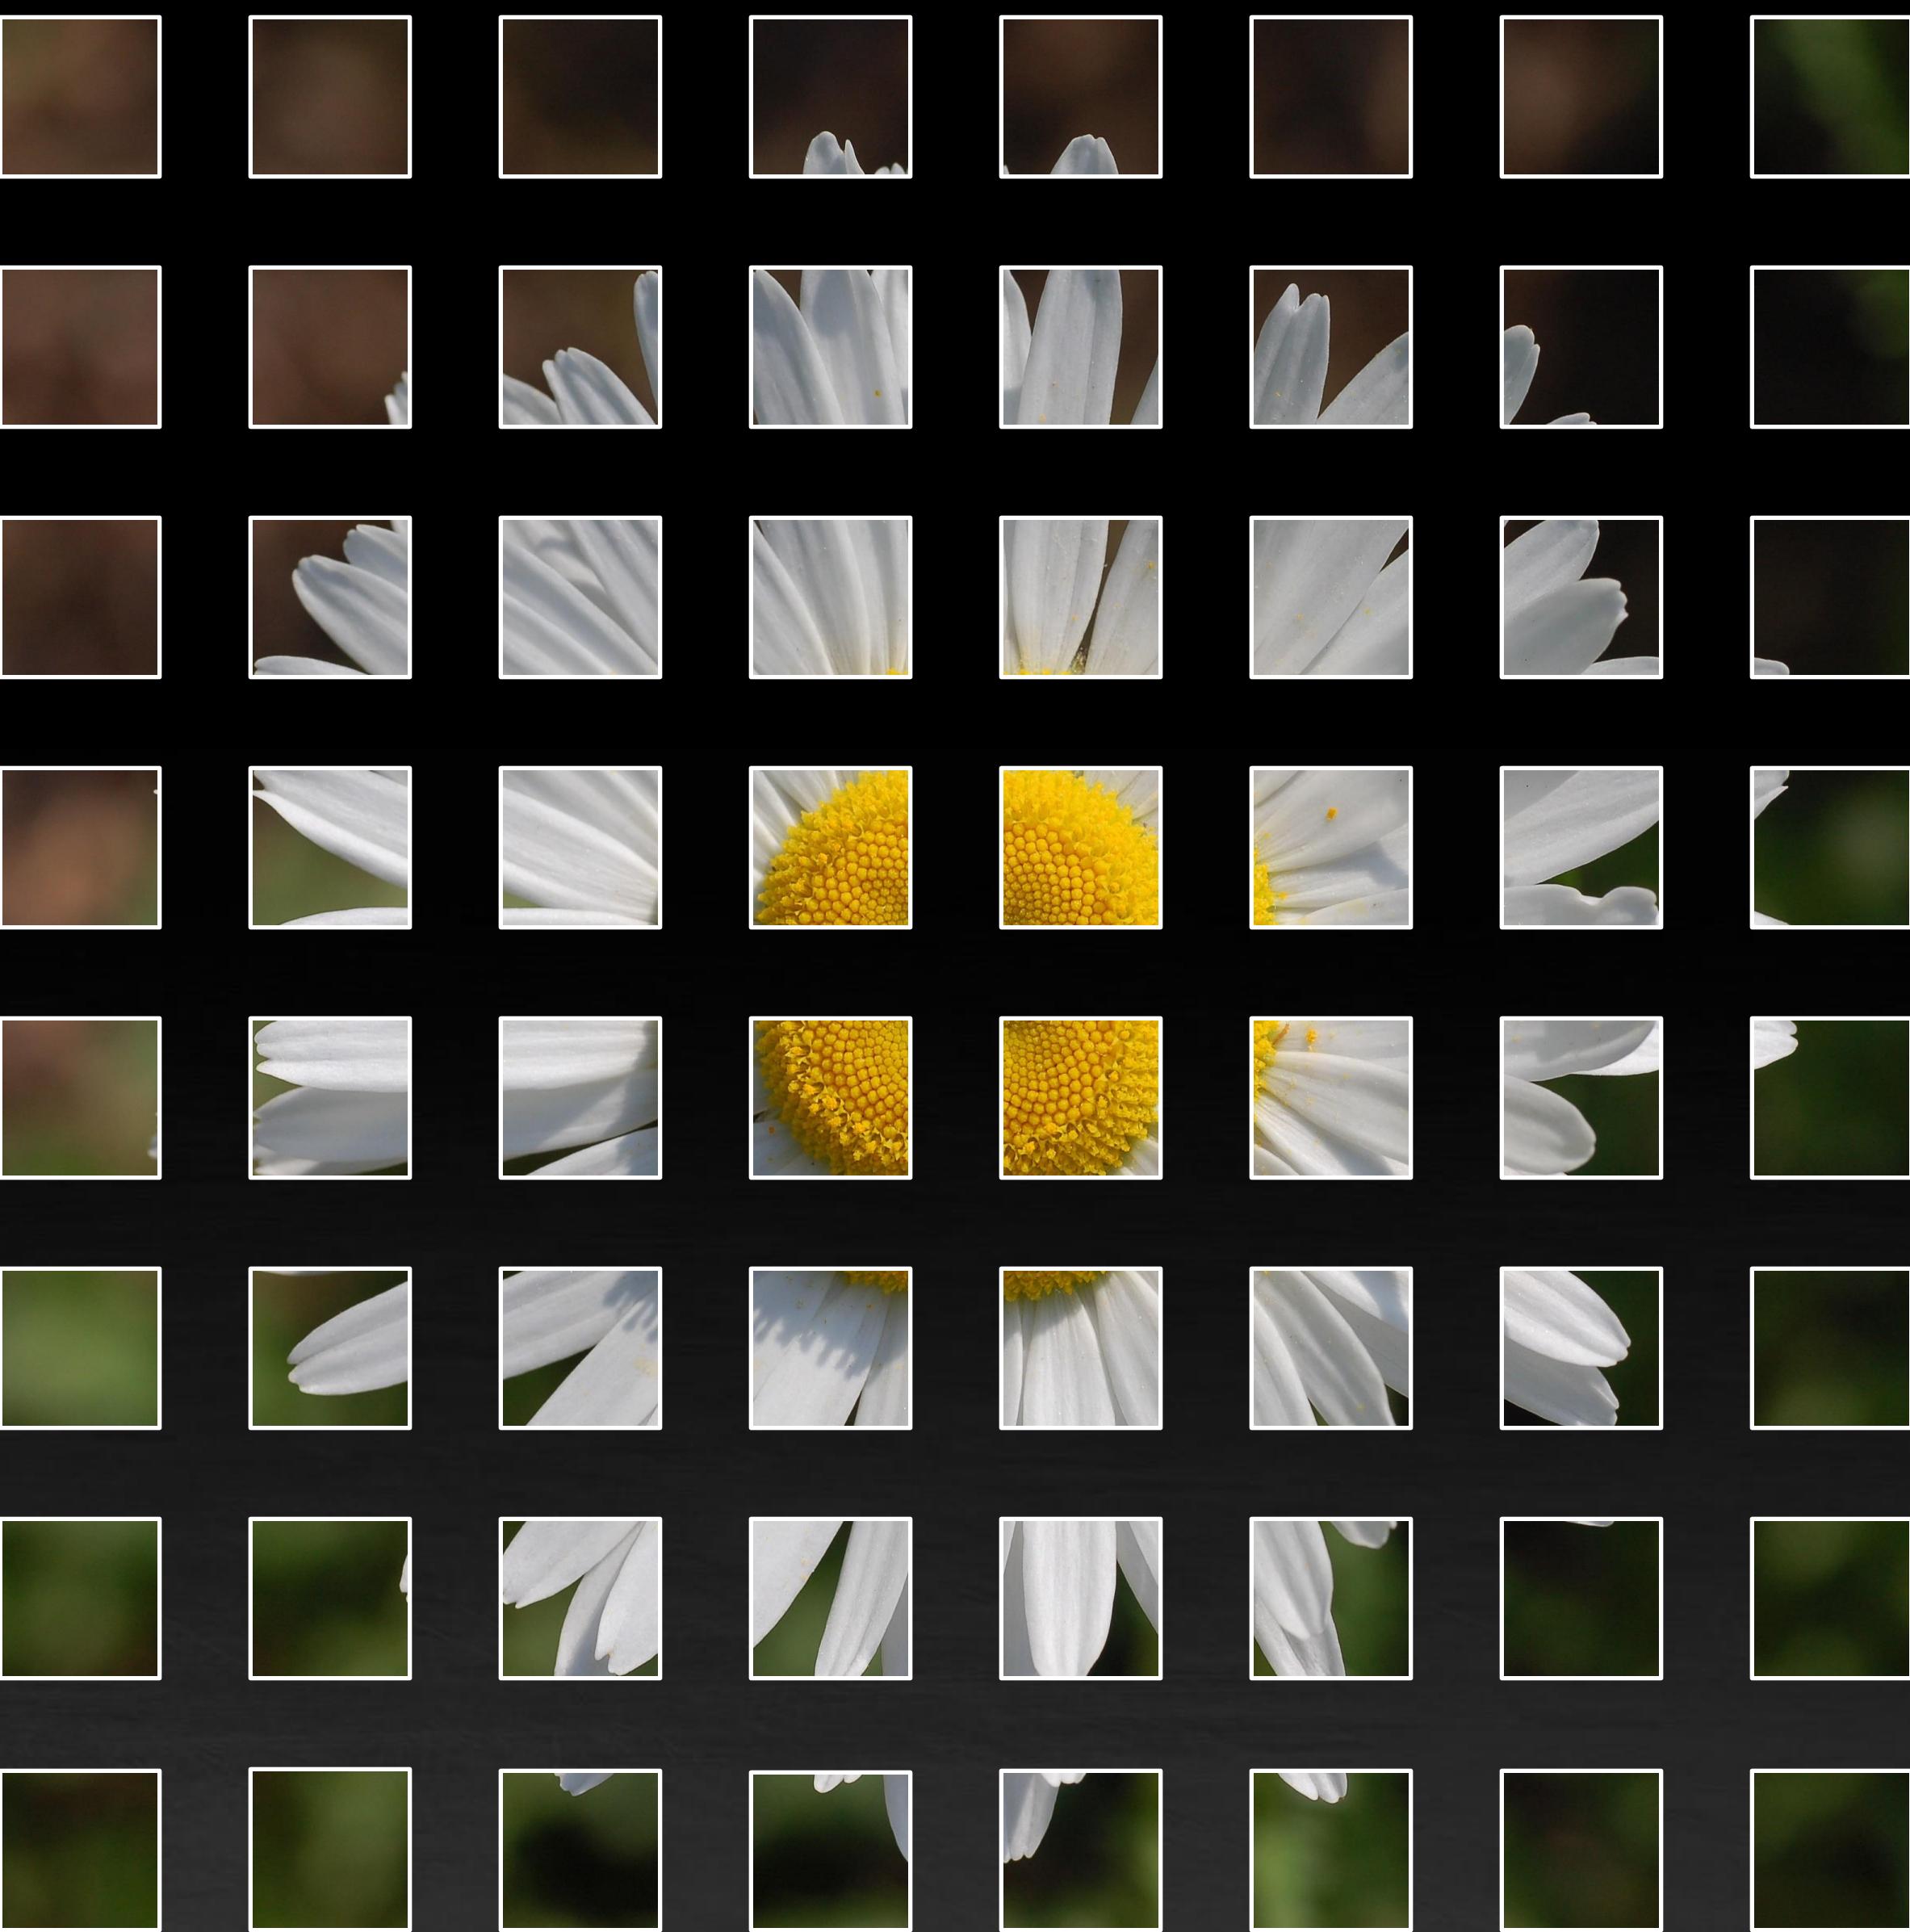


DIVIDE THE WORK INTO A GRID OF EQUAL BLOCKS

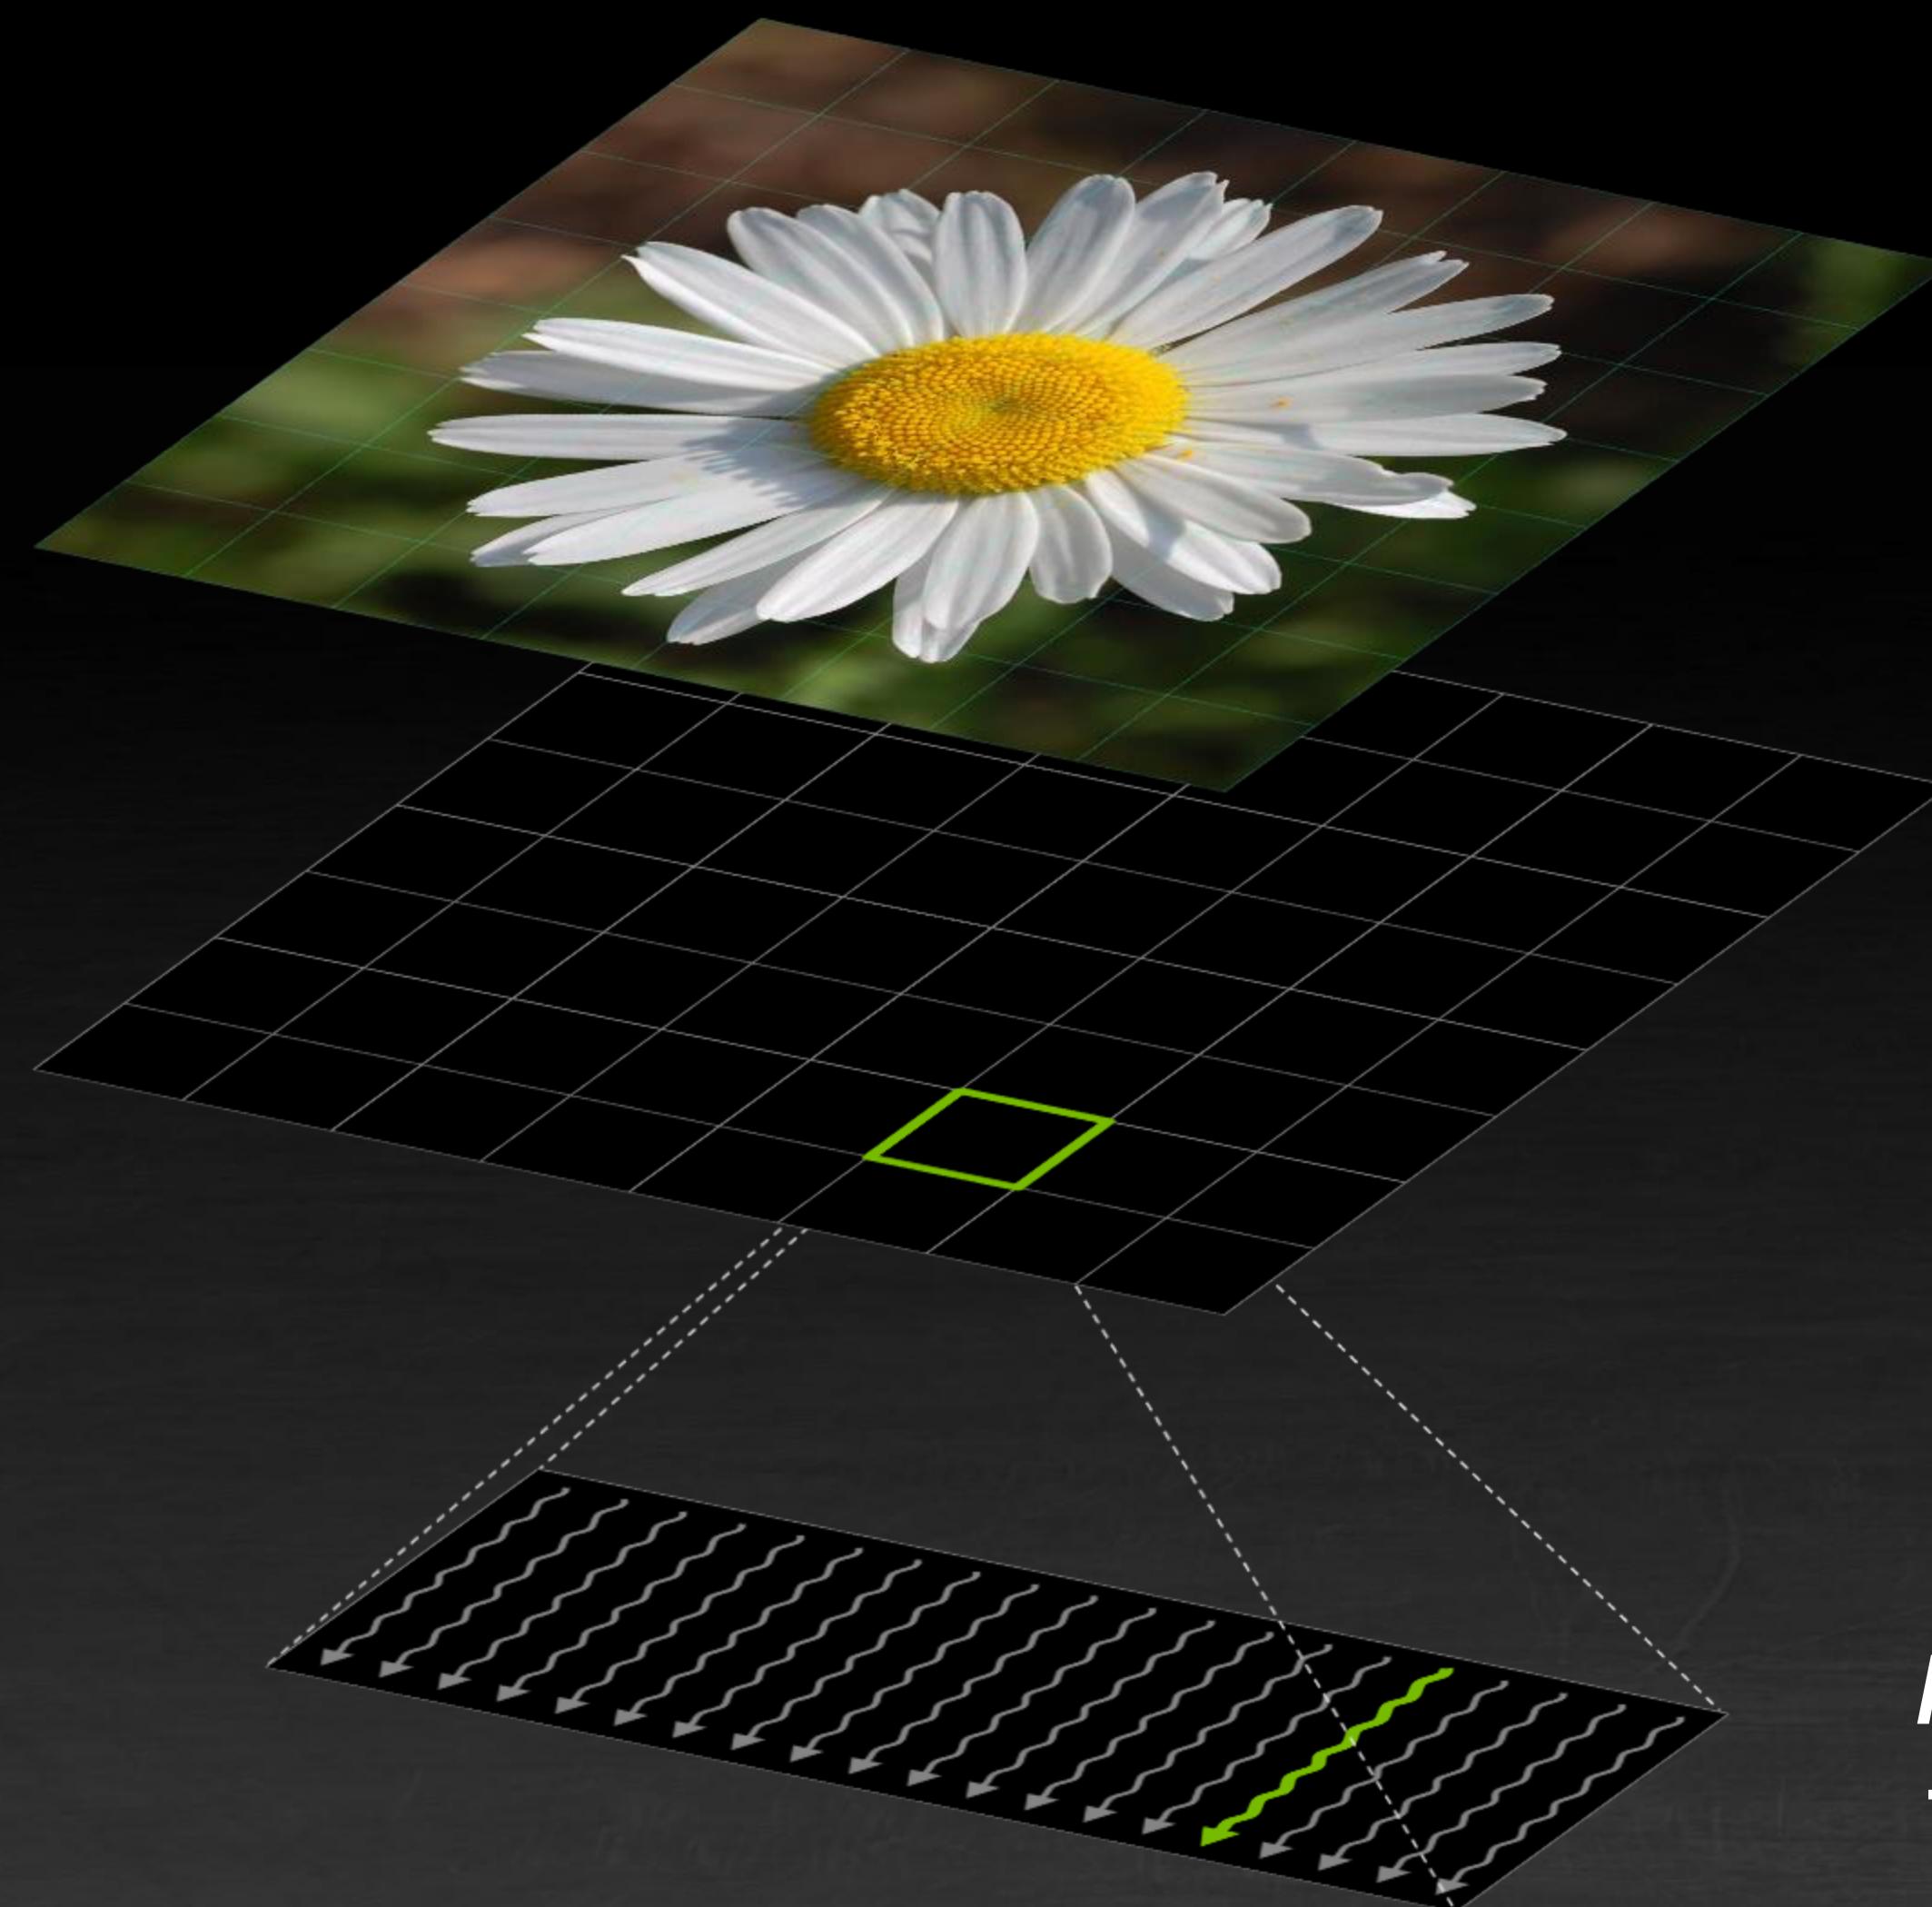


Grid
of work

Divide into
many Blocks



EACH BLOCK RUNS AS IF IT'S AN INDEPENDENT PROGRAM



Grid
of work

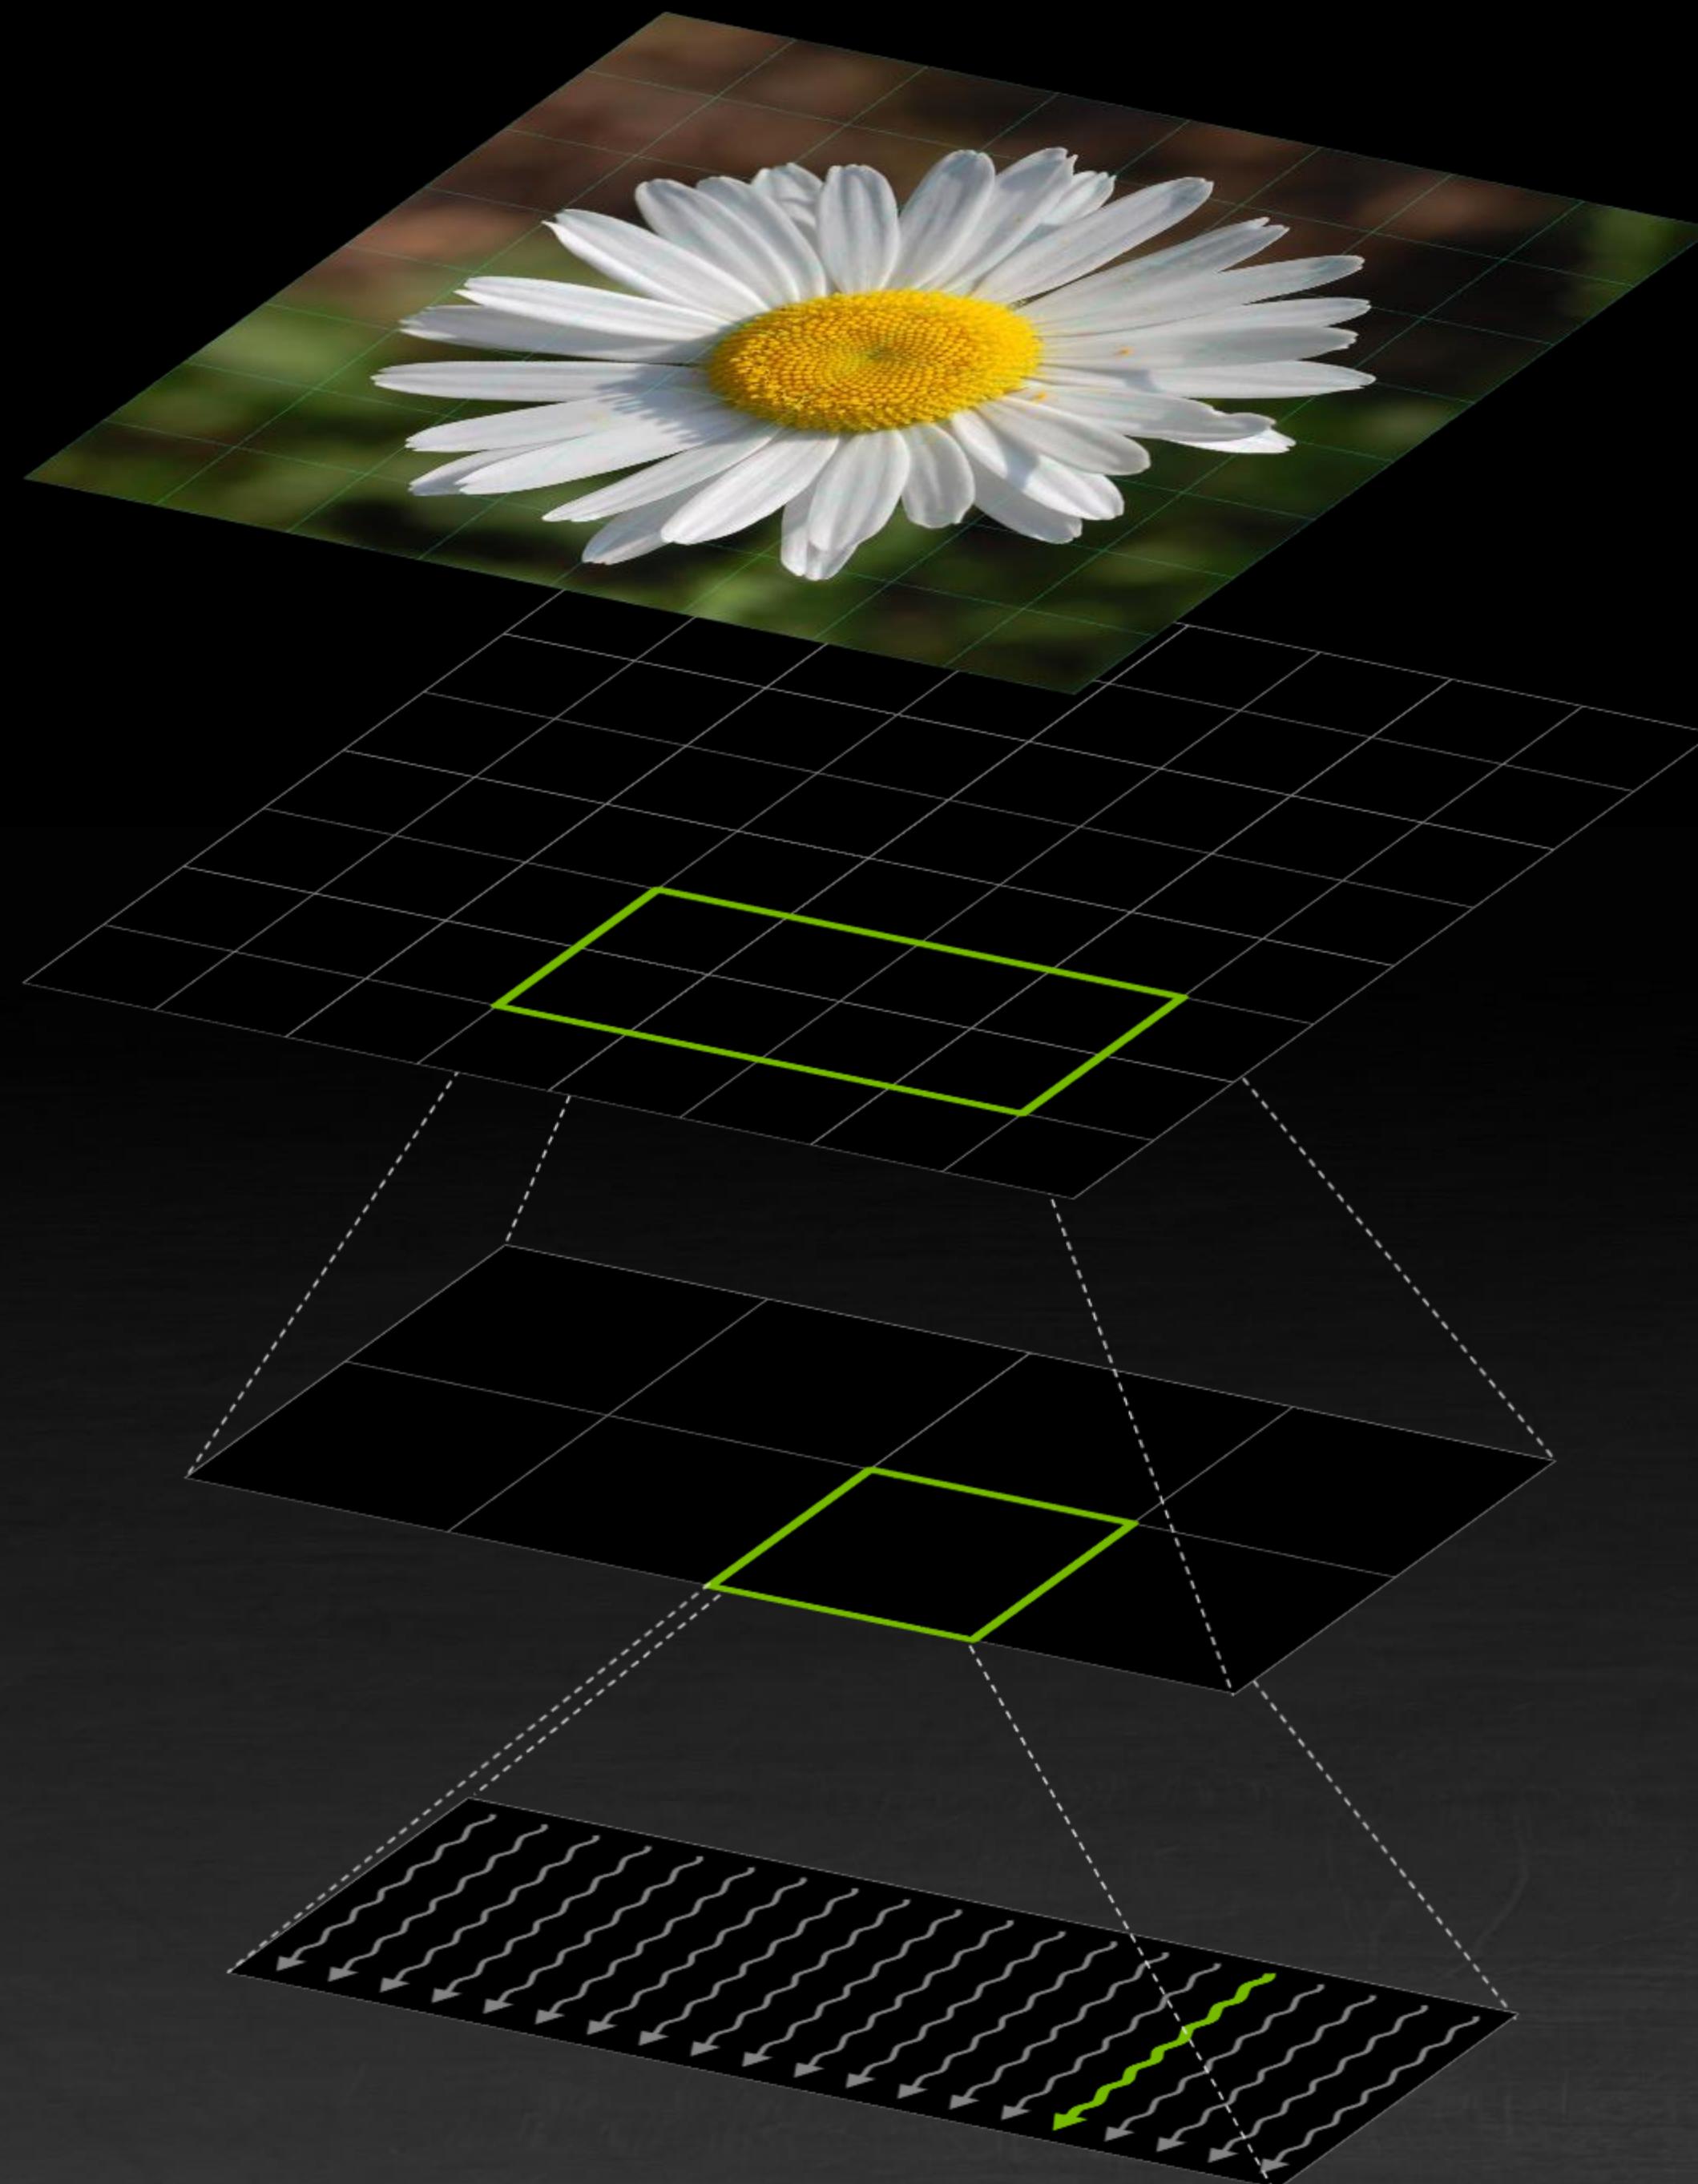
Blocks
of Threads

Many Threads
in each Block



THREAD BLOCK CLUSTER

A collective of blocks, co-scheduled on adjacent multiprocessors

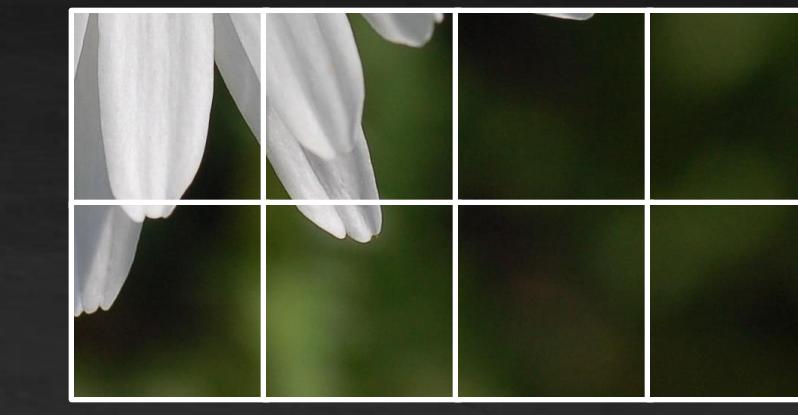
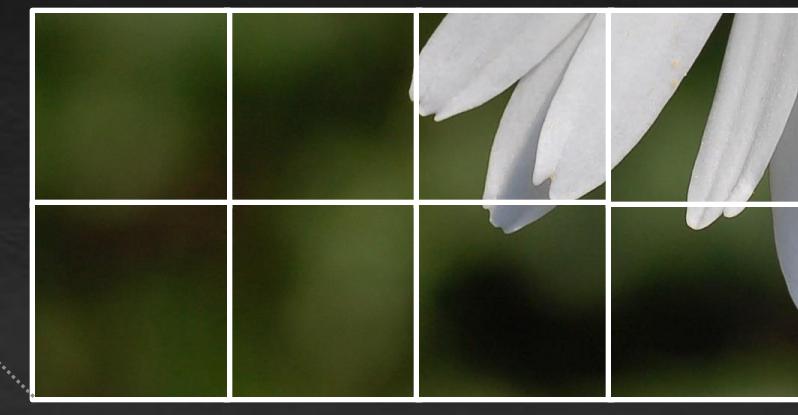
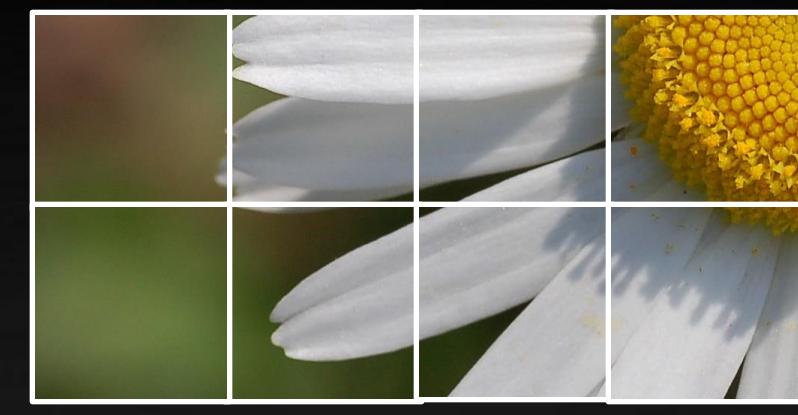
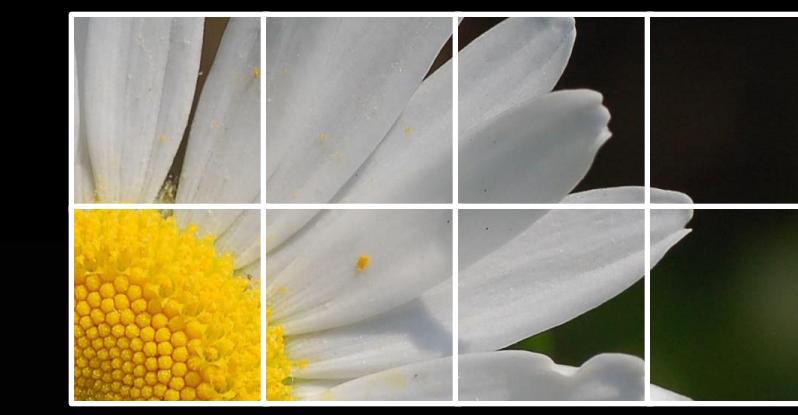
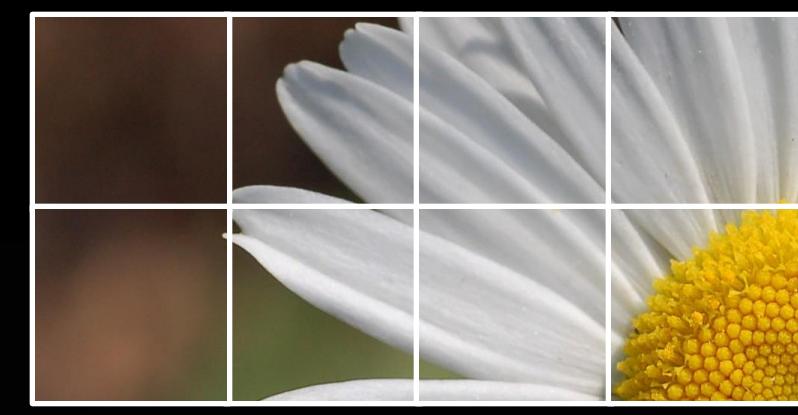
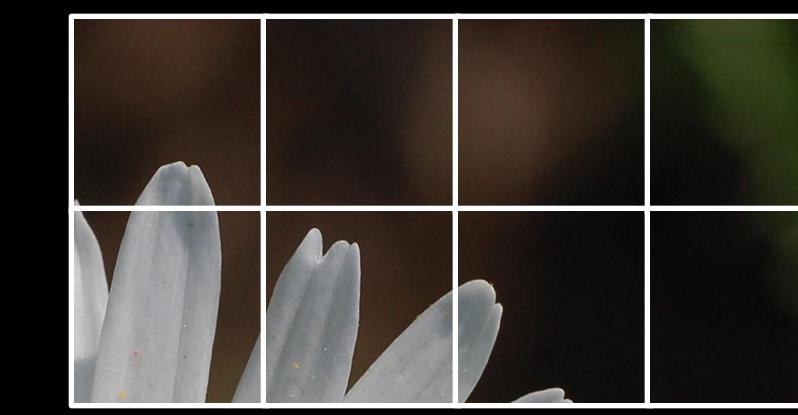
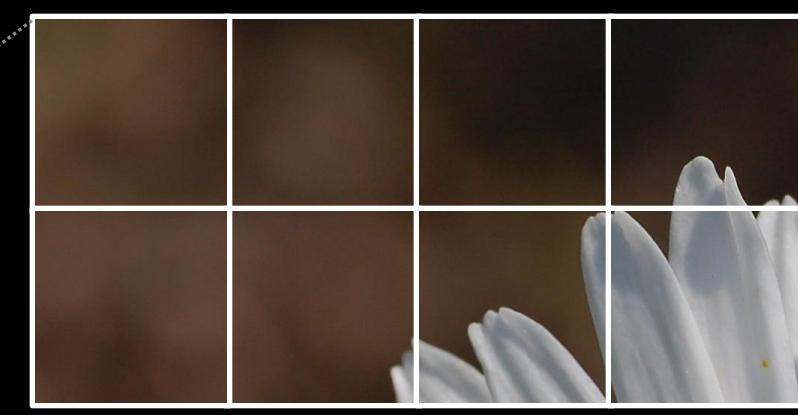


Grid
of work

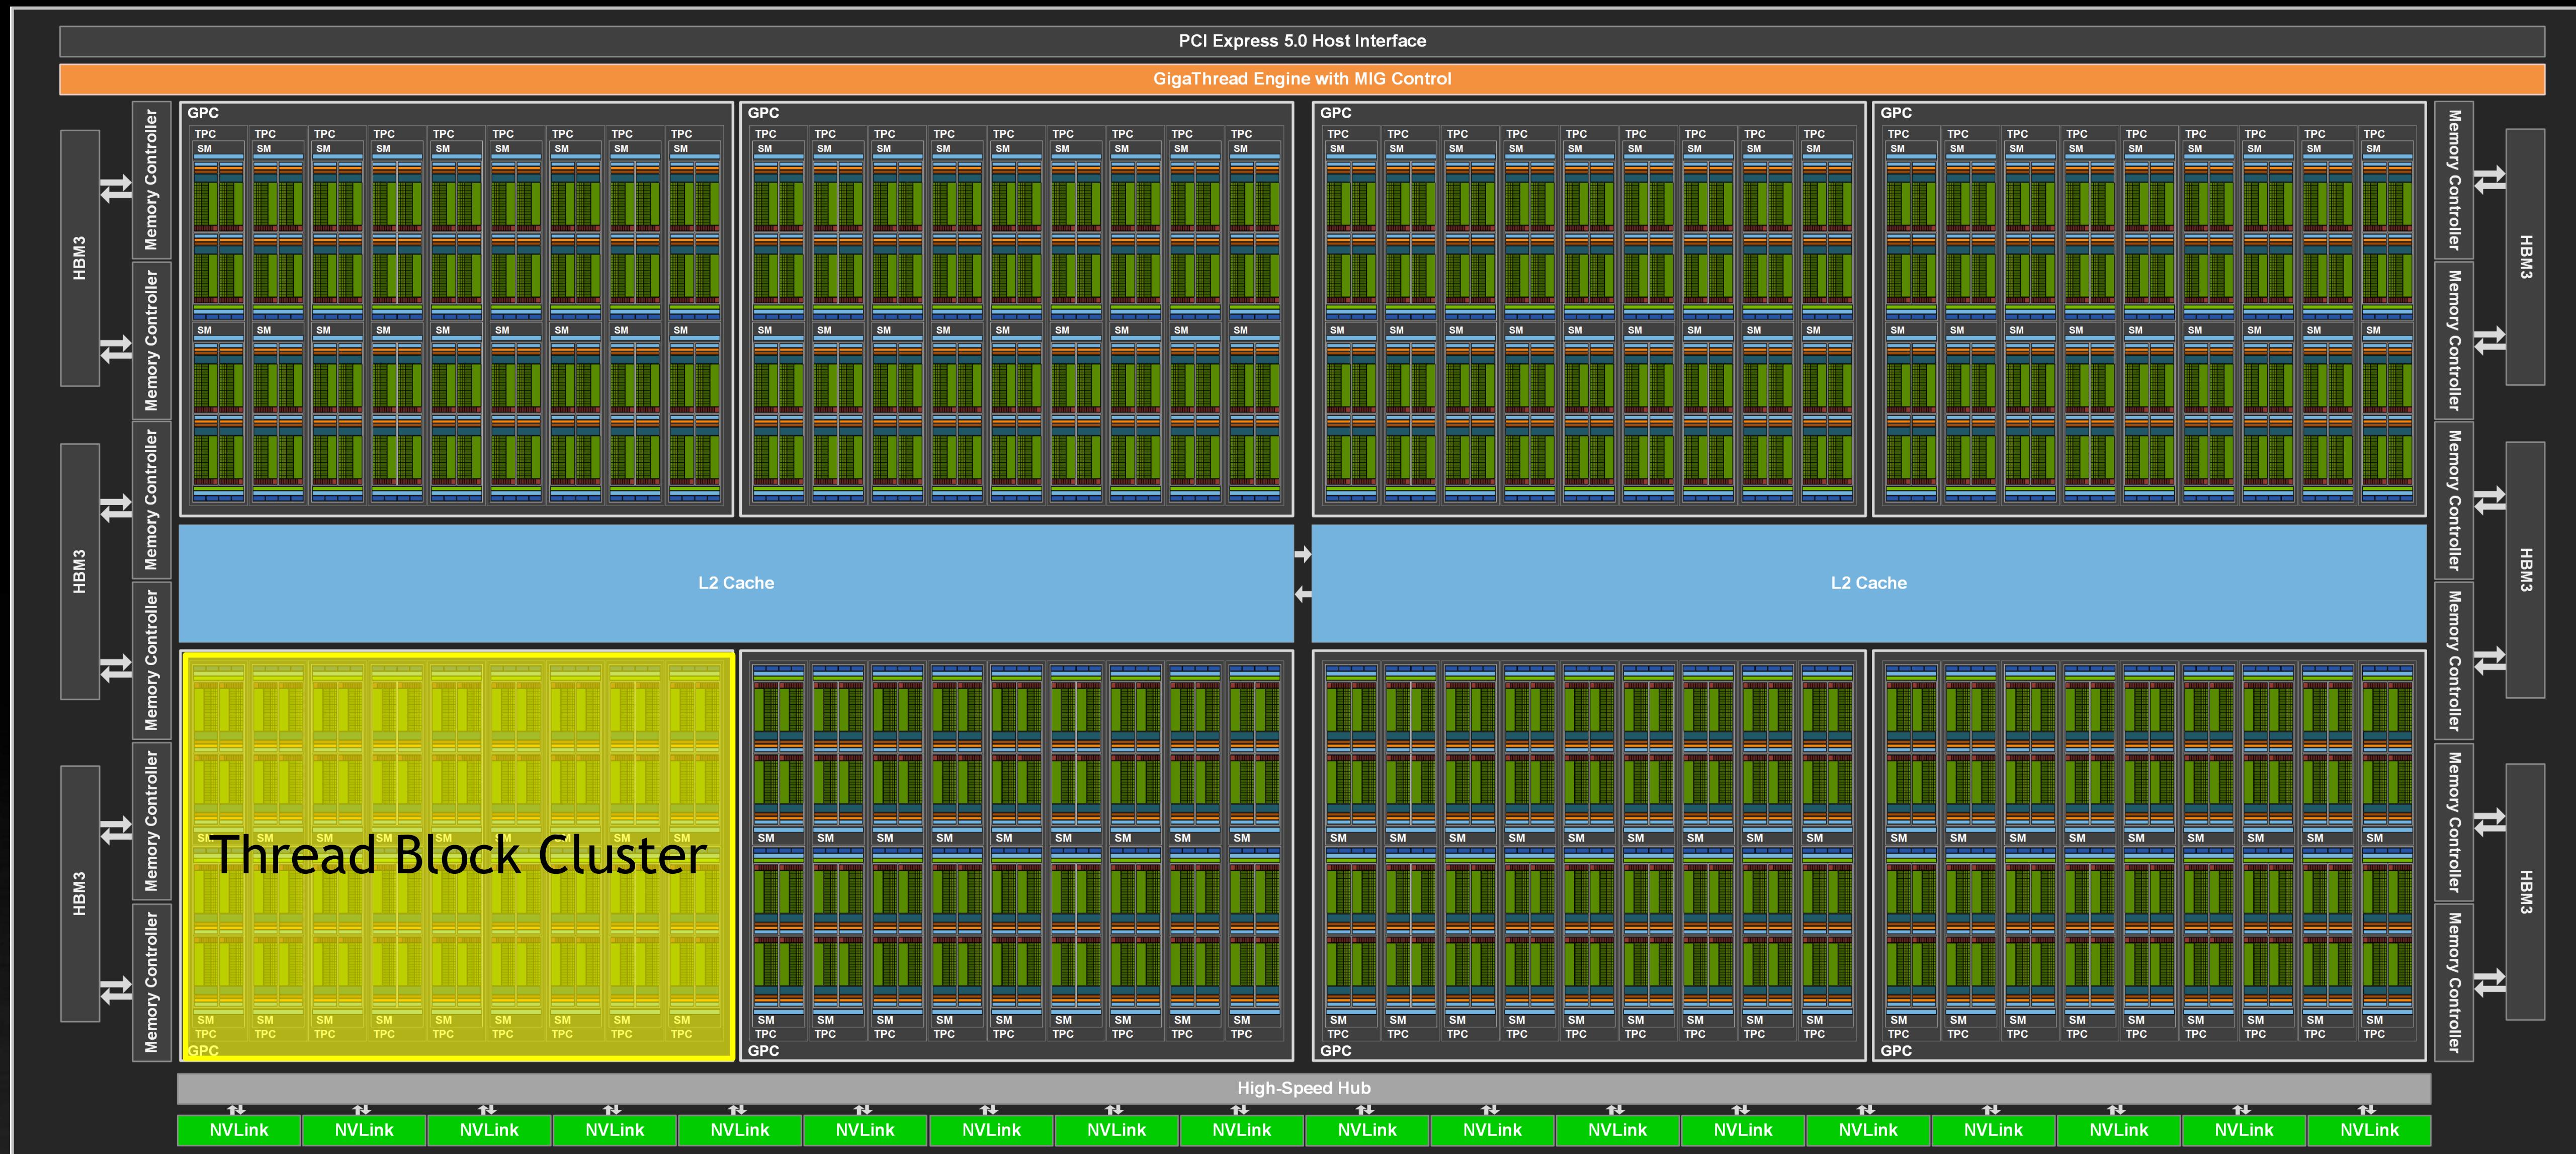
Cluster
of Blocks

Blocks
of Threads

Threads



TAKING ADVANTAGE OF LOCALITY AT A GPU SCALE



Guaranteed co-located blocks
New tier of guaranteed concurrency
Fast data exchange & sync

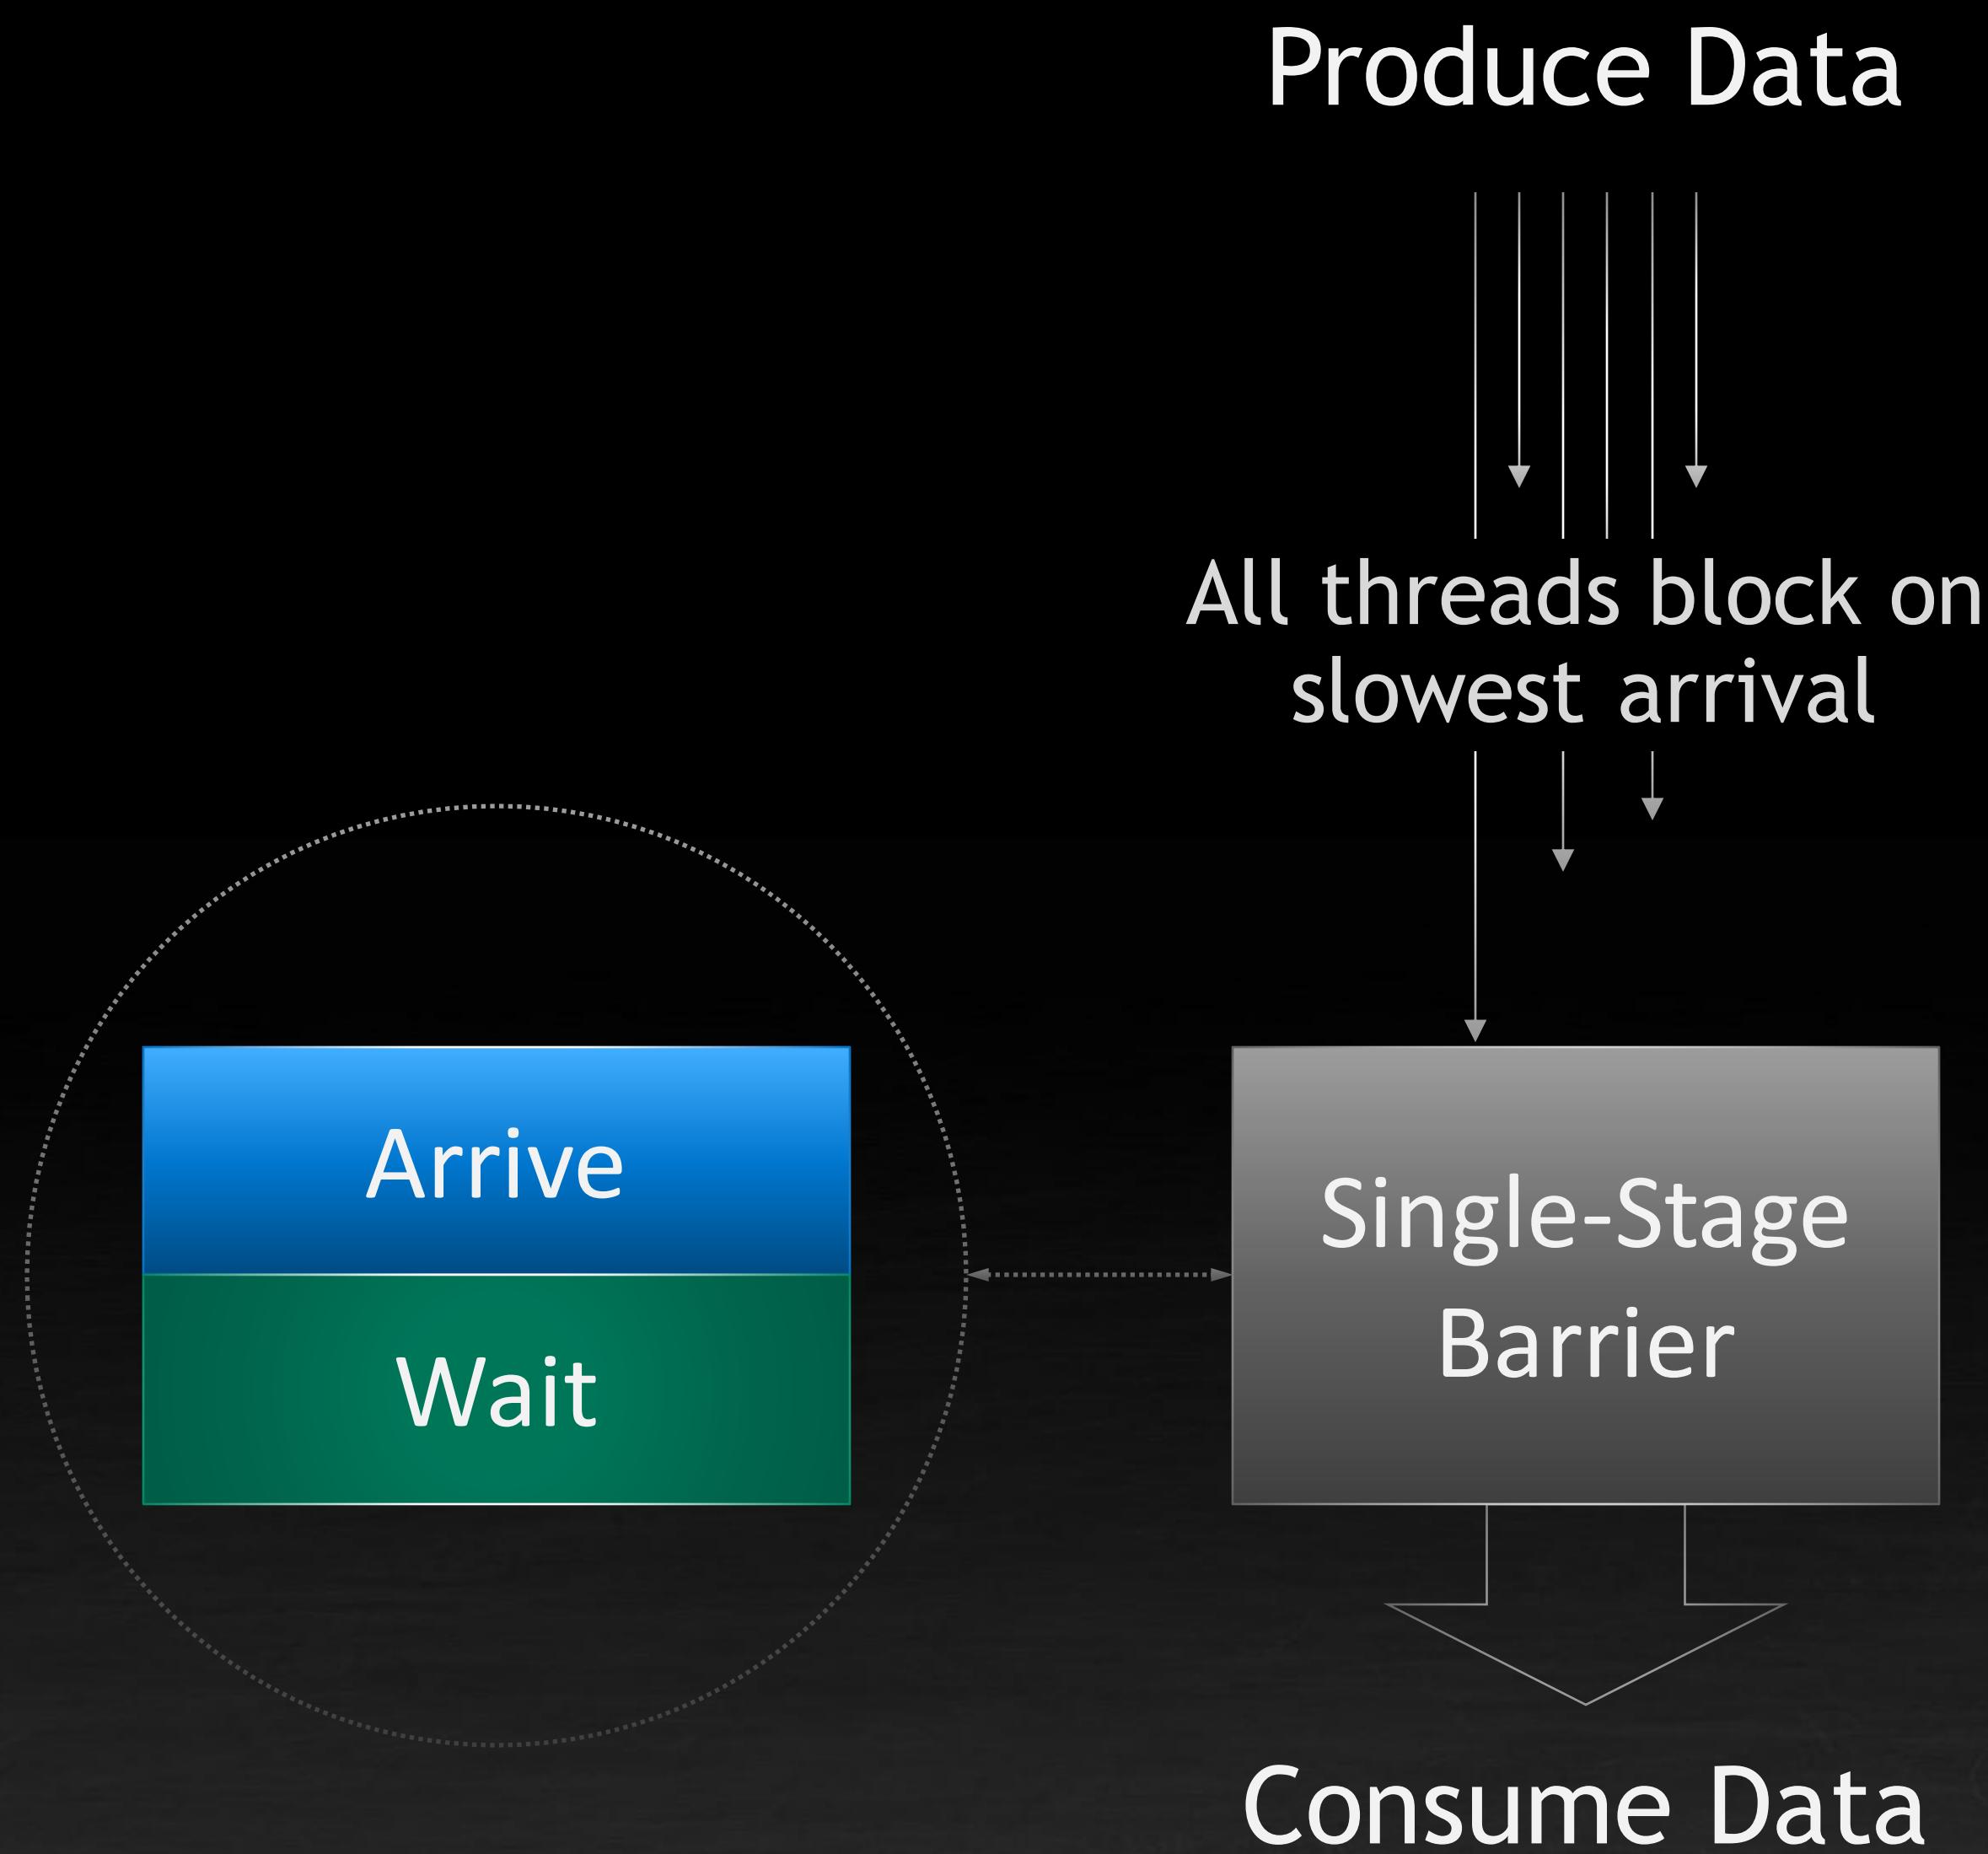
CLUSTER DISTRIBUTED SHARED MEMORY (DSMEM)

Blocks within a cluster are able to access each others' shared memory directly

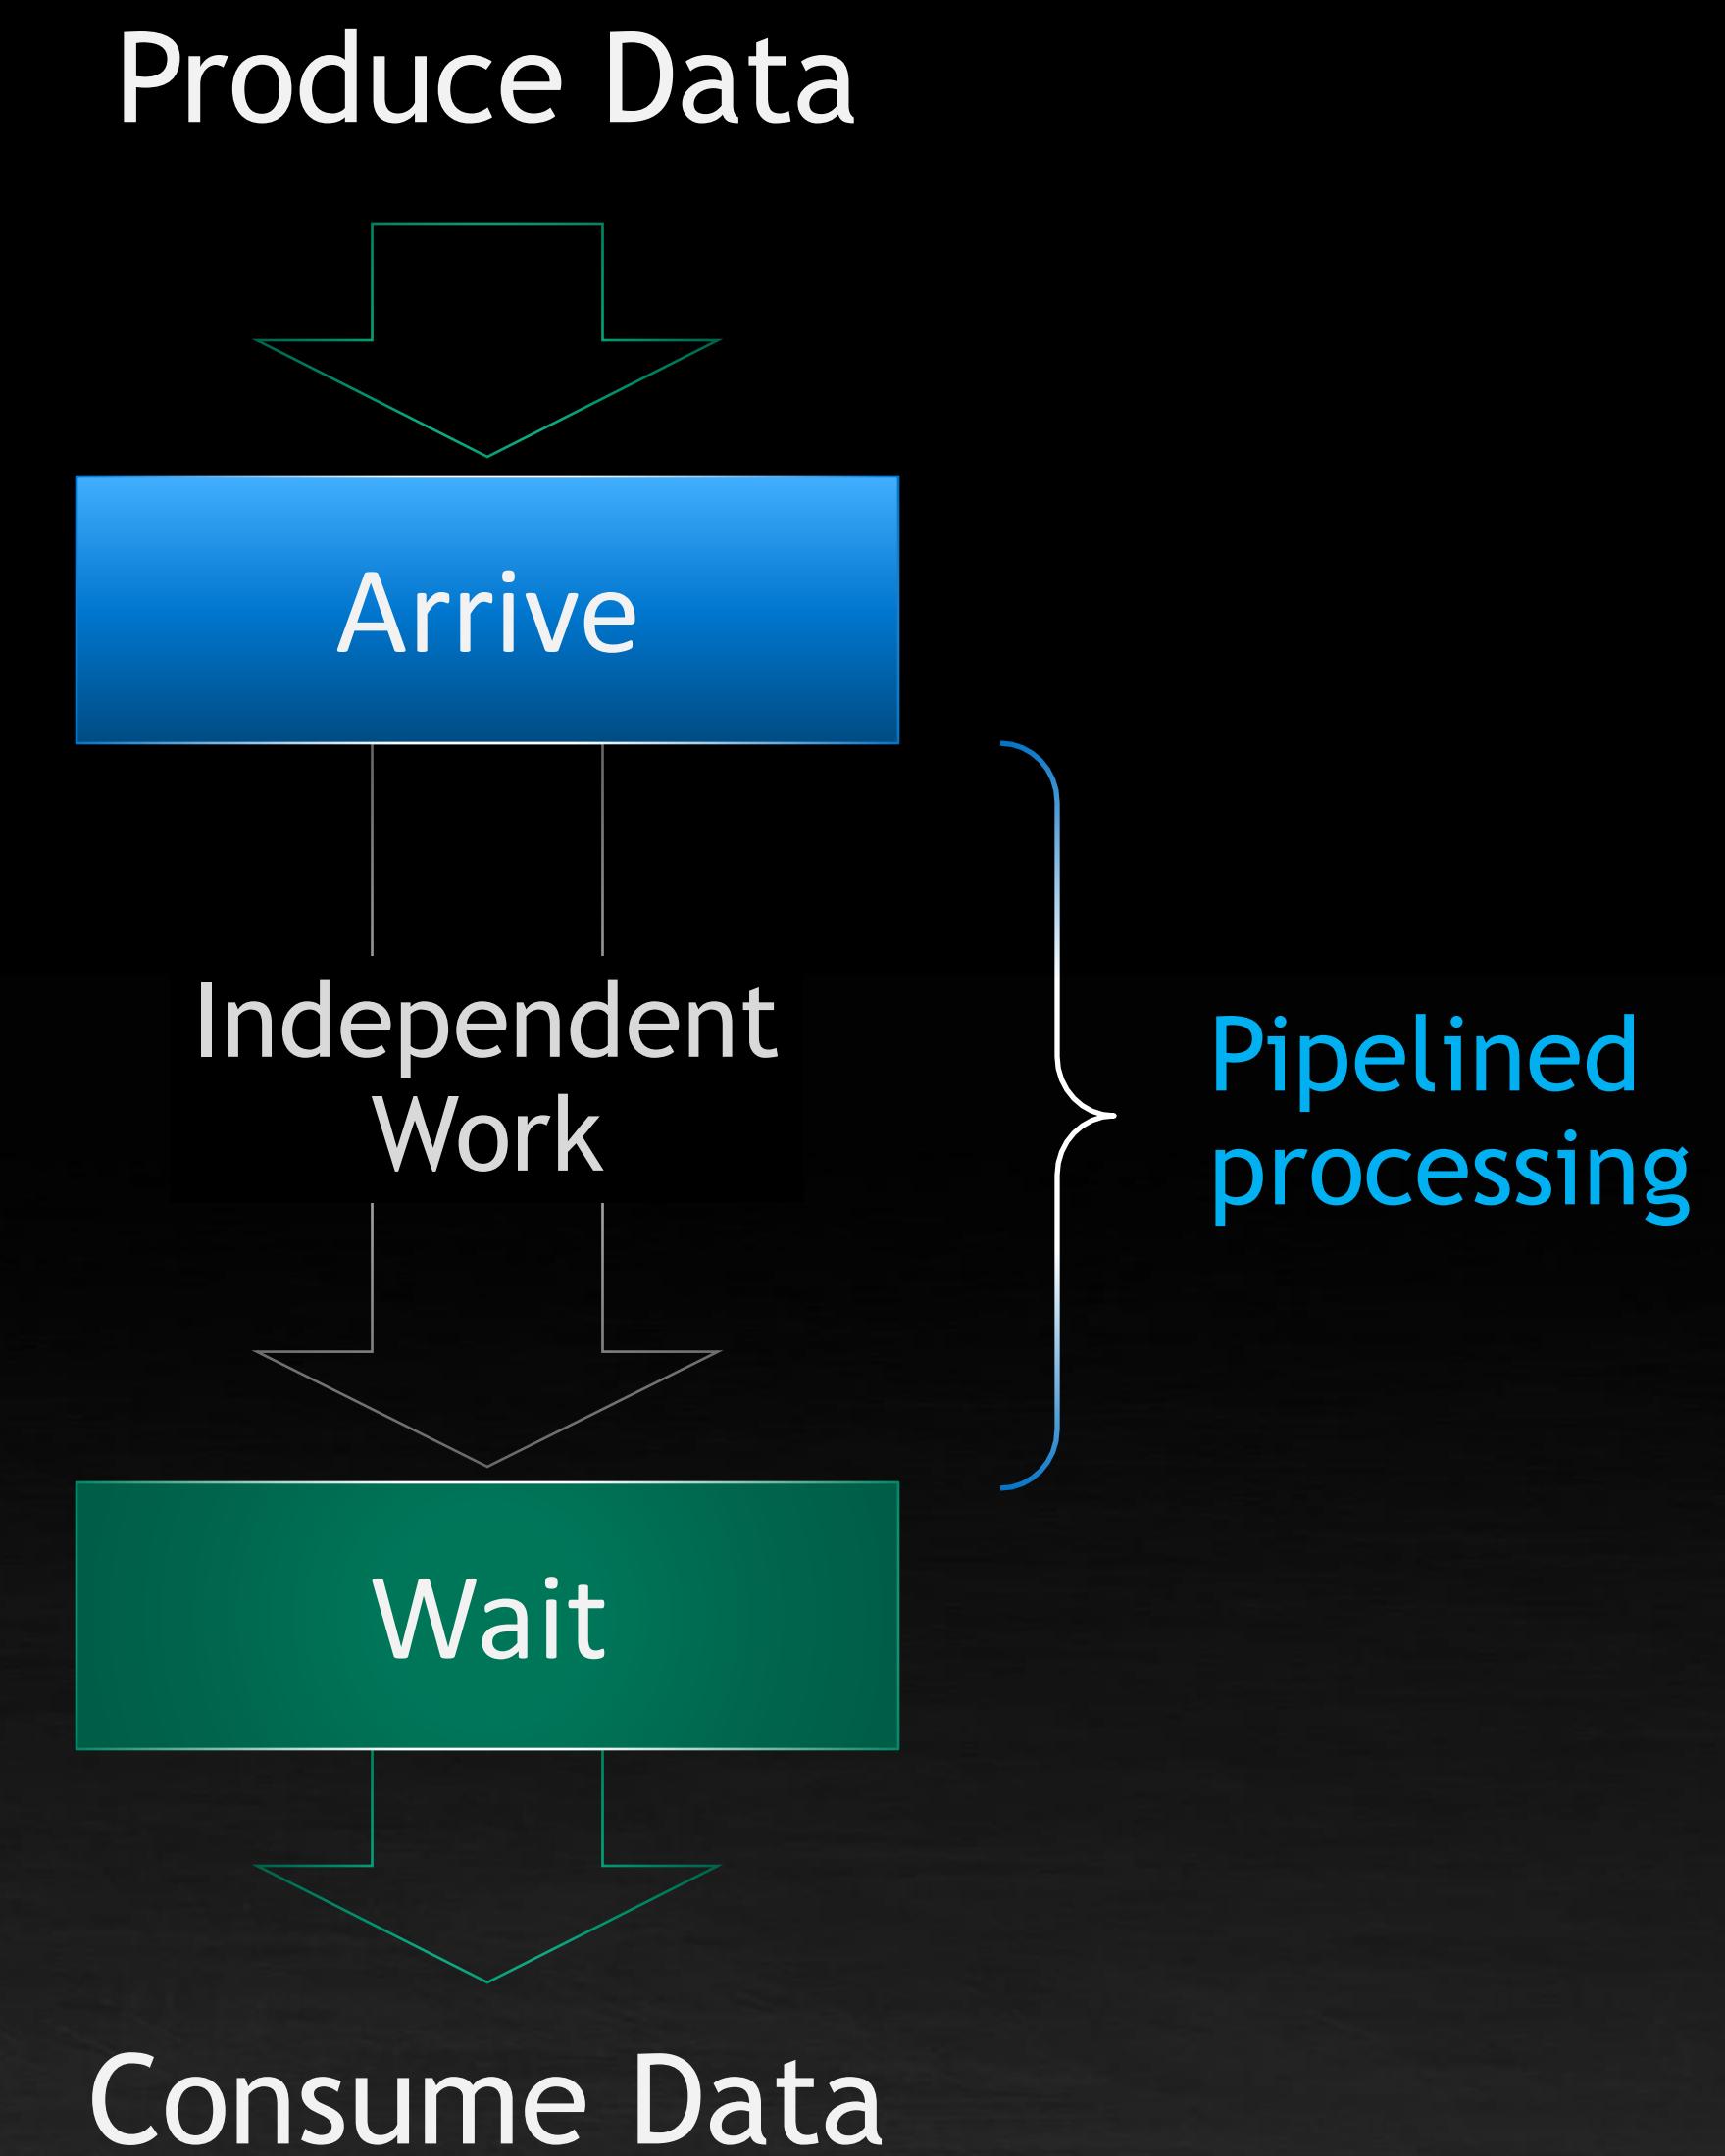


Full load/store/atomic access to all shared memory
between blocks within a cluster

ASYNCHRONOUS BARRIERS



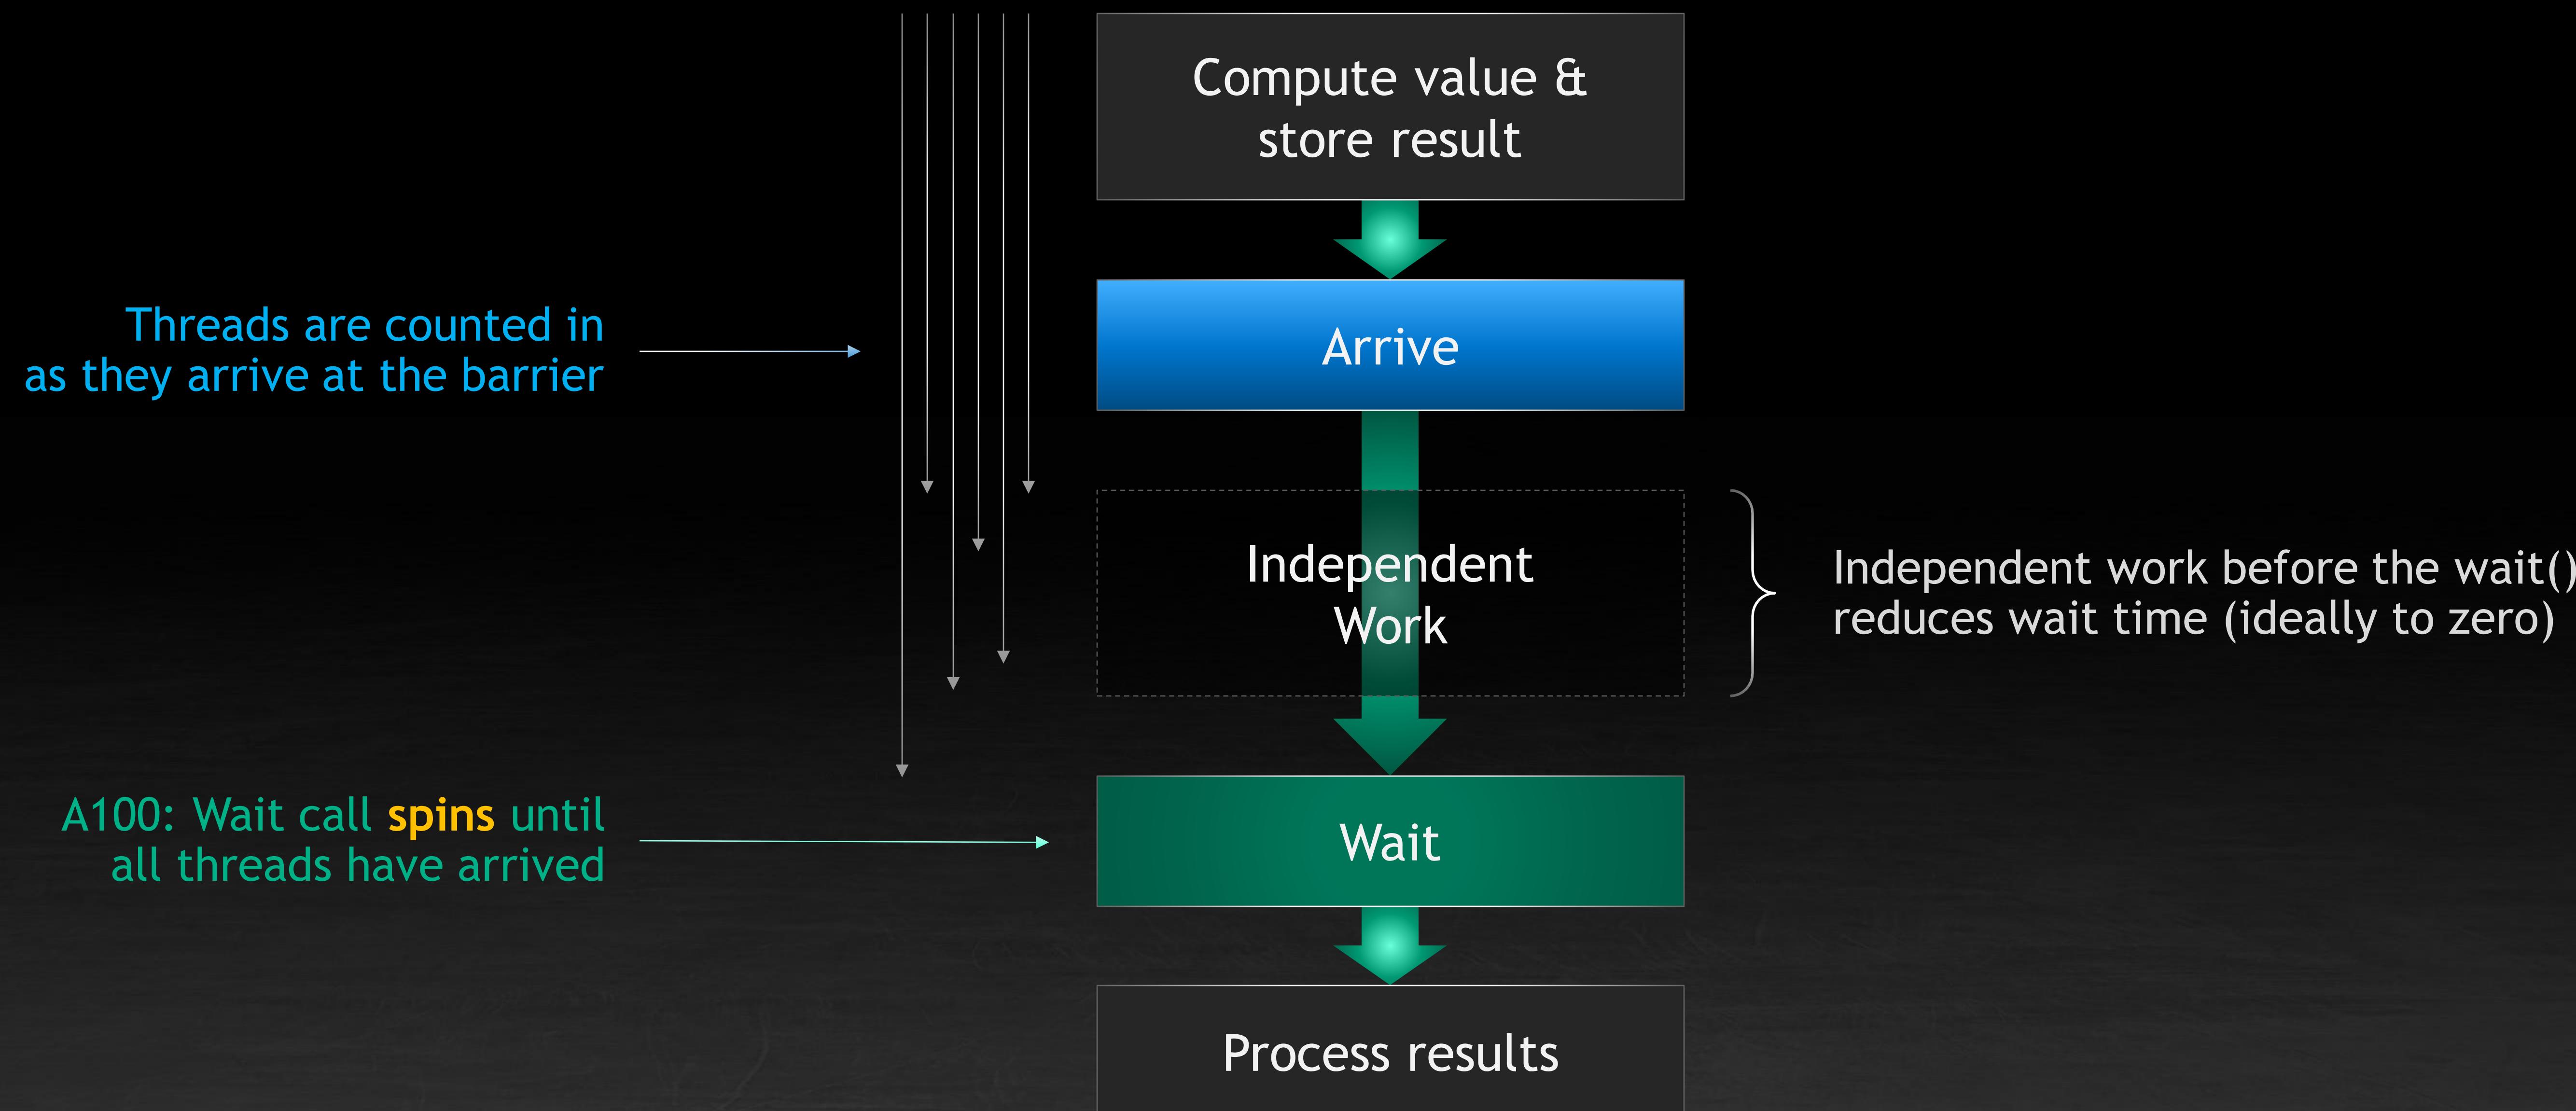
Single-Stage barriers combine back-to-back arrive & wait



Asynchronous barriers enable pipelined processing

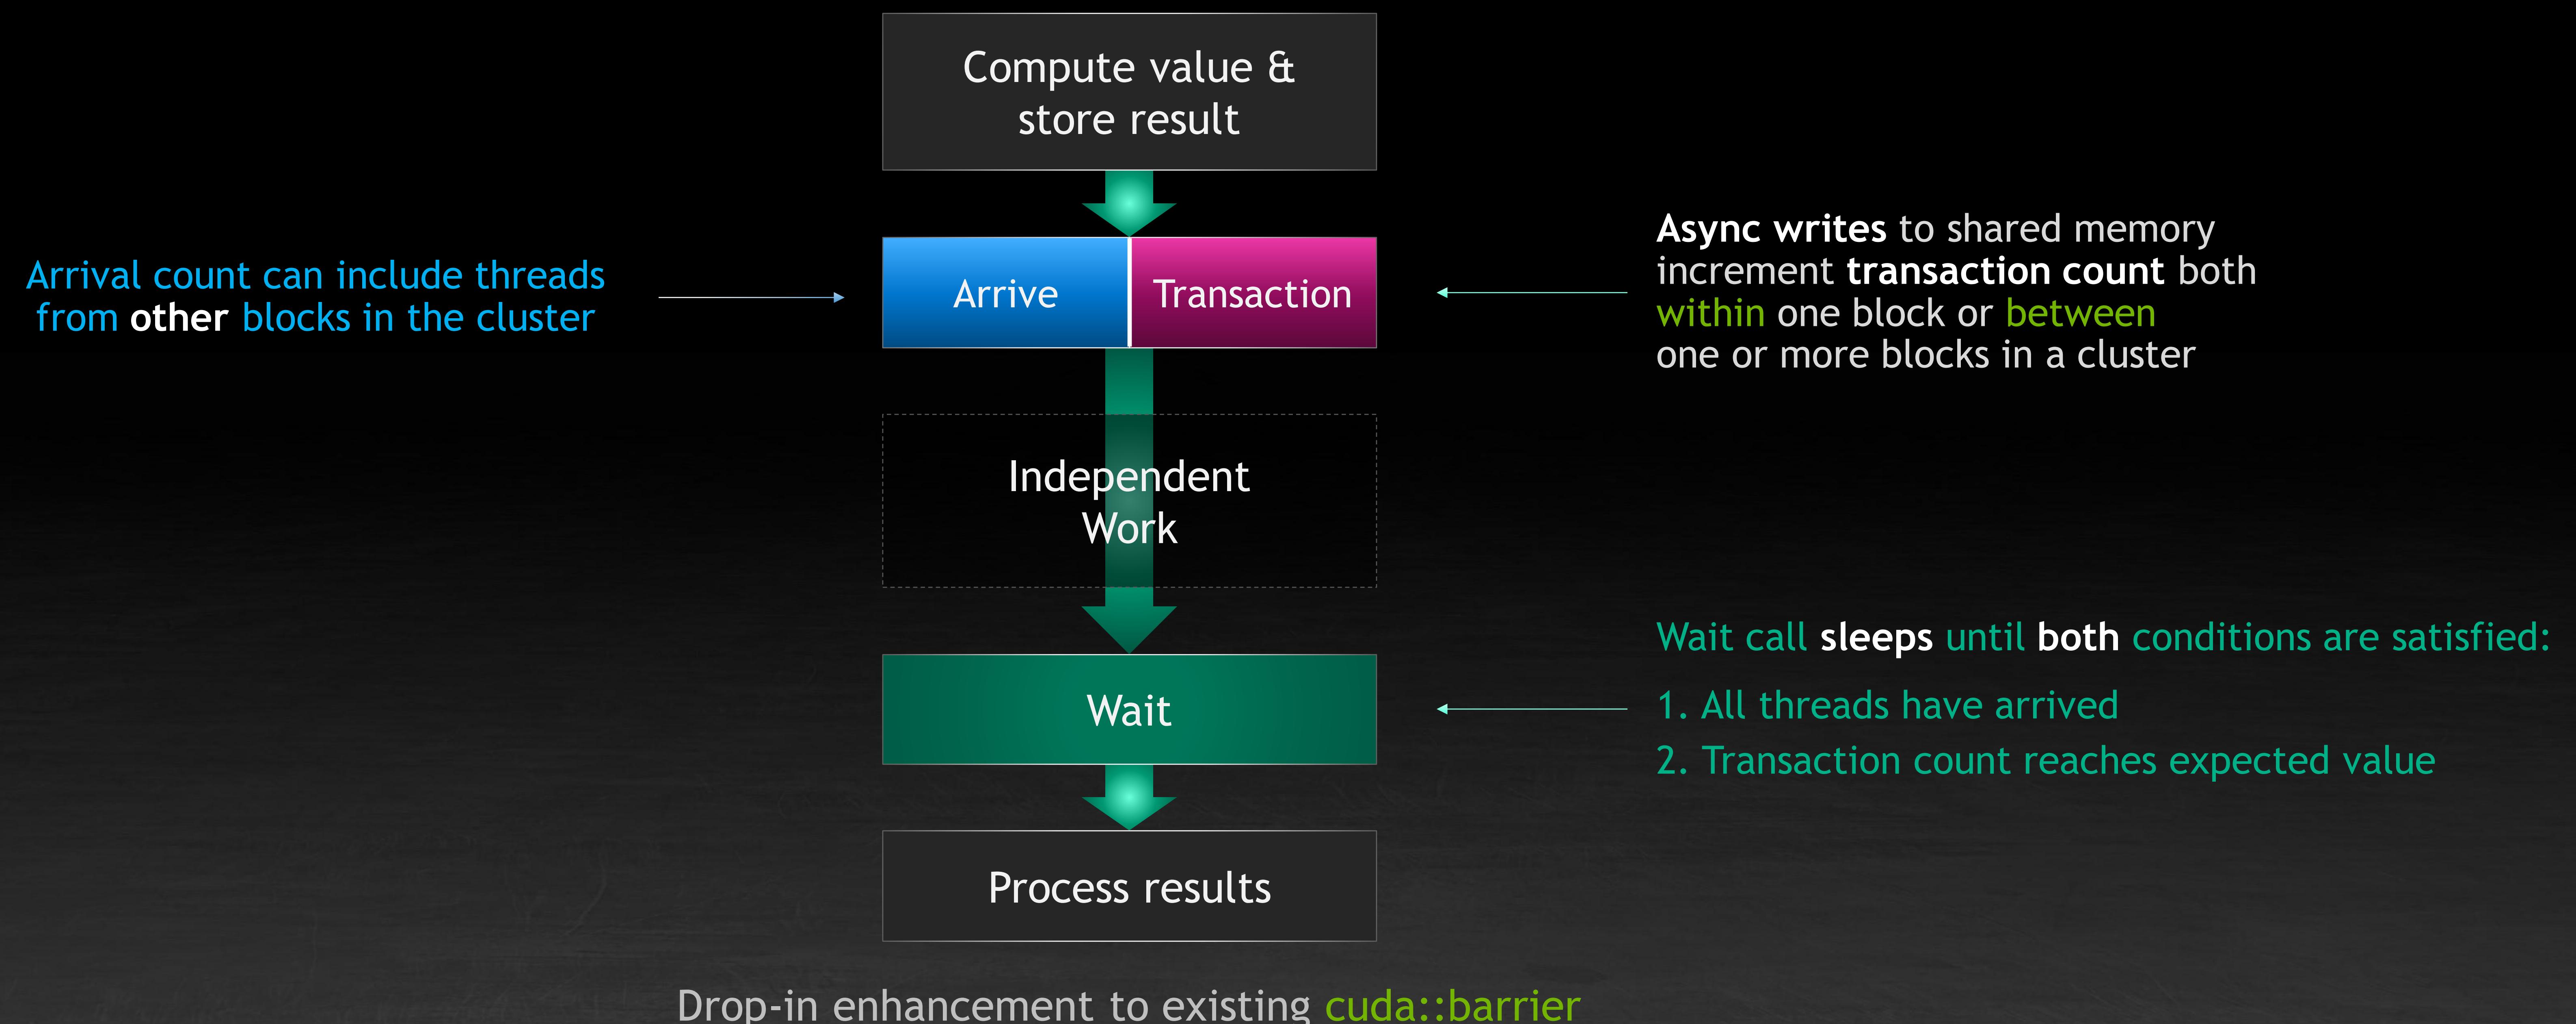
ASYNCHRONOUS BARRIERS

cuda::barrier



ASYNCHRONOUS TRANSACTION BARRIERS

Trigger barrier with both transactions and threads



TENSOR MEMORY ACCELERATOR UNIT (TMA) FOR ASYNC DATA MOVEMENT

Tensor Memory Accelerator Unit

Hardware-accelerated **bi-directional** bulk copy

Global \leftrightarrow Shared Memory

DSMEM \leftrightarrow DSMEM within cluster

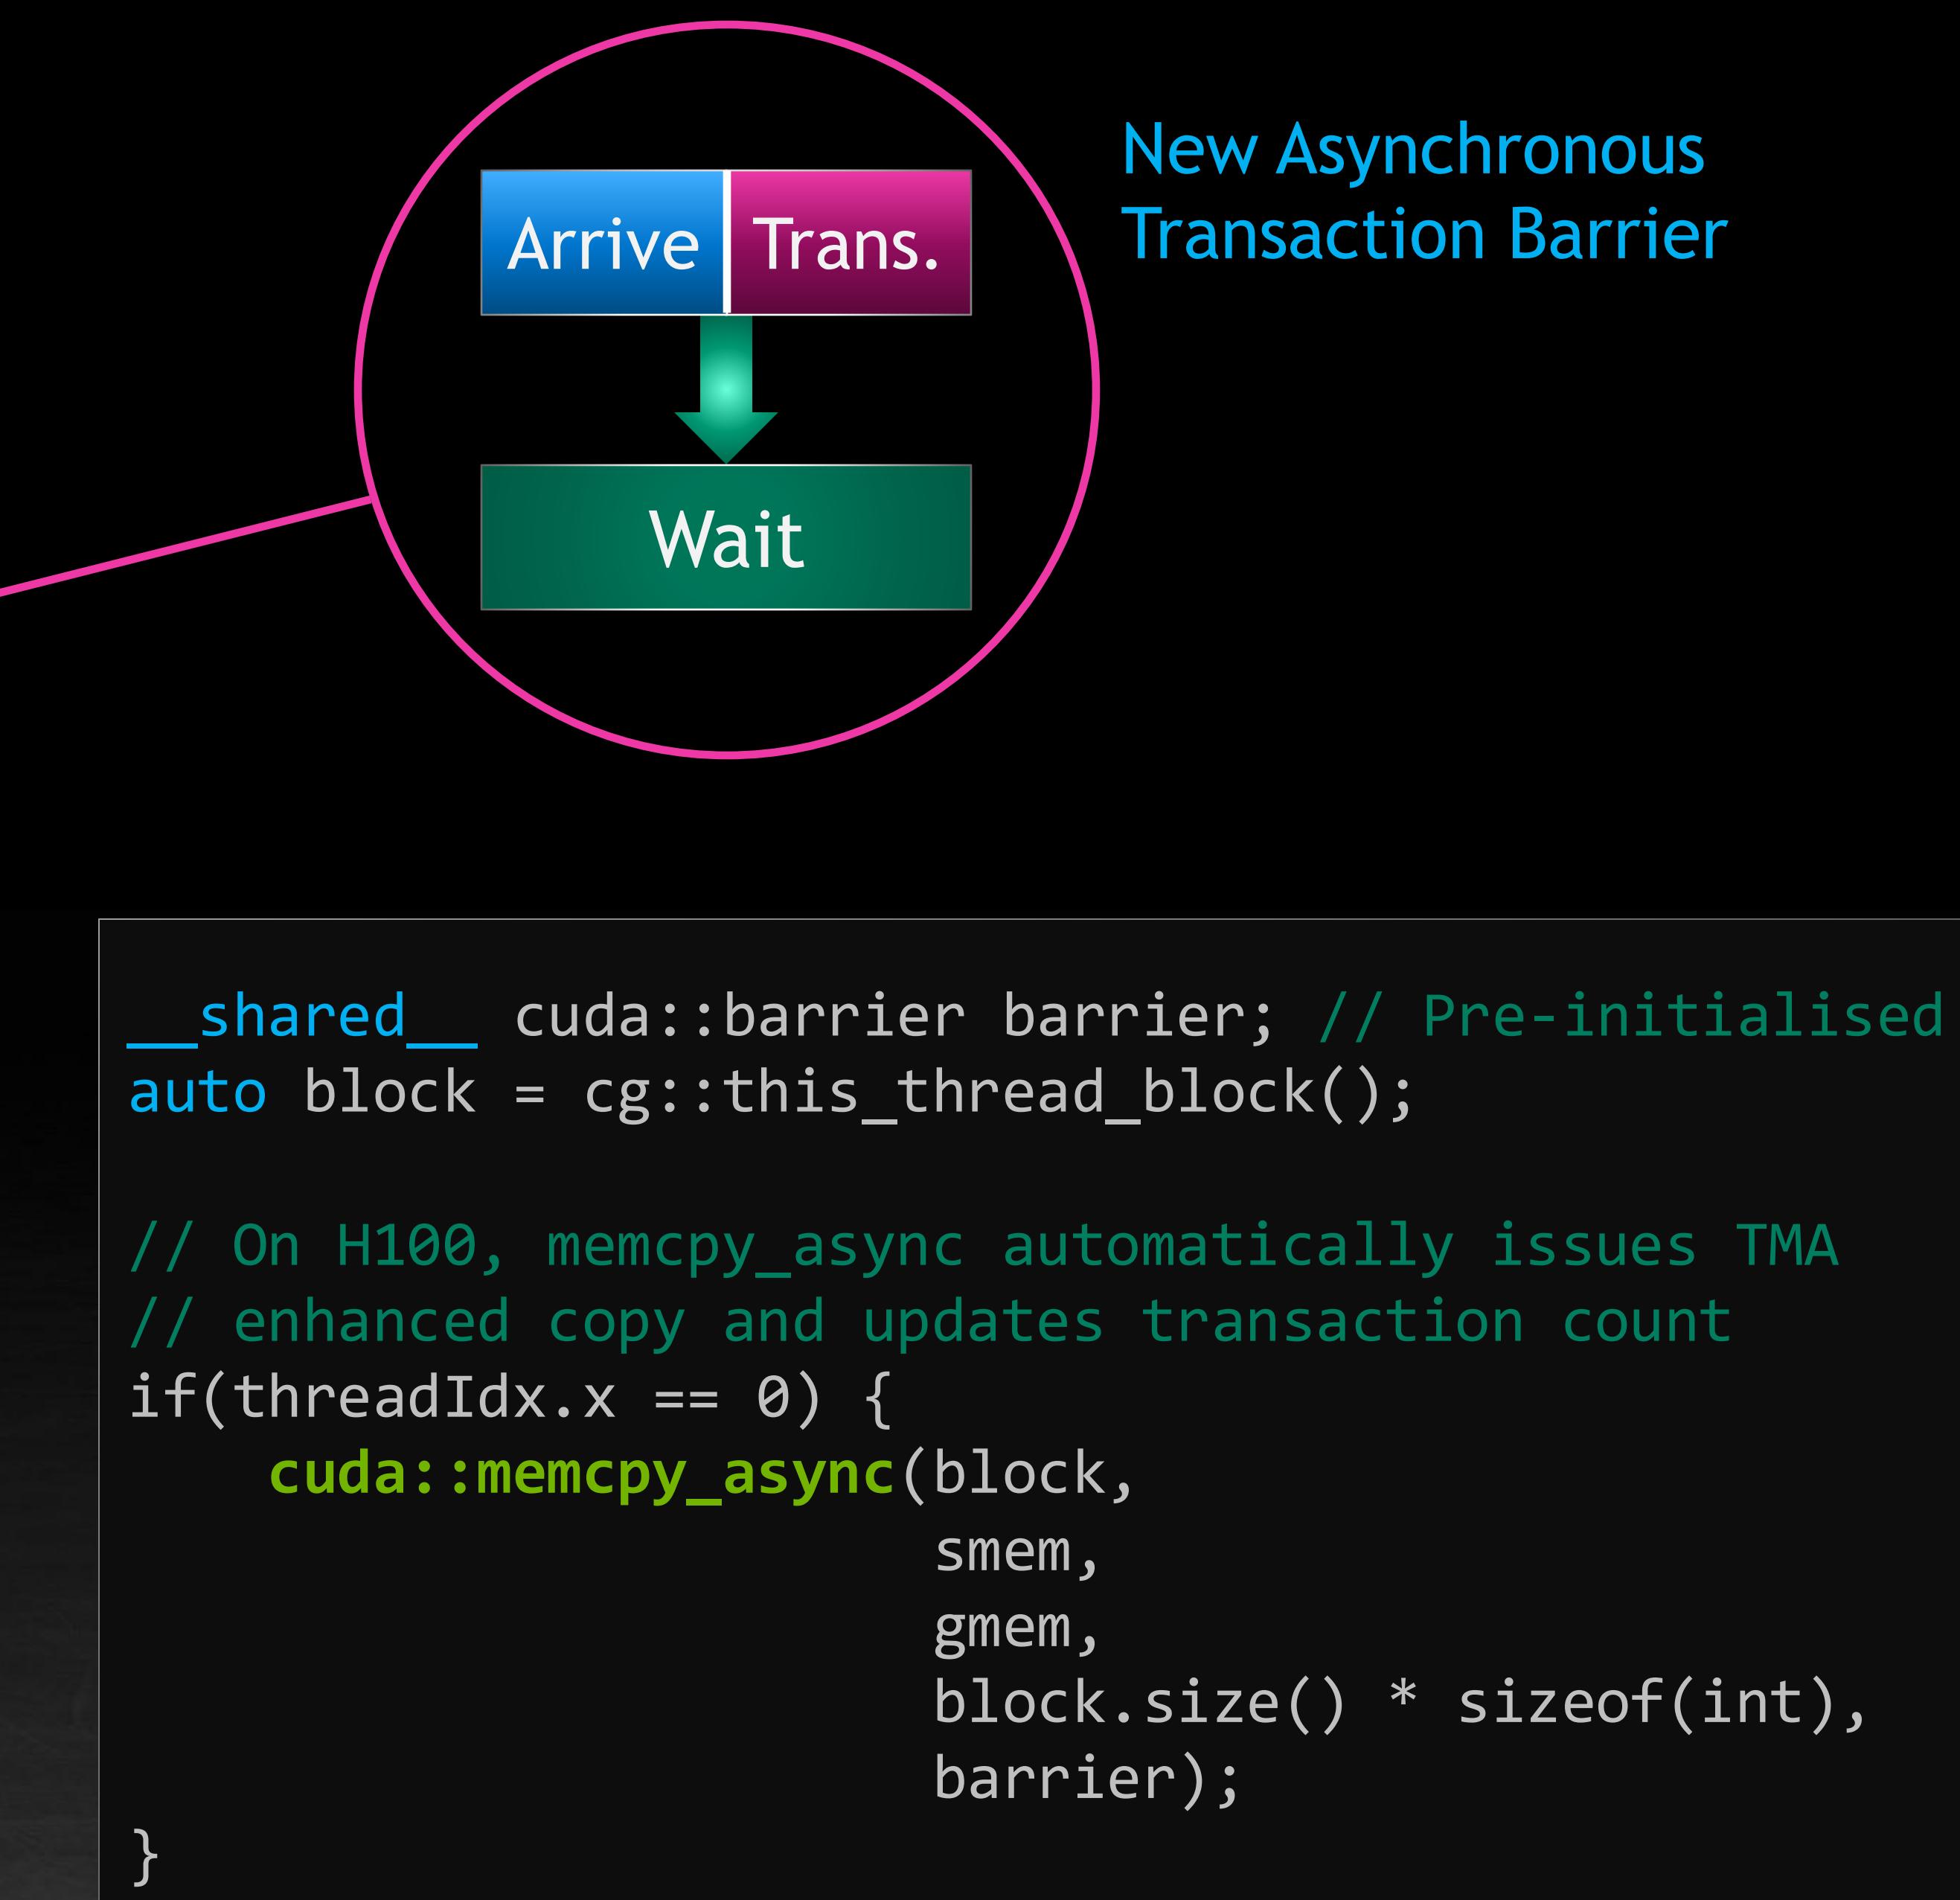
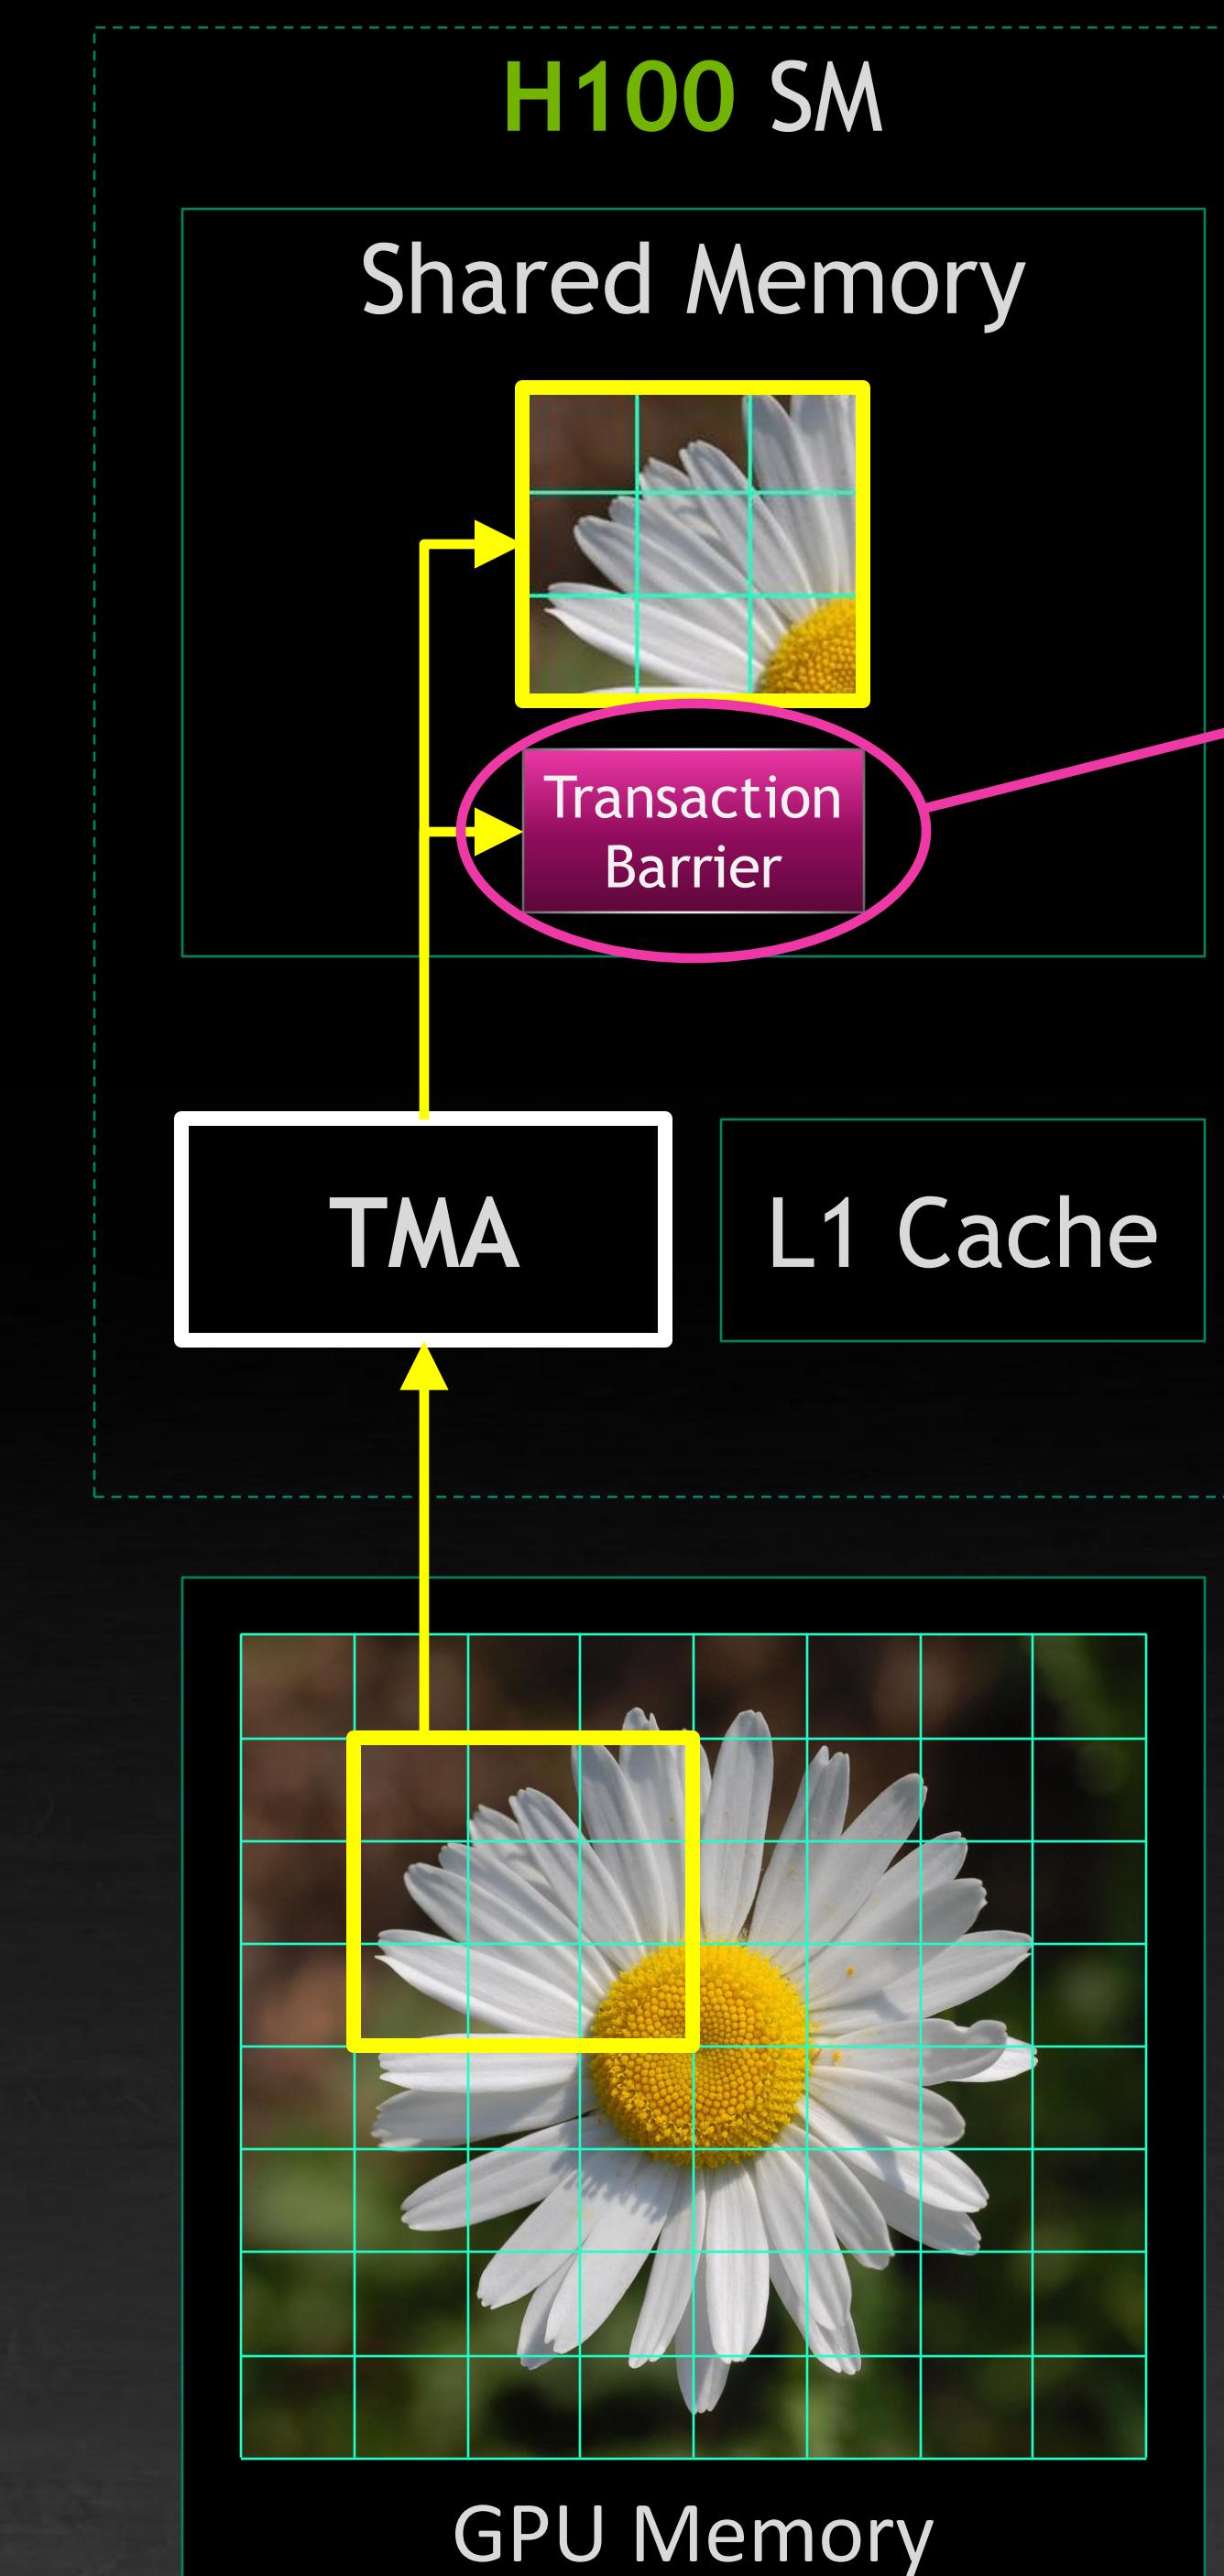
Uses **Asynchronous Transaction Barrier** to track completion

Faster And More Flexible Async Copying

Drop-in enhancement to existing
`cuda::memcpy_async()` function

Single thread triggers any-size copy - no need for
looping or collective copying

Multiple copy operations can contribute to a
single cluster-wide async transaction barrier



Same copy syntax becomes faster and more flexible
New syntax to enable TMA capabilities

HOPPER TENSOR CORES

NVIDIA - H100

New Faster Tensor Core Instructions

16b Floating-point Tensor Core operations **16x** and **32x** faster than F32 CUDA Cores

8b Floating-point Tensor Core operations **32x** and **64x** faster than F32 CUDA Cores

Improved Integer Tensor Core operations **32x** and **64x** faster than F32 CUDA Cores

Improved Tensor Float 32 Tensor Cores **8x** and **16x** faster than F32 CUDA Cores

Improved IEEE double-precision Tensor Cores **2x** faster than F64 CUDA Cores

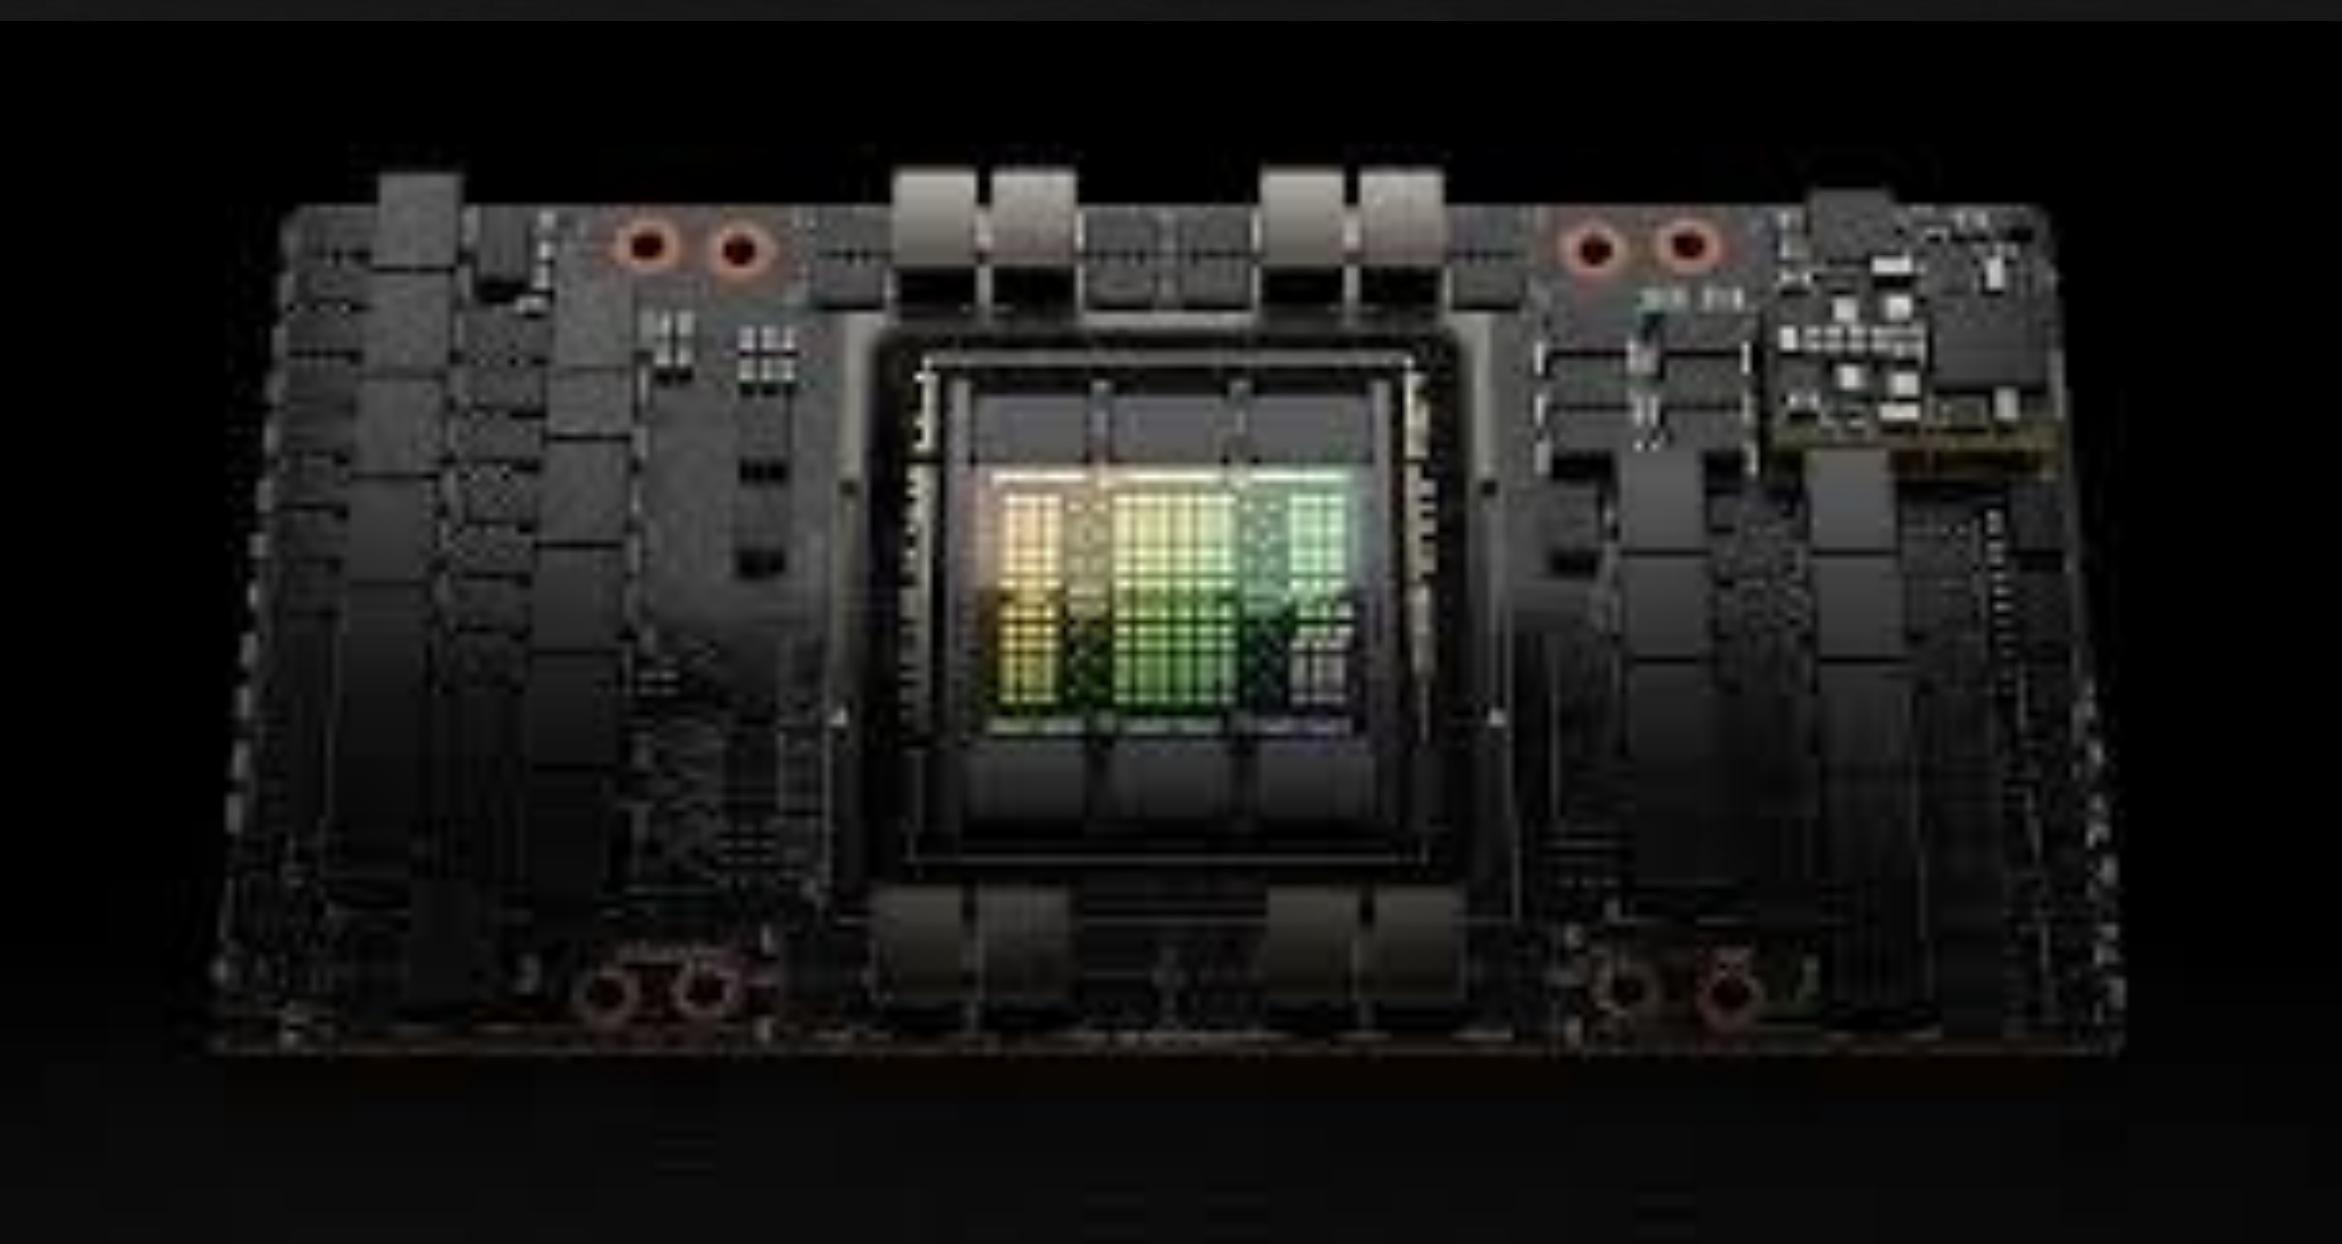


Additional Data Types

8bit Floating point types E5M2 and E4M3

Many additional new features – see "[NVIDIA H100 Tensor Core GPU Architecture](#)" whitepaper

Applicable with using 2:4 sparse computation



NVIDIA HOPPER ARCHITECTURE - TENSOR CORE OPERATIONS

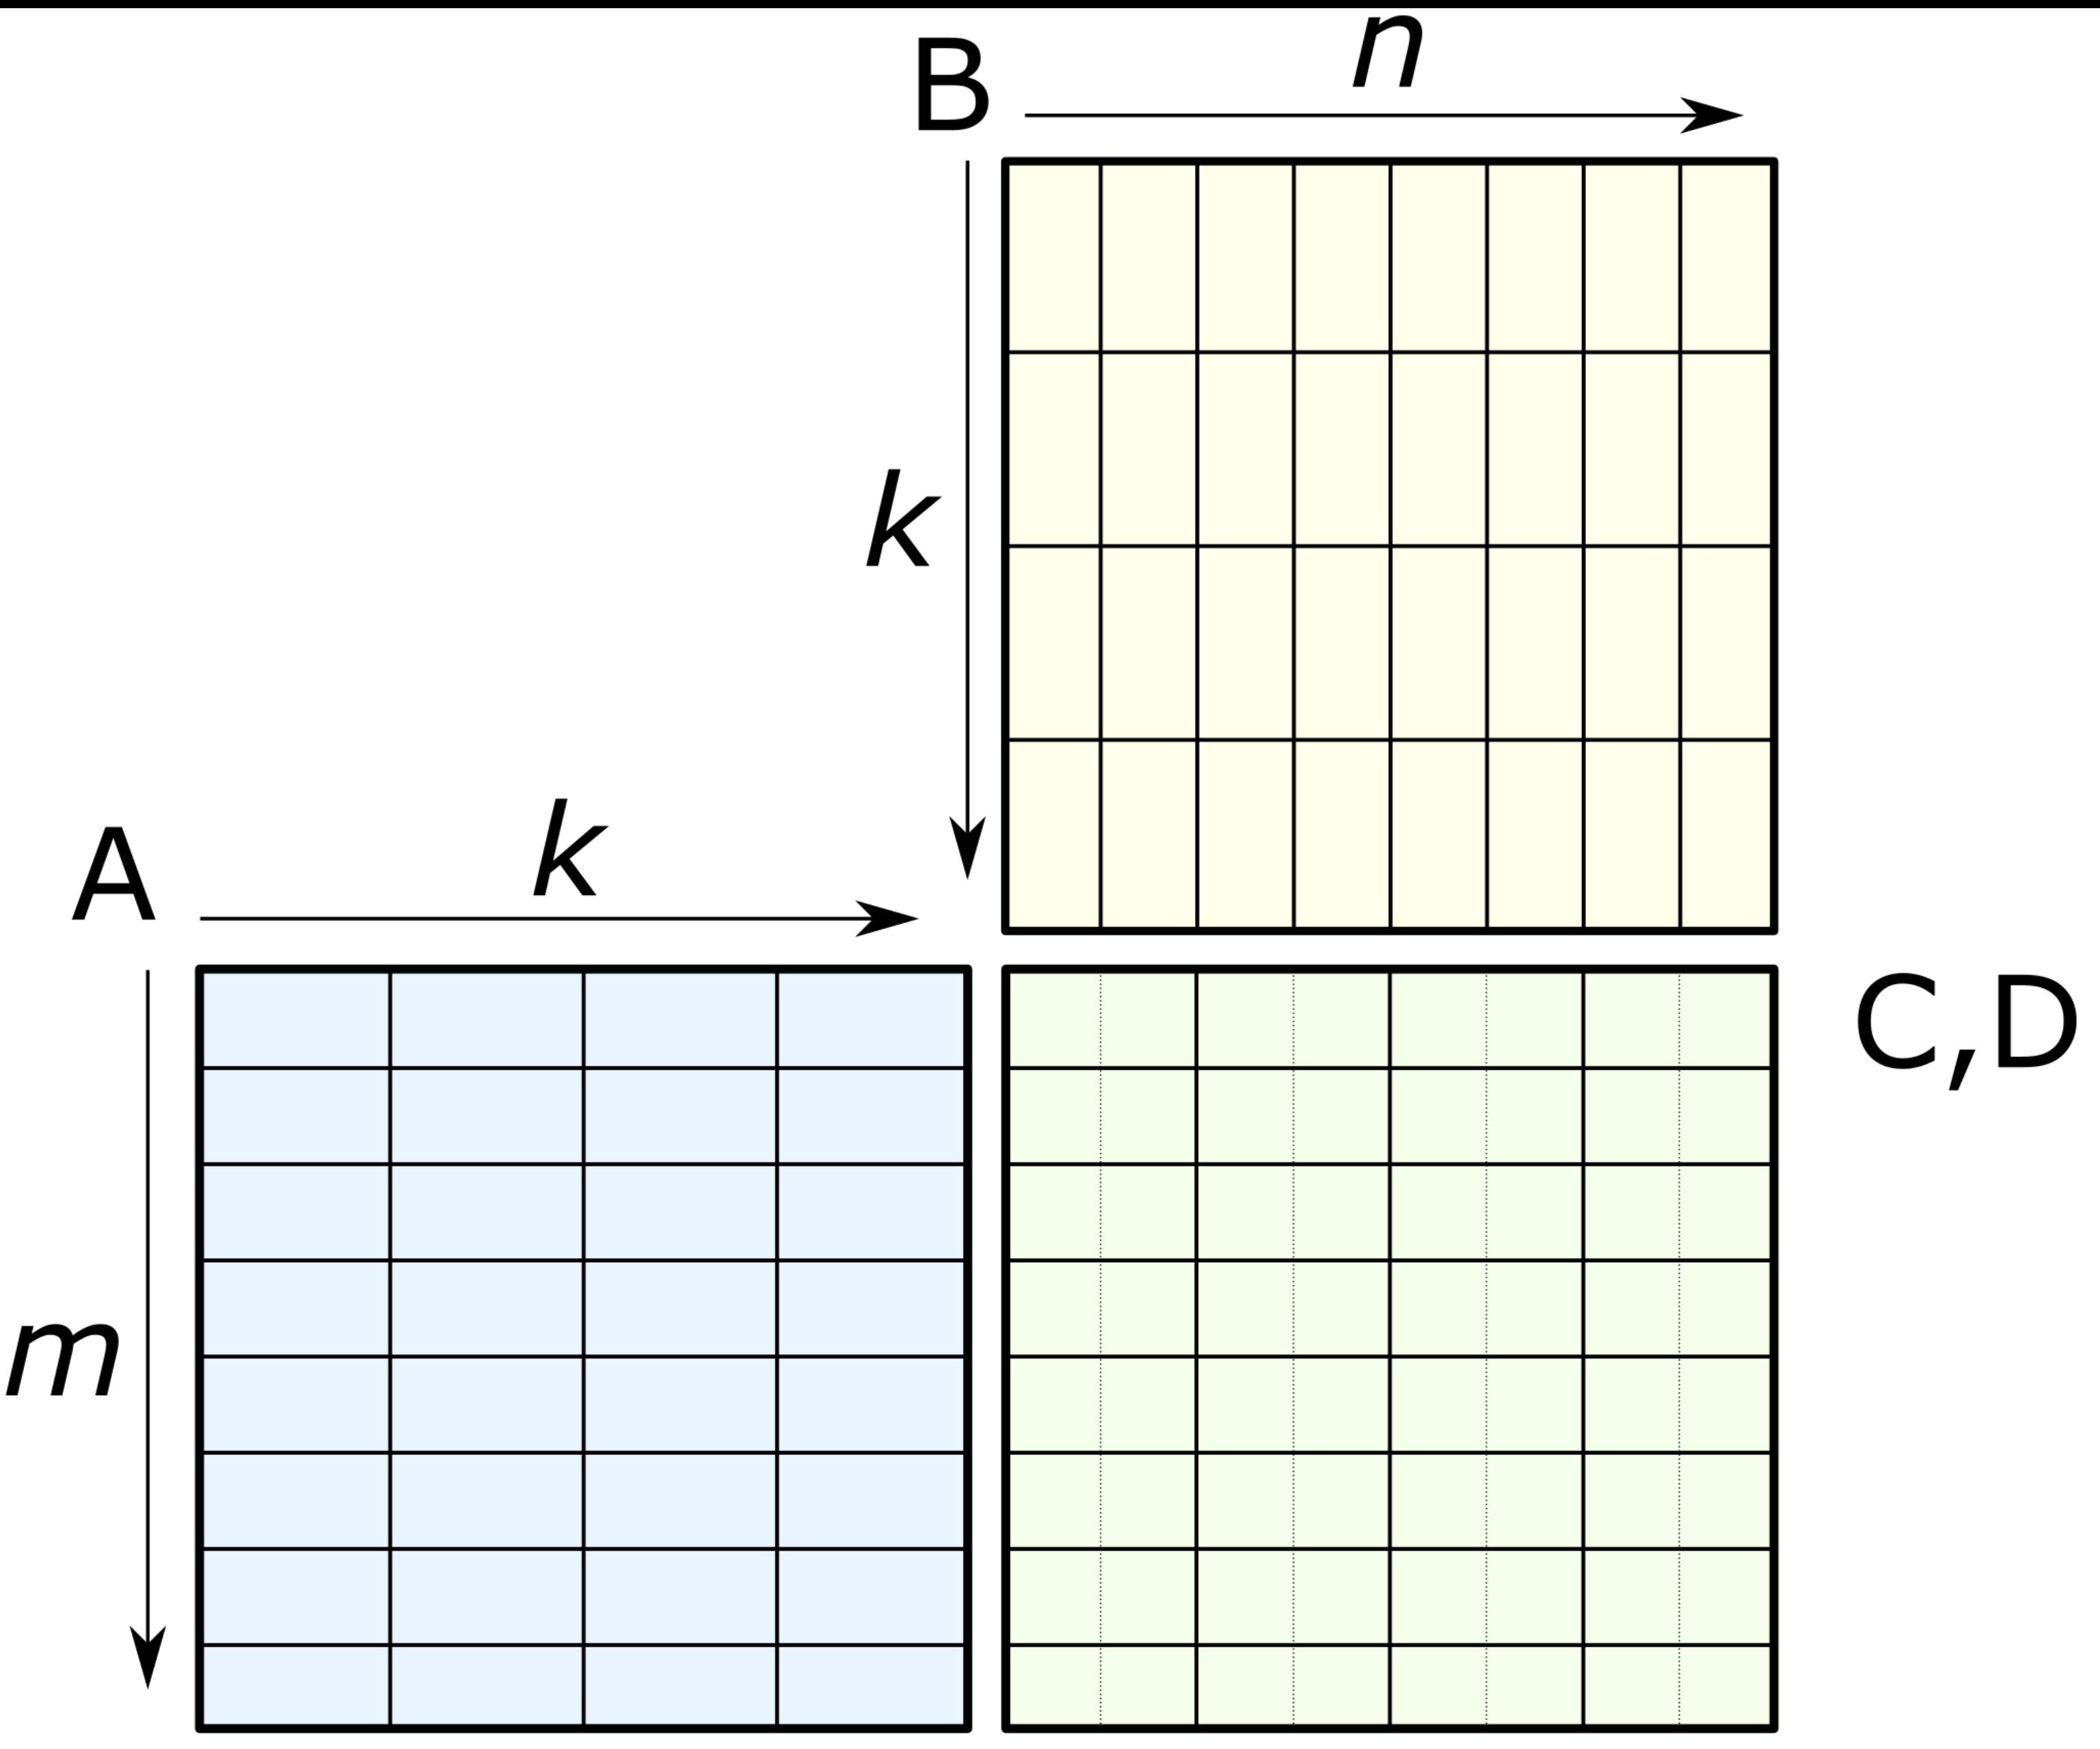
Operation	Data Types (A * B + C)	Shape	TFLOPS Hopper (H100)	TFLOPS Ampere (A100)	TFLOPS Volta (V100)
MMA wgmma, mma_sync	F16 * F16 + F16 F16 * F16 + F32 BF16 * BF16 + F32	64-by-N-by-16	1000	312	125
MMA wgmma, mma_sync	TF32 * TF32 + F32	64-by-N-by-8	500	156	N/A
MMA mma_sync	F64 * F64 + F64	16-by-8-by-4/8/16	60	19	N/A
MMA wgmma, mma_sync	S8 * S8 + S32	64-by-N-by-32	2000	624	N/A
MMA wgmma, mma_sync	F8 * F8 + F32	64-by-N-by-32	2000	N/A	N/A

<https://docs.nvidia.com/cuda/parallel-thread-execution/index.html#asynchronous-warpgroup-level-matrix-instructions>

TENSOR CORES OPERATION : FUNDAMENTAL SHAPE

Matrix operations: $D = \text{op}(A, B) + D$

- Matrix multiply-add
- XOR-POPC
- M-by-N-by-K matrix operation
- Warp-Group-wide, async. collective operation.
- 128 threads within warp-group collectively holds D operand.
- Operands A and B (not both) can optional be loaded from register memory. Loading directly from shared memory using descriptors most optimal.
- Valid value(s) of $M = 64$; $N \in [8, 256]$ and $K \in [8, 16, 32, 256]$ based on data type refer [ISA documentation](#) for details



AMPERE/TURING FP16 MMA

```

template <>
struct MMA_Traits<SM80_16x8x8_F32F16F16F32_TN>
{
    using ElementDVal = float;
    using ElementAVal = half_t;
    using ElementBVal = half_t;
    using ElementCVal = float;

    using Shape_MNK = Shape<_16,_8,_8>;
    using ThrID = Layout<_32>;
    // (T32,V4) -> (M16,K8)
    using ALayout = Layout<Shape <Shape <_4,_8>, Shape <_2,_2>>,
                           Stride<Stride<_32,_1>, Stride<_16,_8>>>;
    // (T32,V2) -> (N8,K8)
    using BLayout = Layout<Shape <Shape <_4,_8>,_2>,
                           Stride<Stride<_16,_1>,_8>>>;
    // (T32,V4) -> (M16,N8)
    using CLayout = Layout<Shape <Shape <_4,_8>, Shape <_2,_2>>,
                           Stride<Stride<_32,_1>, Stride<_16,_8>>>;
};

```

T0 V0	T0 V1	T1 V0	T1 V1	T2 V0	T2 V1	T3 V0	T3 V1
T4 V0	T4 V1	T5 V0	T5 V1	T6 V0	T6 V1	T7 V0	T7 V1
T8 V0	T8 V1	T9 V0	T9 V1	T10 V0	T10 V1	T11 V0	T11 V1
T12 V0	T12 V1	T13 V0	T13 V1	T14 V0	T14 V1	T15 V0	T15 V1
T16 V0	T16 V1	T17 V0	T17 V1	T18 V0	T18 V1	T19 V0	T19 V1
T20 V0	T20 V1	T21 V0	T21 V1	T22 V0	T22 V1	T23 V0	T23 V1
T24 V0	T24 V1	T25 V0	T25 V1	T26 V0	T26 V1	T27 V0	T27 V1
T28 V0	T28 V1	T29 V0	T29 V1	T30 V0	T30 V1	T31 V0	T31 V1
T0 V2	T0 V3	T1 V2	T1 V3	T2 V2	T2 V3	T3 V2	T3 V3
T4 V2	T4 V3	T5 V2	T5 V3	T6 V2	T6 V3	T7 V2	T7 V3
T8 V2	T8 V3	T9 V2	T9 V3	T10 V2	T10 V3	T11 V2	T11 V3
T12 V2	T12 V3	T13 V2	T13 V3	T14 V2	T14 V3	T15 V2	T15 V3
T16 V2	T16 V3	T17 V2	T17 V3	T18 V2	T18 V3	T19 V2	T19 V3
T20 V2	T20 V3	T21 V2	T21 V3	T22 V2	T22 V3	T23 V2	T23 V3
T24 V2	T24 V3	T25 V2	T25 V3	T26 V2	T26 V3	T27 V2	T27 V3
T28 V2	T28 V3	T29 V2	T29 V3	T30 V2	T30 V3	T31 V2	T31 V3

T0 V0	T4 V0	T8 V0	T12 V0	T16 V0	T20 V0	T24 V0	T28 V0
T0 V1	T4 V1	T8 V1	T12 V1	T16 V1	T20 V1	T24 V1	T28 V1
T1 V0	T5 V0	T9 V0	T13 V0	T17 V0	T21 V0	T25 V0	T29 V0
T1 V1	T5 V1	T9 V1	T13 V1	T17 V1	T21 V1	T25 V1	T29 V1
T2 V0	T6 V0	T10 V0	T14 V0	T18 V0	T22 V0	T26 V0	T30 V0
T2 V1	T6 V1	T10 V1	T14 V1	T18 V1	T22 V1	T26 V1	T30 V1
T3 V0	T7 V0	T11 V0	T15 V0	T19 V0	T23 V0	T27 V0	T31 V0
T3 V1	T7 V1	T11 V1	T15 V1	T19 V1	T23 V1	T27 V1	T31 V1

T0 V0	T0 V1	T1 V0	T1 V1	T2 V0	T2 V1	T3 V0	T3 V1
T4 V0	T4 V1	T5 V0	T5 V1	T6 V0	T6 V1	T7 V0	T7 V1
T8 V0	T8 V1	T9 V0	T9 V1	T10 V0	T10 V1	T11 V0	T11 V1
T12 V0	T12 V1	T13 V0	T13 V1	T14 V0	T14 V1	T15 V0	T15 V1
T16 V0	T16 V1	T17 V0	T17 V1	T18 V0	T18 V1	T19 V0	T19 V1
T20 V0	T20 V1	T21 V0	T21 V1	T22 V0	T22 V1	T23 V0	T23 V1
T24 V0	T24 V1	T25 V0	T25 V1	T26 V0	T26 V1	T27 V0	T27 V1
T28 V0	T28 V1	T29 V0	T29 V1	T30 V0	T30 V1	T31 V0	T31 V1
T0 V2	T0 V3	T1 V2	T1 V3	T2 V2	T2 V3	T3 V2	T3 V3
T4 V2	T4 V3	T5 V2	T5 V3	T6 V2	T6 V3	T7 V2	T7 V3
T8 V2	T8 V3	T9 V2	T9 V3	T10 V2	T10 V3	T11 V2	T11 V3
T12 V2	T12 V3	T13 V2	T13 V3	T14 V2	T14 V3	T15 V2	T15 V3
T16 V2	T16 V3	T17 V2	T17 V3	T18 V2	T18 V3	T19 V2	T19 V3
T20 V2	T20 V3	T21 V2	T21 V3	T22 V2	T22 V3	T23 V2	T23 V3
T24 V2	T24 V3	T25 V2	T25 V3	T26 V2	T26 V3	T27 V2	T27 V3
T28 V2	T28 V3	T29 V2	T29 V3	T30 V2	T30 V3	T31 V2	T31 V3

T0 V0	T0 V1	T1 V0	T1 V1	T2 V0	T2 V1	T3 V0	T3 V1
T4 V0	T4 V1	T5 V0	T5 V1	T6 V0	T6 V1	T7 V0	T7 V1
T8 V0	T8 V1	T9 V0	T9 V1	T10 V0	T10 V1	T11 V0	T11 V1
T12 V0	T12 V1	T13 V0	T13 V1	T14 V0	T14 V1	T15 V0	T15 V1
T16 V0	T16 V1	T17 V0	T17 V1	T18 V0	T18 V1	T19 V0	T19 V1
T20 V0	T20 V1	T21 V0	T21 V1	T22 V0	T22 V1	T23 V0	T23 V1
T24 V0	T24 V1	T25 V0	T25 V1	T26 V0	T26 V1	T27 V0	T27 V1
T28 V0	T28 V1	T29 V0	T29 V1	T30 V0	T30 V1	T31 V0	T31 V1
T0 V2	T0 V3	T1 V2	T1 V3	T2 V2	T2 V3	T3 V2	T3 V3
T4 V2	T4 V3	T5 V2	T5 V3	T6 V2	T6 V3	T7 V2	T7 V3
T8 V2	T8 V3	T9 V2	T9 V3	T10 V2	T10 V3	T11 V2	T11 V3
T12 V2	T12 V3	T13 V2	T13 V3	T14 V2	T14 V3	T15 V2	T15 V3
T16 V2	T16 V3	T17 V2	T17 V3	T18 V2	T18 V3	T19 V2	T19 V3
T20 V2	T20 V3	T21 V2	T21 V3	T22 V2	T22 V3	T23 V2	T23 V3
T24 V2	T24 V3	T25 V2	T25 V3	T26 V2	T26 V3	T27 V2	T27 V3
T28 V2	T28 V3	T29 V2	T29 V3	T30 V2	T30 V3	T31 V2	T31 V3

SM80_16x8x8_F32F16F16F32_TN



HOPPER FP16 MMA

```

template <GMMA::Major tnspA, GMMA::Major tnspB, GMMA::ScaleIn scaleA,
GMMA::ScaleIn scaleB>
struct MMA_Traits<SM90_64x16x16_F16F16F16_SS<tnspA, tnspB, scaleA, scaleB>
{
    using ValTypeD = half_t;
    using ValTypeA = half_t;
    using ValTypeB = half_t;
    using ValTypeC = half_t;

    using Shape_MNK = Shape<-64,-16,-16>;
    using ThrID = Layout<-128>;
    using ALayout = GMMA::ABLayout< 64, 16>;
    using BLayout = GMMA::ABLayout< 16, 16>;;

    using CLayout = Layout<Shape < -4,-8, -4>, Shape < -2,-2, -2>,
Stride<Stride<-128,-1,-16>, Stride<-64,-8,-512>>>;
};

    
```

```

    
```

Hopper MMA :

F16 * F16 + F32

64-by-N-by-16

- **wmma.mma_async** issues a warp-group wide asynchronous MxNxK matrix multiply and accumulate operation,

$$D = A * B + D$$

Where A matrix is MxK, B matrix is KxN, and D matrix is MxN.

- Where $N \in [8, 256]$, and $M = 64$.

- A and B matrices are loaded directly from Shared Memory using descriptors.
- Transposition and scaling with certain immediate values supported as part of the instruction.

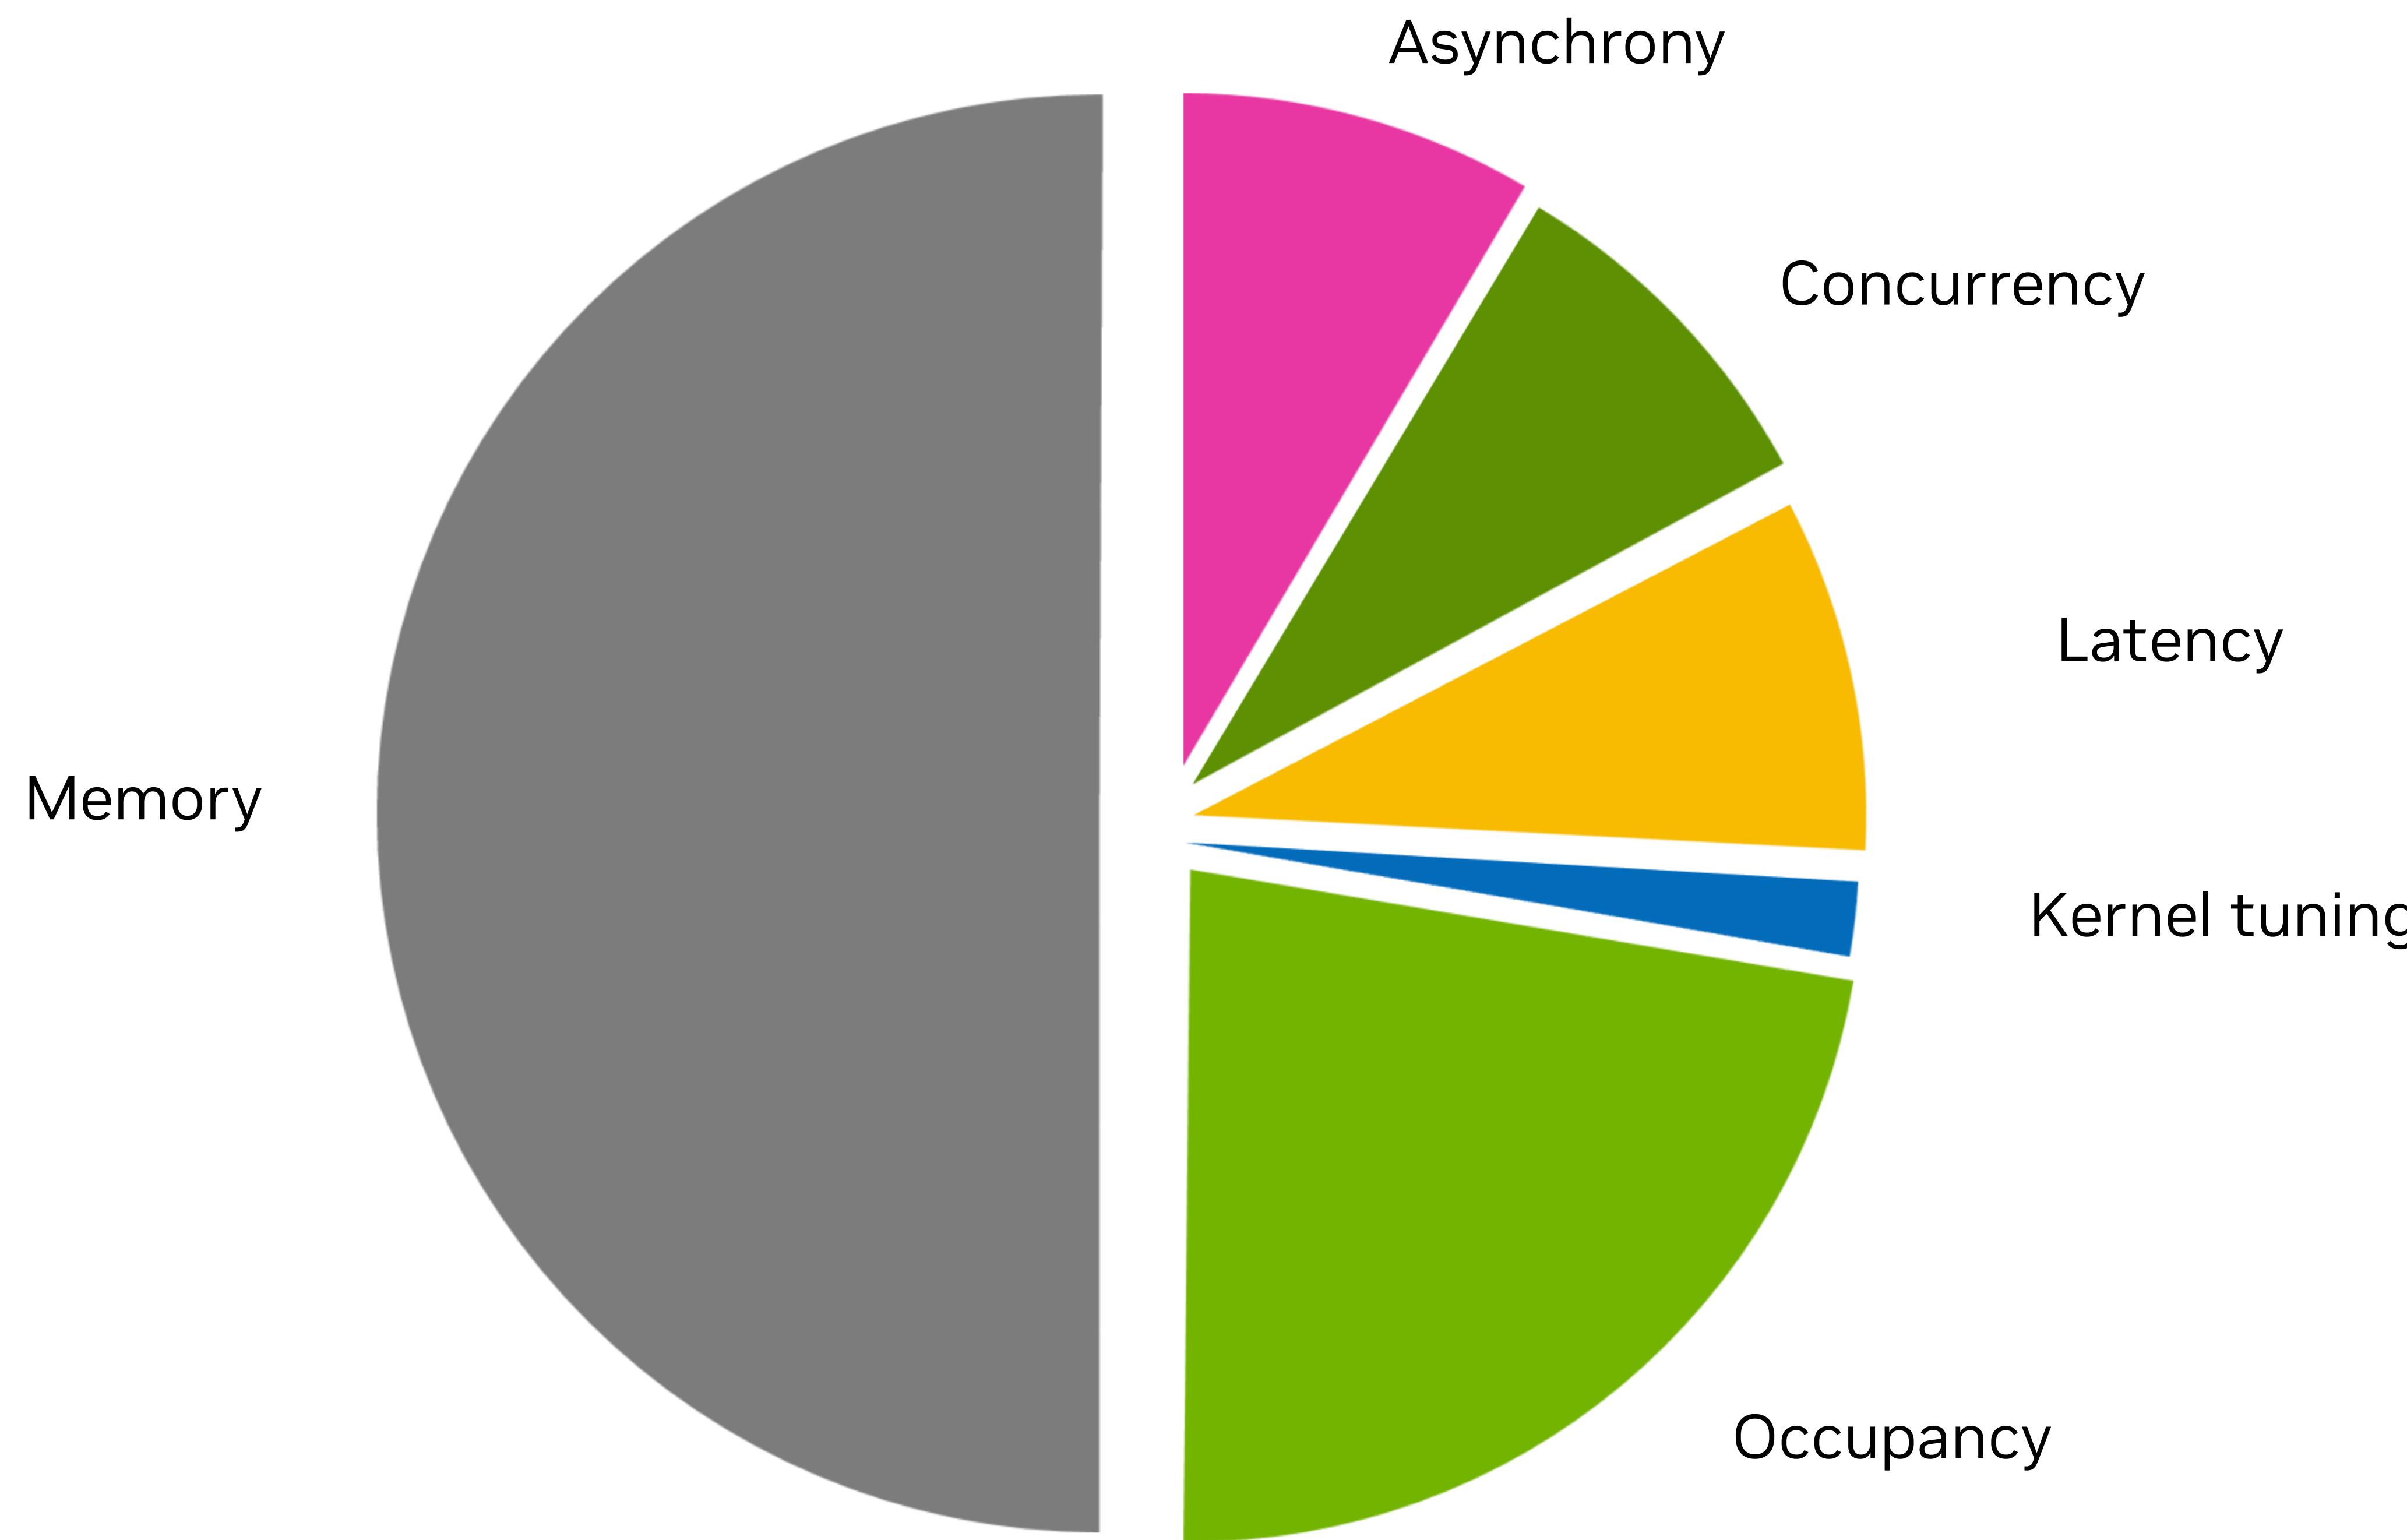
```
uint64_t descriptor_a, descriptor_b;
float D[64];
constexpr int scale_a, scale_b;          // -1 or +1
constexpr uint32_t scale_d;              // 0 or 1
constexpr uint32_t trans_a, trans_b;     // 0 or 1

asm (
    "wgmma.mma_async.sync.aligned.m64n64k16.f32.f16."
f16 "
    "%0, %1, %2, %3, %4, %5, %6, %7,
    %8, %9, %10, %11, %12, %13, %14, %15,
    ...
    ...
    %56, %57, %58, %59, %60, %61, %62, %63},
    %64,
    %65,
    %66, %67, %68, %69, %70;\n"
    : "+f"(D[0]), "+f"(D[1]), "+f"(D[2]), "+f"(D[3]),
    "+f"(D[4]), "+f"(D[5]), "+f"(D[6]), "+f"(D[7]),
    ...
    ...
    "+f"(D[60]), "+f"(D[61]), "+f"(D[62]), "+f"(D[63])
    :
    "l"(descriptor_a), "l"(descriptor_b),
    "n"(scale_d), "n"(scale_a), "n"(scale_b),
    "n"(trans_a), "n"(trans_b)
);
```

Ref. [cute::SM90_64x128x16_F32F16F16_SS](#)

Anatomy of a High Performance GEMM

What I Care About When Programming The GPU

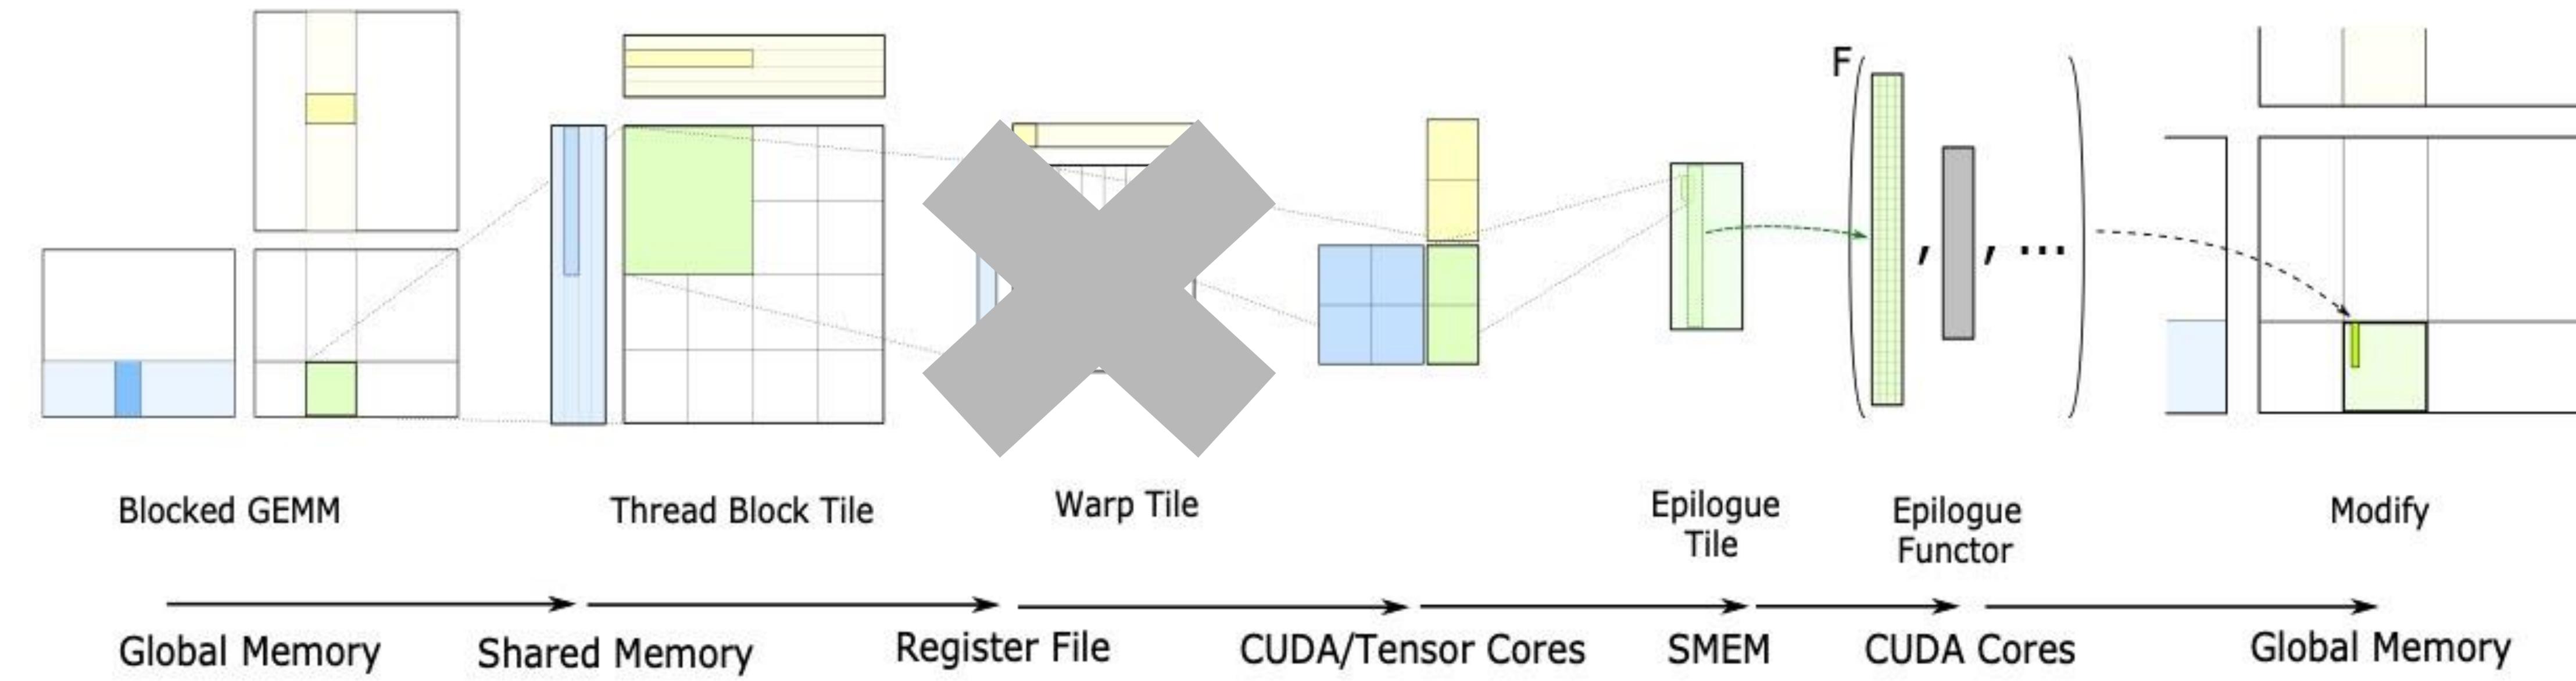


Anatomy of a High Performance Hopper GEMM

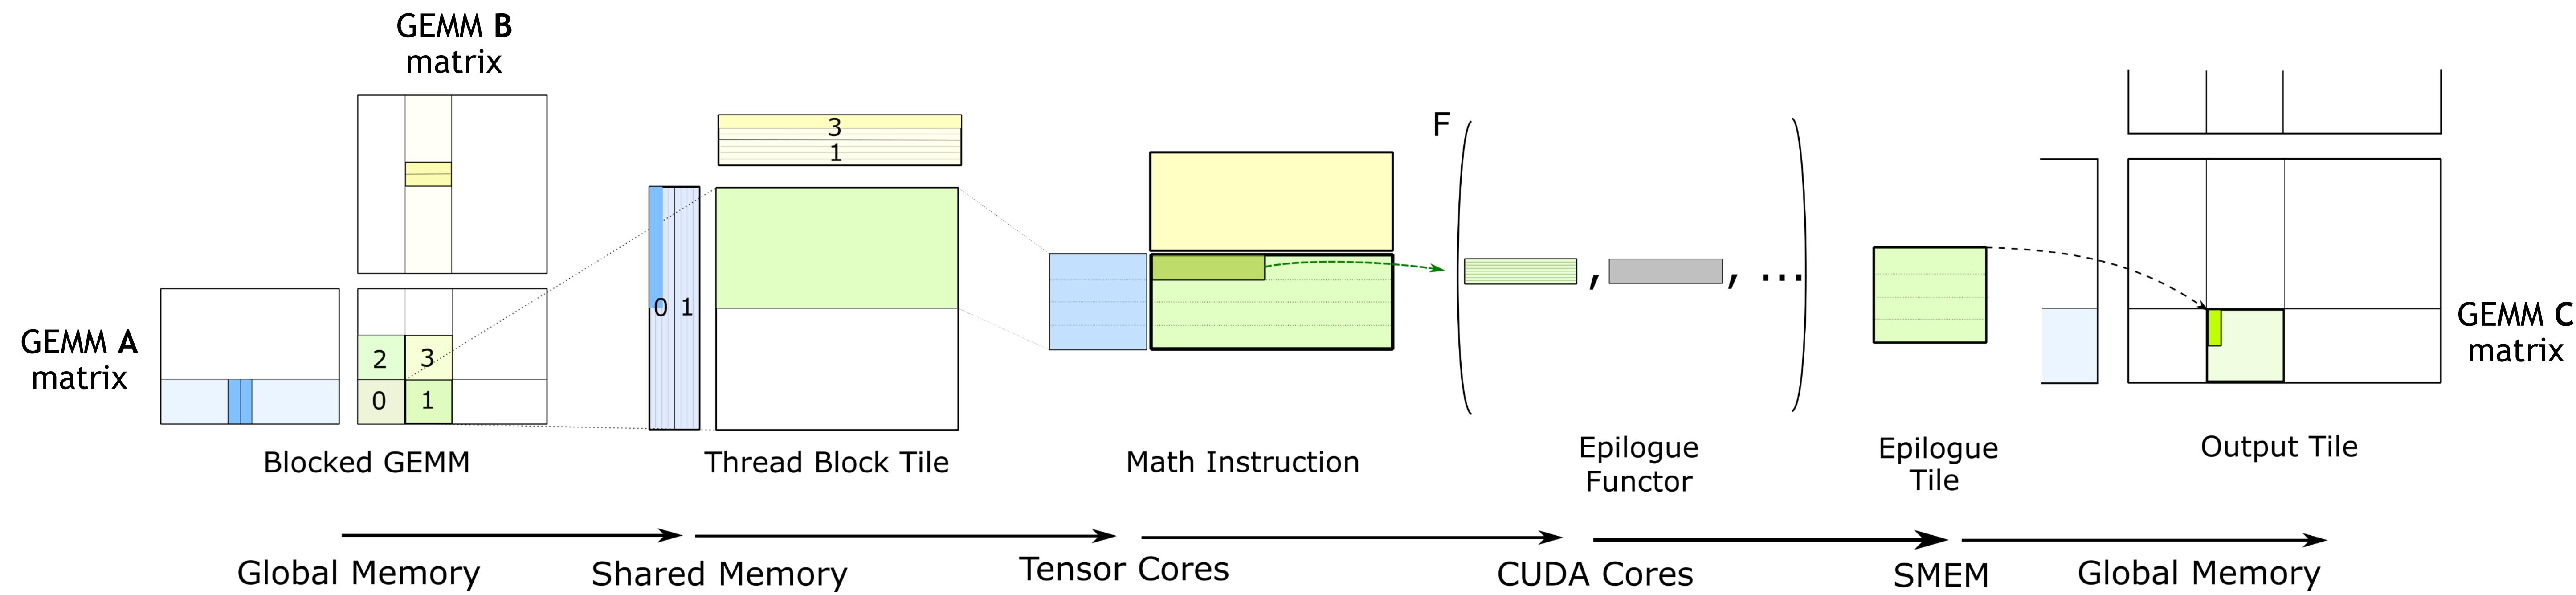
101 Level Optimizations

- Global memory accesses via TMA
 - HW supported Out-Of-Bounds determination a.k.a predication
 - Reduced register pressure and ALU ops
- Threadblock clusters used in concert with TMA to allow for programmatic multicast.
 - Amplification of L2 bandwidth via exploitation of temporal / spatial locality at the cost of quantization, occupancy.
- Deep software pipelining of loads via shared memory
 - Better Latency coverage at the cost of fewer ThreadBlocks running concurrently
- Warp-Group MMA operation
 - Multiple asynchronous MMA instructions in flight
- Mbarrier based synchronization between Async. TMA producers & MMA consumers
- Swizzling of data in shared memory to avoid bank conflicts
 - Prescribed together by MMA, TMA the tile size and the data-type

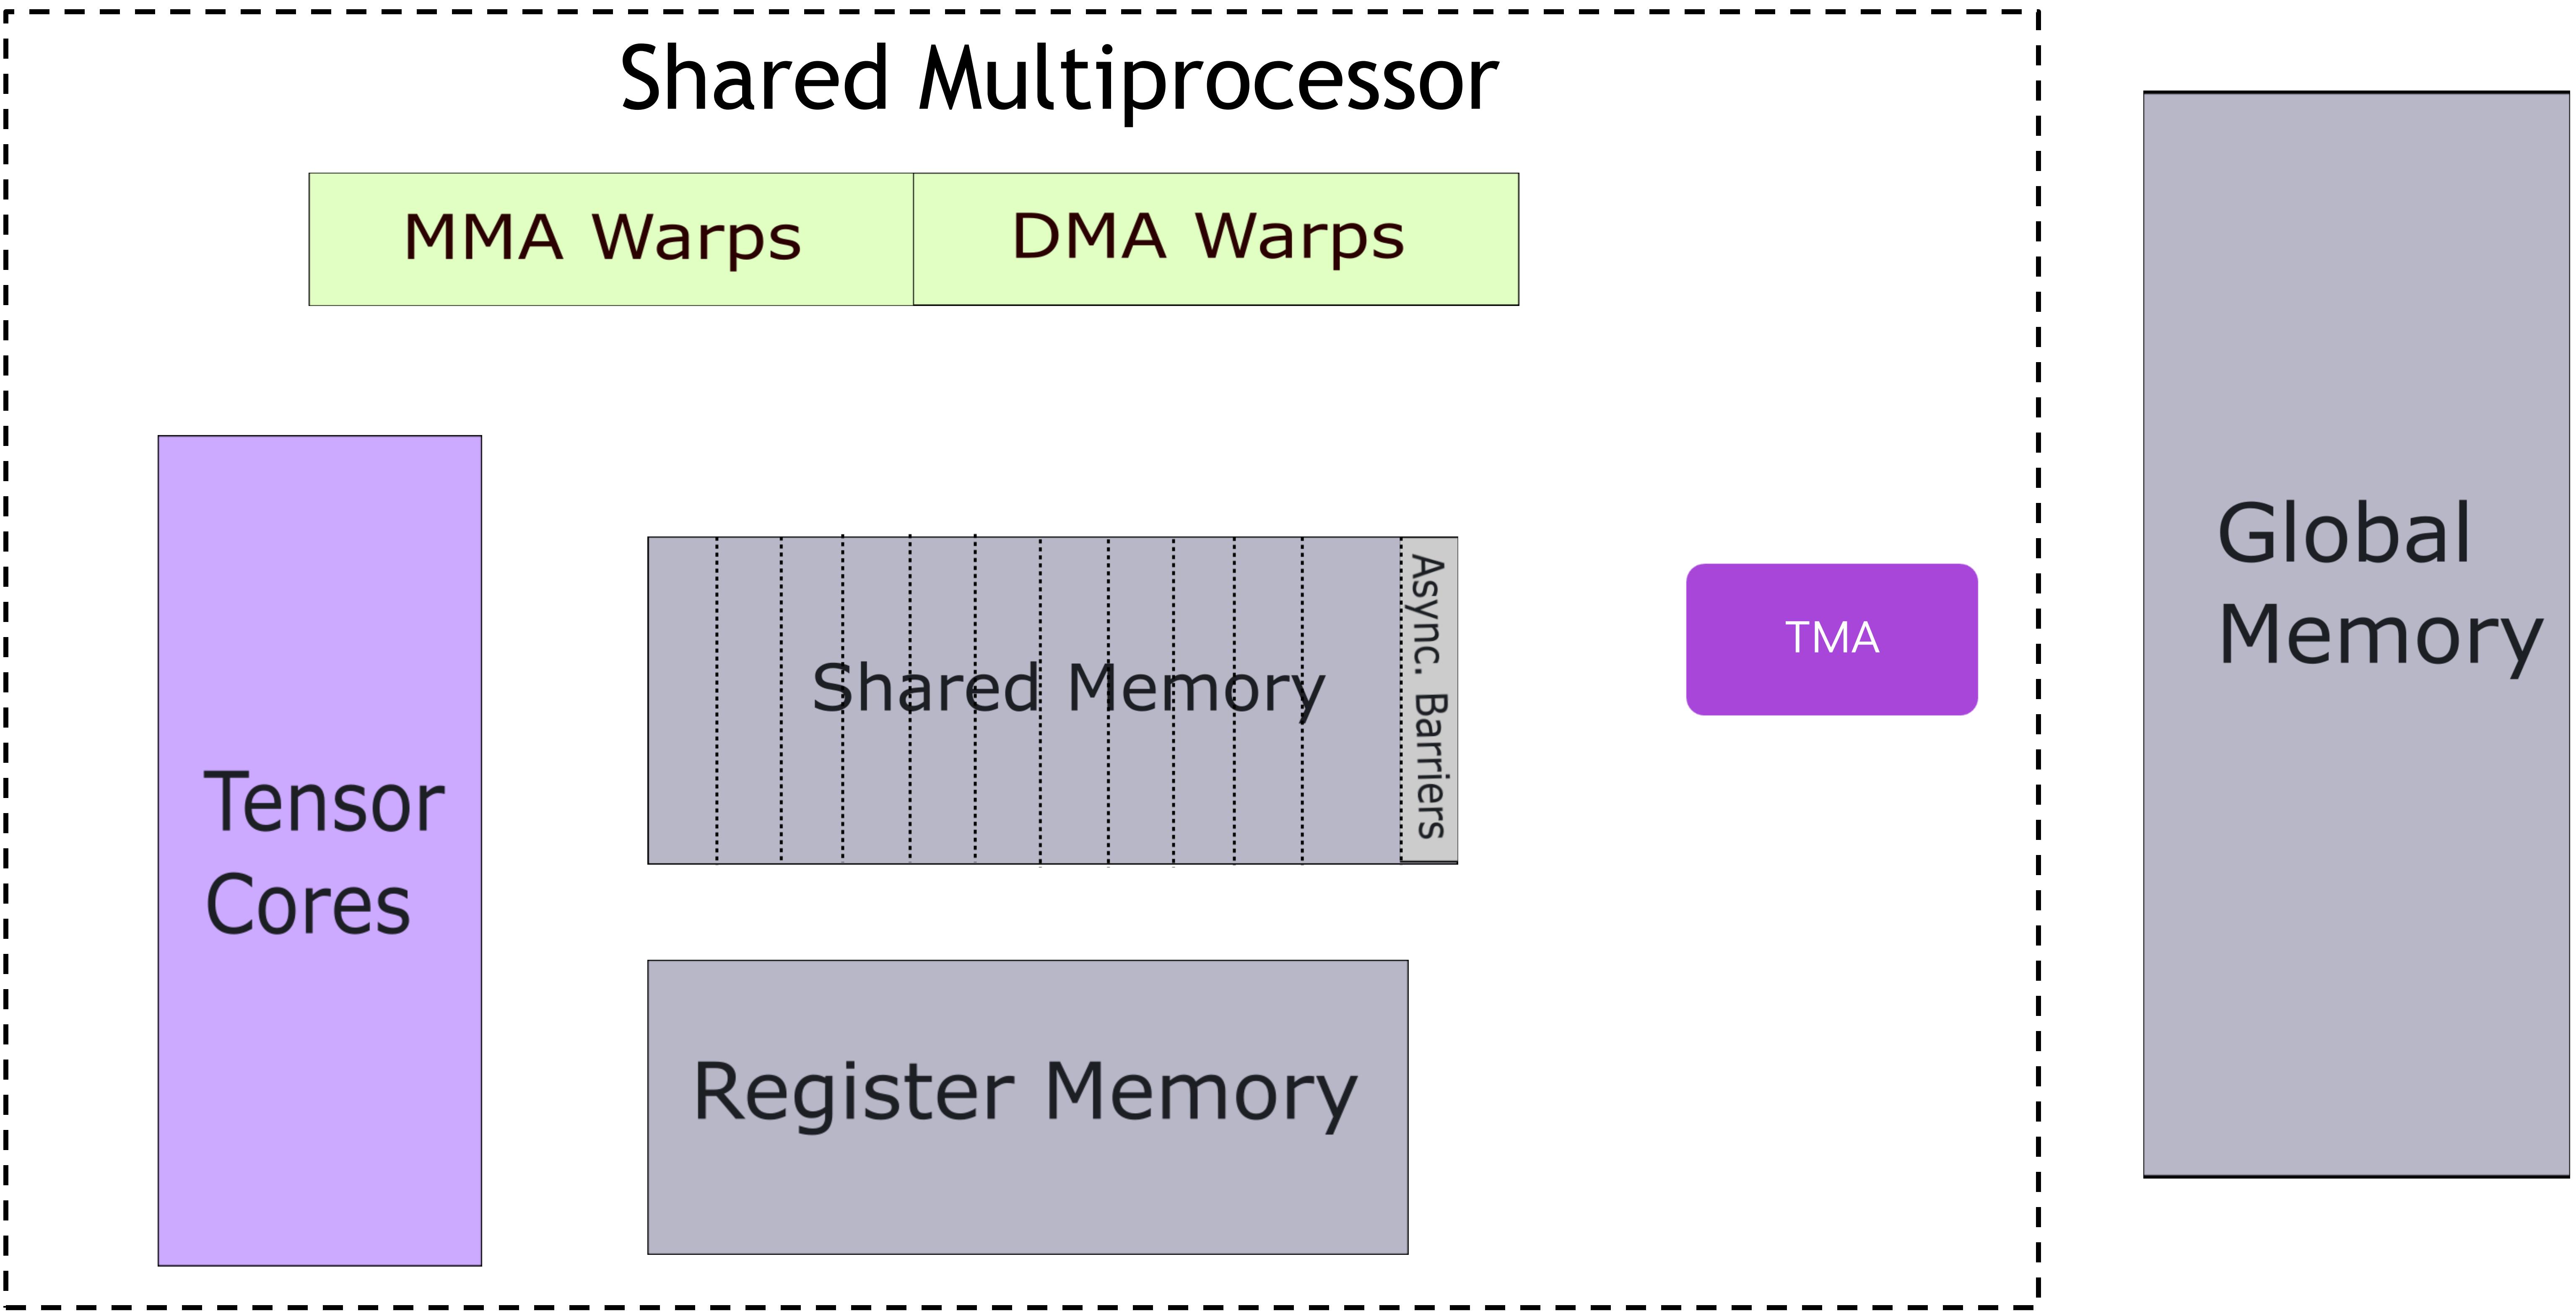
GEMM Data movement

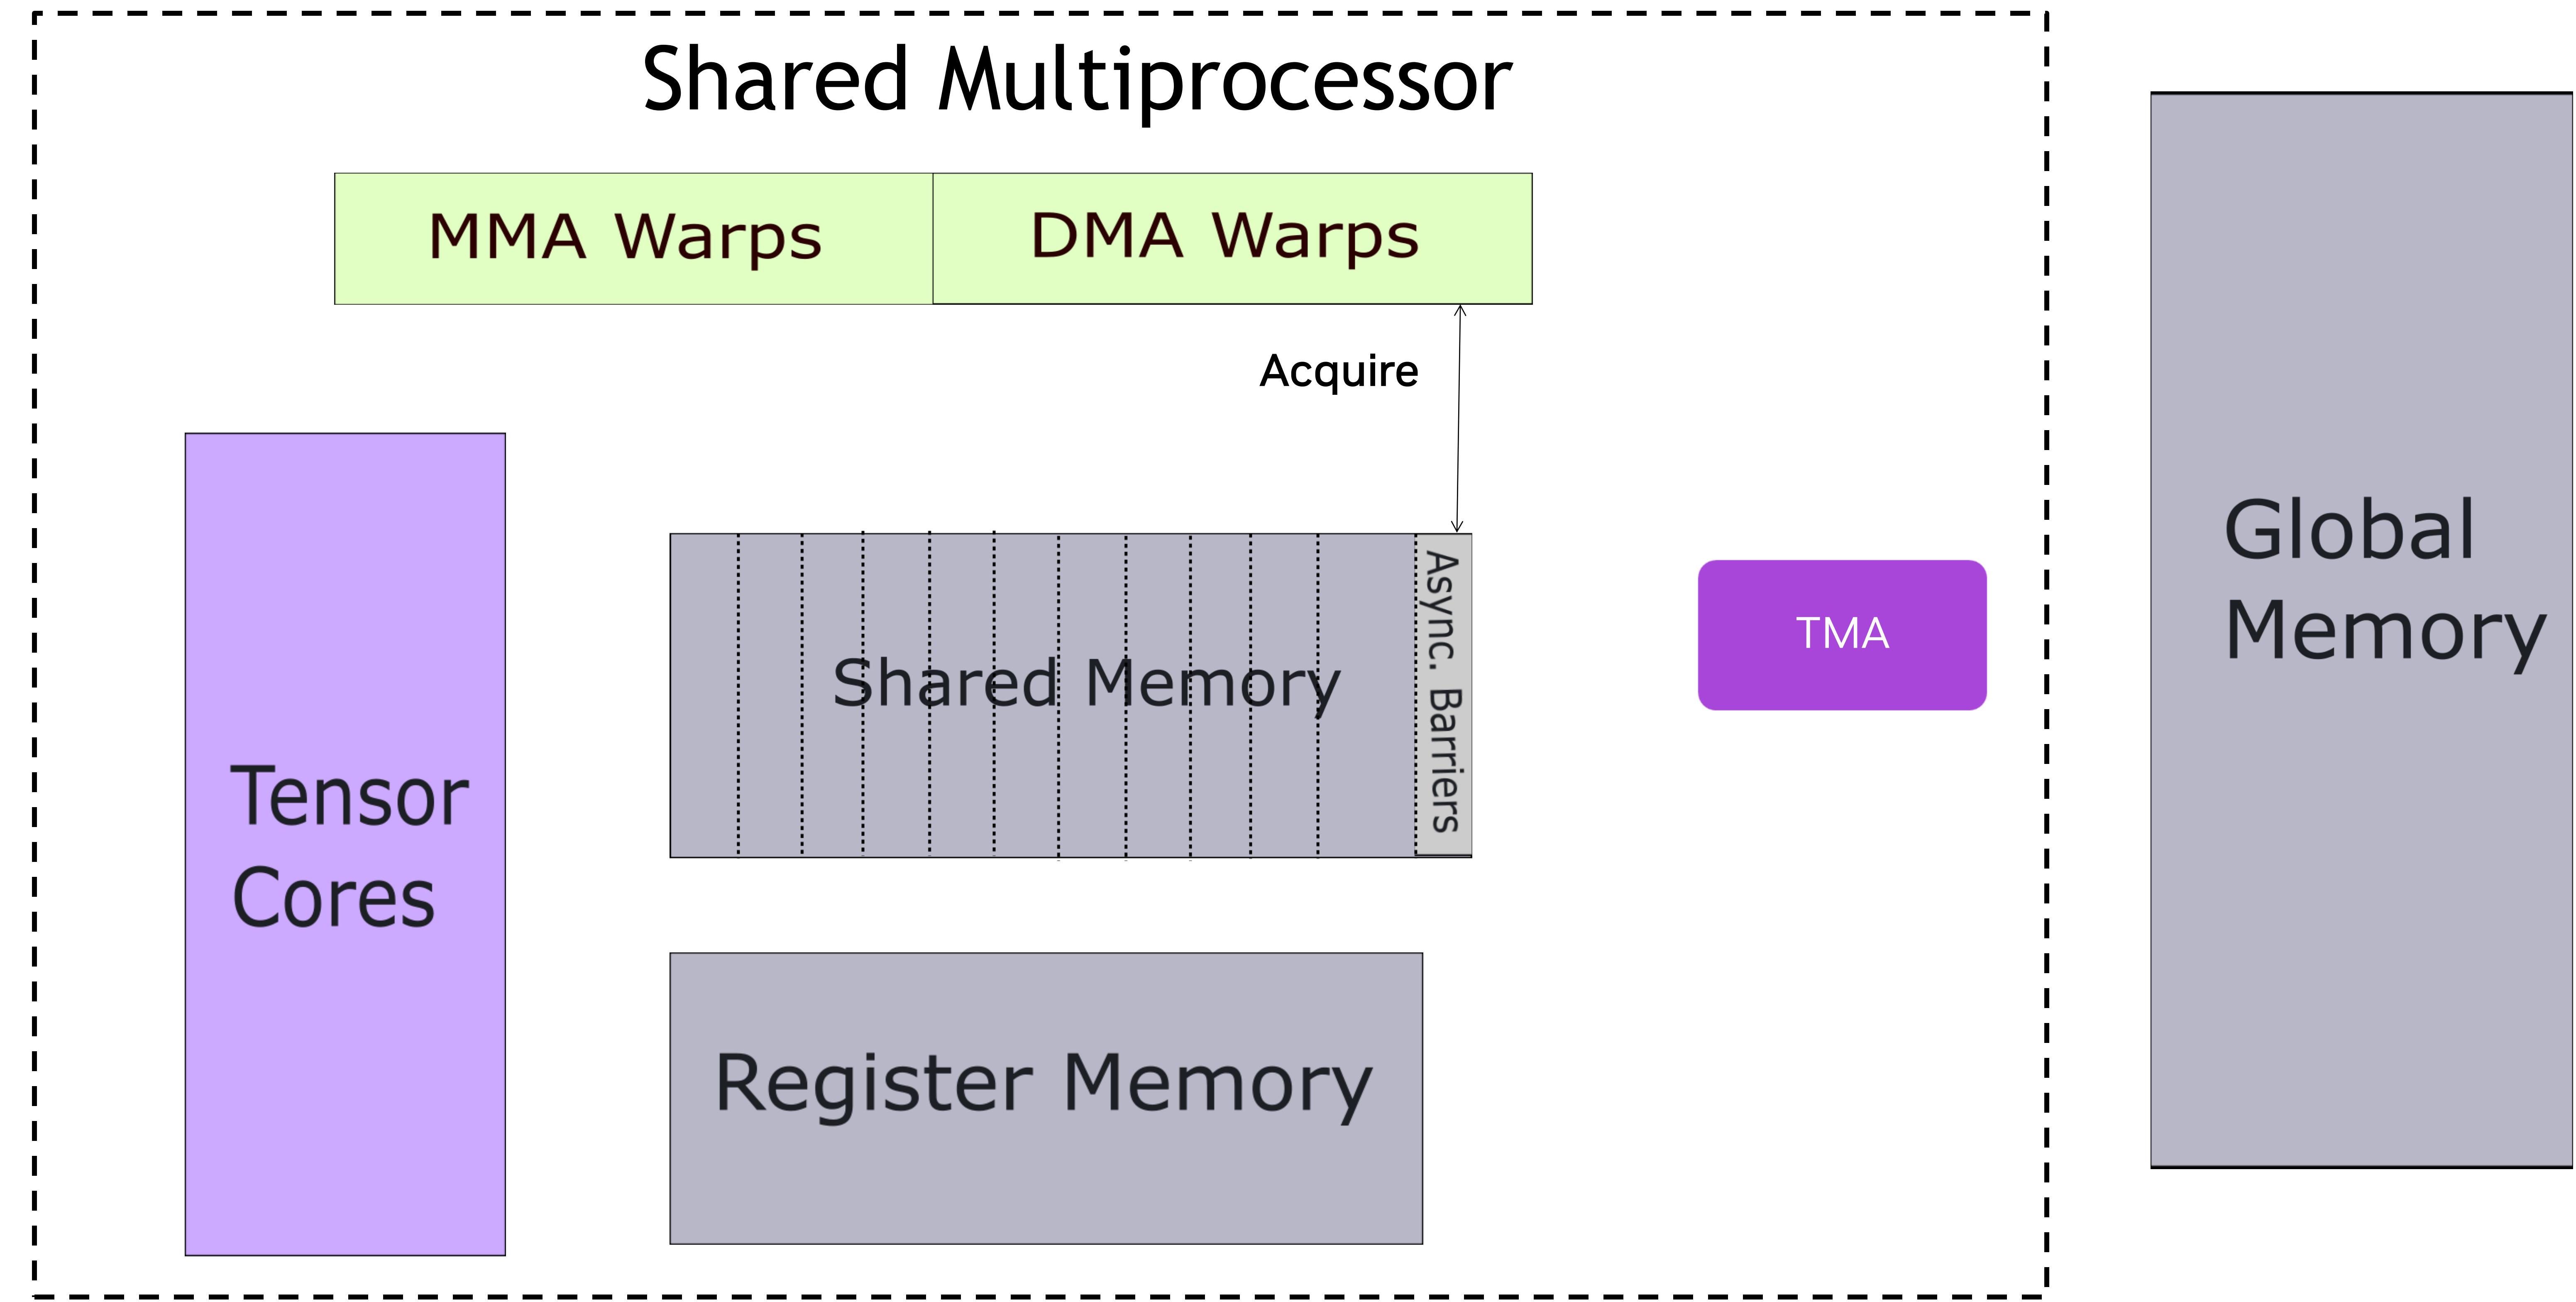


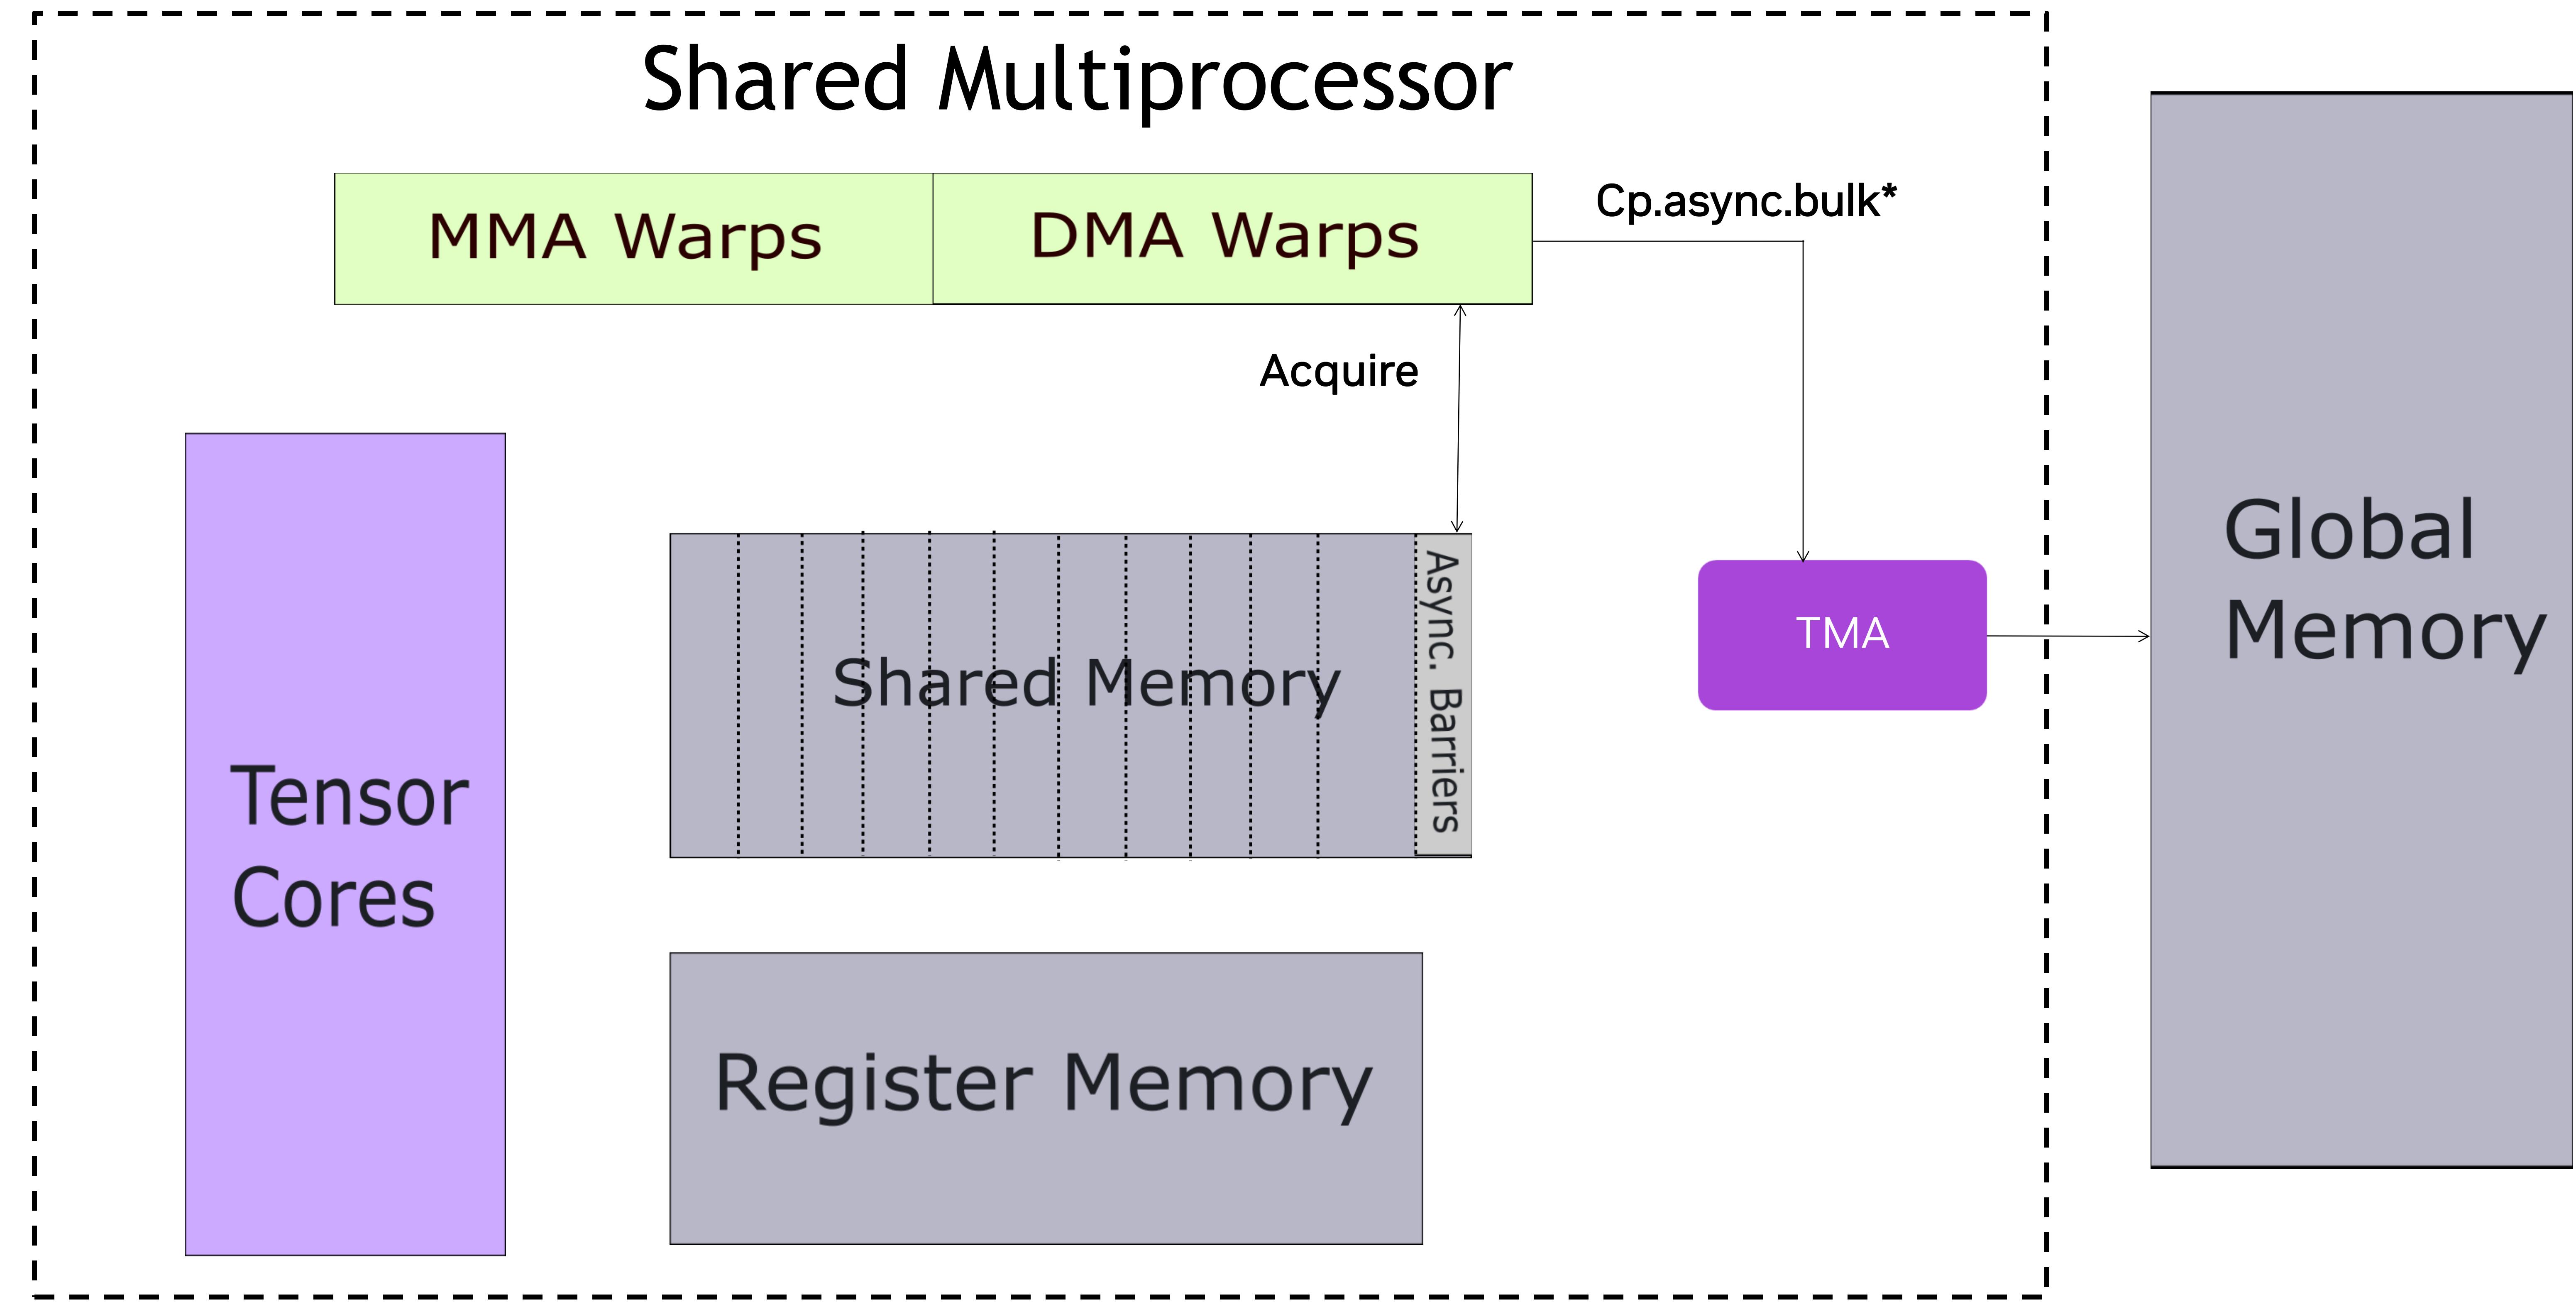
Tensor cores reuse data directly from in shared memory (if possible)

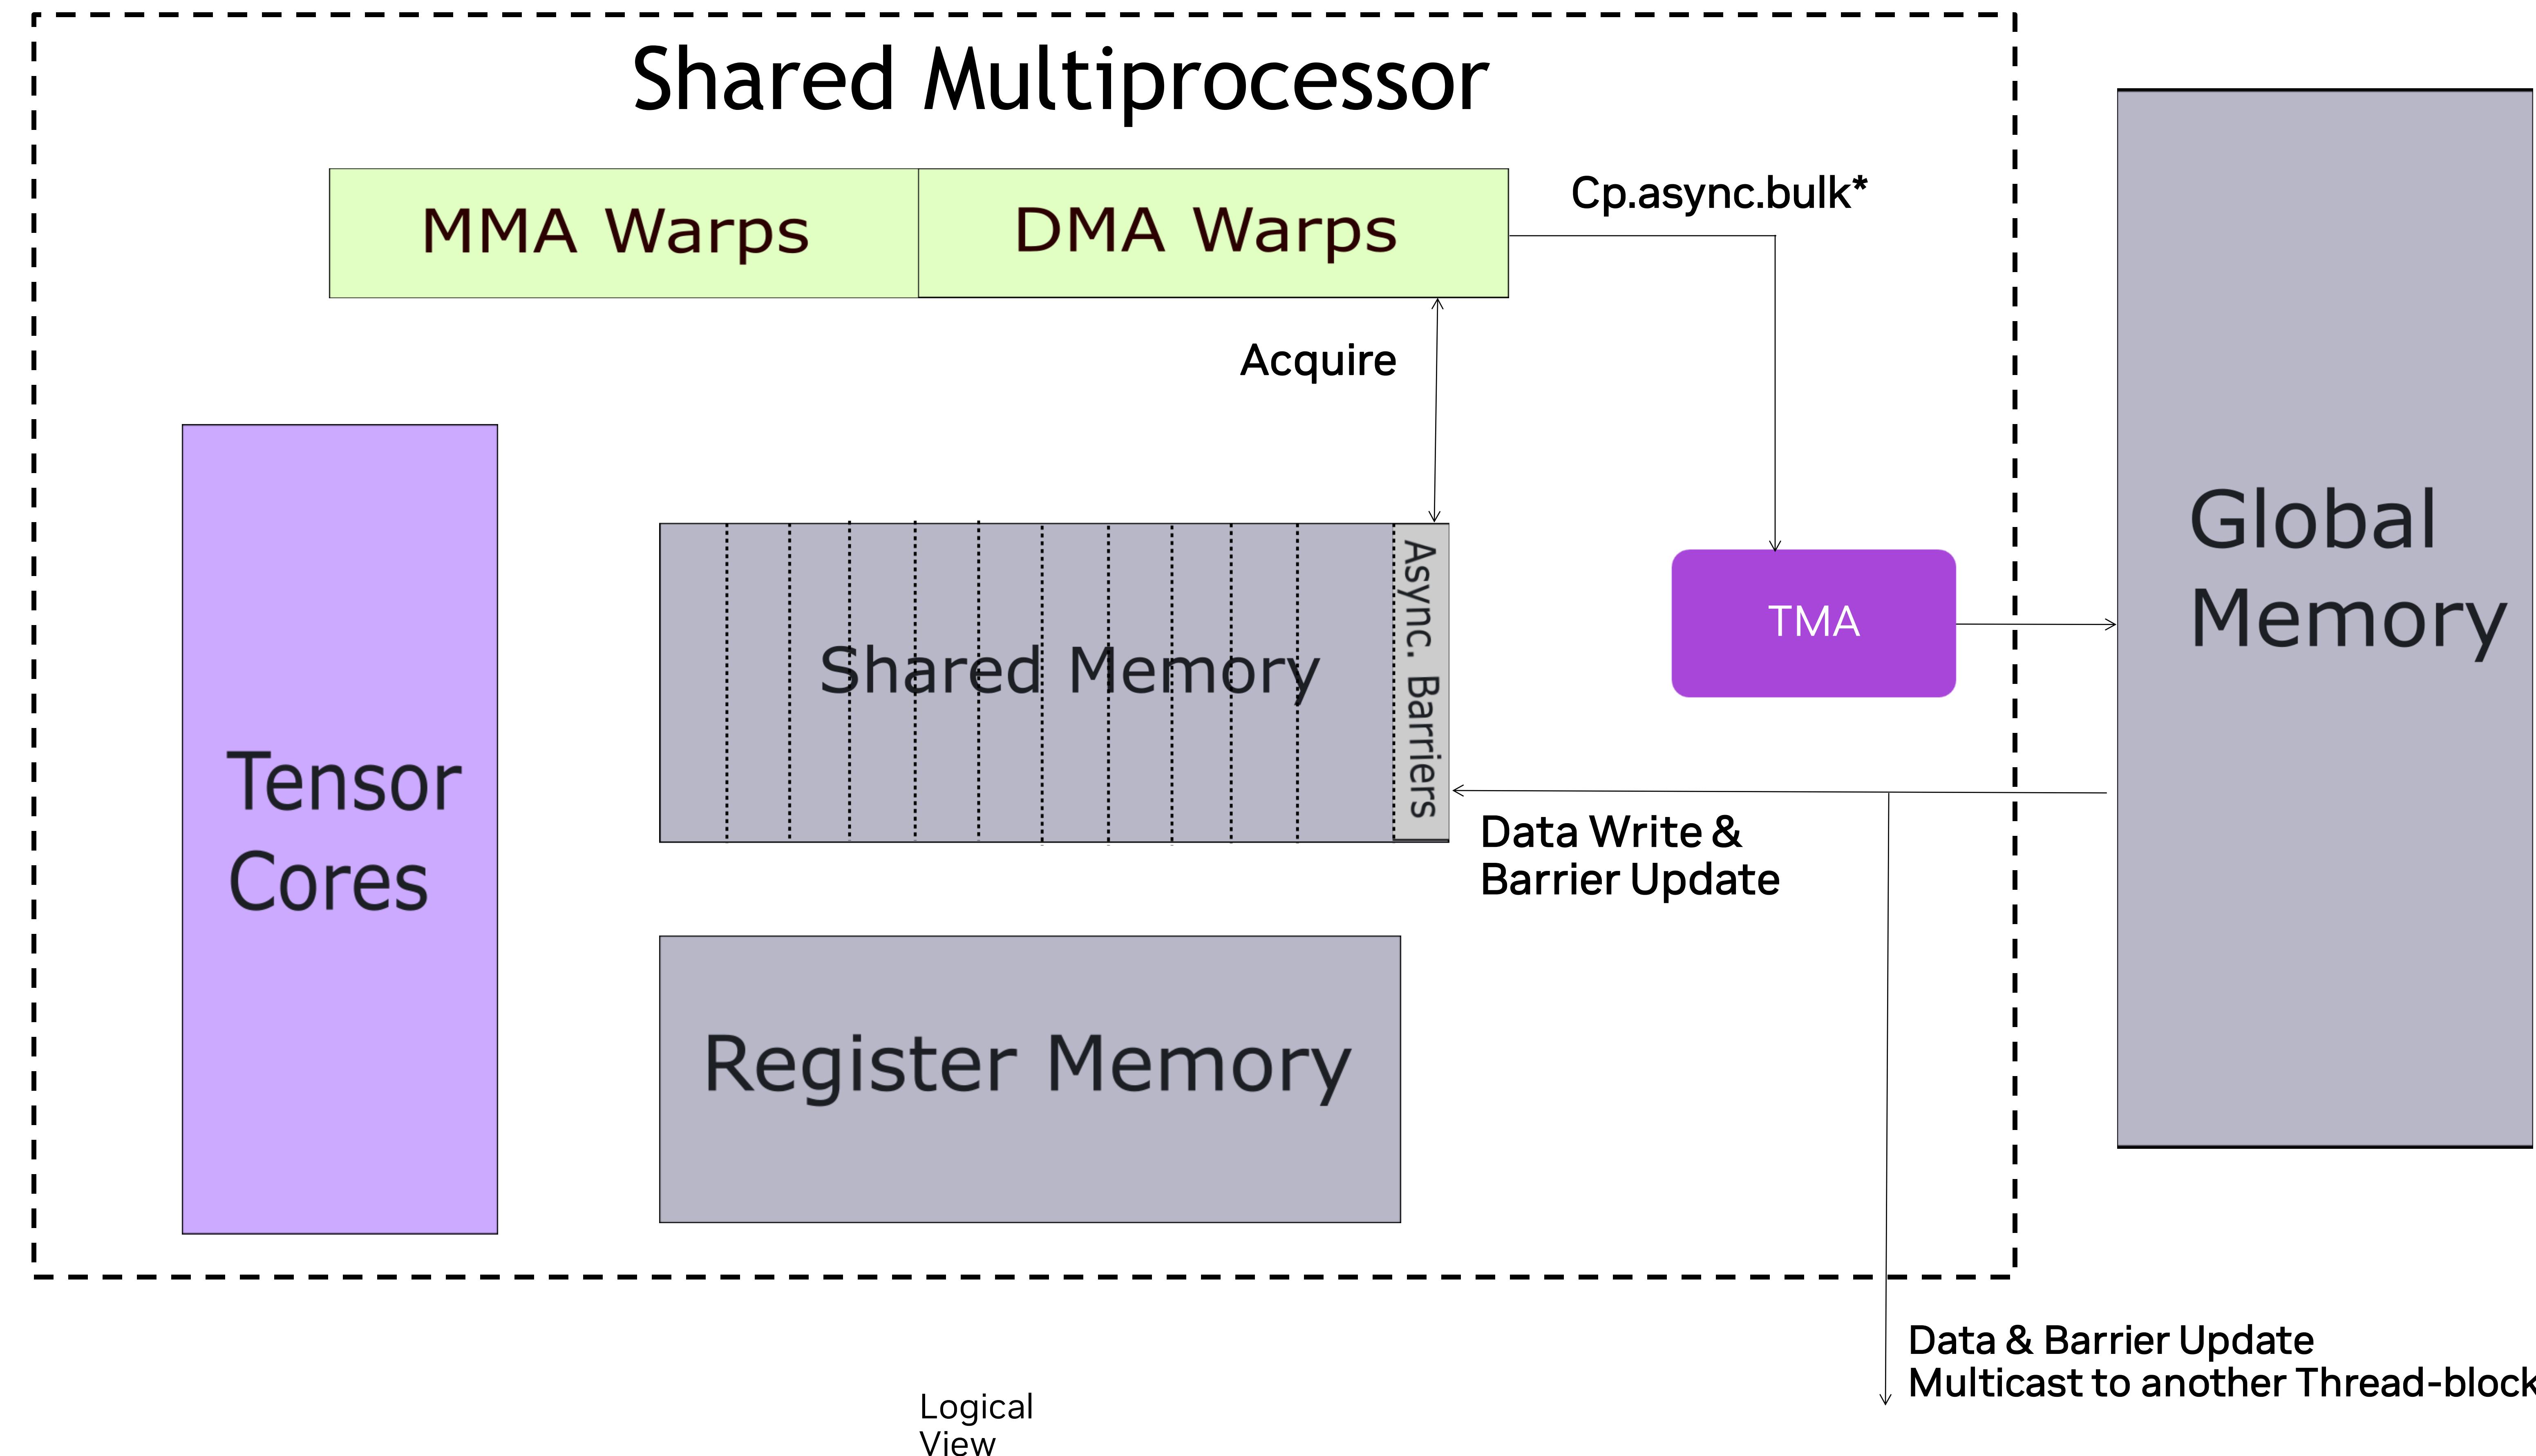


Shared Multiprocessor

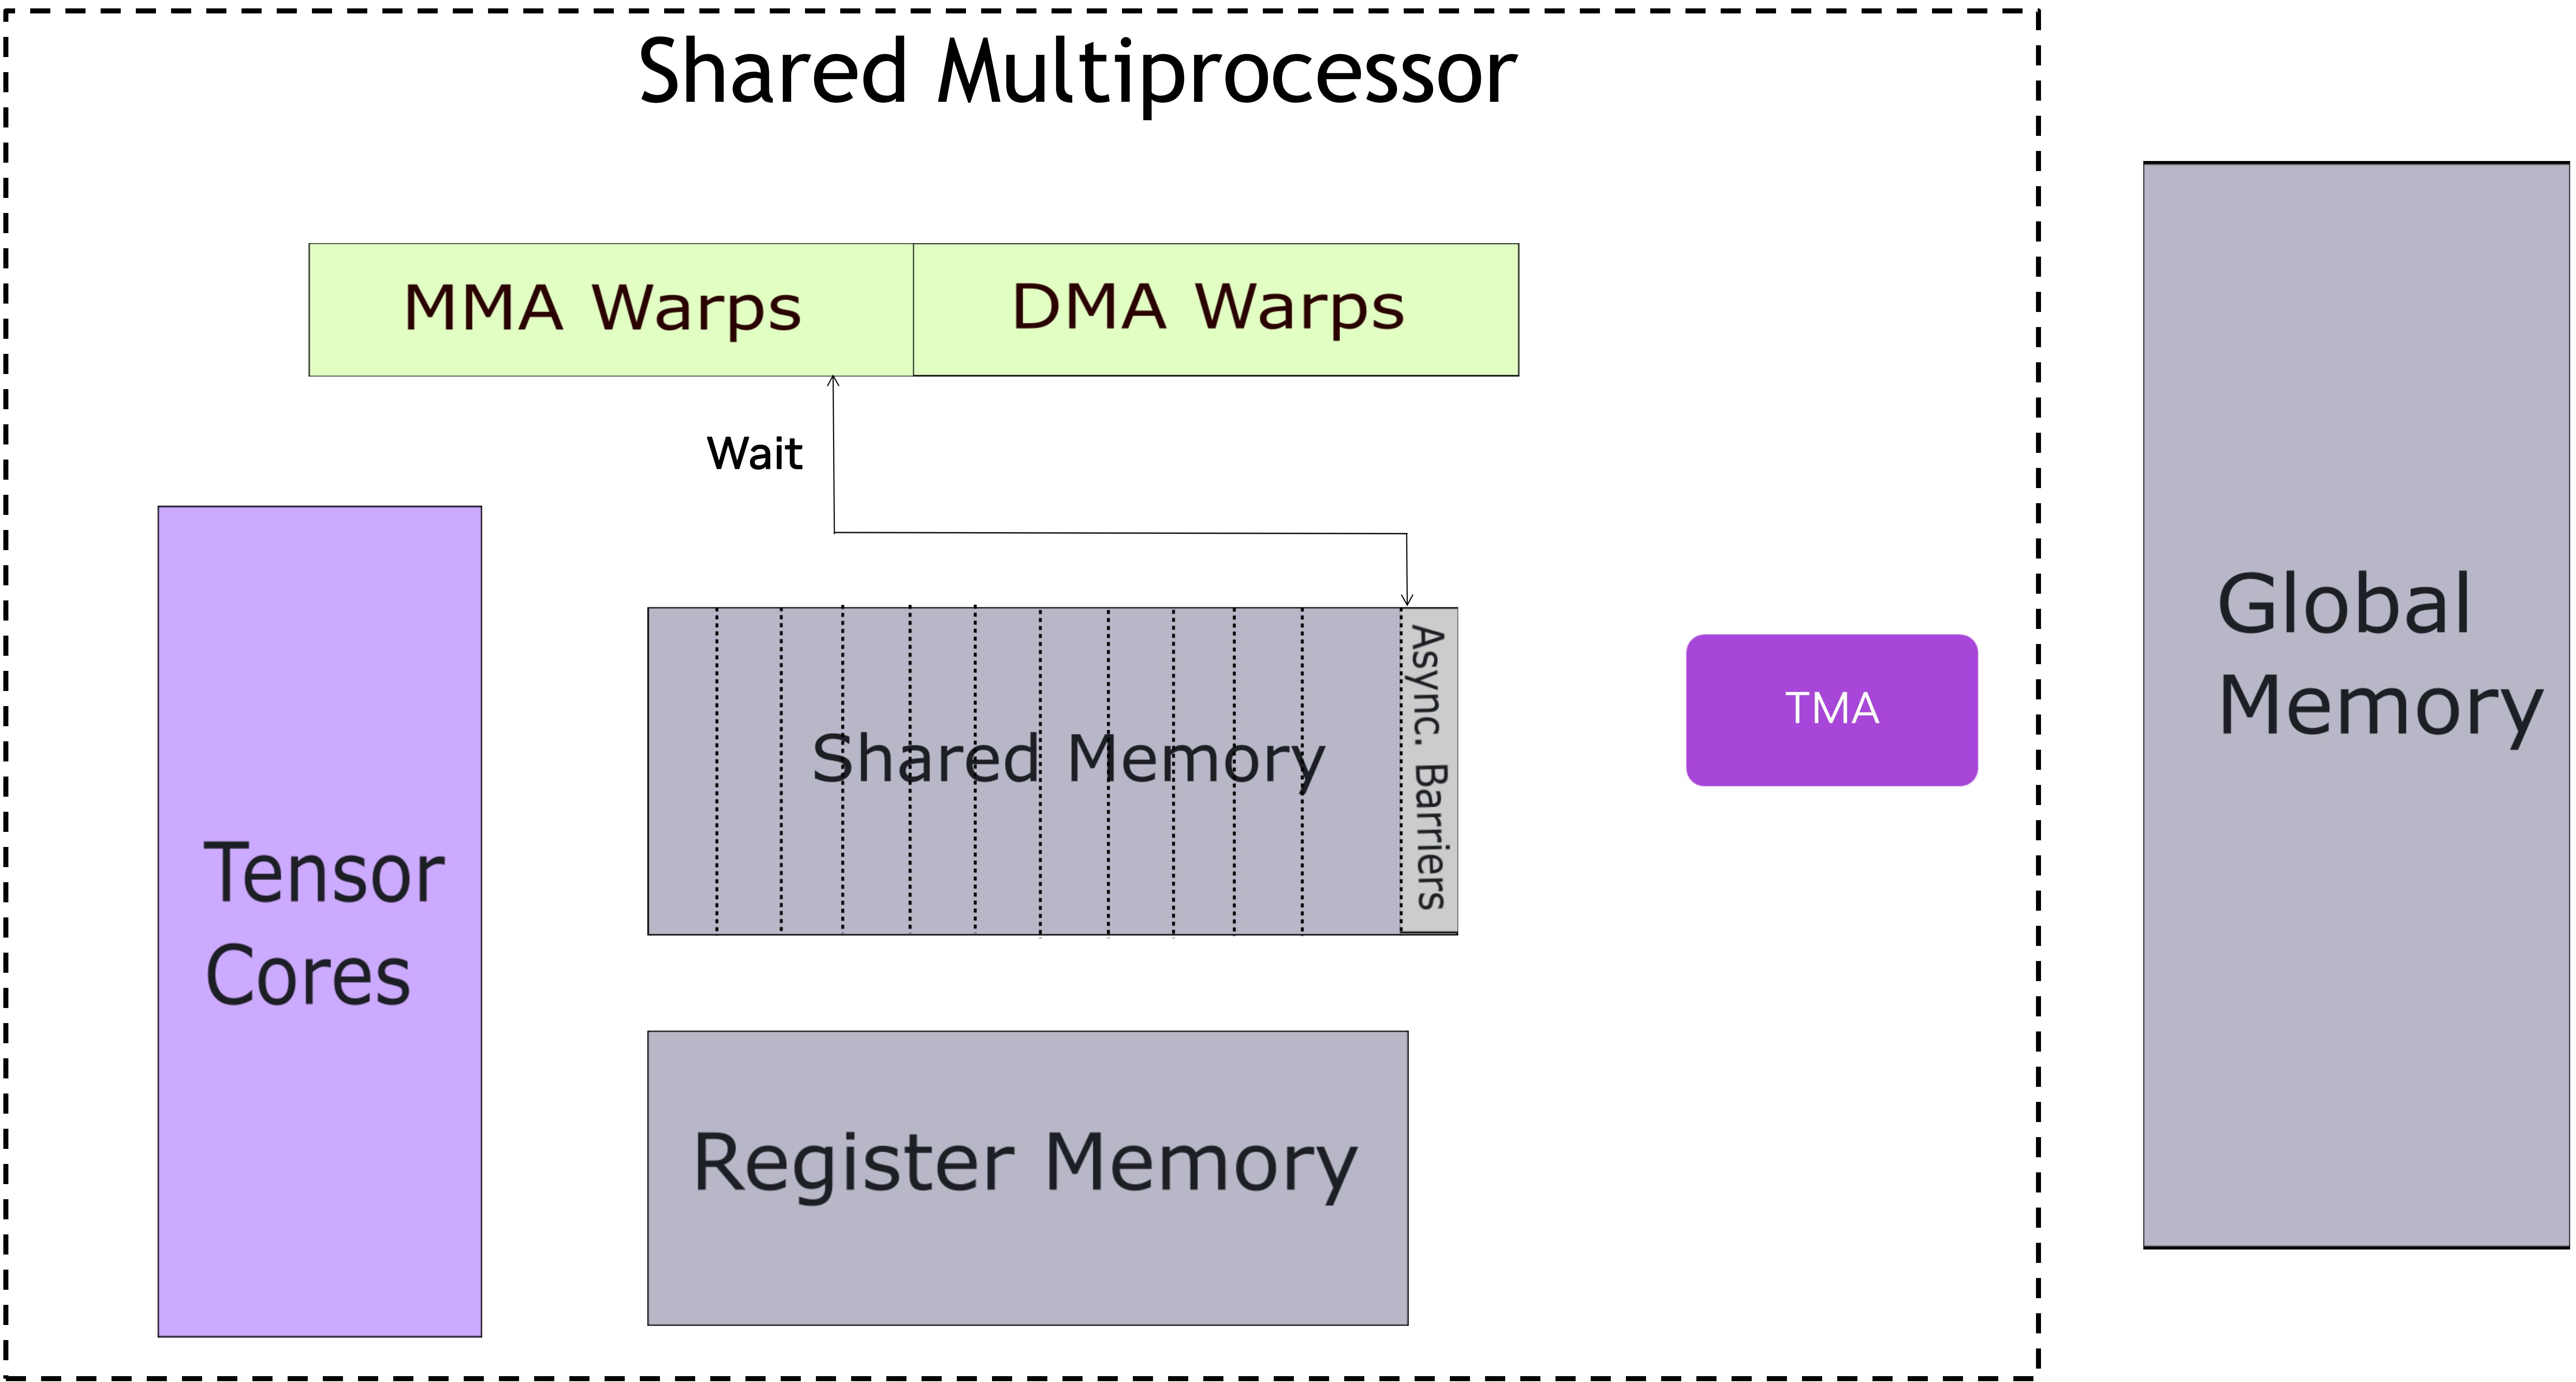




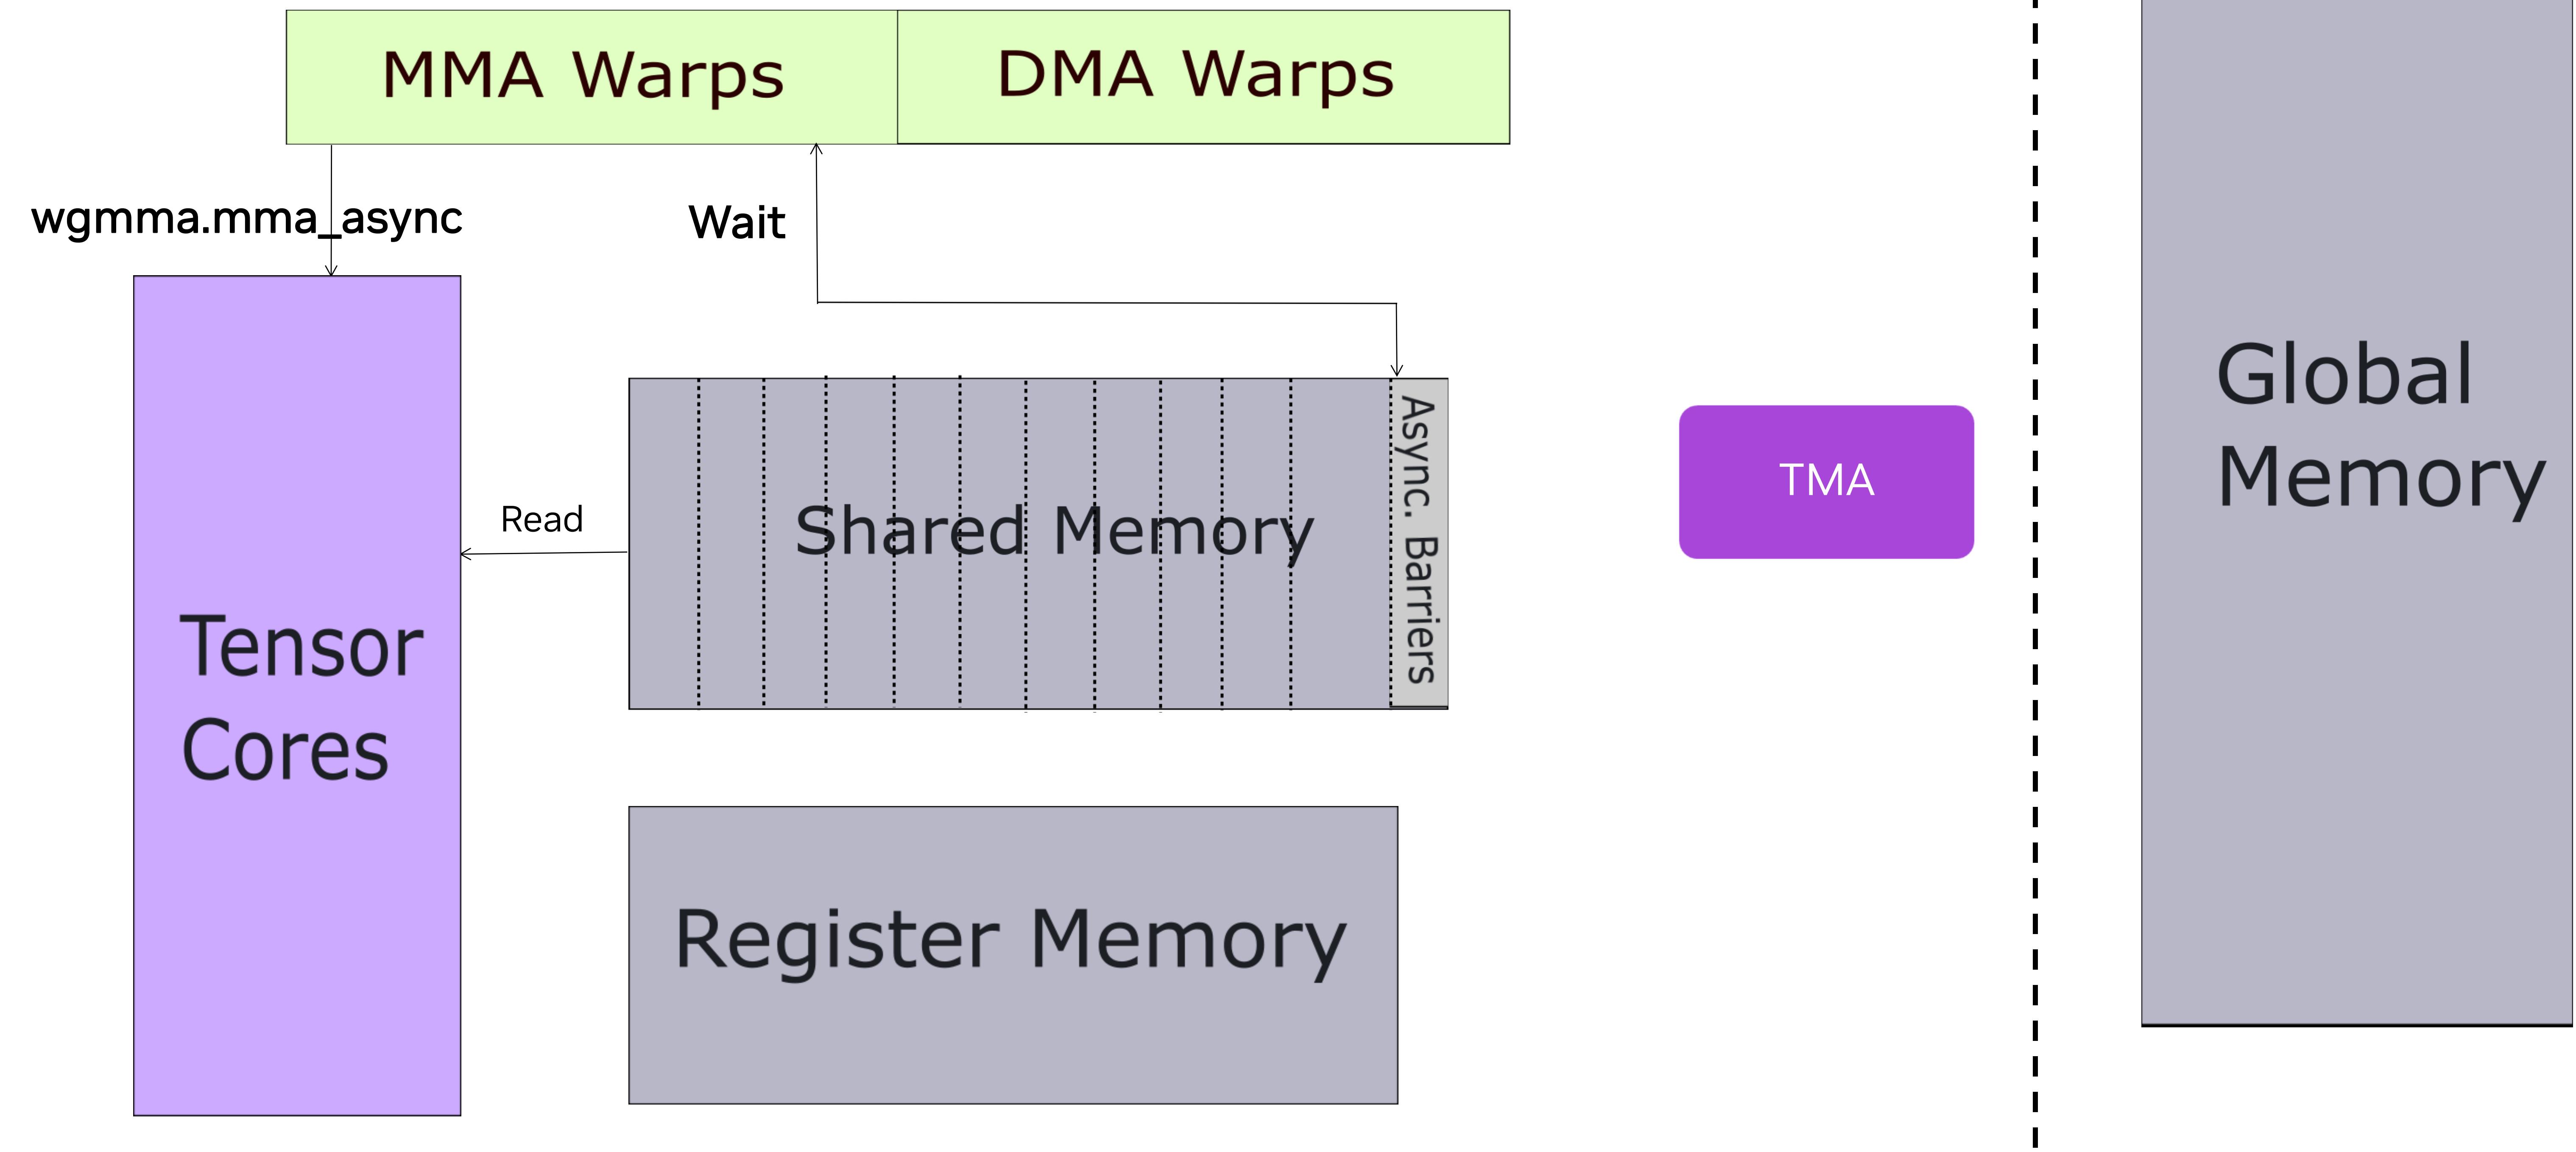




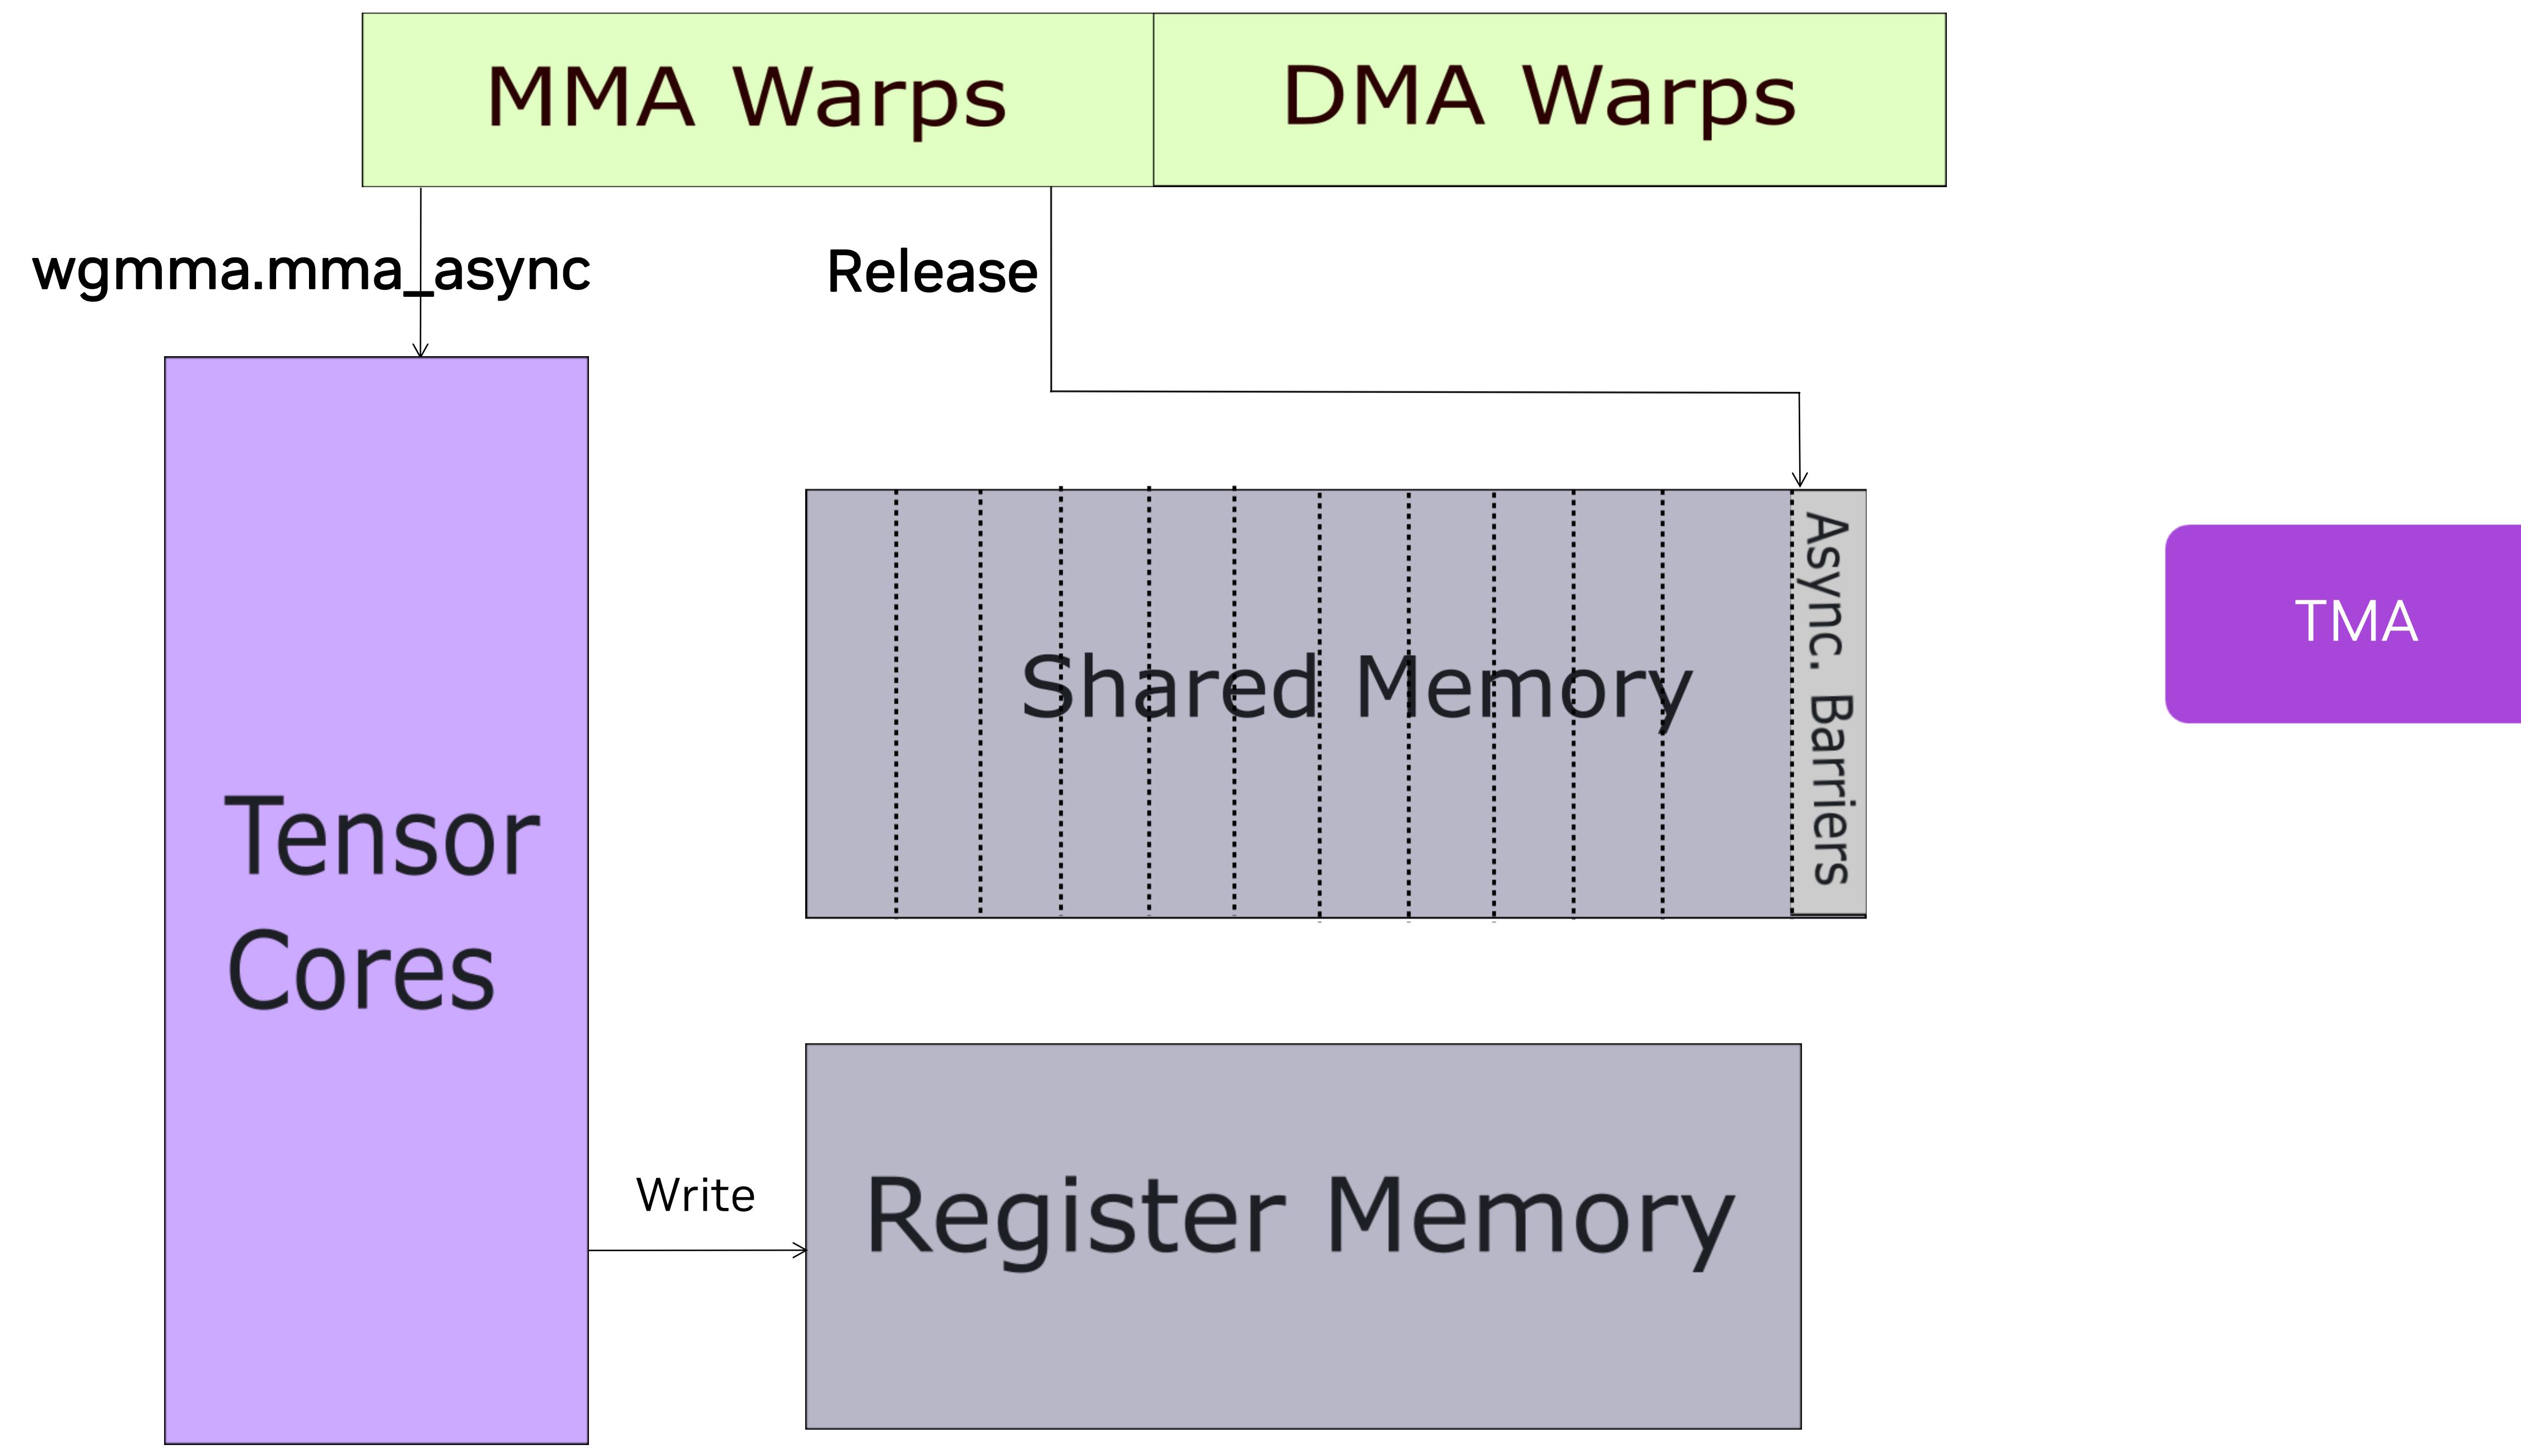
Shared Multiprocessor



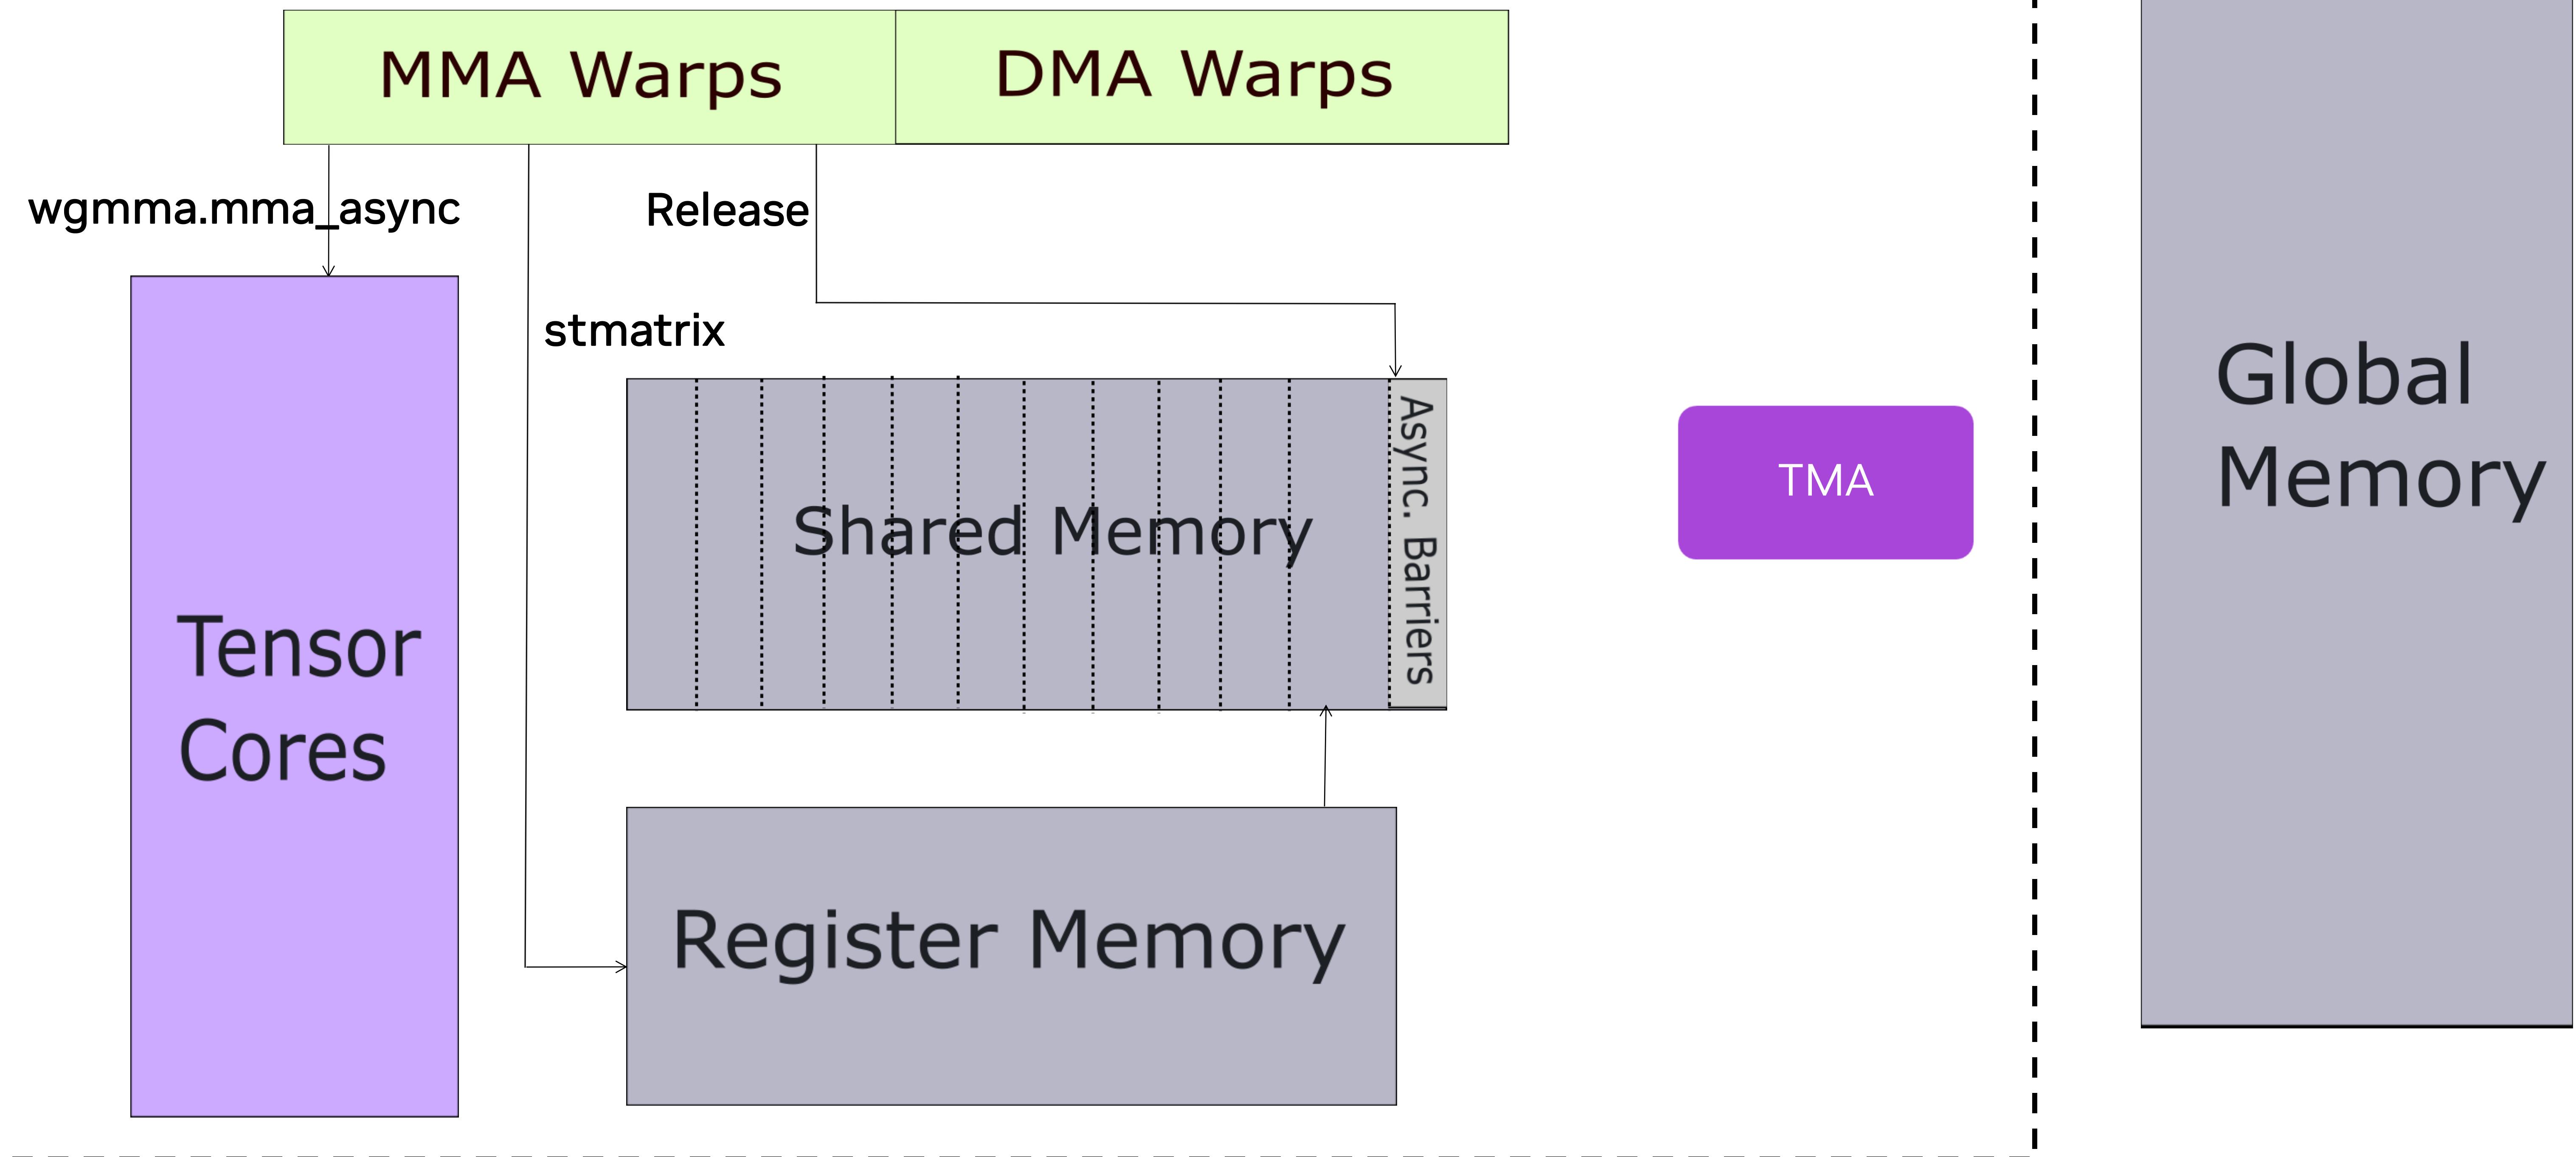
Shared Multiprocessor

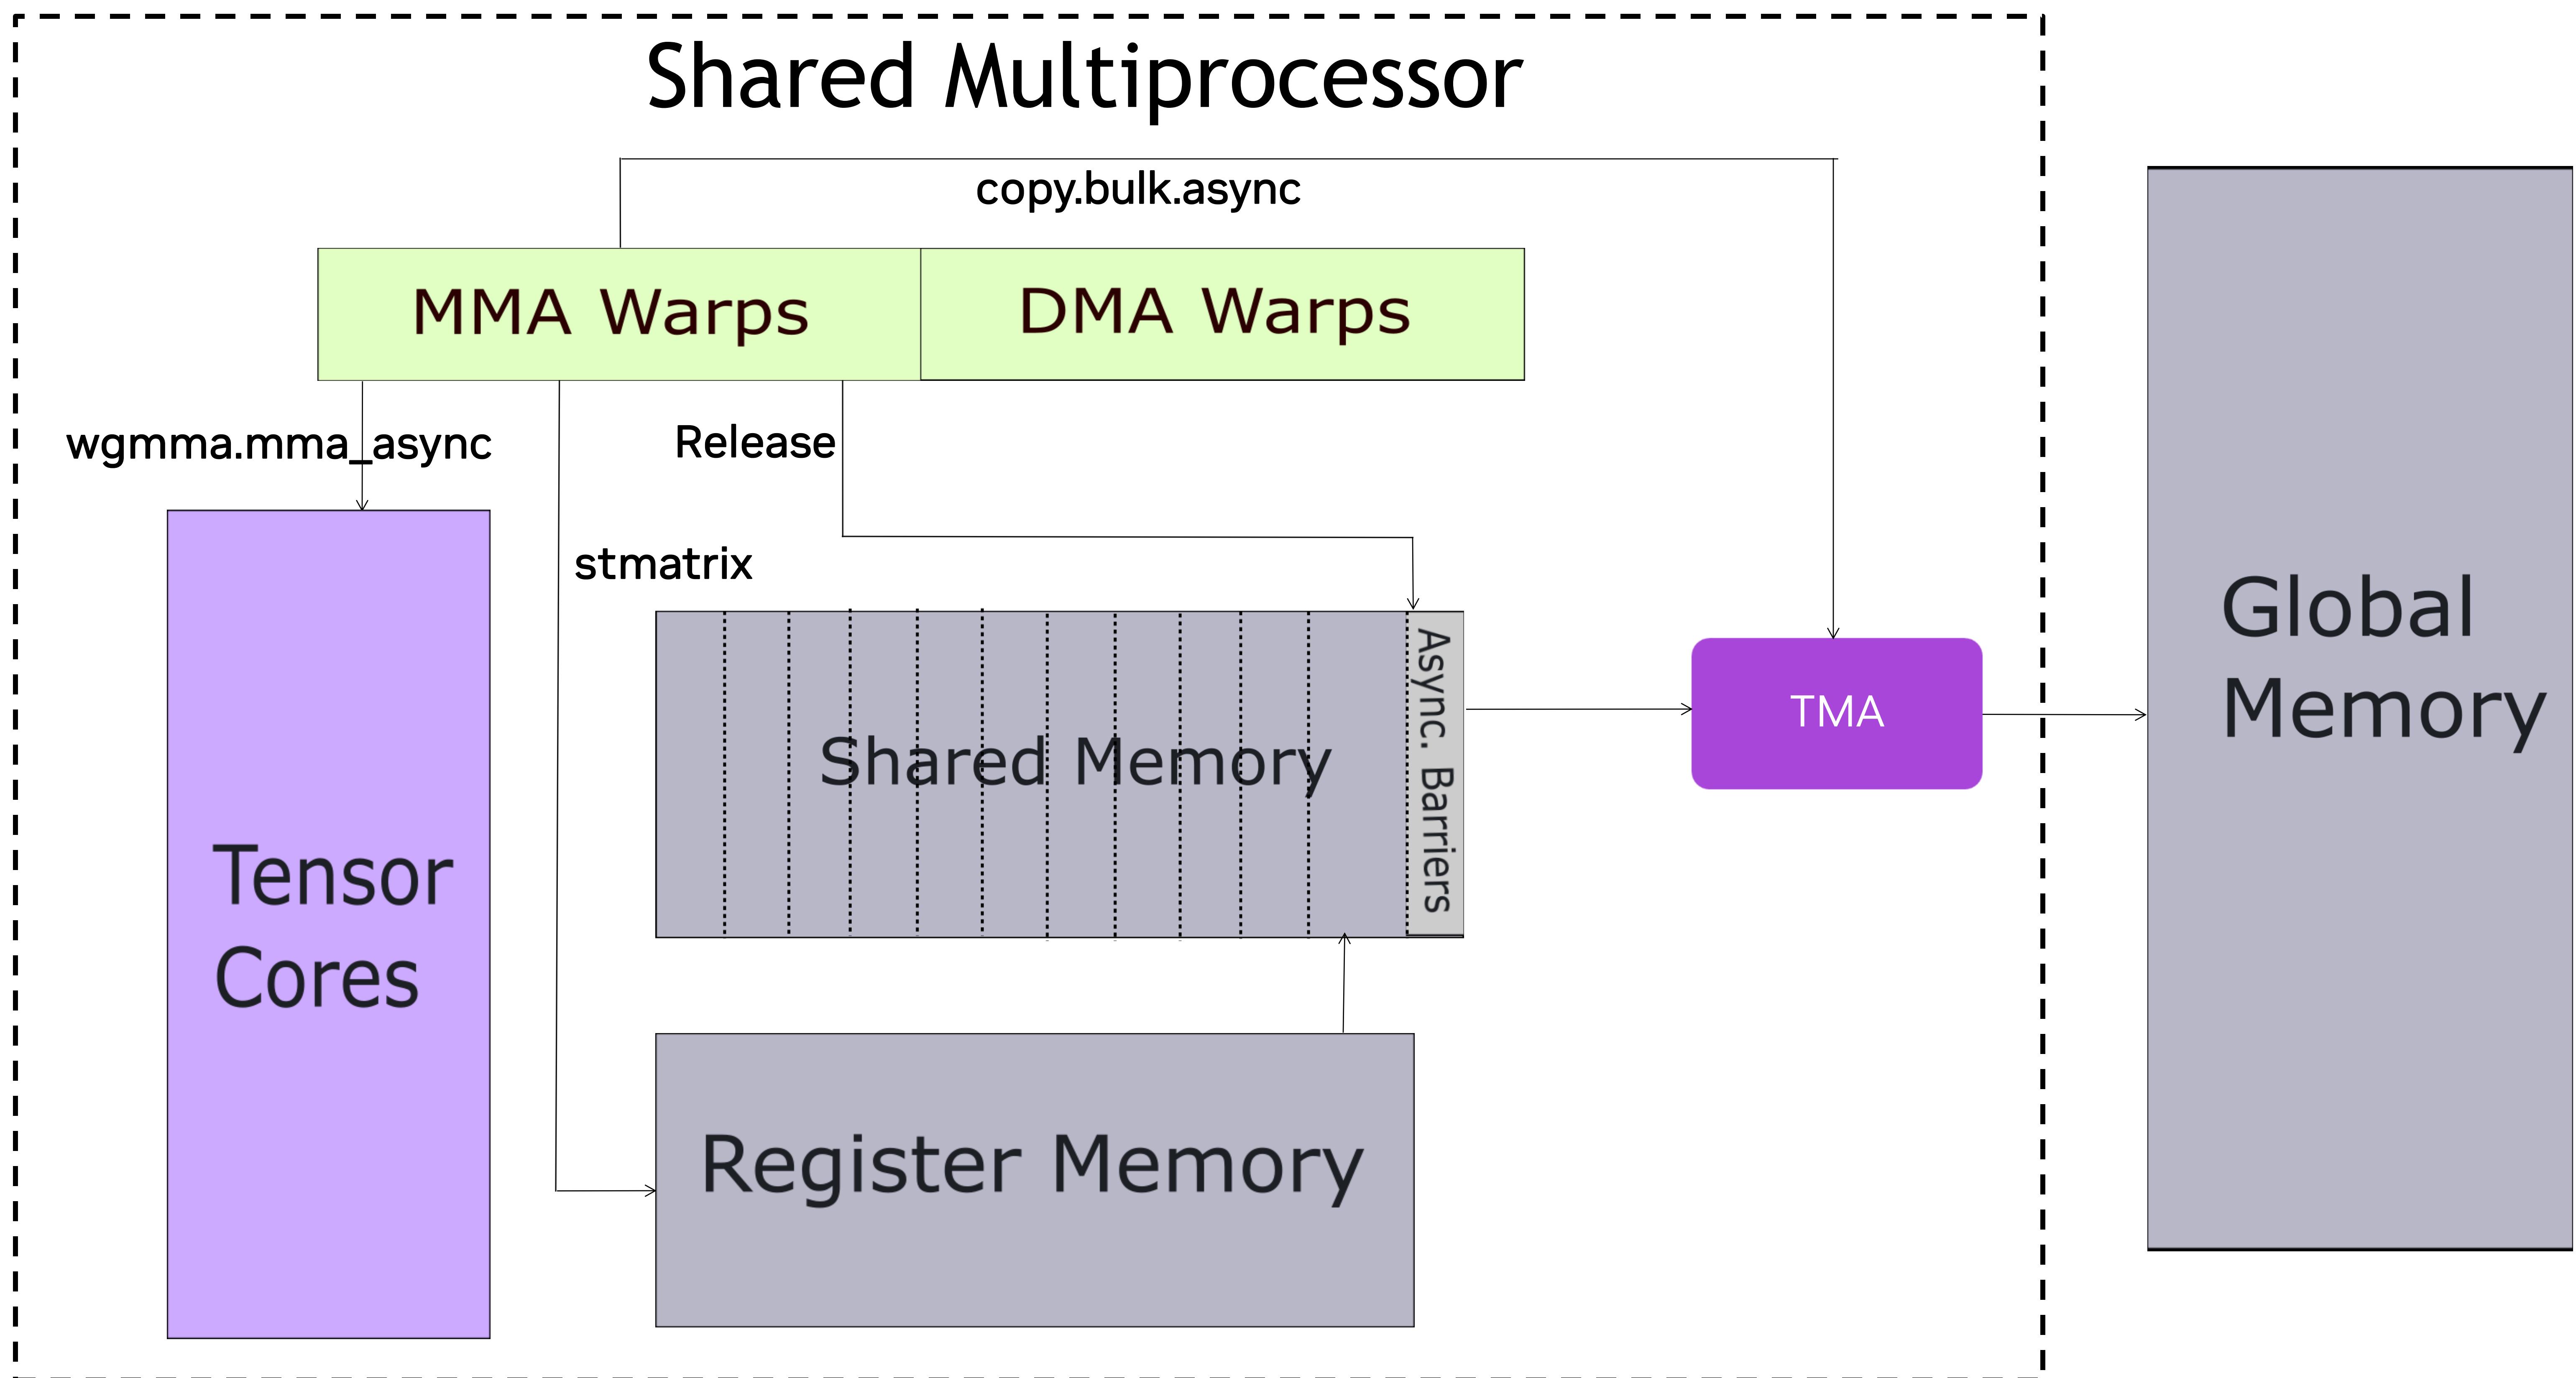


Shared Multiprocessor

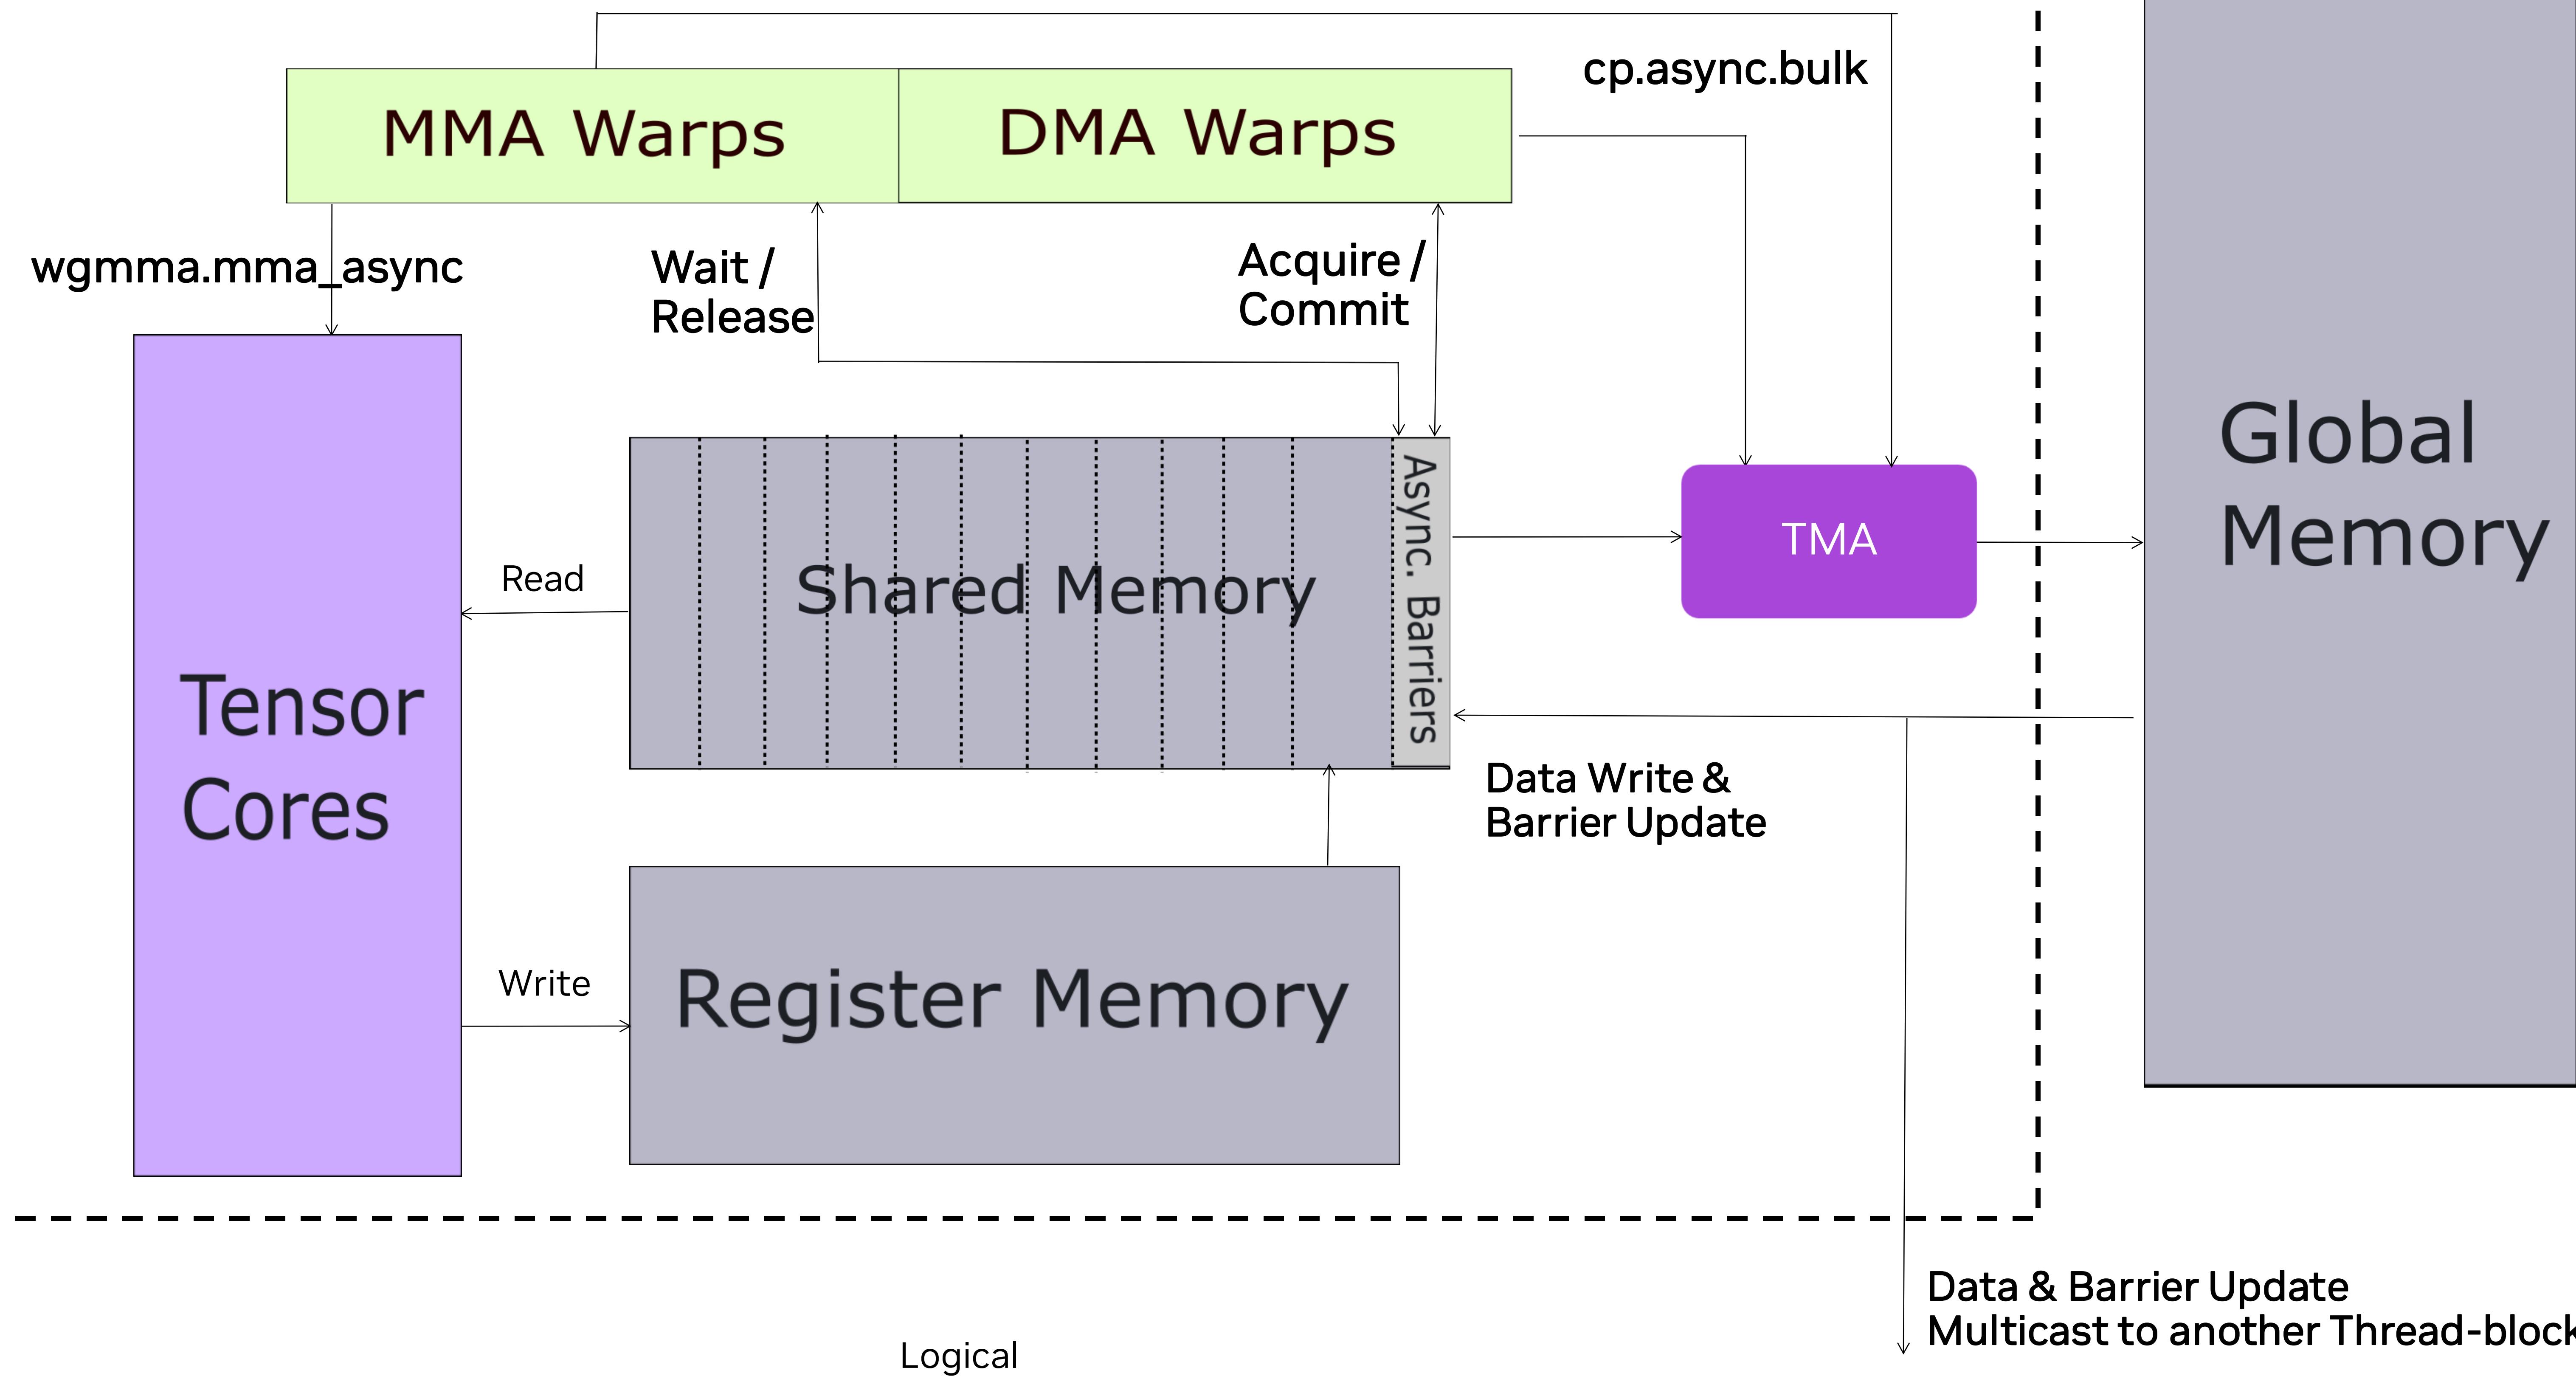


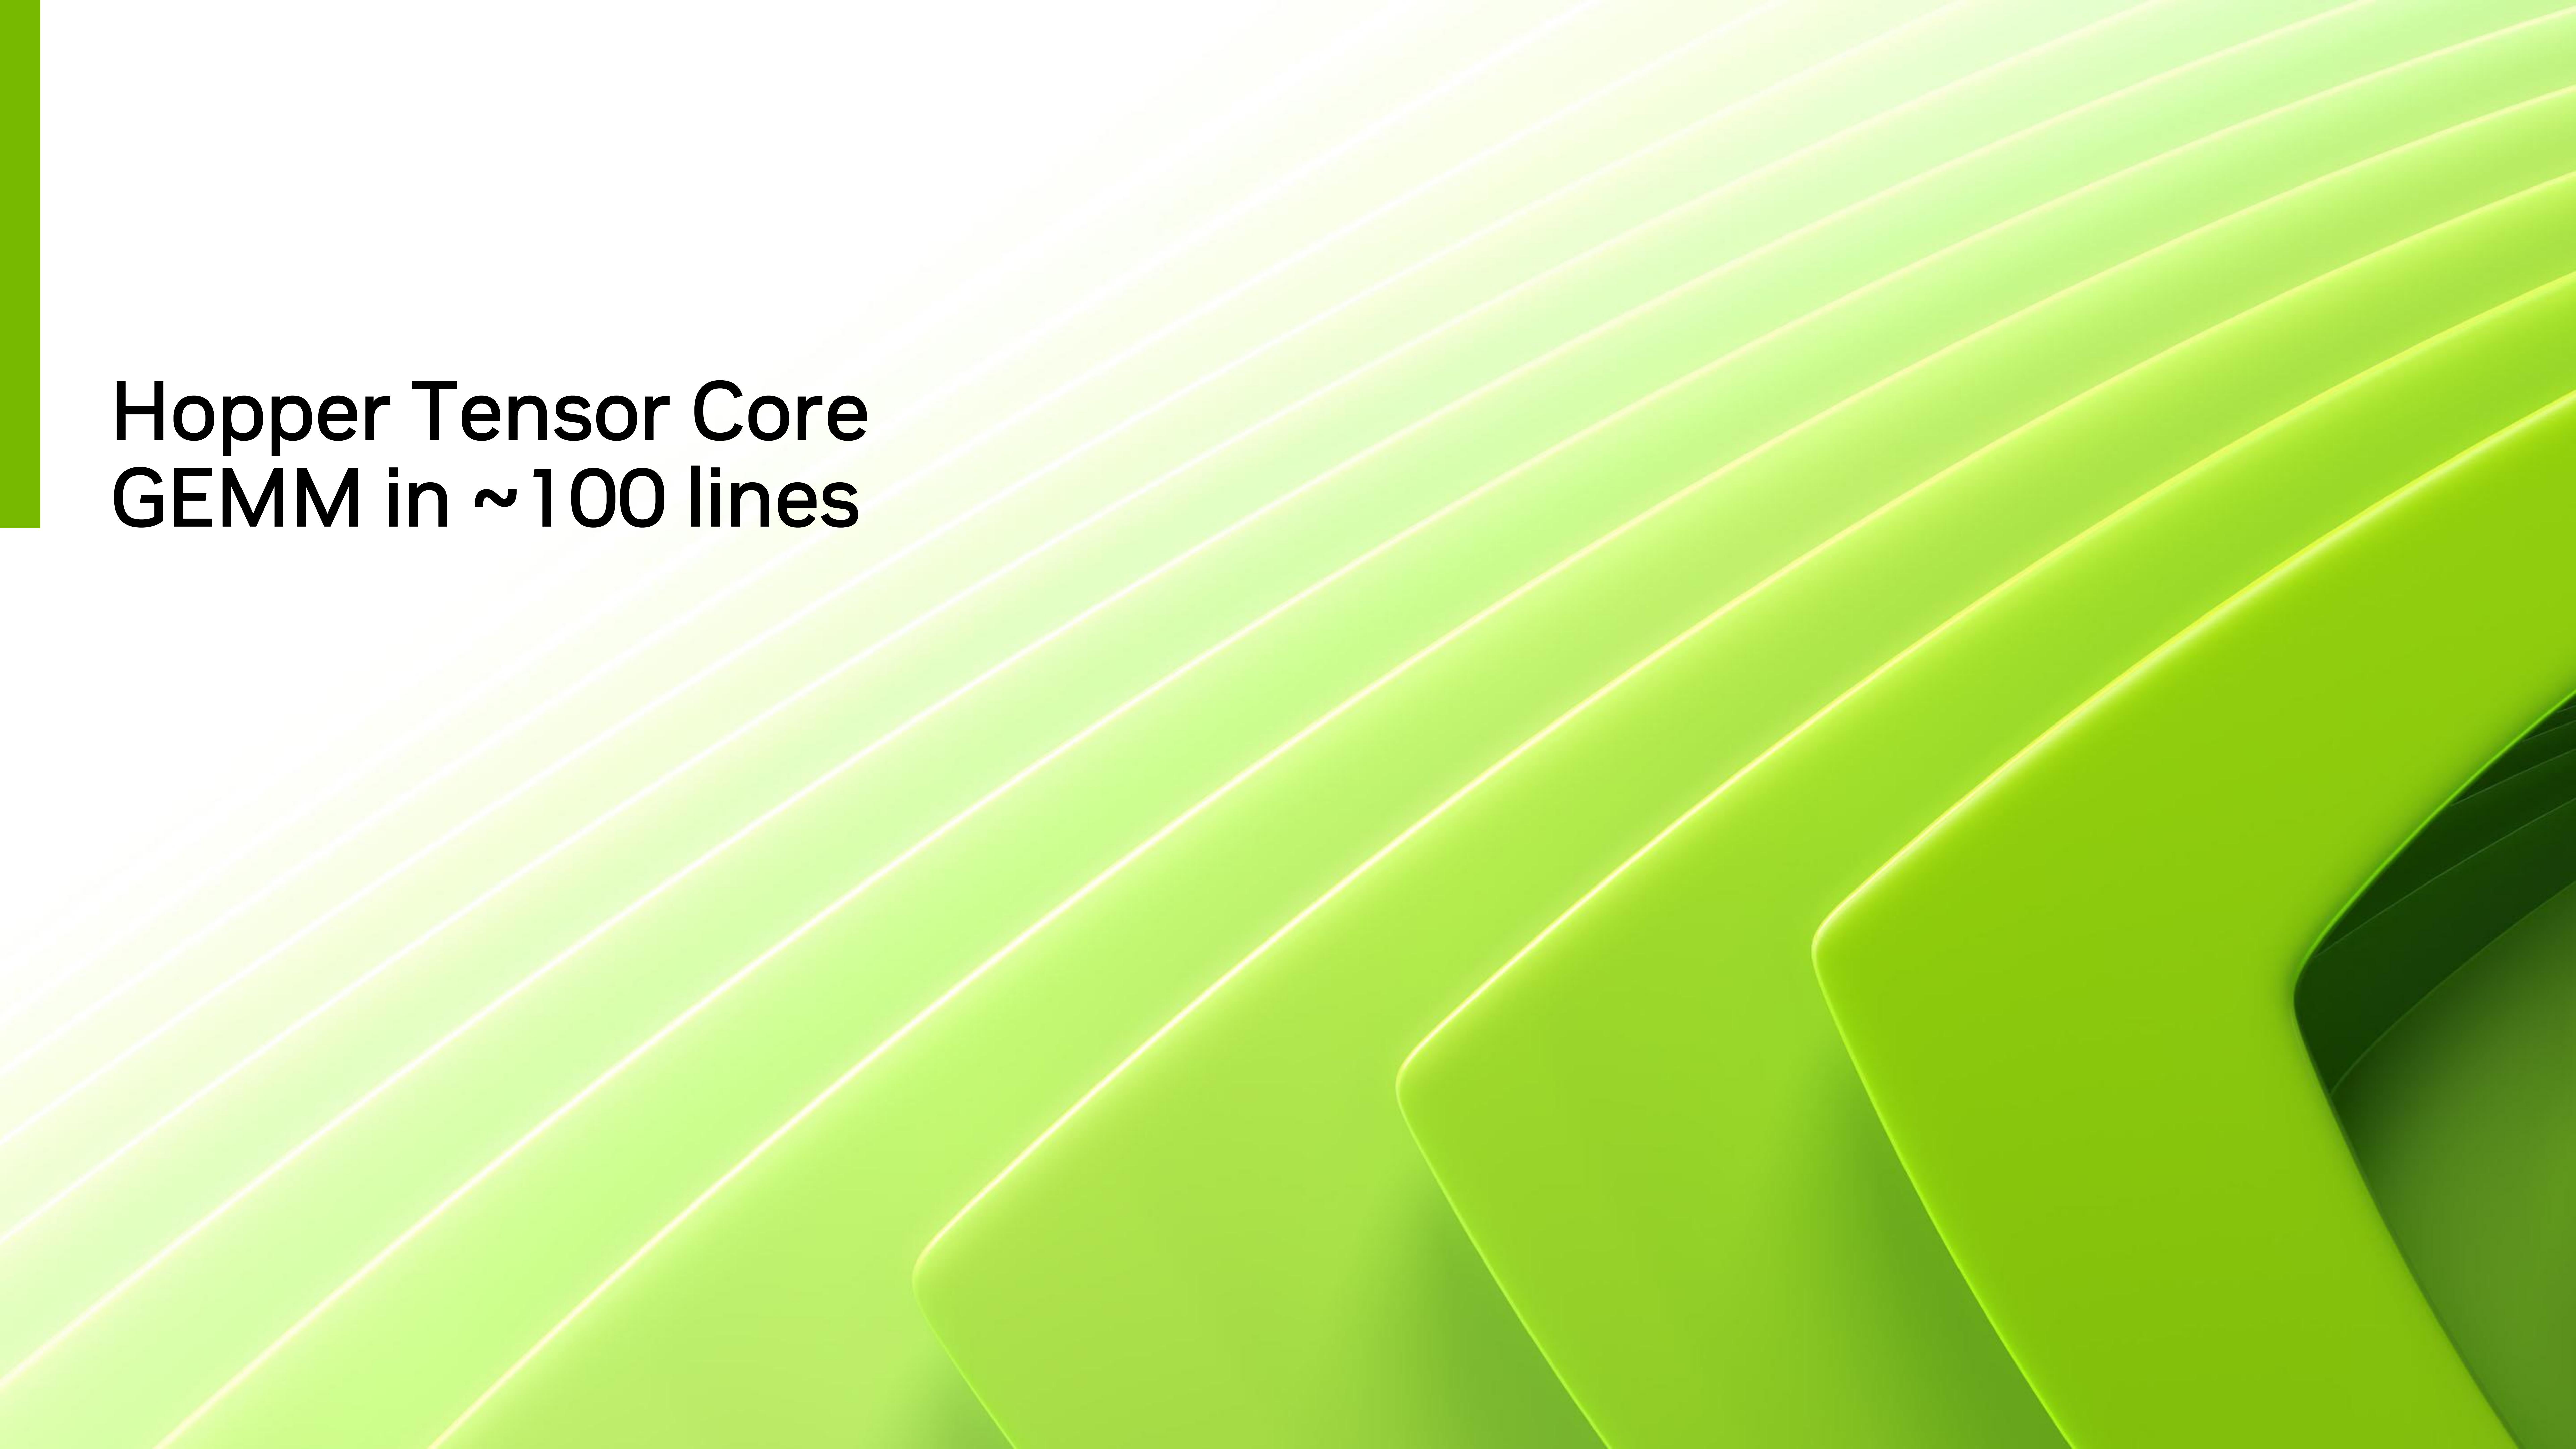
Shared Multiprocessor





Shared Multiprocessor





Hopper Tensor Core
GEMM in ~100 lines

Putting it all together

A basic GEMM kernel in ~100 lines

Code Walk Through

Reading material :

- <https://github.com/NVIDIA/cutlass/blob/main/media/docs/cute/>
- <https://www.nvidia.com/en-us/on-demand/session/gtcspring23-s51413/>
<https://www.nvidia.com/en-us/on-demand/session/gtc24-s61198/>
- <https://github.com/NVIDIA/cutlass/pull/1578>
-

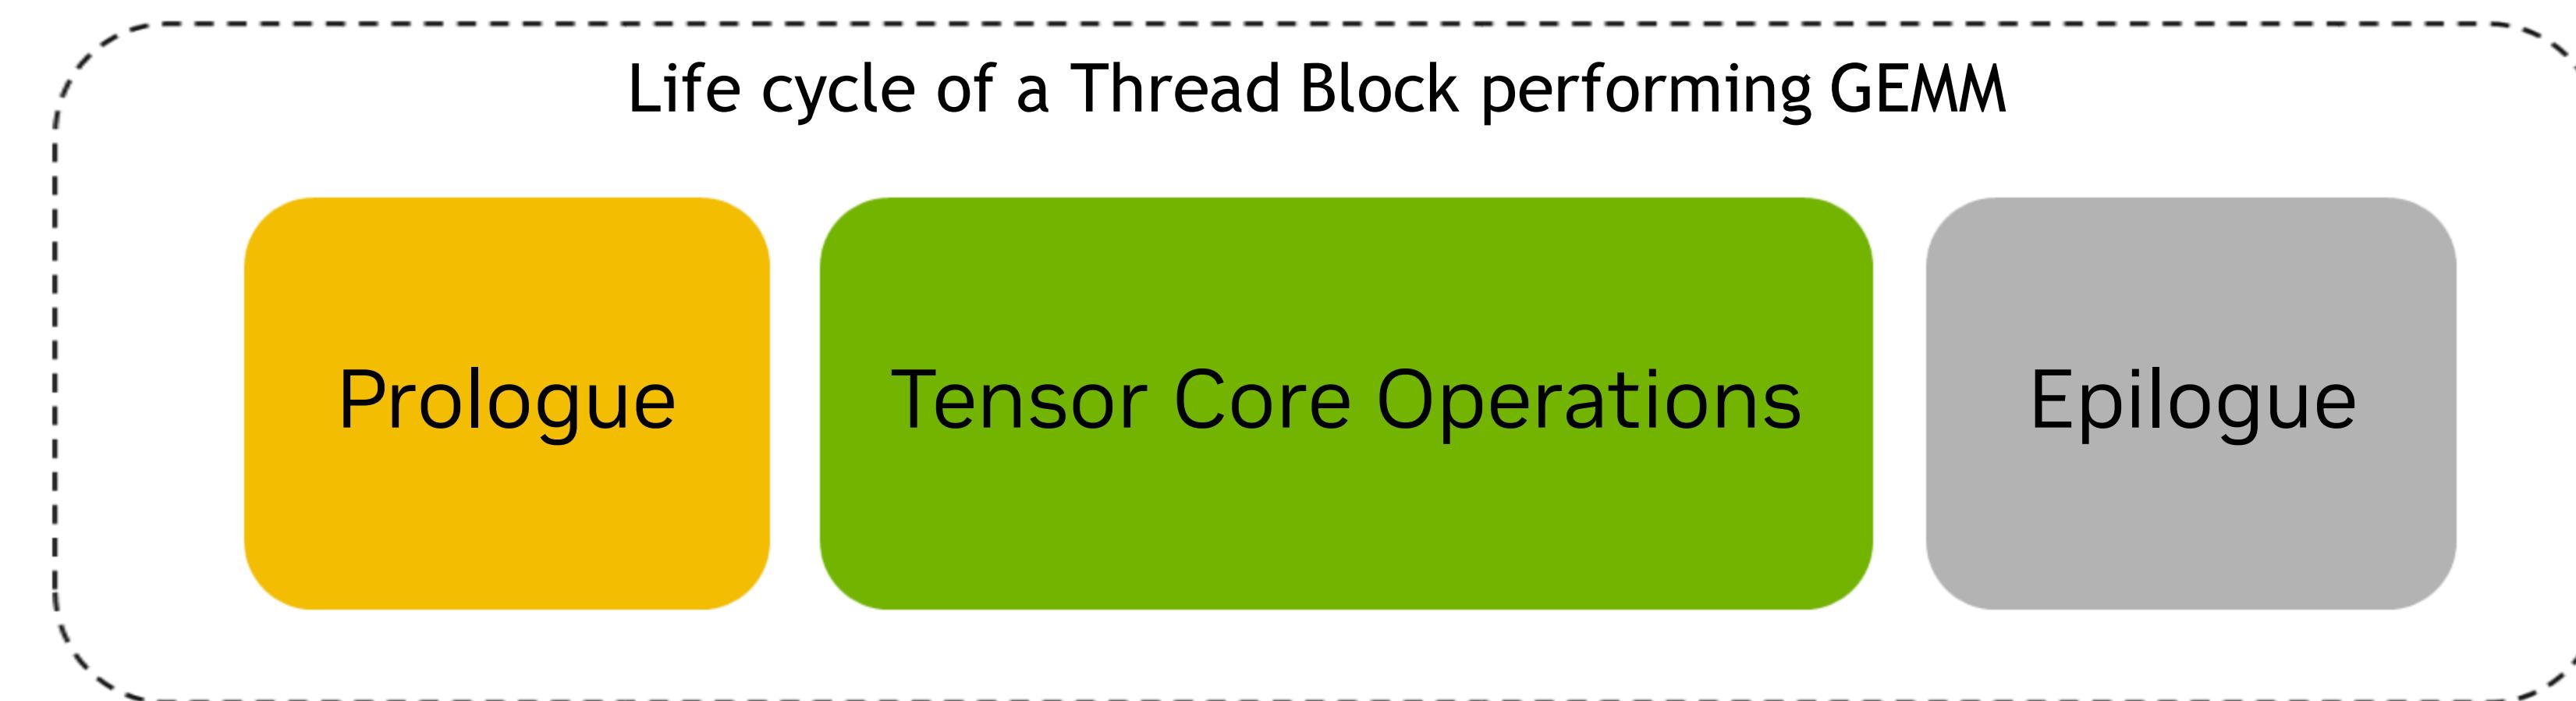
GEMM Advance optimizations

Anatomy of a High Performance Hopper GEMM

201 Level Optimizations

- Warp Specialization
 - Optimal memory and compute throughput across a wide variety of scenarios.
- Persistent schedules
 - Amortizes hardware and software pipelining fixed costs in a kernel
- Epilogue hiding in the shadow of MMAs (a.k.a Ping-Pong Schedule)
 - Critical for small yet compute bound GEMMs
- Maximizing tile size to improve efficiency (a.k.a Cooperative Schedule)
 - Better handling of register pressure
- Efficient usage of synchronization primitives

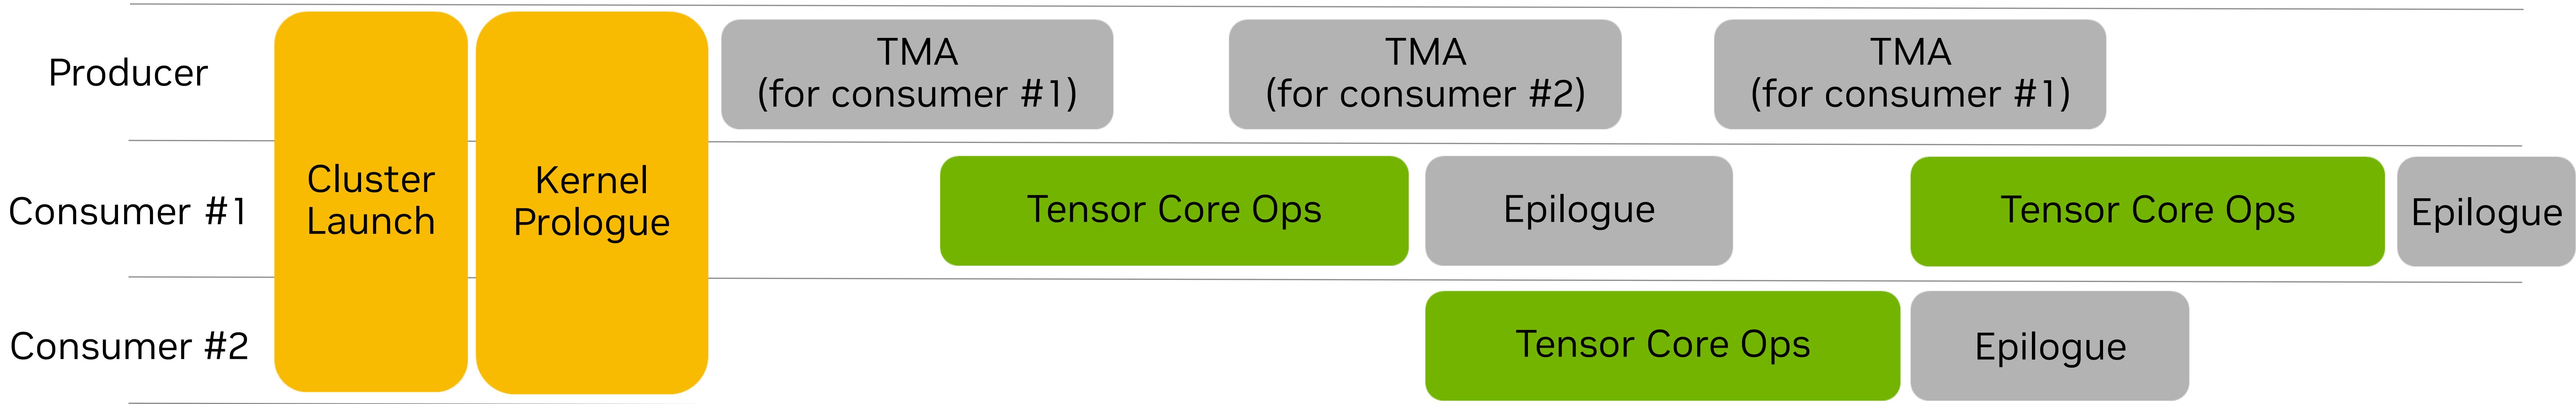
Life of a GEMM kernel



- Prologue and Epilogue are components of the GEMM kernel which involve non-tensor core operations and often latency or bandwidth bound.
- Typically hidden via multiple concurrently running ThreadBlock / SM
- With deep software pipelines – it becomes a tricky problem (due to lack of shared memory capacity)

Warp-Specialized - Persistent Ping-Pong Kernels

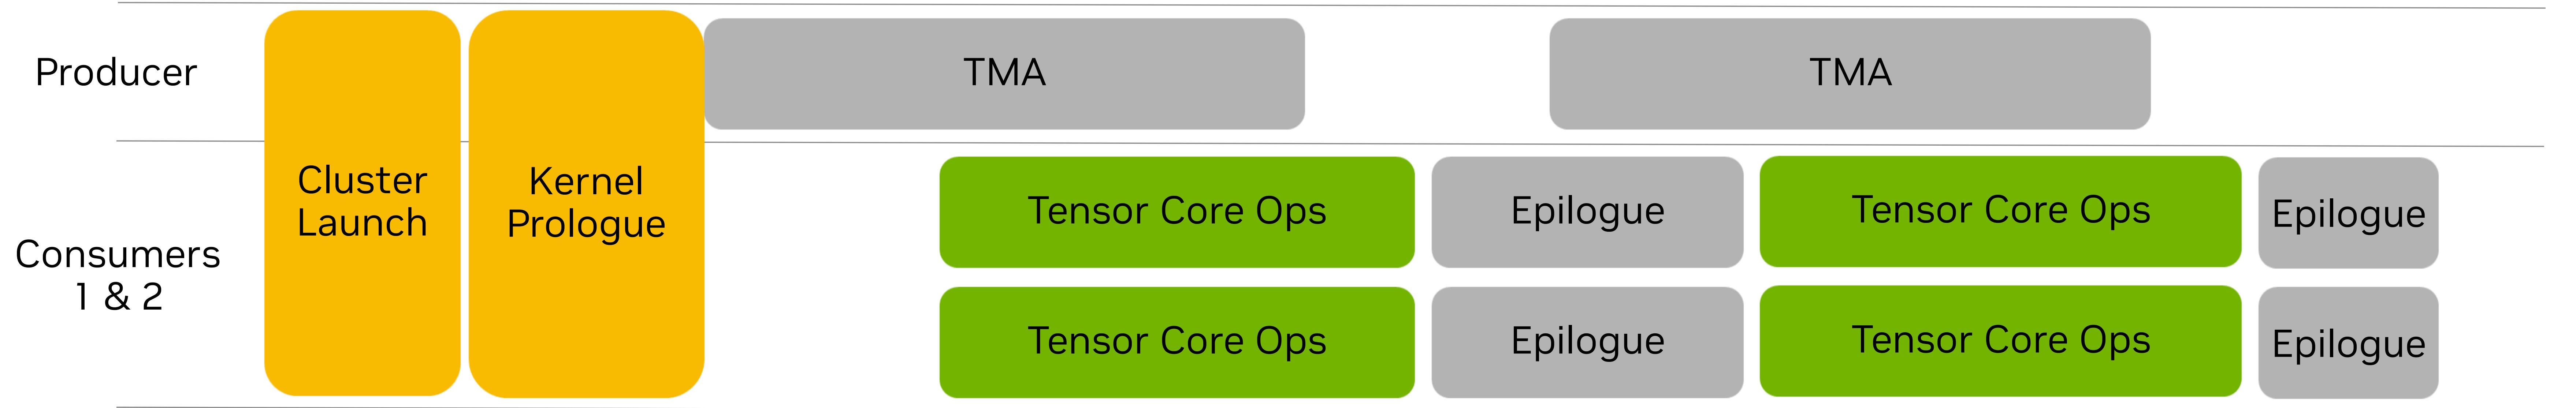
- 2 consumers warp-groups *work on two different work tiles* but ping-pong mainloop & epilogue
- Leads to guaranteed overlapping of epilogue of one consumer with the MMAs of the other
- Always keeps the tensor-cores busy – but at the cost of higher register pressure.



Producer	Consumer #1	Consumer #2
PersistentTileSchedulerSm90 scheduler(problem_shape, blk_shape, cluster_shape) // Data in via TMA while (work_tile_info.is_valid_tile) { collective_mainloop.dma() scheduler.advance_to_next_work() work_tile_info = scheduler.get_current_work() }	// Mainloop, epilogue, and data out while (work_tile_info.is_valid_tile) { collective_mainloop.mma() scheduler.advance_to_next_work(NumConsumers) work_tile_info = scheduler.get_current_work() }	

Warp-Specialized - Persistent Cooperative Kernels

- 2 Consumers each cooperatively each work on *their half of the same work tile*
- 2 consumer warp-groups run in concert and so math and epilogue sections execute ~together
- Reduced register pressure in each consumer – which can optionally be utilized for larger tile sizes



Producer	Consumer #1	Consumer #2
PersistentTileSchedulerSm90 scheduler(problem_shape, blk_shape, cluster_shape)		
// Data in via TMA while (work_tile_info. is_valid_tile) { collective_mainloop. dma () scheduler. advance_to_next_work () work_tile_info = scheduler. get_current_work () }	// Mainloop, epilogue, and data out while (work_tile_info. is_valid_tile) { collective_mainloop. mma () scheduler. advance_to_next_work (NumConsumers) work_tile_info = scheduler. get_current_work () }	

Anatomy of a High Performance Hopper GEMM

501 Level Optimizations

- Optimal ThreadBlock rasterization & swizzling
 - Encourages exploitation of locality
 - https://github.com/NVIDIA/cutlass/blob/main/media/docs/efficient_gemm.md#threadblock-rasterization
- Stream-K scheduling
 - Finding the optimal Trade-off between occupancy and efficiency
 - <https://arxiv.org/abs/2301.03598>
- Efficient Input transformations prior to MMA
 - Optimizing for register usage
 - Pipelined RS kernels
 - https://github.com/NVIDIA/cutlass/blob/main/include/cutlass/gemm/collective/sm90_mma_tma_gmma_rs_warp_specialized.hpp
- Optimal instruction sequence generation
 - Prefetching, Cache management

When to use which recipe?

*Caveat: For well behaved kernels

- Persistent kernels are great! Almost always a win
- “Small” K shape with large MN shape -> persistent ping pong schedule to hide prologue/epilogue overheads
 - Heavy epilogue fusions even for large K shape? Pingpong persistent
- “Large” K shape with large MN shape -> Cooperative persistent
 - Multistage persistent for even larger acc tiles
 - Naïve direct store epilogues to give mainloop even more smem capacity can be a win
- Small MN and large K shape -> cooperative with stream K load balancing
- HBM b/w bound MN shape? -> cleverer tile rasterization for L2 locality
- For fp8 training, follow the CUTLASS FP8 recipe that does higher precision accumulation
- For fp8 quant/dequant, see CUTLASS epilogues fusing amax and aux tensors

Other rules of thumb

- Design around 1 CTA / SM occupancy to start with – only go to higher as a last resort
- Cluster sizes of 2 are almost always a win. Start there with TMA multicast on the largest tile dim
- For larger tile sizes, prefer larger WGMMA N shapes over tiling over M
- Use larger swizzles for WGMMA and TMA if possible (a function of tile shapes)
- Source both A and B from smem if you can
- For non-16b input operands that are not K major, prefer the following in order:
 - Use WGMMA variants that source A from RMEM with or without swap AB only
 - Use WGMMA variants that source A from RMEM with swap AB *and* with pipelined and fused transpose for B
- For epilogue fusions, you have a lot more leeway
- For mainloop fusions
 - If possible, fuse into the epilogue of the previous kernel 
 - Use WGMMA variants that source A from RMEM
 - Avoid fusing things on WGMMA B input as that is a round trip to smem

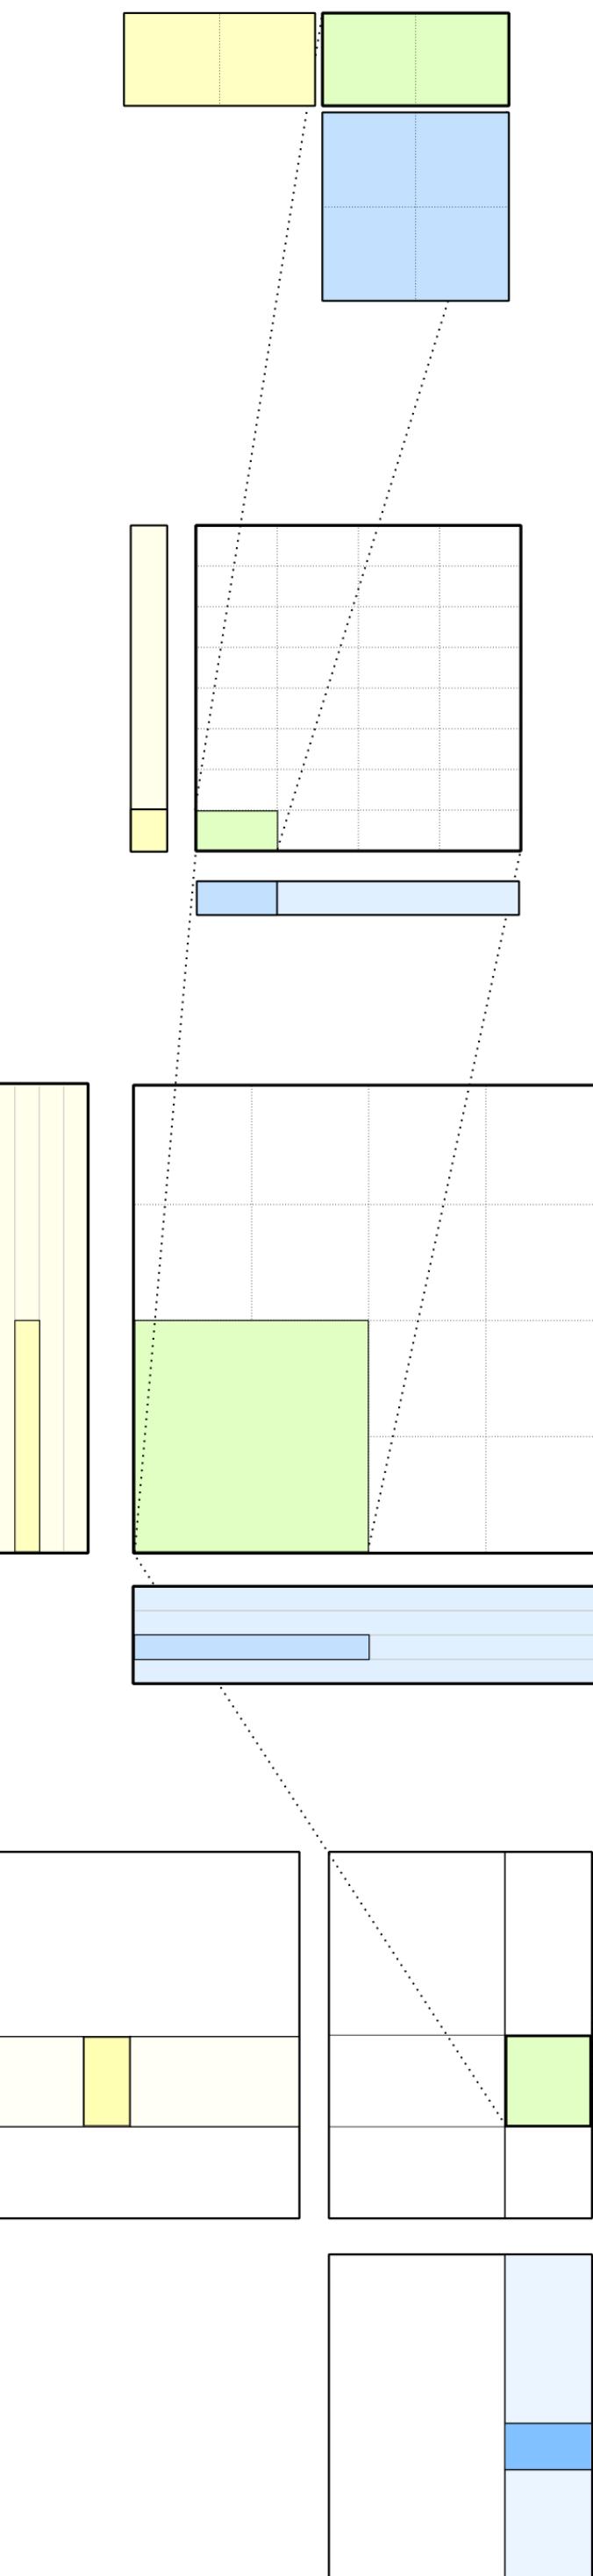
CUTLASS

CUDA C++ Template Library for High Performance Linear Algebra



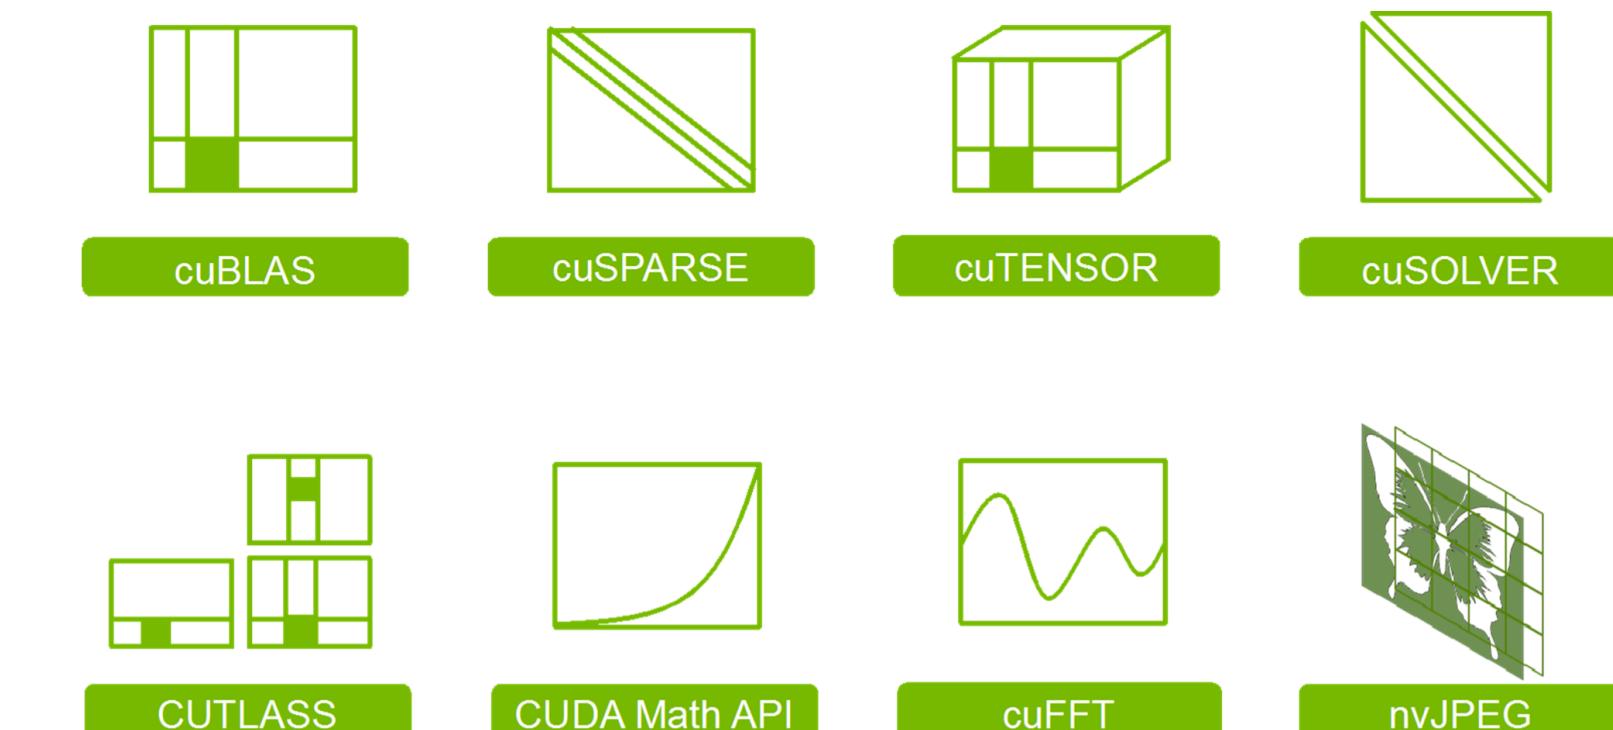
CUTLASS: tensor computations at all scopes and scales, decomposed into their “moving parts”

Device	{ GEMM, Convolution, Reductions , BLAS3 } x { all data types } x { SIMT, Tensor Cores } x { all architectures }
Kernel	GEMM, Batched GEMM, Convolution, Reduction, Fused output operations, Fused input operations
Collective	CUTLASS temporal micro-kernels (async producer/consumer pipelines orchestrating spatial micro-kernels)
Atom	CuTe spatial micro-kernels (Tiled MMA / Copy)
Thread	Numeric conversion, <functional> operators on arrays, complex<T>, fast math algorithms
Architecture intrinsic	Templates wrapping architecture-specific PTX instructions (e.g. mma, cp.async, Idmatrix, cvt)



Open source: <https://github.com/NVIDIA/cutlass>

- 4.7K stars, 3M clones/month, 100+ contributors, and many active users
- Latest revision: CUTLASS 3.5
- Documentation: <https://github.com/NVIDIA/cutlass#documentation>
- Presented: [GTC'18](#), [GTC'19](#), [GTC'20](#), [GTC'21](#), [GTC'22](#), [GTC'22](#), [GTC'23](#)





Thanks !

Q & A